

¹ **Structure-Preserving Generative Models:**
² **Extreme Values, Order Statistics, and**
³ **Diffusion Fine-Tuning**

⁴ December 18, 2025

5 Contents

6	List of Figures	vii
7	List of Tables	viii
8	Introduction	1
9	0.1 Context and Motivation	1
10	0.1.1 Structure in Statistical Data	1
11	0.1.2 Modern Generative Modeling	1
12	0.1.3 The Problem: When Structure Matters	2
13	0.1.4 This Thesis: Theory-Guided Generative Methods	2
14	0.2 Three Structural Regimes	3
15	0.2.1 Extreme Values and Tail Dependence	3
16	0.2.2 Order Statistics and Ranking Constraints	5
17	0.2.3 Reward-Tilted Distributions and Model Fine-Tuning	6
18	0.3 Generative Modeling Background	7
19	0.3.1 The Problem of Generative Modeling	7
20	0.3.2 Neural Networks	8
21	0.3.3 Neural Network Approximation Theory	9
22	0.3.4 Generative Adversarial Networks	9
23	0.3.5 Diffusion and Score-Based Models	10
24	0.4 State of the Art and Open Problems	11
25	0.4.1 Extreme Value Generation	11
26	0.4.2 Order Statistics Generation	12
27	0.4.3 Reward-Tilted Generation	13
28	0.4.4 Objectives of This Thesis	14
29	0.5 Contributions of This Thesis	14
30	0.5.1 Heavy-Tailed GANs for Extreme Value Generation	14
31	0.5.2 Generative Neural Order Statistics	15
32	0.5.3 Gradient-Free Fine-Tuning of Diffusion Models	16
33	0.6 Organization of the Manuscript	16
34	Introduction (Version française)	18
35	0.7 Contexte et motivation	18

36	0.7.1	La structure dans les données statistiques	18
37	0.7.2	La modélisation générative moderne	18
38	0.7.3	Le problème : quand la structure compte	19
39	0.7.4	Cette thèse : méthodes génératives guidées par la théorie	19
40	0.8	Trois régimes structurels	20
41	0.8.1	Valeurs extrêmes et dépendance de queue	20
42	0.8.2	Statistiques d'ordre et contraintes de classement	23
43	0.8.3	Distributions inclinées par récompense et ajustement fin de modèles	23
44	0.9	Contexte de la modélisation générative	25
45	0.9.1	Le problème de la modélisation générative	25
46	0.9.2	Réseaux de neurones	26
47	0.9.3	Théorie d'approximation des réseaux de neurones	26
48	0.9.4	Réseaux antagonistes génératifs	27
49	0.9.5	Modèles de diffusion et basés sur le score	28
50	0.10	État de l'art et problèmes ouverts	28
51	0.10.1	Génération de valeurs extrêmes	28
52	0.10.2	Génération de statistiques d'ordre	30
53	0.10.3	Génération inclinée par récompense	31
54	0.10.4	Objectifs de cette thèse	32
55	0.11	Contributions de cette thèse	32
56	0.11.1	GAN à queue lourde pour la génération de valeurs extrêmes	32
57	0.11.2	Statistiques d'ordre neurales génératives	33
58	0.11.3	Ajustement fin sans gradient des modèles de diffusion	34
59	0.12	Organisation du manuscrit	35
60	1	Heavy-Tailed GANs for Extreme Value Generation	36
61	1.1	Introduction	36
62	1.2	Reproducing dependence in extreme regions with Generative Adversarial Networks: Theory	40
63	1.2.1	Extreme-value theory	40
64	1.2.2	Measuring dependence in extremes	41
65	1.2.3	Generative Adversarial Networks	43
66	1.2.4	New approximation results for the stdf	44
67	1.3	Numerical experiments	45
68	1.3.1	Algorithms and implementation	46
69	1.3.2	Experiments on simulated data	48
70	1.3.3	Experiments on real data	56
71	1.4	Conclusion	58
72	Appendix	60
73	3.A	Theoretical background	60
74	3.B	Extreme-value theory	60
75	3.C	D-norms	61

77	3.D	Quantization	63
78	3.E	Copulas, measures of dependence	64
79	3.F	Proofs of claims	65
80	3.G	Proof of Proposition 1.4	66
81	3.H	Proof of Theorem 1.3	68
82	3.I	Additional experiments	70
83	3.J	Additional plots for experiments on Real data	70
84	2	Generative Neural Order Statistics	77
85	2.1	Introduction	77
86	3.A	Notations	80
87	2.2	A theoretical analysis of generative modeling of order statistics with Neural Networks	80
88	3.A	Representational results for order statistics	81
89	3.B	Generating order statistics: approximation results	81
90	2.3	Experiments	84
91	3.A	Synthetic data	84
92	3.B	Benchmarking on synthetic data	88
93	2.4	Conclusion	90
94	Appendix	92
95	2.A	Additional Figures	92
96	2.B	Some theory of concomitants	92
97	2.C	Concomitants of order statistics from a single distribution	95
98	2.D	Some elements on Functional approximation with Neural Network	98
99	2.E	Proofs	99
100	2.F	Proof of Theorem 2.2	102
101	2.G	Additional experimental details	105
102	2.H	Hyperparameters	106
103	2.I	Metrics	106
104	3	Gradient-Free Fine-Tuning of Diffusion Models	109
105	3.1	Introduction	109
106	3.2	Preliminaries on Denoising Diffusion Models and Fine-Tuning	110
107	3.A	Diffusion models	110
108	3.B	Fine-tuning formulations	113
109	3.3	Iterative Tilting Method for Fine-Tuning Diffusion Models	114
110	3.A	Score of the tilted distribution	115
111	3.B	Neural approximation and algorithm	117
112	3.4	Experiments	117
113	3.A	Gaussian data with quadratic rewards	117
114	3.B	Implementation and results	120
115	3.5	Conclusion	122

117	References	124
-----	-------------------	------------

¹¹⁸ List of Figures

119	1	Light-tailed vs heavy-tailed distributions	4
120	2	Tail dependence with Pareto margins	5
121	3	Reward tilting of a Gaussian distribution	7
122	4	Feedforward neural network architecture	8
123	5	GAN framework as adversarial game	9
124	6	Diffusion model framework on a Gaussian mixture	11
125	7	Distributions à queue légère vs à queue lourde	21
126	8	Dépendance de queue avec marges de Pareto	22
127	9	Inclinaison par récompense d'une distribution gaussienne	24
128	10	Architecture d'un réseau de neurones feedforward	26
129	11	Le cadre GAN comme jeu adversarial	27
130	12	Cadre des modèles de diffusion sur un mélange de gaussiennes	29
131	1.1	Illustration in a two-dimensional setting on simulated data of size $n_{\text{data}} = 10,000$ with Pareto margins and tail parameter $\alpha = 2.0$. Left panel: sample from a Gumbel copula with dependence parameter $\beta = 4/3$. Right panel: sample with two independent margins.	49
132	1.2	Illustration of real and generated data in a two-dimensional setting for data from a Gumbel copula with $\beta = 4/3$ and Pareto margins ($\alpha = 2.0$). For each subfigure (a)-(f), on the left: scatter plot of data points; on the right: histogram of angles $\arccos(X_1 / \sqrt{X_1^2 + X_2^2})$ for the 10% largest Euclidean norms. Subfigures (a)-(e): $n_{\text{gen}} = 50,000$ generated data after training with $n_{\text{data}} = 10,000$ points. Subfigure (f): real data.	50
133	1.3	Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} = 10,000$ from a Gumbel copula (with dependence parameter β) and Pareto margins (with tail parameter $\alpha = 1.5$). Each group of 2×3 plots corresponds to one experimental setting, <i>i.e.</i> a specification of α and β . Figures are averaged over dimensions $d \in \{2, 5, 10, 20, 50\}$. In each of these three groups: the first row plots the specified metric of HTGAN vs LTGAN in log scale for a given hyperparameter configuration. The second row is a histogram of the difference of the log of the metric for HTGAN vs LHTGAN for each parametrization. The legend corresponds to the proportion of cases where HTGAN performs better.	53

151	1.4 Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} =$	
152	10,000 from a Gumbel copula (with dependence parameter β) and Pareto	
153	margins (with tail parameter $\alpha = 2.5$), $d \in \{2, 5, 10, 20, 50\}$. See Figure 1.3	
154	for further details.	54
155	1.5 Estimated Kendall's τ for every pair of index considered. Left: every	
156	S&P500 ticker. On the top right: Tickers for the Financials sector. On	
157	the bottom right: Tickers for the Utilities sector. A comprehensive list of	
158	sectors and subsectors is given in Paragraph 1.3.3.	57
159	1.6 Results of the performance of HTGAN vs LTGAN on the Utilities ($d = 27$)	
160	and Financials subsectors ($d = 57$) of the S&P500 dataset. Results are	
161	aggregated with respect to the dimension d for plot. Each two groups of 2×3	
162	plots corresponds to one experimental setting, <i>i.e.</i> $\alpha = 1.5$ (top) and $\alpha = 2.5$	
163	(bottom). In each of these two groups: the first row plots the specified	
164	metric of HTGAN vs LTGAN in log scale for a given hyperparametrization	
165	configuration. The second row is a histogram of the difference of the log of	
166	the metric for HTGAN vs the metric for LTGAN for each parametrization.	
167	The legend corresponds to the proportion of cases where HTGAN performs	
168	better.	59
169	1.7 Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} =$	
170	10,000 from a Gaussian copula and Pareto margins, $d \in \{2, 5, 10, 20, 50\}$.	
171	Results obtained with the SWD metric. See Figure 1.3 for further details. .	71
172	1.8 Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} =$	
173	10,000 from a Gaussian copula and Pareto margins, $d \in \{2, 5, 10, 20, 50\}$.	
174	Results obtained with the AKE metric. See Figure 1.3 for further details. .	72
175	1.9 Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} =$	
176	10,000 from a Hüsler-ReiSS copula and Pareto margins, $d \in \{2, 5, 10, 20, 50\}$.	
177	Results obtained with the SWD metric. See Figure 1.3 for further details. .	73
178	1.10 Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} =$	
179	10,000 from a Hüsler-ReiSS copula and Pareto margins, $d \in \{2, 5, 10, 20, 50\}$.	
180	Results obtained with the AKE metric. See Figure 1.3 for further details. .	74
181	1.11 (a)-(d): histograms of real vs synthetic marginals across dimensions one to	
182	four. (e)-(f): Scatter plot AMF vs BEN and BK vs BLK tickers. (g)-(h):	
183	Histogram of θ after selecting largest values, at level 80% and 90%.	76
184	2.1 Real and generated data after min-max renormalization. <code>n_samples=1,000</code> . .	86
185	2.2 Comparison of real and generated 2D order statistics under different nor-	
186	malization schemes, with an additional overlay view. The red dashed lines	
187	indicate $x = y$; the mean-scale normalization preserves the sorted structure	
188	(no line crossing), while min-max normalization does not.	87
189	2.3 Meanscale comparison across dimensions	92

190	2.4 Generated vs. real means for order statistics in dimension 10. Every di-	
191	dimension is scattered against one another. Real data points are in blue and	
192	generated data points are in orange.	93
193	3.1 Iterative tilting on the 2-D GMM. Each row shows fixed- N settings (top to	
194	bottom: $N = 20, 50, 100, 200$) at iterations $N/2$ (left) and N (right). The	
195	rightmost plot in each row corresponds to the target distribution.	121
196	3.2 Evolution of the score error during iterative tilting for different N . We plot	
197	RMSE ($\sqrt{\text{MSE}}$) versus tilt iteration for $N \in \{20, 50, 100, 200\}$	122

¹⁹⁸ List of Tables

199	1	Tail probabilities comparison	4
200	2	Comparaison des probabilités de queue	21
201	1.1	Results on simulated data from a Gumbel copula. Percentage (%) of parametriza-	
202		tions for which HTGAN is better than a LTGAN for the six considered	
203		metrics. Results are given with precision $\pm 0.1\%$. Values are averaged over	
204		data dimensions $d \in \{2, 5, 10, 20, 50\}$	52
205	1.2	Results on simulated data from a Gumbel copula. Statistics (mean, min,	
206		max) on three on the top 5% best performing runs of HTGAN (best perform-	
207		ing w.r.t. the swd_90 metric). Statistics are averaged over $\beta \in \{4/3, 2, 4\}$	55
208	1.3	Results on simulated data from Hüsler-ReiSS and Gaussian copulas. Per-	
209		centage (%) of parametrizations for which HTGAN is better than a LTGAN	
210		for the six considered metrics. Results are given with precision $\pm 0.1\%$. Val-	
211		ues are averaged over data dimensions $d \in \{2, 5, 10, 20, 50\}$	56
212	2.1	Hyperparameters used for training our vanilla model.	86
213	2.2	Summary of metric performances for various models across dimensions.	89
214	3.1	Running times and metrics for iterative tilting on the 2-D GMM. Runtime	
215		columns report total sampling and training time in seconds. Metrics report	
216		negative log-likelihood and mean squared error on the mean.	123

²¹⁷ Introduction

²¹⁸ 0.1 Context and Motivation

²¹⁹ 0.1.1 Structure in Statistical Data

²²⁰ The term *statistics* derives from the German *Statistik*, coined by Gottfried Achenwall in
²²¹ 1748 to denote the “science of the state”: governing required organized information. Data
²²² is the raw material of statistics. A scientist measures temperature at several locations; a
²²³ bank records daily returns of assets in a portfolio; a hospital collects patient vital signs.
²²⁴ In each case, we have multiple observations, and each observation may consist of several
²²⁵ related quantities. These quantities are organized in specific ways, governed by underlying
²²⁶ regularities. Uncovering these regularities is the task of statistical science.

²²⁷ This process admits a clear mathematical formulation. We denote n observations by
²²⁸ X_1, \dots, X_n . Each X_i is a *vector*: a list of numbers representing the measured quantities.
²²⁹ We assume these observations are governed by an underlying law. Mathematically, this
²³⁰ law is a *probability distribution*, which we denote P . The distribution specifies how likely
²³¹ each possible outcome is. It is unknown; learning about it is the central problem.

²³² The goal of *statistics* is to infer properties of P from the observations: estimate its
²³³ parameters, test hypotheses about its form, quantify uncertainty. The goal of *generative*
²³⁴ *modeling* is more ambitious: to produce new synthetic data that could plausibly have come
²³⁵ from the same source, indistinguishable from genuine observations.

²³⁶ But data carries *structure*: dependencies between variables, constraints on their values,
²³⁷ properties that define the phenomenon itself. A faithful generative model must respect this
²³⁸ structure, not merely approximate the overall law.

²³⁹ 0.1.2 Modern Generative Modeling

²⁴⁰ Given data, can we learn to generate more of it? Not just compute averages, but produce
²⁴¹ new data that look as if they came from the same source? This is the question of *generative*
²⁴² *modeling*.

²⁴³ When observations have many components (temperatures at hundreds of locations,
²⁴⁴ returns of thousands of assets), directly describing the law P becomes impractical. Modern
²⁴⁵ generative modeling takes a different path: rather than describing P explicitly, we learn

²⁴⁶ to produce data from it. The core idea is:

$$G(Z) \sim P.$$

²⁴⁷ Here Z is a source of randomness, like rolling dice or drawing from a bell curve. The
²⁴⁸ transformation G turns this randomness into something that resembles the data. All the
²⁴⁹ complexity of the unknown law P is encoded in G .

²⁵⁰ This echoes a classical idea. A statistician assumes the data follow a known family
²⁵¹ of laws (bell curves, exponentials, etc.) and estimates a few numbers to pin down which
²⁵² one. A generative modeler makes a similar bet: instead of assuming the law has a certain
²⁵³ shape, we assume the transformation G does. When G is a neural network, we bet that
²⁵⁴ transforming randomness through layers of computations can mimic the data.

²⁵⁵ The breakthrough of the past decade is that neural networks can learn G from examples,
²⁵⁶ producing images indistinguishable from photographs, text that reads as human-written.
²⁵⁷ But the catch: neural networks learn G by minimizing a global measure of error. This
²⁵⁸ process has no mechanism to preserve specific structural properties.

²⁵⁹ 0.1.3 The Problem: When Structure Matters

²⁶⁰ What does it mean for synthetic data to be *faithful* to the phenomenon it imitates?

²⁶¹ Consider a bank generating fake crisis scenarios to test its investments. The generated
²⁶² scenarios look realistic on average, yet the worst losses occur in isolation: when stocks
²⁶³ crash, bonds stay calm. The synthetic data capture each variable separately but miss how
²⁶⁴ they move together. In a real crisis, losses are correlated. Or consider generating synthetic
²⁶⁵ rankings: the third-best item sometimes scores higher than the first-best. The ranking is
²⁶⁶ scrambled; the simulation is useless.

²⁶⁷ These failures reflect a mismatch between what generative models optimize and what
²⁶⁸ applications require. Standard models reward getting the overall shape right but provide
²⁶⁹ no guarantees on specific structural properties.

²⁷⁰ 0.1.4 This Thesis: Theory-Guided Generative Methods

²⁷¹ This manuscript starts from real-world problems where structure matters. We identify the
²⁷² key property that must be preserved (joint crashes, rankings, preferences), formulate it as
²⁷³ a mathematical object, and analyze how well neural networks can approximate it. This is
²⁷⁴ the core methodology: from applied problem, to mathematical abstraction, to theoretical
²⁷⁵ analysis.

²⁷⁶ We study three structural regimes: tail dependence, order statistics, and reward tilting.
²⁷⁷ For each, we design generative methods that *provably* respect the structure, providing
²⁷⁸ guarantees on approximation accuracy.

279 0.2 Three Structural Regimes

280 We now describe the three structural regimes addressed in this thesis. For each regime, we
 281 present the real-world problem that motivates our work and formalize it mathematically.

282 0.2.1 Extreme Values and Tail Dependence

283 **The problem: when extremes drive everything.** In many domains, extreme events
 284 dominate outcomes despite their rarity. In portfolio management, a handful of large losses
 285 can wipe out years of steady gains: the distribution of returns is heavy-tailed, and risk
 286 is concentrated in the extremes (Rama Cont 2001). Value-at-Risk, stress testing, and
 287 regulatory capital requirements all hinge on the probability of rare but catastrophic losses.
 288 The 2008 financial crisis demonstrated that models calibrated to normal conditions fail
 289 when extremes occur (A. J. McNeil et al. 2015).

290 In insurance, extreme events are the core business. Reinsurers must price the risk of
 291 simultaneous catastrophes: a hurricane hitting Florida while an earthquake strikes Califor-
 292 nia. The joint occurrence of such events, not their individual probabilities, determines
 293 solvency requirements. In climate science, compound extremes (concurrent heatwaves and
 294 droughts, sequential flooding events) pose the greatest risks to ecosystems and infrastruc-
 295 ture (Jakob Zscheischler et al. 2018; Bevacqua et al. 2021). Climate models must capture
 296 not just the frequency of individual extremes, but their tendency to cluster in space and
 297 time.

298 The common thread is *tail dependence*: the tendency for extremes to occur together.
 299 Two random variables can have zero correlation under normal conditions yet exhibit strong
 300 tail dependence: when one takes an extreme value, the other is likely to as well (see
 301 Figure 2). For a portfolio manager, this means diversification fails precisely when it is
 302 needed most. For an insurer, it means correlated claims. For climate risk, it means
 303 compound disasters.

304 **Two distinct challenges.** Generating realistic extreme value data poses two fundamen-
 305 tally different challenges.

306 The first is *marginal heavy-tailedness*. A distribution is heavy-tailed when extreme
 307 values are far more likely than light-tailed models would predict. Figure 1 compares three
 308 distributions: the absolute Gaussian $|X|$, the exponential (both light-tailed), and the
 309 Pareto (heavy-tailed). Near zero they look similar, but their tails behave very differently.
 310 Table 1 shows the probability of exceeding various thresholds. At $t = 8$, the Gaussian
 311 assigns probability 10^{-15} , the exponential 10^{-4} , but the Pareto still gives 3.7%. Heavy-
 312 tailed phenomena produce extreme events that light-tailed models dismiss as virtually
 313 impossible.

314 The second challenge is *extremal dependence*: when multiple variables are involved,
 315 their joint behavior in the tails can be entirely different from their dependence in normal
 316 conditions. Figure 2 illustrates this: two samples with identical marginal distributions

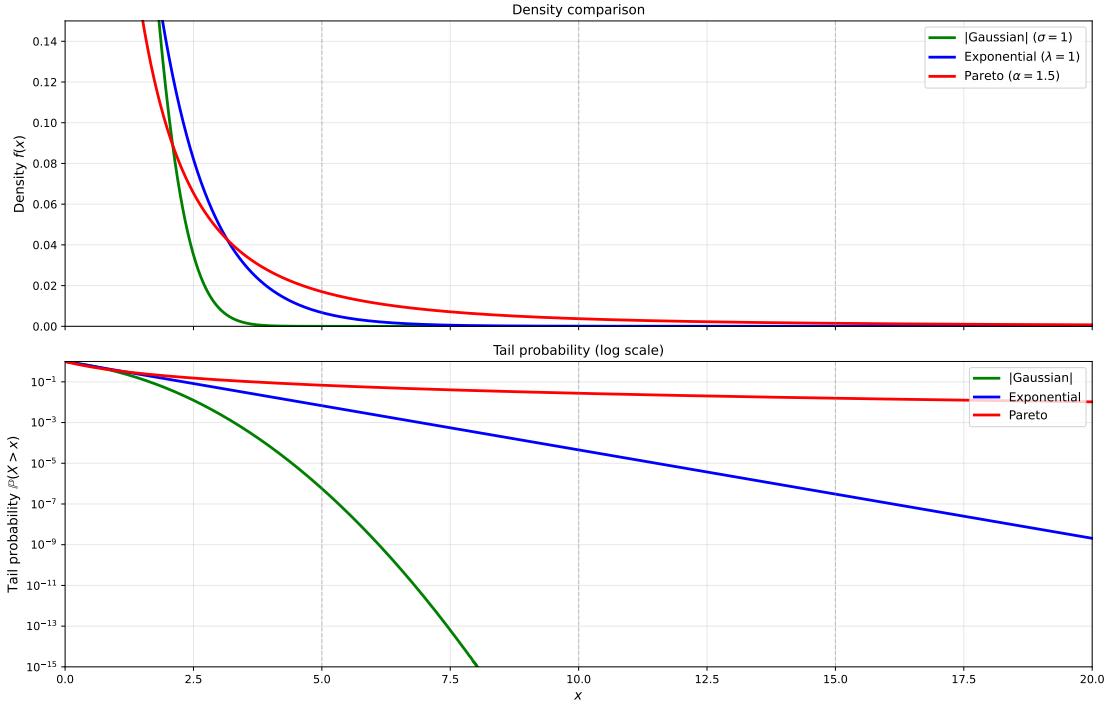


Figure 1: Light-tailed vs heavy-tailed distributions. Left: density comparison for $|Gaussian|$, exponential, and Pareto. Right: tail probability $\mathbb{P}(X > t)$ on log scale. The Gaussian and exponential tails decay rapidly; the Pareto tail decays as a power law, remaining substantial even for large t .

Table 1: Probability of exceeding threshold t for $|Gaussian|$ ($\sigma = 1$), exponential ($\lambda = 1$), and Pareto ($\alpha = 1.5$).

t	$ Gaussian $	Exponential	Pareto
3	2.7×10^{-3}	5.0×10^{-2}	1.3×10^{-1}
5	5.7×10^{-7}	6.7×10^{-3}	6.8×10^{-2}
8	1.3×10^{-15}	3.4×10^{-4}	3.7×10^{-2}

317 behave very differently in the tails depending on whether extremes occur together or in
318 independently.

319 As we shall see in Section 0.4.1, standard generative models fail on both counts.

320 **Mathematical formulation: the stable tail dependence function.** In classical
321 statistics, *correlation* measures how two variables move together on average. Correlation
322 summarizes dependence across the entire distribution but does not fully characterize tail
323 behavior. Two variables can have zero correlation yet crash together in a crisis; conversely,
324 variables with the same correlation can exhibit very different tail dependence structures.

325 The *stable tail dependence function* (stdf) plays for extremes the role that correlation
326 plays for the bulk of the distribution. Just as correlation captures average co-movement,
327 the stdf captures extreme co-movement: when one variable takes a very large value, how
328 likely is the other to be large as well? For a random vector $\mathbf{X} = (X_1, \dots, X_d)$, the stdf ℓ

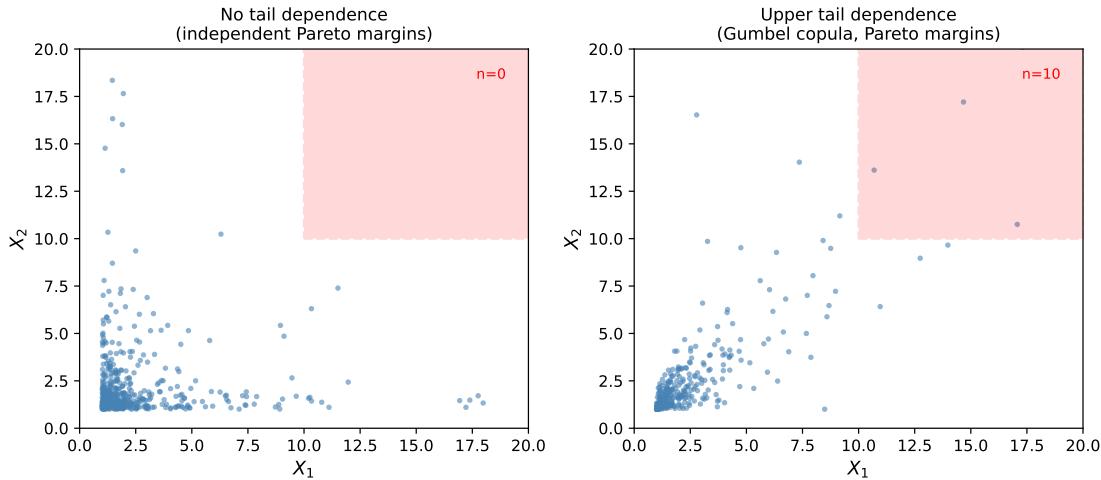


Figure 2: Tail dependence illustrated with Pareto margins ($\alpha = 1.5$). Left: independent margins; extreme values of X_1 and X_2 occur independently, so the upper tail region (red box) is sparse. Right: Gumbel copula ($\theta = 2.5$) with upper tail dependence; when X_1 is extreme, X_2 tends to be extreme as well, creating clustering in the upper-right corner. Both samples have identical marginal distributions, yet behave very differently in the tails.

329 is defined by

$$\ell(\mathbf{x}) = \lim_{t \rightarrow \infty} t \mathbb{P} \left(\bigcup_{j=1}^d \{X_j > tx_j\} \right). \quad (1)$$

330 This limit captures the asymptotic probability of joint exceedances as we move deeper into
 331 the tail. The stdf completely characterizes extremal dependence: two distributions with
 332 the same stdf exhibit identical behavior in extremes, regardless of how they behave in the
 333 bulk. Knowing the stdf is essential for understanding diversification when it matters most:
 334 during extreme events.

335 0.2.2 Order Statistics and Ranking Constraints

336 **The problem: when ranking defines the data.** Many phenomena are defined by
 337 their rank structure. In auctions, what matters is the distribution of the highest bid, or
 338 the gap between the two highest bids. In sports, we care about the performance of the
 339 top-ranked player, not the average. In extreme value statistics, the k largest observations
 340 are used to estimate tail behavior. In each case, the ordering is not incidental: it is the
 341 defining structure of the data.

342 Consider impact investing, where assets are ranked by environmental, social, or gover-
 343 nance (ESG) scores (Lo and R. Zhang 2021). A fund manager backtesting an ESG strategy
 344 needs to simulate the returns of the best-ranked asset, the second-best, and so on. These
 345 are *induced order statistics*: the returns $\Theta_{[1:N]}, \dots, \Theta_{[N:N]}$ associated with assets ranked
 346 by their ESG scores $X_{1:N} \leq \dots \leq X_{N:N}$. Synthetic data must preserve both the marginal
 347 distribution at each rank and the joint dependence across ranks.

348 **Mathematical formulation: order statistics and their joint density.** Given a
349 sample X_1, \dots, X_n from a continuous distribution F with density f , the *order statistics*
350 are the sorted values $X_{1:n} \leq X_{2:n} \leq \dots \leq X_{n:n}$. Their joint density is

$$f_{1:n, \dots, n:n}(x_1, \dots, x_n) = n! \prod_{i=1}^n f(x_i) \cdot \mathbf{1}\{x_1 \leq \dots \leq x_n\}. \quad (2)$$

351 The factor $n!$ counts the number of arrangements, and the indicator $\mathbf{1}\{\cdot\}$ enforces the
352 ordering constraint. Classical representations, such as those of Sukhatme and Schucany,
353 decompose order statistics into tractable building blocks that have enabled exact simula-
354 tion for decades. As we shall see in Section 0.5.2, these representations also provide the
355 foundation for neural network approximation with provable complexity bounds.

356 0.2.3 Reward-Tilted Distributions and Model Fine-Tuning

357 **The problem: aligning generative models with human preferences.** A pre-
358 trained generative model produces samples from a distribution learned from training data.
359 But for many applications, we want samples that are not merely typical, but *good* according
360 to some criterion: images that are aesthetically pleasing, text that is helpful and harm-
361 less, molecules that bind to a target protein. This is the *alignment problem*: adjusting a
362 generative model to favor high-quality outputs while maintaining diversity and coherence.

363 Consider a text-to-image diffusion model trained on web-scraped images. The model
364 generates realistic images, but many are poorly composed, contain artifacts, or depict
365 undesirable content. We have access to a reward model (trained on human preferences)
366 that scores images by aesthetic quality and safety. The goal is to fine-tune the diffusion
367 model so that it preferentially generates high-reward images, without collapsing to a single
368 mode or forgetting how to generate diverse content.

369 **Mathematical formulation: exponential tilting.** The natural mathematical formu-
370 lation is *exponential tilting*. Given a base distribution p_0 (the pre-trained model) and a
371 reward function $r : \mathcal{X} \rightarrow \mathbb{R}$, the target is the *tilted distribution*

$$p(x) \propto p_0(x) \exp(\lambda r(x)), \quad (3)$$

372 where $\lambda > 0$ controls the strength of the tilting. This distribution is the unique solution
373 to the KL-regularized reward maximization problem:

$$p = \arg \max_q \left\{ \mathbb{E}_q[r(x)] - \frac{1}{\lambda} \text{KL}(q \| p_0) \right\}.$$

374 The tilted distribution balances reward maximization against staying close to the pre-
375 trained model, with λ controlling the trade-off. Large λ favors high reward; small λ stays
376 close to p_0 .

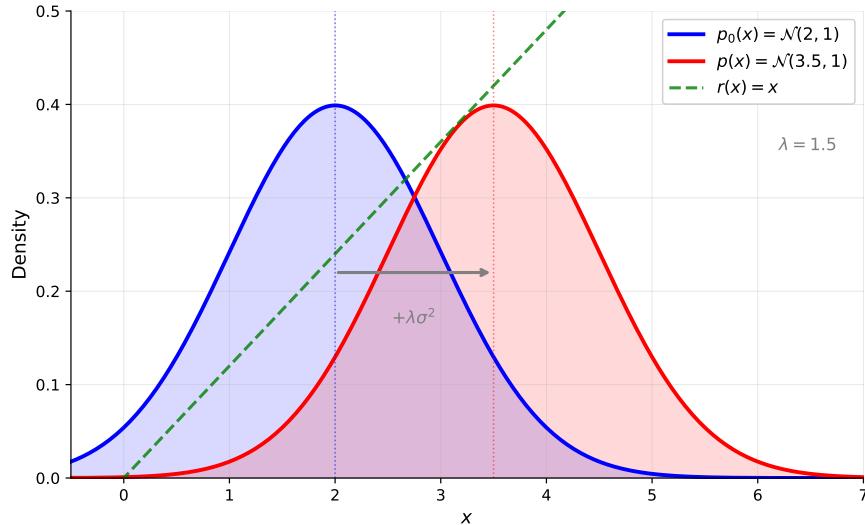


Figure 3: Reward tilting shifts the distribution toward high-reward regions. With $p_0 = \mathcal{N}(2, 1)$, linear reward $r(x) = x$, and tilting strength $\lambda = 1.5$, the tilted distribution $p(x) \propto p_0(x) \exp(\lambda r(x))$ is $\mathcal{N}(3.5, 1)$: the mean shifts by $\lambda\sigma^2 = 1.5$ toward high reward, while the variance is unchanged.

³⁷⁷ 0.3 Generative Modeling Background

³⁷⁸ Section 0.1 introduced the goal of learning from data while respecting structure. Section 0.2
³⁷⁹ presented three regimes where structure matters: tail dependence, order statistics, and
³⁸⁰ reward tilting. To address these challenges, we turn to generative modeling: the task
³⁸¹ of producing synthetic data that follows the same law as the observations. This section
³⁸² provides the technical background on generative models and neural network approximation
³⁸³ theory necessary to understand the contributions of this thesis.

³⁸⁴ 0.3.1 The Problem of Generative Modeling

³⁸⁵ The fundamental problem of generative modeling can be stated as follows: given samples
³⁸⁶ X_1, \dots, X_n drawn independently from an unknown distribution P on \mathbb{R}^d , construct a
³⁸⁷ procedure for generating new samples that are distributed according to P , or a close
³⁸⁸ approximation thereof.

³⁸⁹ This problem has a long history in statistics and probability. Classical approaches
³⁹⁰ include *density estimation* (Silverman 1986; Scott 2015), where one estimates the density
³⁹¹ $p = dP/d\lambda$ and then samples from the estimated density, and *Markov Chain Monte*
³⁹² *Carlo* (MCMC) (Metropolis et al. 1953; Hastings 1970; C. P. Robert and Casella 2004),
³⁹³ where one constructs a Markov chain with P as its stationary distribution and samples
³⁹⁴ by running the chain to equilibrium. Both approaches face the curse of dimensionality
³⁹⁵ (Bellman 1961): density estimation becomes statistically intractable in high dimensions,
³⁹⁶ while MCMC mixing times can grow exponentially with dimension.

³⁹⁷ The modern approach to generative modeling, enabled by deep learning, takes a dif-
³⁹⁸ ferent path. Rather than explicitly estimating the density, we learn a *transport map* G

399 that transforms samples from a simple reference distribution (typically standard Gaussian)
 400 into samples from the target distribution. If Z is Gaussian noise, then $G(Z)$ should be
 401 distributed approximately as P . The map G is parameterized by a neural network, and
 402 the parameters are learned from data.

403 This transport-based view unifies several modern generative paradigms. In *normalizing*
 404 *flows* (Rezende and Mohamed 2015; Dinh et al. 2017), the map G is constrained to be in-
 405 vertible with tractable Jacobian. In *variational autoencoders* (VAEs) (Diederik P. Kingma
 406 and Max Welling 2014), the map is learned jointly with an approximate inverse. In *gen-*
 407 *erative adversarial networks* (GANs) (I. J. Goodfellow et al. 2014), the map is learned via
 408 an adversarial objective. In *diffusion models* (Ho, Jain, et al. 2020; Y. Song, Jascha Sohl-
 409 Dickstein, et al. 2021), the map is implicitly defined by the reverse of a noising process.
 410 We describe GANs and diffusion models in detail below, as they are the foundations of our
 411 contributions.

412 0.3.2 Neural Networks

413 A *feedforward neural network* with L hidden layers transforms an input x through succe-
 414 sive layers (Figure 4). Denoting $h_0 = x$, the hidden representations are

$$h_\ell = \sigma(W_\ell h_{\ell-1} + b_\ell) \quad \text{for } \ell = 1, 2, \dots, L, \quad (4)$$

415 where W_ℓ is a weight matrix, b_ℓ a bias vector, and σ a nonlinear activation function applied
 416 coordinate-wise. The output is $g(x) = W_{L+1}h_L + b_{L+1}$. Common activations include the
 417 sigmoid, tanh, and ReLU $\sigma(t) = \max(0, t)$, which dominates modern practice.

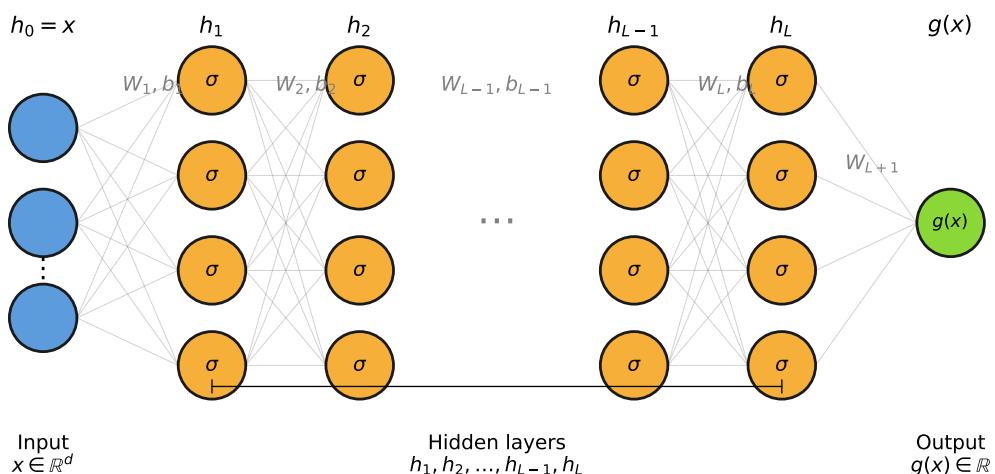


Figure 4: A feedforward neural network. The input $h_0 = x$ passes through L hidden layers, each applying an affine transformation followed by activation σ . The output $g(x)$ is a linear function of h_L .

418 0.3.3 Neural Network Approximation Theory

419 The *universal approximation theorem* (Cybenko 1989; Hornik et al. 1989) states that neural
 420 networks with a single hidden layer can approximate any continuous function on a compact
 421 domain to arbitrary accuracy, given sufficient width. This extends to ReLU networks
 422 (Leshno et al. 1993).

423 **Quantitative bounds.** Quantitative bounds relate approximation error to network size
 424 (Barron 1993; Yarotsky 2017). For Lipschitz functions on $[0, 1]^d$, achieving error ε requires
 425 $O(\varepsilon^{-d})$ parameters, exhibiting the curse of dimensionality. For smoother functions (Hölder
 426 class with smoothness s), networks with $O(\varepsilon^{-d/s})$ parameters suffice (Yarotsky 2017).

427 **Limitations.** These theorems concern pointwise function approximation, not preserva-
 428 tion of distributional structure. A network approximating a transport map G uniformly
 429 may still fail to preserve tail behavior, ordering constraints, or dependence structure. More-
 430 over, the theorems assume optimal parameters are known; in practice, gradient descent
 431 may not find them. These limitations motivate the structure-specific theory developed in
 432 this thesis.

433 0.3.4 Generative Adversarial Networks

434 Generative Adversarial Networks (I. J. Goodfellow et al. 2014) frame generative modeling
 435 as a two-player game. The intuition is best understood through an analogy: imagine a
 436 counterfeiter trying to produce fake banknotes, and a detective trying to detect them.
 437 The counterfeiter improves by studying which fakes get caught; the detective improves
 438 by studying the counterfeiter’s latest techniques. Over time, both become more sophisti-
 439 cated, until eventually the counterfeiter produces notes so convincing that even the expert
 440 detective cannot tell them apart from genuine currency.

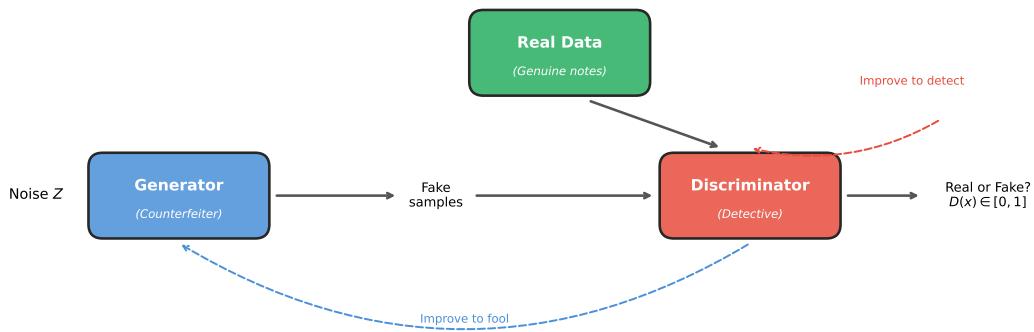


Figure 5: The GAN framework as an adversarial game. The generator (counterfeiter) transforms noise into fake samples; the discriminator (detective) tries to distinguish fakes from real data. Both networks improve through competition until equilibrium, when generated samples are indistinguishable from real data.

441 In mathematical terms, a *generator* $G_\theta : \mathcal{Z} \rightarrow \mathcal{X}$ maps latent noise $Z \sim p_Z$ to
442 synthetic samples (the counterfeiter’s output). A *discriminator* $D_\phi : \mathcal{X} \rightarrow [0, 1]$ attempts
443 to distinguish real samples from generated ones (the detective’s verdict). The networks
444 are trained by solving a minimax problem:

$$\min_{\theta} \max_{\phi} \mathbb{E}_{X \sim p_{\text{data}}} [\log D_\phi(X)] + \mathbb{E}_{Z \sim p_Z} [\log(1 - D_\phi(G_\theta(Z)))].$$

445 The generator minimizes this objective (trying to fool the discriminator), while the dis-
446 criminator maximizes it (trying to correctly classify real vs. fake). At equilibrium, the
447 generator produces samples from the data distribution, and the discriminator outputs 1/2
448 for all inputs, unable to distinguish real from fake.

449 The Wasserstein GAN (WGAN) (Arjovsky et al. 2017) replaces the Jensen-Shannon
450 divergence with the Wasserstein-1 distance:

$$W_1(p_{\text{data}}, p_\theta) = \sup_{\|f\|_L \leq 1} \mathbb{E}_{X \sim p_{\text{data}}} [f(X)] - \mathbb{E}_{Z \sim p_Z} [f(G_\theta(Z))],$$

451 where the supremum is over 1-Lipschitz functions. The WGAN with gradient penalty
452 (WGAN-GP) (Gulrajani et al. 2017) enforces the Lipschitz constraint via a penalty term,
453 improving training stability.

454 0.3.5 Diffusion and Score-Based Models

455 Diffusion models (Ho, Jain, et al. 2020; Y. Song, Jascha Sohl-Dickstein, et al. 2021) define
456 a generative process by constructing a continuous path $(p_t)_{t \in [0, 1]}$ between the data distri-
457 bution p_0 and a tractable reference $p_1 = N(0, I_d)$. The path is defined via the *forward*
458 *interpolation*

$$X_t = \alpha_t X_0 + \sigma_t X_1, \quad X_0 \sim p_0, \quad X_1 \sim N(0, I_d), \quad (5)$$

459 where (α_t, σ_t) are deterministic schedules with $(\alpha_0, \sigma_0) = (1, 0)$ and $(\alpha_1, \sigma_1) = (0, 1)$. At
460 $t = 0$, we have clean data; at $t = 1$, pure noise. Standard choices include the variance-
461 preserving schedule ($\alpha_t^2 + \sigma_t^2 = 1$) and the linear schedule ($\alpha_t = 1 - t$, $\sigma_t = t$).

462 To generate samples, we reverse this process. The key quantity is the *score function*
463 $\nabla_x \log p_t$, which points toward regions of high density at each noise level. For the Gaussian
464 forward kernel, the score satisfies $\nabla_x \log p_t(x_t) = -\hat{x}_1(x_t, t)/\sigma_t$, where $\hat{x}_1(x_t, t) = \mathbb{E}[X_1 |$
465 $X_t = x_t]$ is the conditional expectation of the noise given the noised sample.

466 The score is learned via *denoising score matching* (Hyvärinen 2005; Vincent 2011): a
467 network $s_\theta(x, t)$ is trained to predict the noise X_1 from the noised sample X_t , which is
468 equivalent to learning $\nabla_x \log p_t$. Starting from $X_1 \sim N(0, I_d)$ and iteratively denoising
469 using the learned score yields samples from p_0 .

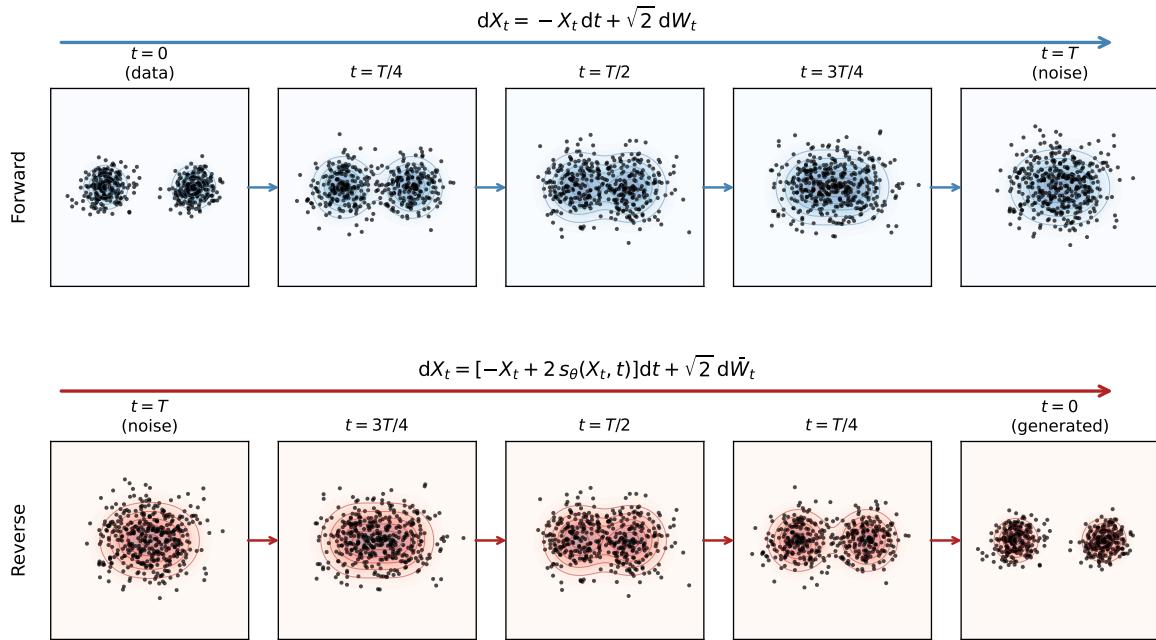


Figure 6: The diffusion model framework on a 2D Gaussian mixture. Forward process (top row): data from p_0 is progressively corrupted by adding noise until it becomes indistinguishable from the Gaussian reference p_1 at $t = 1$. Reverse process (bottom row): starting from pure noise, a learned score network s_θ guides the denoising, gradually recovering the bimodal structure of p_0 at $t = 0$.

470 0.4 State of the Art and Open Problems

471 We now review the state of the art for each structural regime, identify the gaps in the
 472 existing literature, and formulate the objectives of this thesis.

473 0.4.1 Extreme Value Generation

474 **Classical extreme value theory.** The statistical theory of multivariate extremes is
 475 well-developed (Paul Embrechts et al. 1997; L. d. Haan and Ferreira 2006; Sidney I. Resnick
 476 1987). For univariate extremes, the Fisher-Tippett-Gnedenko theorem establishes that
 477 properly normalized maxima converge to one of three distributions: Gumbel, Fréchet, or
 478 Weibull, unified by the Generalized Extreme Value (GEV) distribution. In the multivari-
 479 ate setting, the stable tail dependence function (stdf) provides a complete characterization
 480 of extremal dependence, describing how the probability of joint exceedances scales as one
 481 moves into the tail. Numerous parametric models have been proposed: the Gumbel (logis-
 482 tic) copula (Gumbel 1960), the Hüsler-Reiss model arising from Gaussian processes (Hüsler
 483 and Reiss 1989a), and the extremal- t model generalizing the Hüsler-Reiss to Student- t
 484 margins (Nikoloulopoulos et al. 2009). Estimation methods include empirical likelihood
 485 (Einmahl, Krajina, et al. 2012), rank-based M-estimation (Einmahl, L. d. Haan, et al.
 486 2016), and Bayesian approaches (Sabourin et al. 2013).

487 **Why standard GANs fail.** Standard generative models are fundamentally incompatible
488 with heavy-tailed data. The reason is mathematical: (Wiese et al. 2019) proved that
489 if the latent distribution has light tails (e.g., Gaussian) and the generator G is Lipschitz
490 continuous, then the generated distribution $G(Z)$ necessarily has light tails. Specifically,
491 if Z is sub-Gaussian, then $G(Z)$ is sub-Gaussian for any Lipschitz G . Since modern GANs
492 enforce Lipschitz constraints for training stability (spectral normalization, gradient pen-
493 alties), they *provably cannot* generate heavy-tailed outputs from Gaussian noise. This is not
494 a failure of optimization or architecture; it is a fundamental impossibility.

495 **Heavy-tailed generative models.** Recent work has addressed this limitation through
496 several approaches (Michaël Allouche, Emmanuel Gobet, et al. 2024). For GANs: *prepro-
497 cessing methods* such as Quant-GAN (Wiese et al. 2019) and evtGAN (Youssef Boulaguiem
498 et al. 2022) transform data to remove tail heaviness before training; *heavy-tailed latent
499 methods* (Huster et al. 2021) replace Gaussian noise with Pareto distributions; *EVT-
500 parameterized* approaches such as EV-GAN (Michaël Allouche, Girard, et al. 2022) design
501 generators that explicitly encode extreme-value structure. For diffusion models, heavy-
502 tailed extensions replace the Gaussian noise process with α -stable Lévy processes (Yoon
503 et al. 2023), allowing direct generation of heavy-tailed samples. However, all these ap-
504 proaches focus on *marginal* tail behavior. None addresses the *dependence structure* in the
505 tails, characterized by the stdf.

506 **Gap 1.** To our knowledge, no existing generative models provide provable guarantees
507 on tail dependence approximation. Existing methods address marginal tail behavior but
508 ignore the joint structure. We address this gap in Chapter 1, developing a GAN architecture
509 that provably approximates any stdf.

510 0.4.2 Order Statistics Generation

511 The mathematical formulation of order statistics was given in Section 0.2.2. Classical
512 representations (Sukhatme, Schucany) have enabled exact simulation for decades (H. A.
513 David and H. N. Nagaraja 2003; Arnold et al. 2008). We now review the state of the art.

514 **Machine learning approaches.** The machine learning literature has addressed sorting
515 from the perspective of learning-to-rank (T.-Y. Liu 2009), where the goal is to predict
516 orderings rather than generate them. Differentiable sorting operators have been developed
517 to enable gradient-based learning: (Grover et al. 2019) proposed NeuralSort, a continuous
518 relaxation of the argsort operator based on the Sinkhorn algorithm; (Cuturi et al. 2019)
519 developed optimal transport-based differentiable sorting; (Blondel et al. 2020) introduced
520 fast differentiable sorting with $O(n \log n)$ complexity. These methods enable sorting within
521 neural network pipelines but are designed for ranking tasks, not distribution matching.

522 **Generative models and ordering.** Standard generative models (GANs, VAEs, diffu-
523 sion) produce unordered samples. When applied to ranked data, they face two problems.

524 First, generated samples are not sorted: a GAN trained on order statistics will produce
 525 vectors where coordinates are not monotonically ordered. Post-hoc sorting can restore
 526 order but introduces bias in the distribution. Second, standard preprocessing destroys
 527 ordering structure: per-coordinate normalization to $[0, 1]$, common in GAN training, nor-
 528 malizes each coordinate independently based on its marginal range, destroying the relative
 529 ordering. If the k -th order statistic is normalized by its own mean and variance, and the
 530 $(k + 1)$ -th by different values, there is no guarantee that normalized values remain sorted.

531 **Gap 2.** To our knowledge, no existing generative models produce order statistics with
 532 both correct marginal/joint distributions and guaranteed ordering constraints. Existing
 533 differentiable sorting methods address ranking prediction, not generative modeling. We
 534 address this gap in Chapter 2, developing an architecture that exploits classical represen-
 535 tations to guarantee sortedness while matching target distributions.

536 0.4.3 Reward-Tilted Generation

537 The alignment problem (Section 0.2.3) requires sampling from the tilted distribution (3).
 538 We review existing approaches.

539 **Reinforcement learning from human feedback.** The dominant paradigm for align-
 540 ment is reinforcement learning from human feedback (RLHF) (Christiano et al. 2017; Bai
 541 et al. 2022; Ouyang et al. 2022). A reward model is trained on human preference data
 542 (pairwise comparisons of outputs), then the generative model is fine-tuned to maximize
 543 expected reward using policy gradient methods (PPO, REINFORCE). RLHF has been
 544 transformative for large language models, enabling ChatGPT, Claude, and other conver-
 545 sational agents. However, RLHF is notoriously unstable: reward hacking, mode collapse,
 546 and catastrophic forgetting are common failure modes (Casper et al. 2023).

547 **Diffusion model fine-tuning.** For diffusion models, several fine-tuning approaches have
 548 emerged. DRaFT (Clark et al. 2024) backpropagates through the full denoising chain to
 549 maximize reward, but this requires storing activations across all timesteps, with memory
 550 scaling linearly in chain length. DPOK (Fan et al. 2023) combines policy gradients with
 551 KL regularization. ReFL (Jing Xu et al. 2024) uses reward feedback learning to fine-tune
 552 text-to-image diffusion models. Adjoint matching (Domingo-Enrich, Drozdzał, et al. 2025)
 553 provides a theoretically principled approach based on adjoint SDEs, achieving state-of-the-
 554 art results but still requiring reward gradients.

555 **The gradient bottleneck.** Gradient-based fine-tuning methods (DRaFT, adjoint match-
 556 ing) require backpropagating through the reward function, which creates practical limita-
 557 tions: storing activations for large reward models leads to prohibitive memory costs, and
 558 non-differentiable rewards (human feedback, discrete evaluations, simulator-based metrics)
 559 cannot be handled. RL-based approaches (RLHF, DPOK) avoid this by using only reward
 560 values, not gradients, but introduce other challenges: policy gradient estimators have high

561 variance, and careful hyperparameter tuning is needed to balance reward maximization
 562 against mode collapse. Neither approach provides distributional guarantees on the result-
 563 ing fine-tuned model.

564 **Gap 3.** To our knowledge, no existing gradient-free methods fine-tune diffusion models
 565 to sample from reward-tilted distributions. Existing methods either require reward gra-
 566 dients or use reinforcement learning with high-variance estimators. We address this gap
 567 in Chapter 3, developing a method based on Fisher’s identity that requires only forward
 568 evaluations of the reward function.

569 0.4.4 Objectives of This Thesis

570 Based on the identified gaps, we formulate three objectives:

- 571 1. **Extreme value generation.** Develop a generative model for multivariate extremes
 572 with provable guarantees on stdf approximation. The method should use heavy-
 573 tailed latent noise and provide explicit bounds on the approximation error in terms
 574 of network architecture.
- 575 2. **Order statistics generation.** Develop a generative model for order statistics that
 576 guarantees sortedness while matching the target distribution. The method should ex-
 577 ploit classical representations and provide approximation bounds that scale favorably
 578 with dimension.
- 579 3. **Reward-tilted generation.** Develop a gradient-free method for fine-tuning diffu-
 580 sion models to reward-tilted distributions. The method should require only forward
 581 evaluations of the reward function.

582 0.5 Contributions of This Thesis

583 We now summarize the main contributions, organized by the three objectives.

584 0.5.1 Heavy-Tailed GANs for Extreme Value Generation

585 We develop Heavy-Tailed GANs (HTGAN) for generating data with correct tail depen-
 586 dence.

587 **Theoretical contributions.** We replace Gaussian latent noise with heavy-tailed (Pareto
 588 or Fréchet) noise. Because the generator is Lipschitz, the heavy-tail property propagates
 589 to the output. We prove that the stdf of the generated distribution is always discrete, a
 590 maximum of linear functions with finitely many atoms. Our main approximation theorem
 591 establishes that any stdf can be approximated within L^∞ error ε using latent dimension
 592 $N = O(\varepsilon^{-(d-1)})$ and a single hidden layer with ReLU activations. This appears to be the
 593 first approximation guarantee for tail dependence in the generative modeling literature.

594 **Methodological contributions.** The HTGAN algorithm proceeds in three steps: marginal
 595 transformation to unit Fréchet, GAN training with Pareto latent noise, and inverse trans-
 596 formation. This separation of margins and dependence is standard in copula modeling.

597 **Experimental contributions.** We evaluate on synthetic Gumbel copula data and S&P
 598 500 returns, demonstrating that HTGAN correctly captures tail dependence while standard
 599 GANs fail by orders of magnitude.

600 0.5.2 Generative Neural Order Statistics

601 We develop Generative Neural Order Statistics (GENOS) for generating ranked data with
 602 guaranteed ordering.

603 **Theoretical contributions.** We combine classical statistics representations with neu-
 604 ral network approximation theory (Yarotsky 2017) to derive non-asymptotic complexity
 605 bounds. Specifically, we adapt two classical representations of order statistics and analyze
 606 their neural network approximation complexity.

607 *Schucany's representation* (Schucany 1972) constructs order statistics recursively:

$$U_{n-r:n} = U_{n-r+1:n} \cdot U_r^{1/(n-r)}, \quad r = 1, \dots, n-1. \quad (6)$$

608 This requires n different transition functions. We prove that neural approximation of this
 609 scheme requires a number of parameters that is polynomial of degree n in ε^{-1} .

610 *Sukhatme's representation* (Sukhatme 1937) expresses uniform order statistics via ex-
 611 ponential spacings:

$$U_{k:n} = F \left(\sum_{j=1}^k \frac{E_j}{n-j+1} \right), \quad E_j \stackrel{\text{iid}}{\sim} \text{Exp}(1), \quad F(z) = 1 - e^{-z}. \quad (7)$$

612 The cumulative sum has an exponential distribution; applying the exponential CDF F
 613 yields the k -th uniform order statistic $U_{k:n}$. This requires approximating only two universal
 614 functions (F and F^{-1}), independent of n . We prove that the Sukhatme-based approach
 615 achieves

$$\mathbb{E}|U_{k:n} - \hat{U}_{k:n}| \leq O(\varepsilon \log n) \quad \text{with} \quad \mathcal{C}_{\text{Sukhatme}}(\varepsilon) = O(\varepsilon^{-2} \log(1/\varepsilon)) \quad (8)$$

616 parameters. The complexity is quadratic in ε^{-1} , independent of n , versus polynomial of
 617 degree n for Schucany. This makes the Sukhatme-based approach scalable to large sample
 618 sizes.

619 **Methodological contributions.** We introduce global mean-scale normalization, which
 620 preserves ordering unlike per-coordinate normalization. We also propose an order statistics

621 penalty (OSP) that penalizes violations of the ordering constraint:

$$\mathcal{L}_{\text{OSP}} = \mathbb{E}_Z \left[\frac{1}{n-1} \sum_{k=1}^{n-1} (\max\{G(Z)_k - G(Z)_{k+1}, 0\})^2 \right].$$

622 **Experimental contributions.** We evaluate on synthetic order statistics with varying
 623 dimension, comparing GAN, WGAN, and WGAN with OSP. Adding the OSP penalty
 624 improves sortedness metrics while maintaining competitive performance on Wasserstein
 625 distance.

626 0.5.3 Gradient-Free Fine-Tuning of Diffusion Models

627 We develop gradient-free methods for fine-tuning diffusion models to sample from reward-
 628 tilted distributions.

629 **Theoretical contributions.** Fisher’s identity expresses the score of the tilted distribu-
 630 tion as a covariance:

$$\nabla_x \log p(x) = \nabla_x \log p_0(x) + \lambda \cdot \text{Cov}_{q_{0|x}} [\nabla_{x_0} \log q(x_0|x), r(x_0)]. \quad (9)$$

631 This covariance can be estimated via Monte Carlo using only forward evaluations of r ,
 632 without backpropagation through the reward model.

633 **Methodological contributions.** We propose an iterative algorithm that decomposes
 634 a large tilt into N small steps, estimating the score increment at each step via Fisher’s
 635 identity. The algorithm applies even when r is non-differentiable or a black-box.

636 **Experimental contributions.** On 2D Gaussian mixtures with known ground truth, we
 637 validate that the method correctly recovers the tilted score and that increasing the number
 638 of steps N improves accuracy.

639 0.6 Organization of the Manuscript

640 Chapter 1 develops the theory and algorithms for heavy-tailed generative modeling. We
 641 present mathematical background on multivariate extreme value theory and the stdf, an-
 642 alyze why Lipschitz maps of Gaussian noise cannot produce heavy tails, characterize the
 643 discrete stdf induced by generators with Pareto noise, prove our main approximation the-
 644 orem, and present experiments on synthetic and financial data.

645 Chapter 2 develops generative models for order statistics. We present background
 646 on order statistics and the Sukhatme and Schucany representations, prove approximation
 647 bounds for ReLU networks, introduce global mean-scale normalization and the WGAN-
 648 OSP objective, and present experiments on synthetic order statistics.

649 Chapter 3 develops gradient-free fine-tuning for diffusion models. We present back-
650 ground on exponential tilting and existing gradient-based methods, derive the Fisher
651 identity-based score estimator, analyze the approximation error, present the iterative tilt-
652 ing algorithm, and validate on controlled experiments.

653 This thesis is partially based on collaborative work: Chapter 1 on heavy-tailed gen-
654 erative models, Chapter 2 on neural approximation of order statistics, and Chapter 3 on
655 gradient-free diffusion fine-tuning. Specific venues and co-authors are indicated at the
656 beginning of each chapter.

657 **Introduction (Version 658 française)**

659 **0.7 Contexte et motivation**

660 **0.7.1 La structure dans les données statistiques**

661 Le terme *statistique* dérive de l'allemand *Statistik*, forgé par Gottfried Achenwall en 1748
662 pour désigner la « science de l'État » : gouverner nécessitait une information organisée. Les
663 données sont la matière première de la statistique. Un scientifique mesure la température
664 en plusieurs endroits ; une banque enregistre les rendements journaliers des actifs d'un por-
665 tefeuille ; un hôpital collecte les signes vitaux des patients. Dans chaque cas, nous disposons
666 de plusieurs observations, et chaque observation peut comprendre plusieurs quantités liées
667 entre elles. Ces quantités sont organisées de manières spécifiques, régies par des régularités
668 sous-jacentes. Découvrir ces régularités est la tâche de la science statistique. Ce processus
669 admet une formulation mathématique claire. Nous notons n observations par X_1, \dots, X_n .
670 Chaque X_i est un *vecteur* : c'est à dire une liste de nombres représentant les quantités
671 mesurées. Nous supposons que ces observations sont gouvernées par une loi sous-jacente.
672 Mathématiquement, cette loi est une *distribution de probabilité*, que nous notons P . La dis-
673 tribution spécifie la probabilité de chaque résultat possible. Elle est inconnue ; apprendre
674 à la connaître est le problème central.

675 L'objectif de la *statistique* est d'inférer des propriétés de P à partir des observations :
676 estimer ses paramètres, tester des hypothèses sur sa forme, quantifier l'incertitude. L'ob-
677 jectif de la *modélisation générative* est plus ambitieux : produire de nouvelles données
678 synthétiques qui auraient plausiblement pu provenir de la même source, indiscernables des
679 observations authentiques.

680 Mais les données portent une *structure* : des dépendances entre variables, des contraintes
681 sur leurs valeurs, des propriétés qui définissent le phénomène lui-même. Un modèle génératif
682 fidèle doit respecter cette structure, et non simplement approximer la loi globale.

683 **0.7.2 La modélisation générative moderne**

684 À partir de données, pouvons-nous apprendre à en générer davantage ? Non pas simplement
685 calculer des moyennes, mais produire de nouvelles données qui semblent provenir de la

686 même source ? C'est la question de la *modélisation générative*.

687 Lorsque les observations comportent de nombreuses composantes (températures en des
688 centaines de lieux, rendements de milliers d'actifs), décrire directement la loi P devient
689 impraticable. La modélisation générative moderne emprunte une voie différente : plutôt
690 que de décrire P explicitement, nous apprenons à produire des données qui en sont issues.
691 L'idée centrale est :

$$G(Z) \sim P.$$

692 Ici, Z est une source d'aléa, comme lancer des dés ou tirer selon une courbe en cloche. La
693 transformation G convertit cet aléa en quelque chose qui ressemble aux données. Toute la
694 complexité de la loi inconnue P est encodée dans G .

695 Cela fait écho à une idée classique. Un statisticien suppose que les données suivent
696 une famille connue de lois (courbes en cloche, exponentielles, etc.) et estime quelques
697 paramètres pour déterminer laquelle. Un modélisateur génératif fait un pari similaire : au
698 lieu de supposer que la loi a une certaine forme, nous supposons que la transformation G
699 en a une. Quand G est un réseau de neurones, nous parions que transformer de l'aléa à
700 travers des couches de calculs peut imiter les données.

701 La percée de la dernière décennie est que les réseaux de neurones peuvent apprendre
702 G à partir d'exemples, produisant des images indiscernables de photographies, du texte
703 qui se lit comme écrit par un humain. Mais le revers de la médaille : les réseaux de neu-
704 rones apprennent G en minimisant une mesure globale d'erreur. Ce processus n'a pas de
705 mécanisme pour préserver des propriétés structurelles spécifiques.

706 0.7.3 Le problème : quand la structure compte

707 Que signifie pour des données synthétiques d'être *fidèles* au phénomène qu'elles imitent ?

708 Considérons une banque générant de faux scénarios de crise pour tester ses investis-
709 sements. Les scénarios générés semblent réalistes en moyenne, mais les pires pertes sur-
710 viennent isolément : quand les actions s'effondrent, les obligations restent calmes. Les
711 données synthétiques capturent chaque variable séparément mais manquent la façon dont
712 elles évoluent ensemble. Dans une vraie crise, les pertes sont corrélées. Ou considérons la
713 génération de classements synthétiques : le troisième meilleur élément obtient parfois un
714 score supérieur au premier. Le classement est brouillé ; la simulation est inutile.

715 Ces échecs reflètent un décalage entre ce que les modèles génératifs optimisent et ce que
716 les applications requièrent. Les modèles standards récompensent l'obtention de la forme
717 globale correcte mais ne fournissent aucune garantie sur des propriétés structurelles spéci-
718 fiques.

719 0.7.4 Cette thèse : méthodes génératives guidées par la théorie

720 Ce manuscrit part de problèmes réels où la structure compte. Nous identifions la propriété
721 clé qui doit être préservée (effondrements conjoints, classements, préférences), la formulons
722 comme un objet mathématique, et analysons dans quelle mesure les réseaux de neurones

⁷²³ peuvent l'approximer. C'est la méthodologie centrale : du problème appliqué, à l'abstraction
⁷²⁴ mathématique, à l'analyse théorique.

⁷²⁵ Nous étudions trois régimes structurels : la dépendance de queue, les statistiques d'ordre
⁷²⁶ et l'inclinaison par récompense. Pour chacun, nous concevons des méthodes génératives qui
⁷²⁷ respectent *prouvablement* la structure, fournissant des garanties sur la précision d'approxi-
⁷²⁸ mation.

⁷²⁹ 0.8 Trois régimes structurels

⁷³⁰ Nous décrivons maintenant les trois régimes structurels abordés dans cette thèse. Pour
⁷³¹ chaque régime, nous présentons le problème réel qui motive notre travail et le formalisons
⁷³² mathématiquement.

⁷³³ 0.8.1 Valeurs extrêmes et dépendance de queue

⁷³⁴ **Le problème : quand les extrêmes gouvernent tout.** Dans de nombreux domaines,
⁷³⁵ les événements extrêmes dominent les résultats malgré leur rareté. En gestion de por-
⁷³⁶ tefeuille, une poignée de grandes pertes peut effacer des années de gains réguliers : la
⁷³⁷ distribution des rendements est à queue lourde, et le risque est concentré dans les extrêmes
⁷³⁸ (Rama CONT 2001). La Value-at-Risk, les tests de stress et les exigences réglementaires en
⁷³⁹ capital reposent tous sur la probabilité de pertes rares mais catastrophiques. La crise fi-
⁷⁴⁰ nancière de 2008 a démontré que les modèles calibrés sur des conditions normales échouent
⁷⁴¹ lorsque des extrêmes surviennent (A. J. MCNEIL et al. 2015).

⁷⁴² En assurance, les événements extrêmes sont le cœur de métier. Les réassureurs doivent
⁷⁴³ tarifer le risque de catastrophes simultanées : un ouragan frappant la Floride pendant
⁷⁴⁴ qu'un tremblement de terre frappe la Californie. C'est l'occurrence conjointe de tels évé-
⁷⁴⁵ nements, et non leurs probabilités individuelles, qui détermine les exigences de solvabilité.
⁷⁴⁶ En science du climat, les extrêmes composés (vagues de chaleur et sécheresses simultanées,
⁷⁴⁷ événements d'inondation séquentiels) posent les plus grands risques aux écosystèmes et aux
⁷⁴⁸ infrastructures (Jakob ZSCHEISCHLER et al. 2018 ; BEVACQUA et al. 2021). Les modèles
⁷⁴⁹ climatiques doivent capturer non seulement la fréquence des extrêmes individuels, mais
⁷⁵⁰ leur tendance à se regrouper dans l'espace et le temps.

⁷⁵¹ Le fil conducteur est la *dépendance de queue* : la tendance des extrêmes à survenir
⁷⁵² ensemble. Deux variables aléatoires peuvent avoir une corrélation nulle dans des conditions
⁷⁵³ normales tout en exhibant une forte dépendance de queue : quand l'une prend une valeur
⁷⁵⁴ extrême, l'autre a tendance à en faire autant (voir la Figure 8). Pour un gestionnaire de
⁷⁵⁵ portefeuille, cela signifie que la diversification échoue précisément quand elle est le plus
⁷⁵⁶ nécessaire. Pour un assureur, cela signifie des sinistres corrélés. Pour le risque climatique,
⁷⁵⁷ cela signifie des catastrophes composées.

⁷⁵⁸ **Deux défis distincts.** Générer des données de valeurs extrêmes réalistes pose deux défis
⁷⁵⁹ fondamentalement différents.

760 Le premier est la *lourdeur marginale des queues*. Une distribution est à queue lourde
 761 quand les valeurs extrêmes sont bien plus probables que ce que les modèles à queue légère
 762 prédiraient. La Figure 7 compare trois distributions : la gaussienne absolue $|X|$, l'exponen-
 763 tielle (toutes deux à queue légère) et la Pareto (à queue lourde). Près de zéro, elles se
 764 ressemblent, mais leurs queues se comportent très différemment. Le Tableau 2 montre la
 765 probabilité de dépasser différents seuils. À $t = 8$, la gaussienne attribue une probabilité de
 766 10^{-15} , l'exponentielle 10^{-4} , mais la Pareto donne encore 3,7%. Les phénomènes à queue
 767 lourde produisent des événements extrêmes que les modèles à queue légère considèrent
 768 comme virtuellement impossibles.

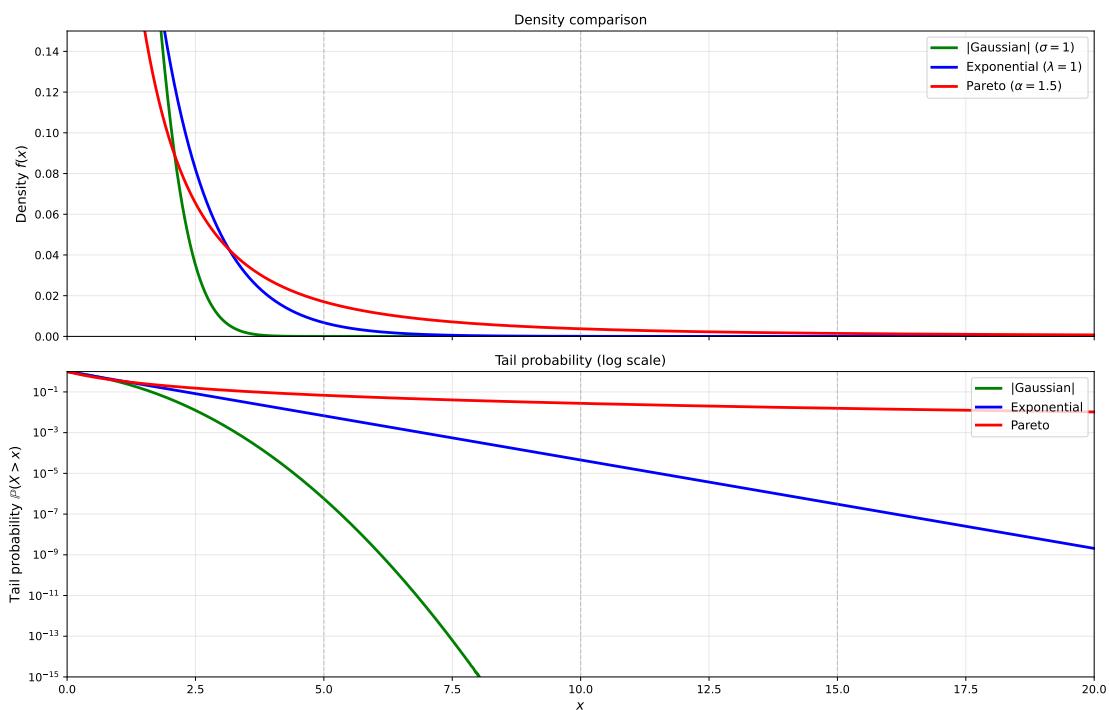


FIGURE 7 : Distributions à queue légère vs à queue lourde. Gauche : comparaison des densités pour |Gaussienne|, exponentielle et Pareto. Droite : probabilité de queue $\mathbb{P}(X > t)$ en échelle logarithmique. Les queues gaussienne et exponentielle décroissent rapidement ; la queue de Pareto décroît comme une loi de puissance, restant substantielle même pour de grandes valeurs de t .

TABLE 2 : Probabilité de dépasser le seuil t pour |Gaussienne| ($\sigma = 1$), exponentielle ($\lambda = 1$) et Pareto ($\alpha = 1,5$).

t	Gaussienne	Exponentielle	Pareto
3	$2,7 \times 10^{-3}$	$5,0 \times 10^{-2}$	$1,3 \times 10^{-1}$
5	$5,7 \times 10^{-7}$	$6,7 \times 10^{-3}$	$6,8 \times 10^{-2}$
8	$1,3 \times 10^{-15}$	$3,4 \times 10^{-4}$	$3,7 \times 10^{-2}$

769 Le second défi est la *dépendance extrémale* : quand plusieurs variables sont impliquées,
 770 leur comportement conjoint dans les queues peut être entièrement différent de leur dépen-
 771 dance dans des conditions normales. La Figure 8 illustre ceci : deux échantillons avec des

⁷⁷² distributions marginales identiques se comportent très différemment dans les queues selon
⁷⁷³ que les extrêmes surviennent ensemble ou indépendamment.

⁷⁷⁴ Comme nous le verrons dans la Section 0.10.1, les modèles génératifs standards échouent
⁷⁷⁵ sur ces deux aspects.

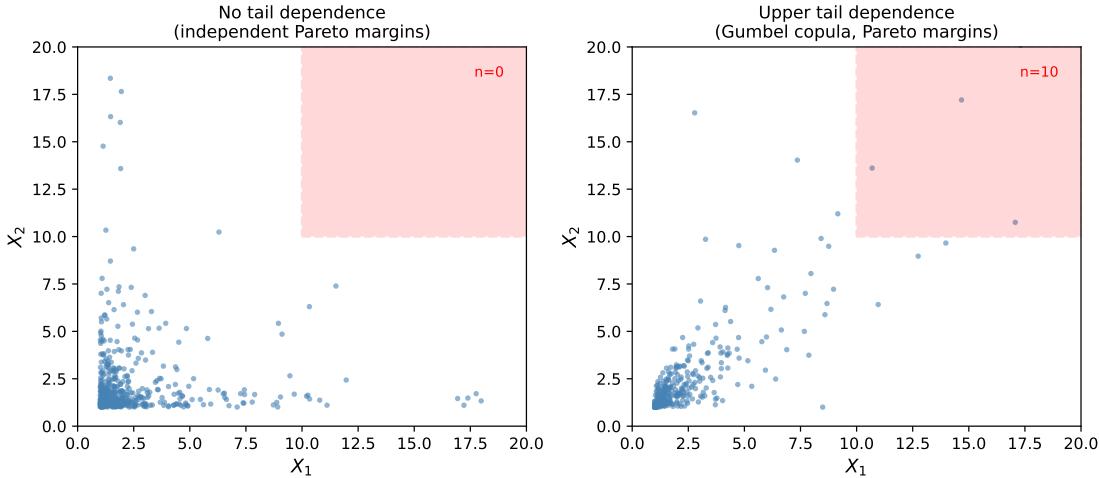


FIGURE 8 : Dépendance de queue illustrée avec des marges de Pareto ($\alpha = 1,5$). Gauche : marges indépendantes ; les valeurs extrêmes de X_1 et X_2 surviennent indépendamment, donc la région de queue supérieure (boîte rouge) est clairsemée. Droite : copule de Gumbel ($\theta = 2,5$) avec dépendance de queue supérieure ; quand X_1 est extrême, X_2 tend à être extrême également, créant un regroupement dans le coin supérieur droit. Les deux échantillons ont des distributions marginales identiques, mais se comportent très différemment dans les queues.

⁷⁷⁶ **Formulation mathématique : la fonction de dépendance de queue stable.** En
⁷⁷⁷ statistique classique, la *corrélation* mesure comment deux variables évoluent ensemble en
⁷⁷⁸ moyenne. La corrélation résume la dépendance sur l'ensemble de la distribution mais ne
⁷⁷⁹ caractérise pas complètement le comportement des queues. Deux variables peuvent avoir
⁷⁸⁰ une corrélation nulle tout en s'effondrant ensemble dans une crise ; inversement, des va-
⁷⁸¹ riabiles avec la même corrélation peuvent exposer des structures de dépendance de queue
⁷⁸² très différentes.

⁷⁸³ La *fonction de dépendance de queue stable* (stdf, de l'anglais *stable tail dependence*
⁷⁸⁴ *function*) joue pour les extrêmes le rôle que la corrélation joue pour le corps de la dis-
⁷⁸⁵ tribution. Tout comme la corrélation capture le co-mouvement moyen, la stdf capture le
⁷⁸⁶ co-mouvement extrême : quand une variable prend une très grande valeur, quelle est la pro-
⁷⁸⁷ babilité que l'autre soit grande également ? Pour un vecteur aléatoire $\mathbf{X} = (X_1, \dots, X_d)$,
⁷⁸⁸ la stdf ℓ est définie par

$$\ell(\mathbf{x}) = \lim_{t \rightarrow \infty} t \mathbb{P} \left(\bigcup_{j=1}^d \{X_j > tx_j\} \right). \quad (10)$$

⁷⁸⁹ Cette limite capture la probabilité asymptotique des dépassements conjoints lorsque l'on
⁷⁹⁰ s'enfonce dans la queue. La stdf caractérise complètement la dépendance extrémale : deux
⁷⁹¹ distributions avec la même stdf exhibent un comportement identique dans les extrêmes,

792 indépendamment de leur comportement dans le corps de la distribution. Connaître la stdf
793 est essentiel pour comprendre la diversification quand elle compte le plus : pendant les
794 événements extrêmes.

795 0.8.2 Statistiques d'ordre et contraintes de classement

796 **Le problème : quand le classement définit les données.** De nombreux phénomènes
797 sont définis par leur structure de rang. Dans les enchères, ce qui compte est la distribution
798 de l'offre la plus élevée, ou l'écart entre les deux offres les plus élevées. En sport, nous nous
799 intéressons à la performance du joueur le mieux classé, pas à la moyenne. En statistique
800 des valeurs extrêmes, les k plus grandes observations sont utilisées pour estimer le com-
801 portement des queues. Dans chaque cas, l'ordonnancement n'est pas accessoire : c'est la
802 structure définissante des données.

803 Considérons l'investissement à impact, où les actifs sont classés par scores environne-
804 mentaux, sociaux ou de gouvernance (ESG) (LO et R. ZHANG 2021). Un gestionnaire de
805 fonds testant rétrospectivement une stratégie ESG doit simuler les rendements du meilleur
806 actif classé, du deuxième meilleur, et ainsi de suite. Ce sont des *statistiques d'ordre induites* : les rendements $\Theta_{[1:N]}, \dots, \Theta_{[N:N]}$ associés aux actifs classés par leurs scores ESG
807 $X_{1:N} \leq \dots \leq X_{N:N}$. Les données synthétiques doivent préserver à la fois la distribution
808 marginale à chaque rang et la dépendance conjointe entre les rangs.

810 **Formulation mathématique : statistiques d'ordre et leur densité conjointe.**
811 Étant donné un échantillon X_1, \dots, X_n d'une distribution continue F de densité f , les
812 *statistiques d'ordre* sont les valeurs triées $X_{1:n} \leq X_{2:n} \leq \dots \leq X_{n:n}$. Leur densité conjointe
813 est

$$f_{1:n, \dots, n:n}(x_1, \dots, x_n) = n! \prod_{i=1}^n f(x_i) \cdot \mathbf{1}\{x_1 \leq \dots \leq x_n\}. \quad (11)$$

814 Le facteur $n!$ compte le nombre d'arrangements, et l'indicatrice $\mathbf{1}\{\cdot\}$ impose la contrainte
815 d'ordonnancement. Les représentations classiques, telles que celles de Sukhatme et Schu-
816 cany, décomposent les statistiques d'ordre en blocs élémentaires tractables qui ont permis
817 la simulation exacte depuis des décennies. Comme nous le verrons dans la Section 0.11.2,
818 ces représentations fournissent également les fondements pour l'approximation par réseaux
819 de neurones avec des bornes de complexité démontrables.

820 0.8.3 Distributions inclinées par récompense et ajustement fin 821 de modèles

822 **Le problème : aligner les modèles génératifs avec les préférences humaines.**
823 Un modèle génératif pré-entraîné produit des échantillons d'une distribution apprise à
824 partir des données d'entraînement. Mais pour de nombreuses applications, nous voulons
825 des échantillons qui ne sont pas simplement typiques, mais *bons* selon un certain critère :
826 des images esthétiquement plaisantes, du texte utile et inoffensif, des molécules qui se lient

à une protéine cible. C'est le *problème d'alignement* : ajuster un modèle génératif pour favoriser les sorties de haute qualité tout en maintenant la diversité et la cohérence.

Considérons un modèle de diffusion texte-vers-image entraîné sur des images collectées sur le web. Le modèle génère des images réalistes, mais beaucoup sont mal composées, contiennent des artefacts ou représentent du contenu indésirable. Nous avons accès à un modèle de récompense (entraîné sur les préférences humaines) qui note les images selon leur qualité esthétique et leur sûreté. L'objectif est d'affiner le modèle de diffusion pour qu'il génère préférentiellement des images à haute récompense, sans s'effondrer vers un mode unique ni oublier comment générer du contenu diversifié.

Formulation mathématique : inclinaison exponentielle. La formulation mathématique naturelle est l'*inclinaison exponentielle*. Étant donné une distribution de base p_0 (le modèle pré-entraîné) et une fonction de récompense $r : \mathcal{X} \rightarrow \mathbb{R}$, la cible est la *distribution inclinée*

$$p(x) \propto p_0(x) \exp(\lambda r(x)), \quad (12)$$

où $\lambda > 0$ contrôle l'intensité de l'inclinaison. Cette distribution est l'unique solution du problème de maximisation de récompense régularisé par KL :

$$p = \arg \max_q \left\{ \mathbb{E}_q[r(x)] - \frac{1}{\lambda} \text{KL}(q \| p_0) \right\}.$$

La distribution inclinée équilibre la maximisation de la récompense contre le fait de rester proche du modèle pré-entraîné, avec λ contrôlant le compromis. Un grand λ favorise la haute récompense ; un petit λ reste proche de p_0 .

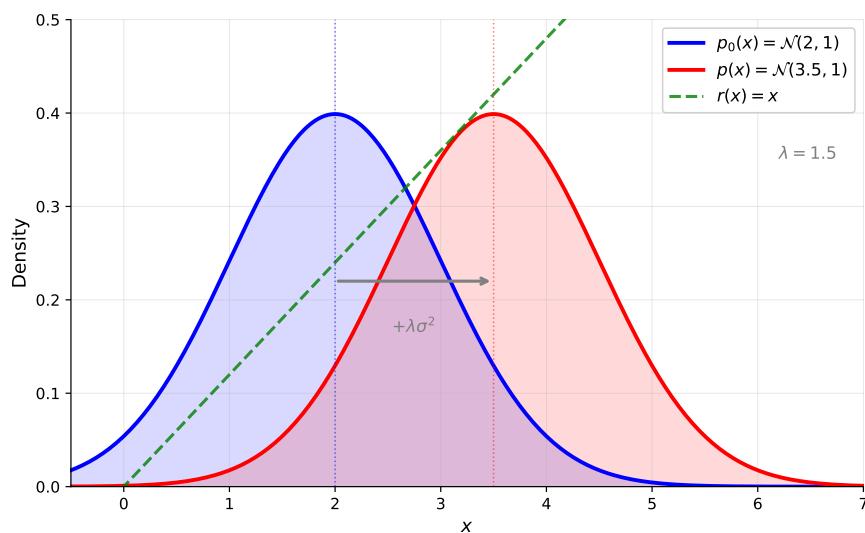


FIGURE 9 : L'inclinaison par récompense déplace la distribution vers les régions à haute récompense. Avec $p_0 = \mathcal{N}(2, 1)$, récompense linéaire $r(x) = x$, et intensité d'inclinaison $\lambda = 1,5$, la distribution inclinée $p(x) \propto p_0(x) \exp(\lambda r(x))$ est $\mathcal{N}(3,5, 1)$: la moyenne se décale de $\lambda\sigma^2 = 1,5$ vers la haute récompense, tandis que la variance reste inchangée.

845 0.9 Contexte de la modélisation générative

846 La Section 0.7 a introduit l'objectif d'apprendre à partir des données tout en respectant la
 847 structure. La Section 0.8 a présenté trois régimes où la structure compte : la dépendance
 848 de queue, les statistiques d'ordre et l'inclinaison par récompense. Pour relever ces défis,
 849 nous nous tournons vers la modélisation générative : la tâche de produire des données
 850 synthétiques qui suivent la même loi que les observations. Cette section fournit le contexte
 851 technique sur les modèles génératifs et la théorie d'approximation par réseaux de neurones
 852 nécessaire pour comprendre les contributions de cette thèse.

853 0.9.1 Le problème de la modélisation générative

854 Le problème fondamental de la modélisation générative peut être énoncé comme suit : étant
 855 donné des échantillons X_1, \dots, X_n tirés indépendamment d'une distribution inconnue P
 856 sur \mathbb{R}^d , construire une procédure pour générer de nouveaux échantillons qui sont distribués
 857 selon P , ou une approximation proche de celle-ci.

858 Ce problème a une longue histoire en statistique et probabilités. Les approches clas-
 859 siques incluent l'*estimation de densité* (SILVERMAN 1986 ; SCOTT 2015), où l'on estime la
 860 densité $p = dP/d\lambda$ puis on échantillonne à partir de la densité estimée, et les *chaînes de*
 861 *Markov Monte Carlo* (MCMC) (METROPOLIS et al. 1953 ; HASTINGS 1970 ; C. P. ROBERT
 862 et CASELLA 2004), où l'on construit une chaîne de Markov ayant P comme distribution
 863 stationnaire et on échantillonne en faisant tourner la chaîne jusqu'à l'équilibre. Les deux
 864 approches font face à la malédiction de la dimension (BELLMAN 1961) : l'estimation de
 865 densité devient statistiquement intractable en haute dimension, tandis que les temps de
 866 mélange MCMC peuvent croître exponentiellement avec la dimension.

867 L'approche moderne de la modélisation générative, rendue possible par l'apprentissage
 868 profond, emprunte une voie différente. Plutôt qu'estimer explicitement la densité, nous
 869 apprenons une *application de transport* G qui transforme des échantillons d'une distribution
 870 de référence simple (typiquement gaussienne standard) en échantillons de la distribution
 871 cible. Si Z est un bruit gaussien, alors $G(Z)$ devrait être distribué approximativement
 872 comme P . L'application G est paramétrée par un réseau de neurones, et les paramètres
 873 sont appris à partir des données.

874 Cette vision basée sur le transport unifie plusieurs paradigmes génératifs modernes.
 875 Dans les *flocs normalisants* (REZENDE et MOHAMED 2015 ; DINH et al. 2017), l'appli-
 876 cation G est contrainte à être inversible avec un jacobien tractable. Dans les *autoencodeurs*
 877 *variationnels* (VAE) (Diederik P. KINGMA et Max WELLING 2014), l'application est ap-
 878 prise conjointement avec une inverse approximative. Dans les *réseaux antagonistes génér-
 879 ratifs* (GAN) (I. J. GOODFELLOW et al. 2014), l'application est apprise via un objectif
 880 adversarial. Dans les *modèles de diffusion* (HO, JAIN et al. 2020 ; Y. SONG, Jascha SOHL-
 881 DICKSTEIN et al. 2021), l'application est implicitement définie par l'inverse d'un processus
 882 de bruitage. Nous décrivons les GAN et les modèles de diffusion en détail ci-dessous, car
 883 ils sont les fondements de nos contributions.

884 0.9.2 Réseaux de neurones

885 Un *réseau de neurones feedforward* à L couches cachées transforme une entrée x à travers
 886 des couches successives (Figure 10). En notant $h_0 = x$, les représentations cachées sont

$$h_\ell = \sigma(W_\ell h_{\ell-1} + b_\ell) \quad \text{pour } \ell = 1, 2, \dots, L, \quad (13)$$

887 où W_ℓ est une matrice de poids, b_ℓ un vecteur de biais, et σ une fonction d'activation non
 888 linéaire appliquée coordonnée par coordonnée. La sortie est $g(x) = W_{L+1}h_L + b_{L+1}$. Les
 889 activations courantes incluent la sigmoïde, la tangente hyperbolique et la ReLU $\sigma(t) =$
 890 $\max(0, t)$, qui domine la pratique moderne.

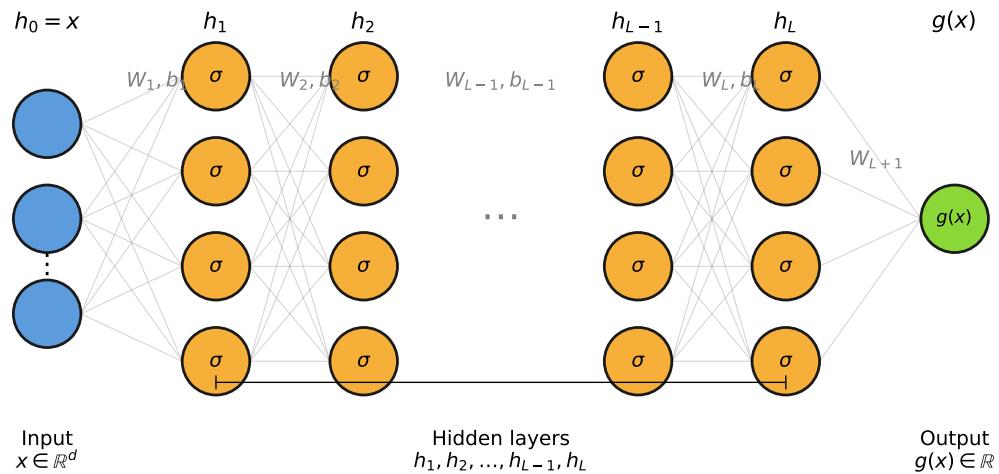


FIGURE 10 : Un réseau de neurones feedforward. L'entrée $h_0 = x$ passe à travers L couches cachées, chacune appliquant une transformation affine suivie de l'activation σ . La sortie $g(x)$ est une fonction linéaire de h_L .

891 0.9.3 Théorie d'approximation des réseaux de neurones

892 Le *théorème d'approximation universelle* (CYBENKO 1989; HORNIK et al. 1989) établit
 893 que les réseaux de neurones à une seule couche cachée peuvent approximer toute fonction
 894 continue sur un domaine compact avec une précision arbitraire, étant donnée une largeur
 895 suffisante. Cela s'étend aux réseaux ReLU (LESHNO et al. 1993).

896 **Bornes quantitatives.** Les bornes quantitatives relient l'erreur d'approximation à la
 897 taille du réseau (BARRON 1993; YAROTSKY 2017). Pour les fonctions lipschitziennes sur
 898 $[0, 1]^d$, atteindre une erreur ε nécessite $O(\varepsilon^{-d})$ paramètres, exhibant la malédiction de
 899 la dimension. Pour des fonctions plus régulières (classe de Hölder avec régularité s), des
 900 réseaux avec $O(\varepsilon^{-d/s})$ paramètres suffisent (YAROTSKY 2017).

901 **Limitations.** Ces théorèmes concernent l'approximation ponctuelle de fonctions, pas
 902 la préservation de structure distributionnelle. Un réseau approximant une application

903 de transport G uniformément peut néanmoins échouer à préserver le comportement des
 904 queues, les contraintes d'ordonnancement ou la structure de dépendance. De plus, les théo-
 905 rèmes supposent que les paramètres optimaux sont connus ; en pratique, la descente de gra-
 906 dient peut ne pas les trouver. Ces limitations motivent la théorie spécifique à la structure
 907 développée dans cette thèse.

908 0.9.4 Réseaux antagonistes génératifs

909 Les réseaux antagonistes génératifs (I. J. GOODFELLOW et al. 2014) formulent la modéli-
 910 sation générative comme un jeu à deux joueurs. L'intuition se comprend mieux à travers
 911 une analogie : imaginez un faussaire essayant de produire de faux billets de banque, et
 912 un détective essayant de les détecter. Le faussaire s'améliore en étudiant quels faux sont
 913 détectés ; le détective s'améliore en étudiant les dernières techniques du faussaire. Au fil
 914 du temps, les deux deviennent plus sophistiqués, jusqu'à ce que finalement le faussaire
 915 produise des billets si convaincants que même le détective expert ne peut les distinguer de
 916 la vraie monnaie.

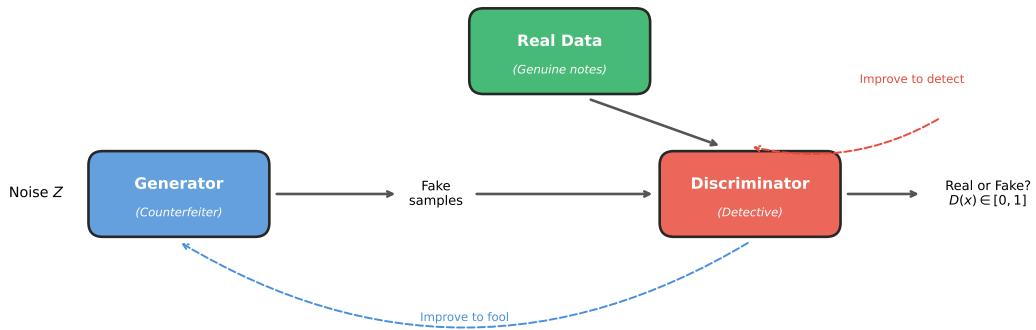


FIGURE 11 : Le cadre GAN comme jeu adversarial. Le générateur (faussaire) transforme du bruit en faux échantillons ; le discriminateur (détective) tente de distinguer les faux des vraies données. Les deux réseaux s'améliorent par compétition jusqu'à l'équilibre, quand les échantillons générés sont indiscernables des vraies données.

917 En termes mathématiques, un *générateur* $G_\theta : \mathcal{Z} \rightarrow \mathcal{X}$ transforme un bruit latent
 918 $Z \sim p_Z$ en échantillons synthétiques (la production du faussaire). Un *discriminateur* $D_\phi :
 919 \mathcal{X} \rightarrow [0, 1]$ tente de distinguer les échantillons réels des générés (le verdict du détective).
 920 Les réseaux sont entraînés en résolvant un problème minimax :

$$\min_{\theta} \max_{\phi} \mathbb{E}_{X \sim p_{\text{data}}} [\log D_\phi(X)] + \mathbb{E}_{Z \sim p_Z} [\log(1 - D_\phi(G_\theta(Z)))].$$

921 Le générateur minimise cet objectif (essayant de tromper le discriminateur), tandis que le
 922 discriminateur le maximise (essayant de classifier correctement réel vs. faux). À l'équilibre,
 923 le générateur produit des échantillons de la distribution des données, et le discriminateur
 924 sort 1/2 pour toutes les entrées, incapable de distinguer le réel du faux.

925 Le GAN de Wasserstein (WGAN) (ARJOVSKY et al. 2017) remplace la divergence de

926 Jensen-Shannon par la distance de Wasserstein-1 :

$$W_1(p_{\text{data}}, p_{\theta}) = \sup_{\|f\|_{L^1} \leq 1} \mathbb{E}_{X \sim p_{\text{data}}}[f(X)] - \mathbb{E}_{Z \sim p_Z}[f(G_{\theta}(Z))],$$

927 où le supremum est pris sur les fonctions 1-lipschitziennes. Le WGAN avec pénalité de
 928 gradient (WGAN-GP) (GULRAJANI et al. 2017) impose la contrainte de Lipschitz via un
 929 terme de pénalité, améliorant la stabilité de l'entraînement.

930 0.9.5 Modèles de diffusion et basés sur le score

931 Les modèles de diffusion (HO, JAIN et al. 2020 ; Y. SONG, Jascha SOHL-DICKSTEIN et al.
 932 2021) définissent un processus génératif en construisant un chemin continu $(p_t)_{t \in [0,1]}$ entre
 933 la distribution des données p_0 et une référence tractable $p_1 = N(0, I_d)$. Le chemin est défini
 934 via l'*interpolation directe*

$$X_t = \alpha_t X_0 + \sigma_t X_1, \quad X_0 \sim p_0, \quad X_1 \sim N(0, I_d), \quad (14)$$

935 où (α_t, σ_t) sont des programmes déterministes avec $(\alpha_0, \sigma_0) = (1, 0)$ et $(\alpha_1, \sigma_1) = (0, 1)$. À
 936 $t = 0$, nous avons des données propres ; à $t = 1$, du pur bruit. Les choix standards incluent
 937 le programme préservant la variance ($\alpha_t^2 + \sigma_t^2 = 1$) et le programme linéaire ($\alpha_t = 1 - t$,
 938 $\sigma_t = t$).

939 Pour générer des échantillons, nous inversons ce processus. La quantité clé est la *fonction de score* $\nabla_x \log p_t$, qui pointe vers les régions de haute densité à chaque niveau de
 940 bruit. Pour le noyau direct gaussien, le score satisfait $\nabla_x \log p_t(x_t) = -\hat{x}_1(x_t, t)/\sigma_t$, où
 941 $\hat{x}_1(x_t, t) = \mathbb{E}[X_1 | X_t = x_t]$ est l'espérance conditionnelle du bruit sachant l'échantillon
 942 bruité.

943 Le score est appris via l'*appariement de score par débruitage* (HYVÄRINEN 2005 ; VINCENT
 944 2011) : un réseau $s_{\theta}(x, t)$ est entraîné à prédire le bruit X_1 à partir de l'échantillon bruité
 945 X_t , ce qui équivaut à apprendre $\nabla_x \log p_t$. En partant de $X_1 \sim N(0, I_d)$ et en débruitant
 946 itérativement en utilisant le score appris, on obtient des échantillons de p_0 .

948 0.10 État de l'art et problèmes ouverts

949 Nous passons maintenant en revue l'état de l'art pour chaque régime structurel, identifions
 950 les lacunes dans la littérature existante et formulons les objectifs de cette thèse.

951 0.10.1 Génération de valeurs extrêmes

952 **Théorie classique des valeurs extrêmes.** La théorie statistique des extrêmes multi-
 953 variés est bien développée (Paul EMBRECHTS et al. 1997 ; L. d. HAAN et FERREIRA 2006 ;
 954 Sidney I. RESNICK 1987). Pour les extrêmes univariés, le théorème de Fisher-Tippett-
 955 Gnedenko établit que les maxima correctement normalisés convergent vers l'une des trois

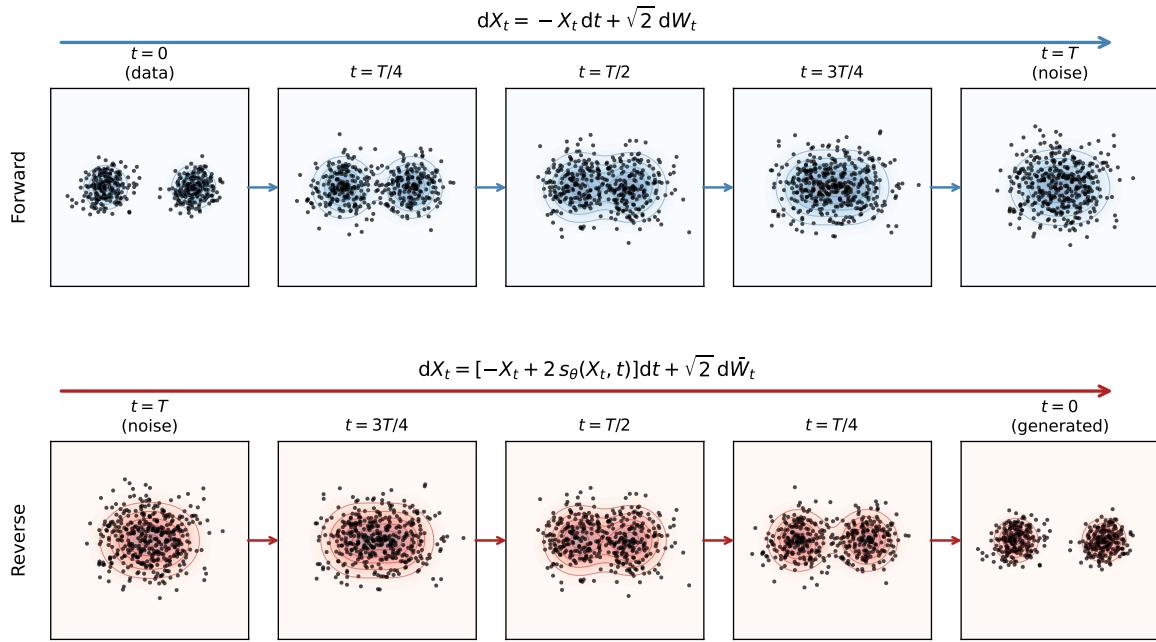


FIGURE 12 : Le cadre des modèles de diffusion sur un mélange de gaussiennes 2D. Processus direct (ligne du haut) : les données de p_0 sont progressivement corrompues par ajout de bruit jusqu'à devenir indiscernables de la référence gaussienne p_1 à $t = 1$. Processus inverse (ligne du bas) : partant du bruit pur, un réseau de score appris s_θ guide le débruitage, récupérant graduellement la structure bimodale de p_0 à $t = 0$.

distributions : Gumbel, Fréchet ou Weibull, unifiées par la distribution des valeurs extrêmes généralisée (GEV). Dans le cadre multivarié, la fonction de dépendance de queue stable (stdf) fournit une caractérisation complète de la dépendance extrémale, décrivant comment la probabilité de dépassements conjoints évolue lorsque l'on s'enfonce dans la queue. De nombreux modèles paramétriques ont été proposés : la copule de Gumbel (logistique) (GUMBEL 1960), le modèle de Hüsler-Reiss provenant des processus gaussiens (HÜSLER et REISS 1989a), et le modèle extrémal- t généralisant Hüsler-Reiss aux marges de Student (NIKOLOUPOULOS et al. 2009). Les méthodes d'estimation incluent la vraisemblance empirique (EINMAHL, KRAJINA et al. 2012), la M-estimation basée sur les rangs (EINMAHL, L. d. HAAN et al. 2016) et les approches bayésiennes (SABOURIN et al. 2013).

Pourquoi les GAN standards échouent. Les modèles génératifs standards sont fondamentalement incompatibles avec les données à queue lourde. La raison est mathématique : (WIESE et al. 2019) ont prouvé que si la distribution latente a des queues légères (par exemple gaussienne) et que le générateur G est lipschitzien continu, alors la distribution générée $G(Z)$ a nécessairement des queues légères. Spécifiquement, si Z est sous-gaussien, alors $G(Z)$ est sous-gaussien pour tout G lipschitzien. Puisque les GAN modernes imposent des contraintes de Lipschitz pour la stabilité de l'entraînement (normalisation spectrale, pénalités de gradient), ils ne peuvent prouvablement pas générer des sorties à queue lourde à partir de bruit gaussien. Ce n'est pas un échec de l'optimisation ou de l'architecture ; c'est une impossibilité fondamentale.

Modèles génératifs à queue lourde. Des travaux récents ont abordé cette limitation à travers plusieurs approches (Michaël ALLOUCHE, Emmanuel GOBET et al. 2024). Pour les GAN : les *méthodes de prétraitement* comme Quant-GAN (WIESE et al. 2019) et evtGAN (Youssef BOULAGUIEM et al. 2022) transforment les données pour supprimer la lourdeur des queues avant l’entraînement ; les méthodes à *latent à queue lourde* (HUSTER et al. 2021) remplacent le bruit gaussien par des distributions de Pareto ; les approches *paramétrées par EVT* comme EV-GAN (Michaël ALLOUCHE, GIRARD et al. 2022) conçoivent des générateurs qui encodent explicitement la structure des valeurs extrêmes. Pour les modèles de diffusion, les extensions à queue lourde remplacent le processus de bruit gaussien par des processus de Lévy α -stables (YOON et al. 2023), permettant la génération directe d’échantillons à queue lourde. Cependant, toutes ces approches se concentrent sur le comportement *marginal* des queues. Aucune n’aborde la *structure de dépendance* dans les queues, caractérisée par la stdf.

Lacune 1. À notre connaissance, aucun modèle génératif existant ne fournit de garanties prouvables sur l’approximation de la dépendance de queue. Les méthodes existantes traitent le comportement marginal des queues mais ignorent la structure conjointe. Nous comblons cette lacune dans le Chapitre 1, en développant une architecture GAN qui approxime prouvablement toute stdf.

0.10.2 Génération de statistiques d’ordre

La formulation mathématique des statistiques d’ordre a été donnée dans la Section 0.8.2. Les représentations classiques (Sukhatme, Schucany) ont permis la simulation exacte depuis des décennies (H. A. DAVID et H. N. NAGARAJA 2003 ; ARNOLD et al. 2008). Nous passons maintenant en revue l’état de l’art.

Approches d’apprentissage automatique. La littérature en apprentissage automatique a abordé le tri du point de vue de l’apprentissage à classer (T.-Y. LIU 2009), où l’objectif est de prédire des ordres plutôt que de les générer. Des opérateurs de tri différentiables ont été développés pour permettre l’apprentissage par gradient : (GROVER et al. 2019) ont proposé NeuralSort, une relaxation continue de l’opérateur argsort basée sur l’algorithme de Sinkhorn ; (CUTURI et al. 2019) ont développé le tri différentiable basé sur le transport optimal ; (BLONDEL et al. 2020) ont introduit le tri différentiable rapide avec une complexité $O(n \log n)$. Ces méthodes permettent le tri dans les pipelines de réseaux de neurones mais sont conçues pour les tâches de classement, pas pour l’appariement de distributions.

Modèles génératifs et ordonnancement. Les modèles génératifs standards (GAN, VAE, diffusion) produisent des échantillons non ordonnés. Appliqués à des données classées, ils font face à deux problèmes. Premièrement, les échantillons générés ne sont pas triés : un GAN entraîné sur des statistiques d’ordre produira des vecteurs dont les coordonnées ne sont pas monotonement ordonnées. Le tri a posteriori peut restaurer l’ordre

1014 mais introduit un biais dans la distribution. Deuxièmement, le prétraitement standard dé-
1015 truit la structure d'ordonnancement : la normalisation par coordonnée vers $[0, 1]$, courante
1016 dans l'entraînement des GAN, normalise chaque coordonnée indépendamment selon son
1017 étendue marginale, détruisant l'ordonnancement relatif. Si la k -ième statistique d'ordre est
1018 normalisée par sa propre moyenne et variance, et la $(k+1)$ -ième par des valeurs différentes,
1019 il n'y a aucune garantie que les valeurs normalisées restent triées.

1020 **Lacune 2.** À notre connaissance, aucun modèle génératif existant ne produit des sta-
1021 tistiques d'ordre avec à la fois les distributions marginales/conjointes correctes et des
1022 contraintes d'ordonnancement garanties. Les méthodes de tri différentiable existantes traitent
1023 la prédiction de classement, pas la modélisation générative. Nous comblons cette lacune
1024 dans le Chapitre 2, en développant une architecture qui exploite les représentations clas-
1025 siques pour garantir l'ordonnancement tout en appariant les distributions cibles.

1026 0.10.3 Génération inclinée par récompense

1027 Le problème d'alignement (Section 0.8.3) requiert d'échantillonner à partir de la distribu-
1028 tion inclinée (12). Nous passons en revue les approches existantes.

1029 **Apprentissage par renforcement à partir de feedback humain.** Le paradigme
1030 dominant pour l'alignement est l'apprentissage par renforcement à partir de feedback hu-
1031 main (RLHF) (CHRISTIANO et al. 2017; BAI et al. 2022; OUYANG et al. 2022). Un mo-
1032 dèle de récompense est entraîné sur des données de préférence humaine (comparaisons par
1033 paires de sorties), puis le modèle génératif est affiné pour maximiser la récompense espé-
1034 rée en utilisant des méthodes de gradient de politique (PPO, REINFORCE). Le RLHF
1035 a été transformateur pour les grands modèles de langage, permettant ChatGPT, Claude
1036 et d'autres agents conversationnels. Cependant, le RLHF est notoirement instable : le pi-
1037 ratage de récompense, l'effondrement de mode et l'oubli catastrophique sont des modes
1038 d'échec courants (CASPER et al. 2023).

1039 **Ajustement fin des modèles de diffusion.** Pour les modèles de diffusion, plusieurs
1040 approches d'ajustement fin ont émergé. DRaFT (CLARK et al. 2024) rétropropage à travers
1041 toute la chaîne de débruitage pour maximiser la récompense, mais cela nécessite de stocker
1042 les activations à travers tous les pas de temps, avec une mémoire croissant linéairement avec
1043 la longueur de la chaîne. DPOK (FAN et al. 2023) combine les gradients de politique avec
1044 la régularisation KL. ReFL (Jing XU et al. 2024) utilise l'apprentissage par feedback de
1045 récompense pour affiner les modèles de diffusion texte-vers-image. L'appariement adjoint
1046 (DOMINGO-ENRICH, DROZDZAL et al. 2025) fournit une approche théoriquement fondée
1047 basée sur les EDS adjointes, atteignant des résultats état de l'art mais nécessitant toujours
1048 les gradients de récompense.

1049 **Le goulot d'étranglement du gradient.** Les méthodes d'ajustement fin basées sur le
1050 gradient (DRaFT, appariement adjoint) nécessitent de rétropropager à travers la fonction

1051 de récompense, ce qui crée des limitations pratiques : stocker les activations pour les grands
 1052 modèles de récompense mène à des coûts mémoire prohibitifs, et les récompenses non
 1053 différentiables (feedback humain, évaluations discrètes, métriques basées sur simulateur)
 1054 ne peuvent pas être traitées. Les approches basées sur l'apprentissage par renforcement
 1055 (RLHF, DPOK) évitent cela en utilisant uniquement les valeurs de récompense, pas les
 1056 gradients, mais introduisent d'autres défis : les estimateurs de gradient de politique ont
 1057 une haute variance, et un réglage fin des hyperparamètres est nécessaire pour équilibrer la
 1058 maximisation de récompense contre l'effondrement de mode. Aucune approche ne fournit
 1059 de garanties distributionnelles sur le modèle affiné résultant.

1060 **Lacune 3.** À notre connaissance, aucune méthode existante sans gradient n'affine les mo-
 1061 dèles de diffusion pour échantillonner à partir de distributions inclinées par récompense.
 1062 Les méthodes existantes nécessitent soit les gradients de récompense, soit utilisent l'ap-
 1063 prentissage par renforcement avec des estimateurs à haute variance. Nous comblons cette
 1064 lacune dans le Chapitre 3, en développant une méthode basée sur l'identité de Fisher qui
 1065 ne requiert que des évaluations directes de la fonction de récompense.

1066 0.10.4 Objectifs de cette thèse

1067 Sur la base des lacunes identifiées, nous formulons trois objectifs :

- 1068 1. **Génération de valeurs extrêmes.** Développer un modèle génératif pour les ex-
 1069 trêmes multivariés avec des garanties prouvables sur l'approximation de la stdf. La
 1070 méthode devrait utiliser un bruit latent à queue lourde et fournir des bornes explicites
 1071 sur l'erreur d'approximation en termes d'architecture de réseau.
- 1072 2. **Génération de statistiques d'ordre.** Développer un modèle génératif pour les
 1073 statistiques d'ordre qui garantit l'ordonnancement tout en appariant la distribution
 1074 cible. La méthode devrait exploiter les représentations classiques et fournir des bornes
 1075 d'approximation qui évoluent favorablement avec la dimension.
- 1076 3. **Génération inclinée par récompense.** Développer une méthode sans gradient
 1077 pour affiner les modèles de diffusion vers des distributions inclinées par récompense. La
 1078 méthode ne devrait requérir que des évaluations directes de la fonction de récompense.

1079 0.11 Contributions de cette thèse

1080 Nous résumons maintenant les contributions principales, organisées selon les trois objectifs.

1081 0.11.1 GAN à queue lourde pour la génération de valeurs ex- 1082 trêmes

1083 Nous développons les GAN à queue lourde (HTGAN) pour générer des données avec une
 1084 dépendance de queue correcte.

1085 **Contributions théoriques.** Nous remplaçons le bruit latent gaussien par un bruit à
1086 queue lourde (Pareto ou Fréchet). Comme le générateur est lipschitzien, la propriété de
1087 queue lourde se propage à la sortie. Nous prouvons que la stdf de la distribution générée
1088 est toujours discrète, un maximum de fonctions linéaires avec un nombre fini d'atomes.
1089 Notre théorème d'approximation principal établit que toute stdf peut être approximée
1090 dans l'erreur $L^\infty \varepsilon$ en utilisant une dimension latente $N = O(\varepsilon^{-(d-1)})$ et une seule couche
1091 cachée avec des activations ReLU. Cela semble être la première garantie d'approximation
1092 pour la dépendance de queue dans la littérature de modélisation générative.

1093 **Contributions méthodologiques.** L'algorithme HTGAN procède en trois étapes :
1094 transformation marginale vers Fréchet unitaire, entraînement GAN avec bruit latent de
1095 Pareto, et transformation inverse. Cette séparation des marges et de la dépendance est
1096 standard en modélisation par copules.

1097 **Contributions expérimentales.** Nous évaluons sur des données synthétiques de copule
1098 de Gumbel et des rendements du S&P 500, démontrant que HTGAN capture correctement
1099 la dépendance de queue tandis que les GAN standards échouent de plusieurs ordres de
1100 grandeur.

1101 0.11.2 Statistiques d'ordre neurales génératives

1102 Nous développons les statistiques d'ordre neurales génératives (GENOS) pour générer des
1103 données classées avec un ordonnancement garanti.

1104 **Contributions théoriques.** Nous combinons les représentations statistiques classiques
1105 avec la théorie d'approximation des réseaux de neurones (YAROTSKY 2017) pour dériver
1106 des bornes de complexité non asymptotiques. Spécifiquement, nous adaptons deux repré-
1107 sentations classiques des statistiques d'ordre et analysons leur complexité d'approximation
1108 par réseaux de neurones.

1109 *La représentation de Schucany* (SCHUCANY 1972) construit les statistiques d'ordre
1110 récursivement :

$$U_{n-r:n} = U_{n-r+1:n} \cdot U_r^{1/(n-r)}, \quad r = 1, \dots, n-1. \quad (15)$$

1111 Cela nécessite n fonctions de transition différentes. Nous prouvons que l'approximation
1112 neurale de ce schéma requiert un nombre de paramètres qui est polynomial de degré n en
1113 ε^{-1} .

1114 *La représentation de Sukhatme* (SUKHATME 1937) exprime les statistiques d'ordre uni-
1115 formes via des espacements exponentiels :

$$U_{k:n} = F \left(\sum_{j=1}^k \frac{E_j}{n-j+1} \right), \quad E_j \stackrel{\text{iid}}{\sim} \text{Exp}(1), \quad F(z) = 1 - e^{-z}. \quad (16)$$

1116 La somme cumulative a une distribution exponentielle ; appliquer la fonction de réparti-
1117 tion exponentielle F donne la k -ième statistique d'ordre uniforme $U_{k:n}$. Cela ne requiert

₁₁₁₈ d'approximer que deux fonctions universelles (F et F^{-1}), indépendamment de n . Nous
₁₁₁₉ prouvons que l'approche basée sur Sukhatme atteint

$$\mathbb{E}|U_{k:n} - \hat{U}_{k:n}| \leq O(\varepsilon \log n) \quad \text{avec} \quad \mathcal{C}_{\text{Sukhatme}}(\varepsilon) = O(\varepsilon^{-2} \log(1/\varepsilon)) \quad (17)$$

₁₁₂₀ paramètres. La complexité est quadratique en ε^{-1} , indépendante de n , contre polynomiale
₁₁₂₁ de degré n pour Schucany. Cela rend l'approche basée sur Sukhatme scalable à de grandes
₁₁₂₂ tailles d'échantillon.

₁₁₂₃ **Contributions méthodologiques.** Nous introduisons la normalisation globale moyenne-
₁₁₂₄ échelle, qui préserve l'ordonnancement contrairement à la normalisation par coordonnée.
₁₁₂₅ Nous proposons également une pénalité de statistiques d'ordre (OSP) qui pénalise les vio-
₁₁₂₆ lations de la contrainte d'ordonnancement :

$$\mathcal{L}_{\text{OSP}} = \mathbb{E}_Z \left[\frac{1}{n-1} \sum_{k=1}^{n-1} (\max\{G(Z)_k - G(Z)_{k+1}, 0\})^2 \right].$$

₁₁₂₇ **Contributions expérimentales.** Nous évaluons sur des statistiques d'ordre synthé-
₁₁₂₈ tiques avec dimension variable, comparant GAN, WGAN et WGAN avec OSP. L'ajout de
₁₁₂₉ la pénalité OSP améliore les métriques d'ordonnancement tout en maintenant des perfor-
₁₁₃₀ mances compétitives sur la distance de Wasserstein.

₁₁₃₁ 0.11.3 Ajustement fin sans gradient des modèles de diffusion

₁₁₃₂ Nous développons des méthodes sans gradient pour affiner les modèles de diffusion afin
₁₁₃₃ d'échantillonner à partir de distributions inclinées par récompense.

₁₁₃₄ **Contributions théoriques.** L'identité de Fisher exprime le score de la distribution
₁₁₃₅ inclinée comme une covariance :

$$\nabla_x \log p(x) = \nabla_x \log p_0(x) + \lambda \cdot \text{Cov}_{q_{0|x}} [\nabla_{x_0} \log q(x_0|x), r(x_0)]. \quad (18)$$

₁₁₃₆ Cette covariance peut être estimée par Monte Carlo en utilisant uniquement des évaluations
₁₁₃₇ directes de r , sans rétropropagation à travers le modèle de récompense.

₁₁₃₈ **Contributions méthodologiques.** Nous proposons un algorithme itératif qui décom-
₁₁₃₉ pose une grande inclinaison en N petites étapes, estimant l'incrément de score à chaque
₁₁₄₀ étape via l'identité de Fisher. L'algorithme s'applique même lorsque r est non différentiable
₁₁₄₁ ou une boîte noire.

₁₁₄₂ **Contributions expérimentales.** Sur des mélanges de gaussiennes 2D avec vérité ter-
₁₁₄₃ rain connue, nous validons que la méthode récupère correctement le score incliné et que
₁₁₄₄ l'augmentation du nombre d'étapes N améliore la précision.

1145 0.12 Organisation du manuscrit

1146 Le Chapitre 1 développe la théorie et les algorithmes pour la modélisation générative à
1147 queue lourde. Nous présentons le contexte mathématique sur la théorie des valeurs extrêmes
1148 multivariées et la stdf, analysons pourquoi les applications lipschitziennes de bruit gaussien
1149 ne peuvent pas produire de queues lourdes, caractérisons la stdf discrète induite par les
1150 générateurs avec bruit de Pareto, prouvons notre théorème d'approximation principal et
1151 présentons des expériences sur des données synthétiques et financières.

1152 Le Chapitre 2 développe les modèles génératifs pour les statistiques d'ordre. Nous
1153 présentons le contexte sur les statistiques d'ordre et les représentations de Sukhatme et
1154 Schucany, prouvons des bornes d'approximation pour les réseaux ReLU, introduisons la
1155 normalisation globale moyenne-échelle et l'objectif WGAN-OSP, et présentons des expé-
1156 riences sur des statistiques d'ordre synthétiques.

1157 Le Chapitre 3 développe l'ajustement fin sans gradient pour les modèles de diffusion.
1158 Nous présentons le contexte sur l'inclinaison exponentielle et les méthodes existantes basées
1159 sur le gradient, dérivons l'estimateur de score basé sur l'identité de Fisher, analysons
1160 l'erreur d'approximation, présentons l'algorithme d'inclinaison itératif et validons sur des
1161 expériences contrôlées.

1162 Cette thèse est partiellement basée sur des travaux collaboratifs : le Chapitre 1 sur les
1163 modèles génératifs à queue lourde, le Chapitre 2 sur l'approximation neurale des statistiques
1164 d'ordre, et le Chapitre 3 sur l'ajustement fin sans gradient des modèles de diffusion. Les
1165 lieux de publication et co-auteurs spécifiques sont indiqués au début de chaque chapitre.

1166 **Chapter 1**

1167 **Heavy-Tailed GANs for
1168 Extreme Value Generation**

1169 **1.1 Introduction**

1170 **Statement of the problem.** Examining extreme events is a critical concern across var-
1171 ious fields such as economics, engineering, and life sciences, with wide-ranging applications
1172 like actuarial and financial risks (Asmussen and Albrecher 2010, Chapters X and XIII)-(P.
1173 Embrechts et al. 1997, Chapters 1 and 6)-(A. McNeil et al. 2015, Chapters 5 and 16, Part
1174 III), communication network reliability (P. Robert 2003), and aircraft safety (Prandini and
1175 Watkins 2005). Extreme events play a crucial role also in the context of climate change
1176 (J. Zscheischler et al. 2020), with the occurrence of more and more severe weather events,
1177 or in the context of cyber-security (Cremer et al. 2022) with the increasing number of
1178 cyber-attacks of private companies or public entities. In recent decades, the importance
1179 of extreme event analysis has surged in financial risk management, which calls for an
1180 ever-increasing number of types of risk (market, credit, operational, reputational, cyber,
1181 climate, and so on and so forth) to be encompassed in order to measure their impacts on
1182 banking activity and the stability of the financial system. In particular, regulators (Euro-
1183 pean Banking Authority 2014) are increasingly imposing stress tests to test the resilience
1184 of the banking system against different risk categories. These stress tests involve numerical
1185 simulations of unfavorable yet plausible extreme scenarios, serving as a primary means to
1186 evaluate the potential impact on risks (see (ECB Banking Supervision 2023) for the 2023
1187 exercises in Europe). Often, extreme events are described as tail events of multivariate
1188 relevant random variables.

1189 The challenge arises in efficiently obtaining relevant samples for assessing these tail risks.
1190 Informally, it writes as finding a generator G and a latent distribution μ_Z which one can
1191 easily sample from, such that:

$$Z \sim \mu_Z, \quad G(Z) \stackrel{d}{=} \text{target distribution.} \quad (1.1)$$

1192 A large family of numerical approaches relies on physics-based methods, and aims to
 1193 efficiently sample the distribution tail with minimal computational effort and to avoid
 1194 the inefficiency of a basic accept-rejection algorithm. These methods are specific to the
 1195 model at hand: among these methods, we mention importance sampling (Bucklew 2004,
 1196 Chapter 4), MCMC with splitting – (E. Gobet and G. Liu 2015), or interacting particles
 1197 system – (Del Moral and Garnier 2005). Another family of simulation schemes is data-
 1198 driven approaches which do not require the knowledge of the (supposedly existing) physical
 1199 model behind the sample generation. These data-driven schemes take as input a data
 1200 set of observations of given size n_{data} and design a generative model able to re-generate
 1201 samples that mimic the empirical distribution of the observed data. This concept aligns
 1202 with recent paradigms in Artificial Intelligence such as Generative Adversarial Networks
 1203 (GANs) initiated by (I. Goodfellow et al. 2014) and Variational Autoencoders (VAEs) by
 1204 (D. P. Kingma and M. Welling 2014), score-based diffusion models (J. Sohl-Dickstein et al.
 1205 2015), normalising flows (Papamakarios et al. 2021), etc.

1206 Our work is in this vein for generating extreme samples. Unlike the classical uses of
 1207 Deep Generative Models as mentioned above, where the number n_{data} of training data is
 1208 generally colossal (millions), in the context of extreme events, n_{data} is small by nature (a few
 1209 hundred at most). Standard training procedures are thus bound to have poor performance
 1210 all the more so as, by construction, their design makes them unable to reproduce heavy
 1211 and dependent tails (M. Allouche et al. 2022). The aim of this work is to design new
 1212 efficient procedures taking advantage of the probabilistic structure behind extreme values.

1213 **State of the art.** Very recently, the machine learning community became aware of the
 1214 difficulties of sampling extremes using Deep Generative Models, requiring to appropriately
 1215 adapt known generative methods to the extreme setting. Most of the newly designed
 1216 algorithms were based on GANs, see an overview in (Michaël Allouche, Emmanuel Gobet,
 1217 et al. 2024). Three directions have been mainly investigated in the literature. First, certain
 1218 works apply some preprocessing to the data to get rid of the tail heaviness, see Quant-GAN
 1219 (Wiese et al. 2020) and evtGAN (Y. Boulaguiem et al. 2022): essentially, it consists in
 1220 suitably transforming each data coordinate in order to retrieve either Gaussian or uniform
 1221 marginal distributions, then performing a usual GAN method and reverting to the original
 1222 space by inverse transform. A second direction is to consider new latent variables Z with a
 1223 heavy-tailed distribution. In (Feder et al. 2020; Huster et al. 2021), a Generalized Pareto
 1224 Distribution is adopted for the latent distribution μ_Z , in contrast to original GANs which
 1225 use light-tailed distributions such as uniform or Gaussian. This idea finds its origins in a
 1226 crucial observation: neural networks are Lipschitz functions (Wiese et al. 2020). As such,
 1227 they cannot transform a random variable with light-tailed to a distribution with heavier
 1228 tail, this is the main result of (Liang et al. 2022). Using heavy-tailed noise however raises
 1229 an issue with the training procedure: if the underlying distribution is too heavy-tailed,
 1230 the gradient of the loss function may not exist. To alleviate this, in (Huster et al. 2021),
 1231 the authors introduce a loss function designed to produce stable training. To be effective,
 1232 these methods require to estimate accurately the tail-index parameters for each marginal

1233 distribution, which is a difficult task because the number n_{data} of training data is usually
1234 small. In (Huster et al. 2021), the authors provide empirical evidence that the Pareto-
1235 GAN method is able to match heavy-tailed distributions embedded in lower dimensional
1236 manifolds in a high dimensional setup. However, they do not provide a quantification of
1237 the power of Pareto-GAN to reproduce asymptotic dependence in the data, which is our
1238 main scope of interest in this paper.

1239 A third approach consists in suitably parameterize the generator to account for extreme-
1240 value theory. This approach has been developed in (M. Allouche et al. 2022) (EV-GAN)
1241 where the renormalized log-quantile (called tail index function by the authors) is learnt
1242 using ReLU neural networks. The authors provide error bounds in uniform norm according
1243 to the complexity of neural networks and the second-order conditions in extreme-value
1244 theory. The latter reference shows that EV-GAN largely outperforms usual GANs, and is
1245 able to accurately sample X in dimension up to 50. For higher dimension, the marginals
1246 of X are still accurately reproduced but the tail dependence structure is under-estimated.
1247 Our work tackles the problem of better learning the dependence between margins with
1248 GANs.

1249 The closest work to ours is presumably (Hasan et al. 2022). The authors endeavour to
1250 approximate the stable tail dependence function (stdf) of the vector X of interest, see Def-
1251 inition 1.1 later, which characterizes the dependence in the extremes, once the marginals
1252 have been standardized to a unit Fréchet scale. Recall that, once the stdf is known (even
1253 approximately), the vector X can be sampled using the spectral representation of a max-
1254 stable process (Hofert et al. 2018, Theorem 2.1). In (Hasan et al. 2022) two learning
1255 techniques of the stdf are introduced. On the one hand, a deterministic approach is pro-
1256 posed using a dedicated architecture of neural networks called “d-Max NN”. The Fréchet
1257 marginalised data are projected onto directions in the simplex leading to exponentially
1258 distributed data. An algorithm is then introduced to maximise the log-likelihood using
1259 randomised directions on the neural network parameters. On the other hand, a stochastic
1260 approach is designed, based on the stochastic representation of the stdf via the spectral
1261 measure (see Proposition 1.1). Note that this approach is calibrated using the previous
1262 deterministic method. At first glance, our approach may look similar to (Hasan et al. 2022)
1263 but it is actually quite different, even though we also work on unitary Fréchet marginals.
1264 Here, we do rely on the spectral representation of max-stable processes, but we rather
1265 directly design a generative model of the vector X using an unitary Fréchet latent noise.
1266 In our opinion, this is much simpler since our approach avoids truncating (at a poorly con-
1267 trolled rank) the max-stable representation. Then, in contrast to Pareto-GAN, we show
1268 a convergence result of the approximated stdf to the true one when the dimension of the
1269 latent noise increases.

1270 Last, we shall mention the Tail-GAN of (R. Cont et al. 2022) in a financial setting,
1271 where a usual GAN is performed but the loss function for learning depends on the risk-
1272 metrics of the scalar quantity of interest (e.g. Value-at-Risk and Expected Shortfall). It is
1273 a way to enforce the generative model to reproduce some specific metrics: note that this
1274 is different to sample a multidimensional extreme distribution.

On the VAE side, it is shown in (Lafon et al. 2023, Corollary 7) that a VAE built with piecewise linear activation functions and Gaussian distributions cannot reproduce heavy-tailed margins. It is moreover established that the spectral measure (see Section 1.2.2) associated with a classical VAE output is necessarily discrete (Lafon et al. 2023, Proposition 8). To overcome these problems, the authors consider a univariate heavy-tailed distribution to sample the radius, and, conditionally on the latter, an angle is sampled from a multivariate normal distribution. The product of the two yields the desired multivariate regularly-varying vector. The appropriate KL-divergences are derived leading to two objective functions: one for the radius VAE and one for the angular VAE. In the recent work (L. Zhang et al. 2024), the authors develop a VAE incorporating the max-id spatial model in order to sample spatial extremes with flexible and non-stationary dependence structures.

Our contributions. Focusing on GANs, we investigate the representation power of neural network based generative models with heavy-tailed input noise. In particular, we study the ability of such models to reproduce multivariate asymptotic dependence, which is quantified by the stdf. To this end, we first provide a modification to the original GAN training to accommodate the specificity of heavy-tailed input noise. Then, the quality of approximation of the stdf by the modified GAN is assessed thanks to the stdf representation in terms of D-norms (see Section 1.2.2). We show that the approximation error is decreasing when the dimension of the latent noise increases (see Theorem 1.3). Finally, we provide a set of experiments to illustrate the behaviour of the algorithm and perform a comparison with a benchmark method on both synthetic and real datasets.

Organization of the paper. Section 1.2 is devoted to the theory on the representation capability of the stdf with GANs based on heavy-tailed input noise. The main theoretical tools are introduced before the statement of our approximation results. Experiments on synthetic and real datasets are summarized in Section 1.3 to illustrate the generative power of heavy-tailed GANs. Some background results and the proofs of main results are postponed to the Supplementary (Appendix 3.A and Appendix 3.F respectively). Extra experiments are available in Appendix 3.I.

Notations.

- $[n] = \{i \in \mathbb{N} : 1 \leq i \leq n\}$ denotes the set of natural numbers from 1 to n . $\lceil \cdot \rceil$ denotes the ceil function.
- The vectors of the canonical basis of \mathbb{R}^d are $\mathbf{e}_i = (0, \dots, 0, 1, 0, \dots, 0)$, $i \in [d]$, with a 1 on the i^{th} coordinate.
- Vectors or matrices are denoted with bold symbols and scalar with roman ones. The i^{th} coordinate of a vector \mathbf{x} is referred to as x_i . Random variables are denoted with capital letters and constant (deterministic) values with small letters. To fix ideas, x is a constant scalar, X is an \mathbb{R} -valued random variable, \mathbf{x} is a constant vector and \mathbf{X}

is a random vector. Vectors with equal scalar coordinates are denoted in bold, such as $\mathbf{0} = (0, \dots, 0)$ and $\mathbf{1} = (1, \dots, 1)$. The dimension will be clear from the context.

- Without further specifications, operations such as multiplication, division, max, exponentiation and boolean operations on vectors are meant componentwise. For instance, for two vectors $\mathbf{x}, \mathbf{y} \in \mathbb{R}^d$, $\mathbf{x}\mathbf{y} := (x_1y_1, \dots, x_dy_d)$. Boolean operations are meant componentwise: $\mathbf{x} \leq \mathbf{y} \Leftrightarrow \forall i \in [d], x_i \leq y_i$. The scalar product is denoted with a dot: $\mathbf{x} \cdot \mathbf{y} = \sum_{i=1}^d x_i y_i$.
- The cumulative distribution function (cdf) of a random variable \mathbf{X} is denoted by $F_{\mathbf{X}}$ or F according to the context. The cdf of the extreme-value distribution is denoted by G and the associated domain of attraction is denoted by $\text{Dom}(G)$.
- $\text{supp}(P)$ denotes the support of P .
- The Pareto distribution is defined by its cdf $1_{x \geq 1} (1 - x^{-1/\gamma})$ with $\gamma > 0$. When introducing a Pareto distribution, we might also parametrize it with $\alpha = 1/\gamma$.
- $\#A$ denotes the cardinal of a set A .

1.2 Reproducing dependence in extreme regions with Generative Adversarial Networks: Theory

Some background on extreme-value theory is briefly given in Section 1.2.1. Section 1.2.2 provides the tools necessary to study dependence in extreme regions of a distribution. Finally, some key elements on Generative Adversarial Networks (GANs) are recalled in Section 1.2.3. Additional material on each topic is also given in the Supplementary (Appendix 3.A). Finally, our theoretical results are presented in Section 1.2.4: we provide a quantization bound on how well GANs with heavy-tailed noise (HTGANs) approximate the stdf of a target distribution.

1.2.1 Extreme-value theory

Extreme-value theory is a branch of statistics concerned with studying the upper tails of probability distributions. We refer to Appendix 3.B in the Supplementary for a brief introduction to this theory. Given a \mathbb{R}^d valued random vector \mathbf{X} with associated cdf $\mathbf{x} \in \mathbb{R}^d \mapsto F(\mathbf{x}) = \mathbb{P}(\mathbf{X} \leq \mathbf{x})$, extreme-value theory establishes the asymptotic distribution of normalized maxima of realizations of \mathbf{X} . Whence normalized to unit Fréchet margins (1.28), the limiting distribution function G_* is a simple max-stable distribution, see Definition 1.6 and Equation (1.27). As illustrated in the next paragraph, G_* encodes the dependence structure in the tails of F .

1345 1.2.2 Measuring dependence in extremes

1346 We first recall the definition of the stable tail dependence function (stdf) associated with
 1347 F .

1348 **Definition 1.1.** (L. d. Haan and Ferreira 2006, Corollary 6.1.4) Let F be a continuous
 1349 cdf on \mathbb{R}^d belonging to the max-domain of attraction of some distribution. Then, one can
 1350 define

$$\forall \mathbf{x} \in (0, \infty)^d : \ell_F(\mathbf{x}) = \lim_{t \rightarrow +\infty} t \mathbb{P} \left(1 - F_1(X_1) \leq \frac{x_1}{t} \text{ or } \dots \text{ or } 1 - F_d(X_d) \leq \frac{x_d}{t} \right), \quad (1.2)$$

1351 where F_j 's are the marginal cdf of \mathbf{X} . The limit ℓ_F is the so-called *stable tail dependence*
 1352 *function* (stdf) of F .

1353 The stdf provides a “normalized” representation of the dependence structure of the
 1354 distribution in the tails (as opposed to the copula function which globally models the
 1355 dependence), in the sense that its marginals are normalized: the stdf does not carry any
 1356 information on the marginal distributions of the base distribution of interest.

1357 The following proposition gives the classical parametrization of the dependence struc-
 1358 ture in terms of spectral measure. A more recent and less known parametrization using
 1359 D-norms is provided in Theorem 1.1 below.

1360 **Proposition 1.1** (Spectral representation of the stdf, (L. d. Haan and Ferreira 2006,
 1361 Remark 6.1.16.)). Consider a cdf F (as in Definition 1.1) and its stdf ℓ_F . Then, there
 1362 exists a probability measure Λ on the simplex $\Delta_{d-1} := \{\boldsymbol{\omega} \geq \mathbf{0} : \sum_{i \in [d]} \omega_i = 1\}$ such that:

$$\forall \mathbf{x} \in (0, \infty)^d : \ell_F(\mathbf{x}) = d \int_{\Delta_{d-1}} \max_{i \in [d]} (\omega_i x_i) \Lambda(d\boldsymbol{\omega}), \quad (1.3)$$

1363 with constraints: $\forall i \in [d] : \int_{\Delta_{d-1}} \omega_i \Lambda(d\boldsymbol{\omega}) = 1/d$.

1364 **Quantifying dependence in extremes with D-norms.** D-norms (Falk 2019) are a
 1365 class of norms in vector space, very convenient for studying dependence in extremes.

1366 **Definition 1.2** (D-norms, (Falk 2019, Lemma 1.1.3)). Let $\boldsymbol{\Gamma} = (\boldsymbol{\Gamma}_1, \dots, \boldsymbol{\Gamma}_d) \in \mathbb{R}^d$ be a
 1367 random vector, whose components satisfy for all $i \in [d]$, $\boldsymbol{\Gamma}_i \geq 0$ a.s. and $\mathbb{E}[\boldsymbol{\Gamma}_i] = 1$. Then,

$$\|\cdot\|_{\boldsymbol{\Gamma}} : \mathbf{x} \mapsto \|\mathbf{x}\|_{\boldsymbol{\Gamma}} = \mathbb{E} \left[\max_{i \in [d]} (|x_i| \boldsymbol{\Gamma}_i) \right] \quad (1.4)$$

1368 defines a norm on \mathbb{R}^d , called a D-norm, and $\boldsymbol{\Gamma}$ is called a generator of the D-norm.

1369 Main properties of the D-norm are recalled in Appendix 3.C. To alleviate the role of
 1370 the generator, we may simply denote the above norm by $\|\cdot\|_D$. The key role of D-norms is
 1371 highlighted in the next result.

1372 **Theorem 1.1.** Let F be a continuous cdf on \mathbb{R}^d belonging to the max-domain of some
1373 distribution. Then, there exists a D-norm that exactly represents the stdf of F :

$$\forall \mathbf{x} \in (0, \infty)^d, \quad \ell_F(\mathbf{x}) = \|\mathbf{x}\|_D. \quad (1.5)$$

1374 This is a consequence of Theorem 1.6 in the Supplementary. Note that the two rep-
1375 resentations in (1.3) and in (1.4)–(1.5) may look similar at first sight, but actually they
1376 are substantially different. While both representations involve a probability measure (the
1377 spectral measure Λ on the one hand, and that of the generator Γ associated with the
1378 D-norm on the other hand), the second representation is more tractable for approximation
1379 purposes: Instead of designing a non-negative measure on the simplex Δ_{d-1} , one *simply*
1380 has to define a random vector $\Gamma \in [0, +\infty)^d$ with unit expectation.

1381 **Definition 1.3.** Let F be a cdf satisfying the assumptions of Theorem 1.1. The stdf ℓ_F
1382 is said to be *discrete*, with N atoms, when its D-norm is related to a generator Γ whose
1383 distribution is *discrete* and made of N atoms.

1384 **Example 1.1.** Let us consider the special case where $d = N$. Note that it can be typically
1385 the case of the input of a generative model with a latent distribution of dimension N . Let F
1386 be a cdf in \mathbb{R}^N with independent margins. Then, let $\ell_F(\mathbf{x}) = \sum_{i \in [N]} x_i, \quad \forall \mathbf{x} \in (0, \infty)^N$.
1387 This stdf is discrete with N atoms: indeed, one can easily check that, if Γ has the discrete
1388 uniform distribution on $\{N \mathbf{e}_1, \dots, N \mathbf{e}_N\}$, then $\mathbb{E}[\Gamma_i] = 1$ for any $i \in [N]$, and its D-norm
1389 is given by

$$\|\mathbf{x}\|_D = \mathbb{E} \left[\max_{i \in [N]} (x_i \Gamma_i) \right] = \frac{1}{N} \sum_{i \in [N]} x_i N = \ell_F(\mathbf{x}), \quad \forall \mathbf{x} \in (0, \infty)^N. \quad (1.6)$$

1390 With similar arguments, it is easily checked that the spectral measure is the discrete
1391 uniform distribution on $\{\mathbf{e}_1, \dots, \mathbf{e}_N\}$.

1392 The concept of Definition 1.3 is at the core of our numerical scheme, since we will
1393 consider approximations based on discrete distributions for the generator Γ in Section 1.2.4.

1394 **Characterization of simple max-stable distribution using D-norms and homo-**
1395 **geneous function.** We aim at giving another criterion to establish the max-domain of
1396 attraction of a simple max-stable distribution. It will be crucially used in the proof of our
1397 main result (Theorem 1.3).

1398 **Definition 1.4** (1-homogeneous function). A function $h : [0, \infty)^d \rightarrow [0, \infty)^{d'}$ is 1-
1399 homogeneous if it satisfies: $\forall \mathbf{x} \in [0, \infty)^d, \forall \lambda \geq 0 : h(\lambda \mathbf{x}) = \lambda h(\mathbf{x})$. The set of 1-
1400 homogeneous continuous functions from \mathcal{X} to \mathcal{Y} is denoted by $\mathbb{H}_1(\mathcal{X}, \mathcal{Y})$.

1401 The next result claims that identifying the stdf (or equivalently the D-norm) of a cdf
1402 F (with asymptotically simple max-stable distribution) is equivalent to analyze tails of
1403 1-homogeneous transforms of the associated random vector.

1404 **Theorem 1.2** ((Falk and Fuller 2021, Theorem 4.4)). Let \mathbf{X} be a random vector in
1405 $[0, +\infty)^d$ with continuous cdf F and let $\|\cdot\|_D$ be a D-norm in \mathbb{R}^d generated by $\boldsymbol{\Gamma}$. The
1406 following equivalence holds:

$$\forall \mathbf{x} \in (0, \infty)^d : F^t(t\mathbf{x}) \xrightarrow[t \rightarrow \infty]{} \exp - \left\| \frac{1}{\mathbf{x}} \right\|_D \quad (1.7)$$

$$\iff \forall h \in \mathbb{H}_1([0, \infty)^d, [0, \infty)) : t\mathbb{P}(h(\mathbf{X}) > t) \xrightarrow[t \rightarrow \infty]{} \mathbb{E}[h(\boldsymbol{\Gamma})]. \quad (1.8)$$

1407 An indirect consequence of the above theorem is that the limiting value $\mathbb{E}[h(\boldsymbol{\Gamma})]$ does
1408 not depend on the chosen generator of the D-norm, when h is a continuous 1-homogeneous
1409 function, see (Falk and Fuller 2021, Theorem 3.5) for a direct proof.

1410 1.2.3 Generative Adversarial Networks

1411 **Neural networks.** We refer to (Murphy 2022, Chapter 13) for a thorough introduction
1412 on neural networks. Let us consider fully connected neural networks, also called Multi
1413 Layer Perceptron (MLP). These neural networks consist of a chain of L operations of the
1414 form:

$$\mathbf{z}_l = f_l(\mathbf{z}_{l-1}) = \varphi_l(\mathbf{b}_l + \mathbf{W}_l \mathbf{z}_{l-1}) \quad (1.9)$$

1415 for $l \in [L]$, where $\mathbf{z}_l \in \mathbb{R}^{q_l}$, $\varphi_l : \mathbb{R}^{q_l} \rightarrow (\mathbb{R}^+)^{q_l}$ is an activation function, q_l is the number
1416 of neurons on the layer l , \mathbf{W}_l is a matrix and \mathbf{b}_l is a bias vector. Note that \mathbf{z}_0 corresponds
1417 to the q_0 -dimensional input of the network, while \mathbf{z}_L is the q_L -dimensional output. In this
1418 work, we focus on neural networks with ReLU activation functions: $\varphi_l(\mathbf{z}) = \max\{\mathbf{z}, \mathbf{0}\}$.
1419 The parameters of the network are:

$$\theta = \{\mathbf{b}_1, \dots, \mathbf{b}_L, \mathbf{W}_1, \dots, \mathbf{W}_L\} \in \Theta \quad (1.10)$$

1420 (activation functions are fixed), where Θ is the space of parameters, the full Euclidean
1421 space without constraints.

1422 **Proposition 1.2.** A neural network parameterized by (1.9) with ReLU activation functions
1423 is a piecewise affine function. It is a continuous 1-homogeneous function when the bias
1424 vectors are zero.

1425 The proof is easy and left to the reader. Combining Proposition 1.2 with Theorem 1.2
1426 shows that a distribution (with simple max-stable distribution) transformed by a ReLU
1427 neural network remains a heavy-tailed distribution with the same tail parameter γ .

1428 **Generative Adversarial Networks.** Generative modeling is the task of approximating
1429 at best a distribution of interest, here written $\mathbf{X} \sim p_{\text{data}} \in \mathbb{R}^d$, and being able to draw
1430 samples from it. One way to do is to find a map \mathcal{G} which takes as input a random variable
1431 $\mathbf{Z} \sim p_{\mathbf{Z}} \in \mathbb{R}^N$ which is easy to sample, such that:

$$\mathcal{G}(\mathbf{Z}) \stackrel{d}{\approx} \mathbf{X}. \quad (1.11)$$

The random vector \mathbf{Z} is commonly referred as the latent noise and $p_{\mathbf{Z}}$ the latent distribution. Generative Adversarial Networks (I. Goodfellow et al. 2014) are widespread generative models. GANs consist of a generator \mathcal{G} and a discriminator \mathcal{D} . GANs play an adversarial game where the generator tries to mimic the true distribution generating the data whereas the discriminator tries to uncover fake data points simulated by the generator. Namely, GANs optimize for:

$$\min_{\mathcal{G}} \max_{\mathcal{D}} V(\mathcal{D}, \mathcal{G}) \quad \text{where} \quad V(\mathcal{D}, \mathcal{G}) = \mathbb{E}_{\mathbf{X} \sim p_{\text{data}}} [\log \mathcal{D}(\mathbf{X})] + \mathbb{E}_{\mathbf{Z} \sim p_{\mathbf{Z}}} [\log(1 - \mathcal{D}(\mathcal{G}(\mathbf{Z})))], \quad (1.12)$$

where p_{data} is the distribution of the data and $p_{\mathbf{Z}}$ is the latent distribution. Usually, p_{data} is chosen among the uniform or the Gaussian distribution. Both \mathcal{G} and \mathcal{D} are neural networks. In this work, we consider neural networks parameterized with MLPs (1.9), and, with the notations of (1.9), note that $q_0 = N$ and that $q_L = d$.

1.2.4 New approximation results for the stdf

Our goal is to assess how well the dependence in the extremes of a target distribution can be approximated with a GAN based on heavy-tailed independent input noise. To this end, our first result establishes the nature of the stdf of a neural network output, when used with heavy-tailed input noise.

Proposition 1.3. *Consider a random vector \mathbf{Z} with a discrete stdf $\ell_{\mathbf{Z}}$ with N atoms. Consider a neural network $G_{\theta, N}$, with ReLU activation functions and $L - 1$ hidden layers, where θ is the parametrization of the network (1.10) and N is the dimension of the input noise ($q_0 = N$). Assume that all the matrix weights $(\mathbf{W}_1, \dots, \mathbf{W}_L)$ are non-negative. Then, the stdf of the output $G_{\theta, N}(\mathbf{Z})$ is discrete with at most N atoms.*

See Section 3.F for the proof. It appears that the output of a neural network with heavy-tailed input noise has a discrete stdf. The next result states that it is possible to attain perfect reproduction of any discrete stdf with MLP neural networks.

Proposition 1.4. *Consider a heavy-tailed random vector \mathbf{X} with a discrete stdf ℓ with N atoms. There exists a neural network $\mathcal{G}_{\theta^*, N}$ with parametrization $\theta^* \in \Theta$ and input noise $\mathbf{Z} \in \mathbb{R}^N$ with i.i.d. unit Fréchet margins such that:*

$$\ell = \ell_{\mathcal{G}_{\theta^*, N}(\mathbf{Z})} \quad (1.13)$$

where $\ell_{\mathcal{G}_{\theta^*, N}(\mathbf{Z})}$ is the stdf of $\mathcal{G}_{\theta^*, N}(\mathbf{Z})$.

In words, any discrete stdf can be perfectly approximated by a neural network with i.i.d. unit Fréchet margins as input noise with sufficiently large dimension.

When the stdf of interest is not discrete, our final result proves that it can be nevertheless approximated using a MLP neural network with an arbitrary precision.

Theorem 1.3. *Let \mathbf{X} be a random vector in \mathbb{R}^d , with continuous cdf F , belonging to the max-domain of some distribution. Denote by ℓ_F its stdf that is related to a D -norm*

1465 (*Theorem 1.1*). Consider using a MLP neural network $\mathcal{G}_{\theta,N}$ with a single layer (i.e. $L = 1$)
1466 and with as latent noise a random vector of N i.i.d. unit Fréchet margins, and denote by
1467 $\ell_{\mathcal{G}_{\theta,N}}$ its stdf. Then,

$$\inf_{\theta \in \Theta} \sup_{\mathbf{x} \in (0, \infty)^d} \frac{|\ell_F(\mathbf{x}) - \ell_{\mathcal{G}_{\theta,N}}(\mathbf{x})|}{\|\mathbf{x}\|_\infty} \leq \epsilon(N), \quad (1.14)$$

$$\epsilon(N) \sim C(d)N^{-1/(d-1)} \text{ as } N \rightarrow \infty, \quad (1.15)$$

1468 where $C(d)$ is a constant depending on d and $\ell_F(\mathbf{1})$. The bound (1.14) is achieved with
1469 one layer neural networks.

1470 The intuition behind the bounds (1.14)-(1.15) can be developed as follows. Approximating
1471 the stdf amounts to approximating a probability measure in dimension d using N
1472 atoms; it is thus natural to expect an error bound of order $N^{-1/d}$. However, the random
1473 variable having this probability measure is subject to a realization or moment constraint
1474 (see the proof for details), which reduces the degrees of freedom by one, leading to an
1475 approximation error of at most order $N^{-1/(d-1)}$.

1476 In this previous theorem, note that the error bound is a worst case scenario. For
1477 example, when the stdf is discrete, the error is zero for N sufficiently large in virtue of
1478 Proposition 1.4. Let us highlight that, in view of the upper bound (1.15), the attainable
1479 precision of the neural network approximation increases with the dimension of the latent
1480 noise. However, the higher the dimension d , the slower is the convergence. One can then
1481 expect that accurate approximations of the dependence in high dimension would require
1482 high dimensional latent noise, which is consistent with the intuition. This will be illustrated
1483 in the Numerical Experiments section.

1484 In other words, the above result states a denseness property. We would like to thank
1485 one of the referees to spot the analogy with the results from (Fougères and Mercadier 2013,
1486 Lemma 9, Theorem 3). As a difference, we use other arguments (see the proof based on
1487 quantization tools) and, consequently, we are able to quantify the effect of both latent
1488 and ambient dimensions N and d . Let us finally note that this denseness property is used
1489 in (JanSSen and Wan 2020) to interpret the proposed k -means procedure as an inference
1490 method for max-linear models.

1491 1.3 Numerical experiments

1492 First, we describe in Section 1.3.1 the proposed algorithms for training and sampling
1493 from a dataset presenting heavy-tailed characteristics. Second, experiments performed
1494 on synthetic datasets generated using the Gumbel copula are detailed in Section 1.3.2 to
1495 compare the proposed method with the standard GAN approach. Finally, experiments on
1496 the stock market index S&P500 are presented in Section 1.3.3.

1497 1.3.1 Algorithms and implementation

1498 **Algorithms**

1499 We accommodate the original GAN algorithm (I. Goodfellow et al. 2014) to have a heavy-
 1500 tailed latent noise. This idea originates from (Huster et al. 2021). Algorithm 1 describes
 1501 the learning phase while Algorithm 2 summarizes the generation phase.

Algorithm 1: Generative modelling: learning algorithm

input :

- Dataset $\{\mathbf{X}^{(1)}, \dots, \mathbf{X}^{(n_{\text{data}})}\}$ of n_{data} i.i.d. sample points in \mathbb{R}^d .
- Fully connected initialized generator \mathcal{G} and discriminator \mathcal{D} .
- Estimators $\hat{F}_1, \dots, \hat{F}_d$ of marginal cdf.
- Tail parameter $\gamma > 0$ for the renormalization step.
- GAN hyperparameters: Network architecture, optimization parameters and callbacks (early stopping, maximum number of iterations ...). The parameters are dependent on the chosen implementation. See Section 3 for more details.

1 begin

1. For any $j \in [d]$, estimate the marginal cdf of X_j on the training set
 $\{X_j^{(1)}, \dots, X_j^{(n_{\text{data}})}\}$ by

$$x \in \mathbb{R} \mapsto \hat{F}_j(x) \in [0, 1]. \quad (1.16)$$

2. Transform the marginals to Fréchet distributions with tail parameter γ :

$$\forall i \in [n_{\text{data}}], j \in [d] : \tilde{X}_j^{(i)} = \left(\frac{1}{1 - \hat{F}_j(X_j^{(i)})} \right)^\gamma. \quad (1.17)$$

3. Train GAN over transformed data $\{\tilde{\mathbf{X}}^{(1)}, \dots, \tilde{\mathbf{X}}^{(n_{\text{data}})}\}$: optimize (1.12) where p_{data} is a componentwise Pareto distribution with tail parameter γ and independent margins.

return the optimal generator \mathcal{G}^* , and the estimators $\hat{F}_1, \dots, \hat{F}_d$ of the margins.

2 end

1502 **Remarks about the marginals.** In Subsection 1.3.3, which focuses on experiments
 1503 with real data, the estimator \hat{F} is taken to be the empirical cdf. While other choices are
 1504 possible – since estimating a one-dimensional cdf is a classical problem in statistics – this
 1505 simple choice suffices to run Algorithm 1 and obtain an approximation of the generator
 1506 \mathcal{G}^* . In doing so, we address the primary objective: reproducing the dependence structure
 1507 in the extremes. It is important to note that Algorithm 2 requires access to the inverse cdf

Algorithm 2: Generative modelling: sampling algorithm

input :

- Number of data points to sample n_{sample} .
- Tail parameter $\gamma > 0$.
- Estimators $\hat{F}_1, \dots, \hat{F}_d$ of the margins.
- Generator \mathcal{G}^* .

1 begin

1. Sample n_{sample} points $(\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(n_{\text{sample}})})$ from p_{data} , a componentwise Pareto distribution with tail parameter γ and independent margins.

2. Transform the noise through the generator:

$$\forall i \in [n_{\text{sample}}] : \tilde{\mathbf{X}}_{\text{sample}}^{(i)} = \mathcal{G}^*(\mathbf{z}^{(i)}). \quad (1.18)$$

3. Transform the sample points $\tilde{\mathbf{X}}_{\text{sample}}$ to the original scale with the inverse of the estimated margins:

$$\forall i \in [n_{\text{sample}}], j \in [d] : \hat{X}_j^{(i)} = \hat{F}_j^{-1}(\tilde{X}_j^{(i)}). \quad (1.19)$$

2 end**3 return** Sampled points $\left\{ \hat{\mathbf{X}}^{(1)}, \dots, \hat{\mathbf{X}}^{(n_{\text{sample}})} \right\}$.

1508 (*i.e.*the quantile function). To ensure an accurate fit of the marginals – while preserving
1509 the learned dependence – we recommend using a quantile estimator that performs well
1510 across the entire distribution, both in the bulk and in the tails. A natural candidate is the
1511 estimator based on eLU-activated neural networks proposed by (M. Allouche et al. 2024).

1512 **Implementation**

1513 **Code.** The implementation is based on a publicly available GitHub repository¹ using the
1514 machine learning library PyTorch². We adapt the source code so that both the generator
1515 and discriminator are fully connected MLPs (1.9). For writing a specific architecture of a
1516 neural network, we use a list that details all hidden layers, that is layers q_l for $l \in [L - 1]$
1517 in (1.9). As an example, a network parametrized by $[100, 200]$ means that it has 2 hidden
1518 layers, $L = 3$, $q_1 = 100$ and $q_2 = 200$. For the activation functions, we use LeakyReLU
1519 (Maas et al. 2013): note that our theoretical result (Theorem 1.3) still holds since a
1520 LeakyReLU is a difference of two ReLUs. For gradient descent, we use the optimizer Adam
1521 (D. P. Kingma and Ba 2014) and we do not tune the base parameters. Early stopping is
1522 used for regularization.

¹<https://github.com/eriklindernoren/PyTorch-GAN?tab=readme-ov-file#gan>

²<https://pytorch.org/>

In these experiments, we allow the networks to have several hidden layers. The proof of Proposition 1.4 only involves networks of depth one to handle the dependence in the extremes, but higher depths may be necessary to reproduce the dependence in the tails or in the bulk of the distribution. In (1.49), it is seen that a one layer transformation produces a heavy-tail term (left term) and a bounded term. Our aim is that on one side, we manage to catch the behaviour of the distribution tail with the left unbounded, 1-homogenous term and, on the other side, we hope to capture the behaviour of the bulk with the right bounded term.

Hardware. All experiments were performed on the cluster Cholesky³ of Ecole Polytechnique. Code was run on Intel Xeon CPU Gold 6230 20 cores @ 2.1 Ghz with 192 GB of memory.

1.3.2 Experiments on simulated data

We first propose in Section 1.3.2 a visual illustration of the behaviour of heavy-tailed latent noise models on a two dimensional dataset. Second, we investigate in Section 1.3.2 the performance of Algorithm 1 on higher dimensional data and perform a comparison with the baseline GAN algorithm using Gaussian input noise. Specifically, we study the model’s ability to reproduce data with both given dependence degree and tail heaviness. To this end, the target distribution is obtained by combining Gumbel, Gaussian and Hüsler-ReiSS copulas with Pareto margins. We refer to Section 3.E for basic material on copulas.

1.3.2 Two-dimensional data: an illustration

Let us first consider a toy dataset of size $n_{\text{data}} = 10,000$ generated by a Gumbel copula in dimension $d = 2$ with identically distributed Pareto margins, see the left panel of Figure 1.1 for an illustration. It appears that, in the tail, data are evenly distributed along the horizontal axis \mathbf{e}_1 and the vertical one \mathbf{e}_2 . As a comparison, the right panel of Figure 1.1 displays a plot of the input noise with independent Pareto random margins. In contrast to the Gumbel training data, the independent Pareto noise whence conditioned to being large in norm, is concentrated in the two directions \mathbf{e}_1 and \mathbf{e}_2 . This is due to the fact that the spectral measure associated with independent Pareto random variable is discrete, with weights on \mathbf{e}_1 and \mathbf{e}_2 .

The left panels of Figure 1.2 show simulated data generated by a GAN trained on the previously described dataset, using $N \in \{2, 3, 5, 10, 50\}$ independent Pareto random variables as input noise. The right panels display the histogram of the angles $\arccos(\hat{\mathbf{X}}_1 / \|\hat{\mathbf{X}}\|)$ for the 10% largest Euclidean norms $\|\hat{\mathbf{X}}\|$. For small latent dimensions ($N \in \{2, 3\}$), spikes are clearly visible in the tail of the distribution: the generated datasets exhibit as many spikes as the dimension of the latent noise. This is expected, since the spectral measure of i.i.d. Pareto noise is discrete, and, in accordance with Proposition 1.3, transforming N -dimensional independent Pareto noise (with a discrete spectral measure of N atoms)

³https://docs.idcs.mesocentre.ip-paris.fr/cholesky/hardware_description/

through a 1-homogeneous mapping yields a distribution whose spectral measure has at most N atoms. For larger dimensions ($N \in \{5, 10\}$), spikes become less pronounced, though still discernible, and the dependence structure appears visually closer to the target. For $N = 50$, the spikes are almost indistinguishable, and the scatter plot of the generated data closely resembles that of the real data (Figure 1.2-f). To sum up, these tests show that the larger N is, the closer the distribution gets to the target distribution, which supports the idea of convergence as N tends to infinity, in line with our theoretical results.

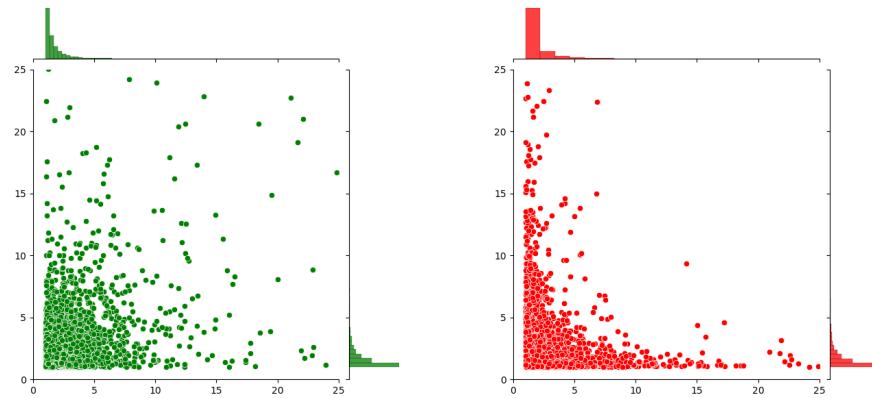


Figure 1.1: Illustration in a two-dimensional setting on simulated data of size $n_{\text{data}} = 10,000$ with Pareto margins and tail parameter $\alpha = 2.0$. Left panel: sample from a Gumbel copula with dependence parameter $\beta = 4/3$. Right panel: sample with two independent margins.

1568 Higher dimensional experiments on Gumbel copula

1569 After a thorough understanding of the behaviour of the model in a two dimensional setting,
 1570 we explore its performance on higher dimensional datasets, still using the Gumbel copula,
 1571 but with a variety of dependence and tail coefficients.

1572 **Experimental setup.** The trained data are sampled from the d -dimensional Gumbel copula (see Section 3.E) with $d \in \{2, 5, 10, 20, 50\}$ and dependence parameter $\beta \in \{4/3, 2, 4\}$, leading to Kendall's $\tau \in \{1/4, 1/2, 3/4\}$. See (1.41) for a definition of
 1573 Kendall's τ . The margins are chosen to be Pareto distributed with shape parameter
 1574 $\alpha \in \{1.5, 2.0, 2.5\}$. In each of the $5 \times 3 \times 3 = 45$ considered situations, $n_{\text{data}} = 10,000$
 1575 data points are simulated for training and 20,000 data points are considered for testing.
 1576 Note that the testing set is larger than the training set in order to assess the ability of
 1577 GANs to extrapolate to extreme regions unseen in training. In this synthetic case, we
 1578 provide the tail-index $\gamma = 1/\alpha$ to Algorithm 1 and focus on the quality of reproduction of
 1579 dependence in extreme regions.

1582 **Metrics.** Two metrics are considered for assessing the performance of the GANS: the
 1583 Absolute Kendall Error (AKE), see (M. Allouche et al. 2022, Section 3.2 and Appendix A)

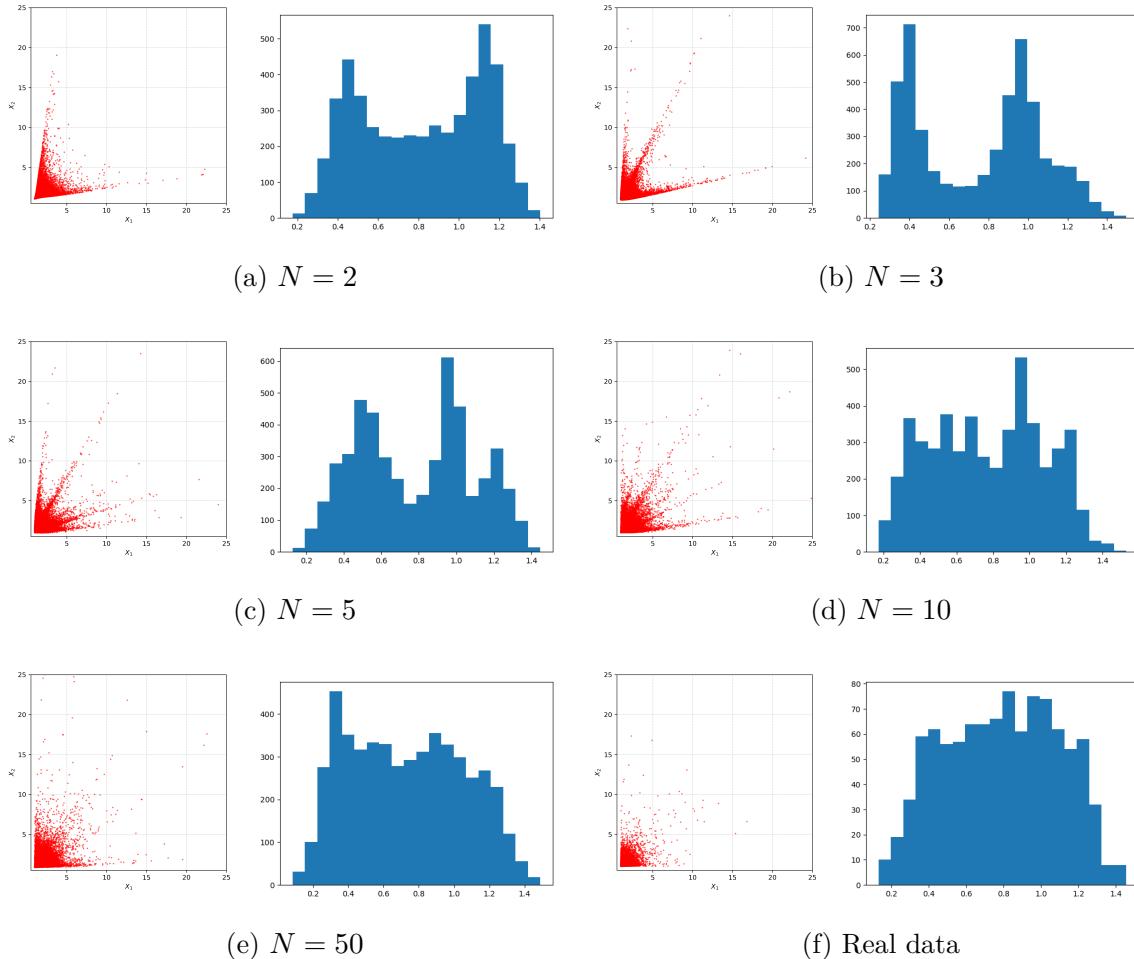


Figure 1.2: Illustration of real and generated data in a two-dimensional setting for data from a Gumbel copula with $\beta = 4/3$ and Pareto margins ($\alpha = 2.0$). For each sub-figure (a)-(f), on the left: scatter plot of data points; on the right: histogram of angles $\arccos(X_1/\sqrt{X_1^2 + X_2^2})$ for the 10% largest Euclidean norms. Subfigures (a)-(e): $n_{\text{gen}} = 50,000$ generated data after training with $n_{\text{data}} = 10,000$ points. Subfigure (f): real data.

for a definition and some background, and the Sliced Wasserstein Distance (SWD), see (1.42). Since the focus is on extreme regions, we compute the Euclidean norm of all data points and select data points whose norms are larger than the upper ξ -quantile, with $\xi \in \{90\%, 95\%, 99\%\}$. Both metrics are then computed (denoted by AKE_ξ and SWD_ξ) on the resulting points. This operation is performed both on the sample generated by GAN methods and on the dataset simulated from the true distribution. These new subsamples are then normalized on the Euclidean sphere, where the Sliced Wasserstein distance is estimated.

1592 Preliminary hyperparameter search. We explore a number of possible parametrizations
1593 of GAN models within a selected range of hyperparameters. For both the generator
1594 and the discriminator, fully connected neural networks are used with the following 15 hid-
1595 den dimensions: [50], [75], [100], [125], [200], [300], [400], [100, 100], [200, 200], [300, 300],

[400, 400], [100, 200, 100], [200, 400, 200], [100, 200, 200, 100] and [200, 400, 400, 200]. For the latent dimension, we explore values N ranging from 1 to 200. For the optimization parameter, values ranging (log-uniformly) from 10^{-6} to 10^{-1} are tested. We investigate the use of four batch sizes: {128, 256, 512, 1024}. A Bayesian search (see (J. Wu et al. 2019)) is implemented using weight and biases' sweep tool⁴. At first, some parametrizations are selected at random and, next, the probability for a configuration to be selected is updated with the use of a Bayesian rule. See the referred online documentation for more details on Bayesian hyperparameter tuning. Considering the top performing models of this research, the following conclusions can be drawn on the choice of hyper parameters: Among tested models, best performing generators (80%+) have a light parametrization (fewer parameters) [100] or [200], while best performing discriminators (80%+) have a heavy parametrization (more parameters) [100, 200, 200, 100] or [200, 400, 400, 200]. Bigger batch sizes yield better performances. Good values for the learning rate include range $[5 \cdot 10^{-6}, 5 \cdot 10^{-4}]$.

Setting up models for comparison. After this hyperparameter exploration, two versions of GAN are compared: the base version of GAN with a Gaussian noise $\mathbf{Z} \sim \mathcal{N}(0, \mathbf{I}_{N \times N})$, referred to as LTGAN (LT for light-tailed), and our version, Algorithm 1, referred to as HTGAN (HT for heavy-tailed). Thanks to the previous insights on the preliminary hyperparameter search, we are able to choose a subset of evaluation hyperparameters to compare both models. Both models are run for each configuration. In order to have an exhaustive comparison, a grid search is performed on the following range of hyperparameters: $N \in \{2, 5, 10, 20, 50, 80\}$, hidden dimensions ranging in $\{[100, 200, 100], [200, 400, 200], [100, 200, 200, 100], [200, 400, 400, 200]\}$ for the discriminator, in $\{[100], [200]\}$ for the generator, $\{10^{-4}, 10^{-5}\}$ for the learning rate.

Results. It appears in Table 1.1 that HTGAN performs better than LTGAN with respect to all six metrics and all considered dependence coefficients $\beta \in \{4/3, 2, 4\}$ in the heavier tail setting $\alpha = 1.5$ (all percentages are above 50%). For $\alpha = 2.5$, it appears that HTGAN fails to outperform the baseline on the $\beta = 4, \alpha = 2.5$ setting. It suggests that in low tail-index settings, light-tailed outputs still manage to provide a good approximation of the target distribution, even though the target is heavy-tailed. For almost all experiment settings, HTGAN is comparatively better than LTGAN for the highest tail coefficient. For the sliced Wasserstein distance, it appears that HTGAN gets comparatively better as both β and $\gamma = 1/\alpha$ increase. This result is coherent with our theoretical development: It is known that (piecewise) linear transformations of Gaussian variables cannot generate (asymptotic) dependence in the tails (Wiese et al. 2020), a case which is met when β gets larger and larger.

Figures 1.3 and 1.4 plot the performance of both methods for $\alpha = 1.5$ and $\alpha = 2.5$, with varying values of β . With $\alpha = 1.5$, the point cloud is well above the median and our

⁴<https://docs.wandb.ai/guides/sweeps>

1634 method is better performing than the baseline. In the case $\alpha = 2.5$, the difference is not
1635 neat and no clear conclusion can be drawn.

Table 1.1: Results on simulated data from a Gumbel copula. Percentage (%) of parametrizations for which HTGAN is better than a LTGAN for the six considered metrics. Results are given with precision $\pm 0.1\%$. Values are averaged over data dimensions $d \in \{2, 5, 10, 20, 50\}$.

α	$\beta = 4/3$			$\beta = 2$			$\beta = 4$		
	1.5	2	2.5	1.5	2	2.5	1.5	2	2.5
AKE_90	65.8	68.1	65.6	80.4	78.4	66.7	88.7	86.6	72.6
AKE_95	65.6	69.5	66.4	79.7	79.4	69.2	87.8	87.2	75.1
AKE_99	62.5	67.4	65.1	77.0	80.1	72.7	86.8	88.0	78.4
SWD_90	81.7	69.1	45.9	76.7	67.7	52.6	63.7	54.7	41.4
SWD_95	84.5	74.3	51.1	76.9	66.7	54.3	62.0	52.2	39.5
SWD_99	86.4	81.5	64.8	77.5	59.5	49.6	56.9	48.3	34.4

1636 **Latent dimension.** Equation (1.14) gives evidence that a higher value of the latent di-
1637 mension N is necessary to obtain a better approximation of the dependence structure. For
1638 each experiment specification (*i.e.*a value for α and for d), we have performed a Bayesian
1639 hyperparameter search (J. Wu et al. 2019) with 400 runs, with the same hyperparameter
1640 ranges as for the previous experiment section. Table 1.2 summarizes statistics on the la-
1641 tent dimension, $\hat{\alpha}$ and SWD_90 for the top 5% run on each experiment. It appears that,
1642 empirically, the theoretical result is verified: generally, for the various values of α , as the
1643 dimension of the data increases, the optimal value of the latent dimension increases. We
1644 also note that, for each value of α , as the dimension of the data increases, the estimate
1645 $\hat{\alpha}$ increases accordingly, which means lower tails. However, this conclusion is to be inter-
1646 preted with caution in view of the great variability associated with the statistics (reflected
1647 by the values taken by min and max). This may be a consequence of the instability in
1648 training generative models with heavy-tailed inputs (Hickling and Prangle 2024).

1649 Other models: Gaussian and Hüsler-ReiSS copulas

1650 We extend our evaluation to higher dimensions using two dependence structures: the
1651 multivariate Hüsler–ReiSS (HR) extreme-value copula and the Gaussian copula. These
1652 represent two qualitatively different tail regimes: HR exhibits non-trivial asymptotic tail
1653 dependence, whereas the Gaussian copula is asymptotically independent.

1654 **Background: multivariate Hüsler–ReiSS.** The HR model is a max-stable (extreme-
1655 value) dependence model that arises as the limit of componentwise maxima of *log-Gaussian*
1656 vectors, see (Hüsler and Reiss 1989b). In dimension d , the dependence is encoded by a
1657 matrix $(\lambda_{ij})_{1 \leq i,j \leq d}$ with $\lambda_{ii} = 0$ and $\lambda_{ij} > 0$ for $i \neq j$. Smaller λ_{ij} means stronger extremal
1658 dependence between components i and j ; larger λ_{ij} weakens it. In the bivariate case, the

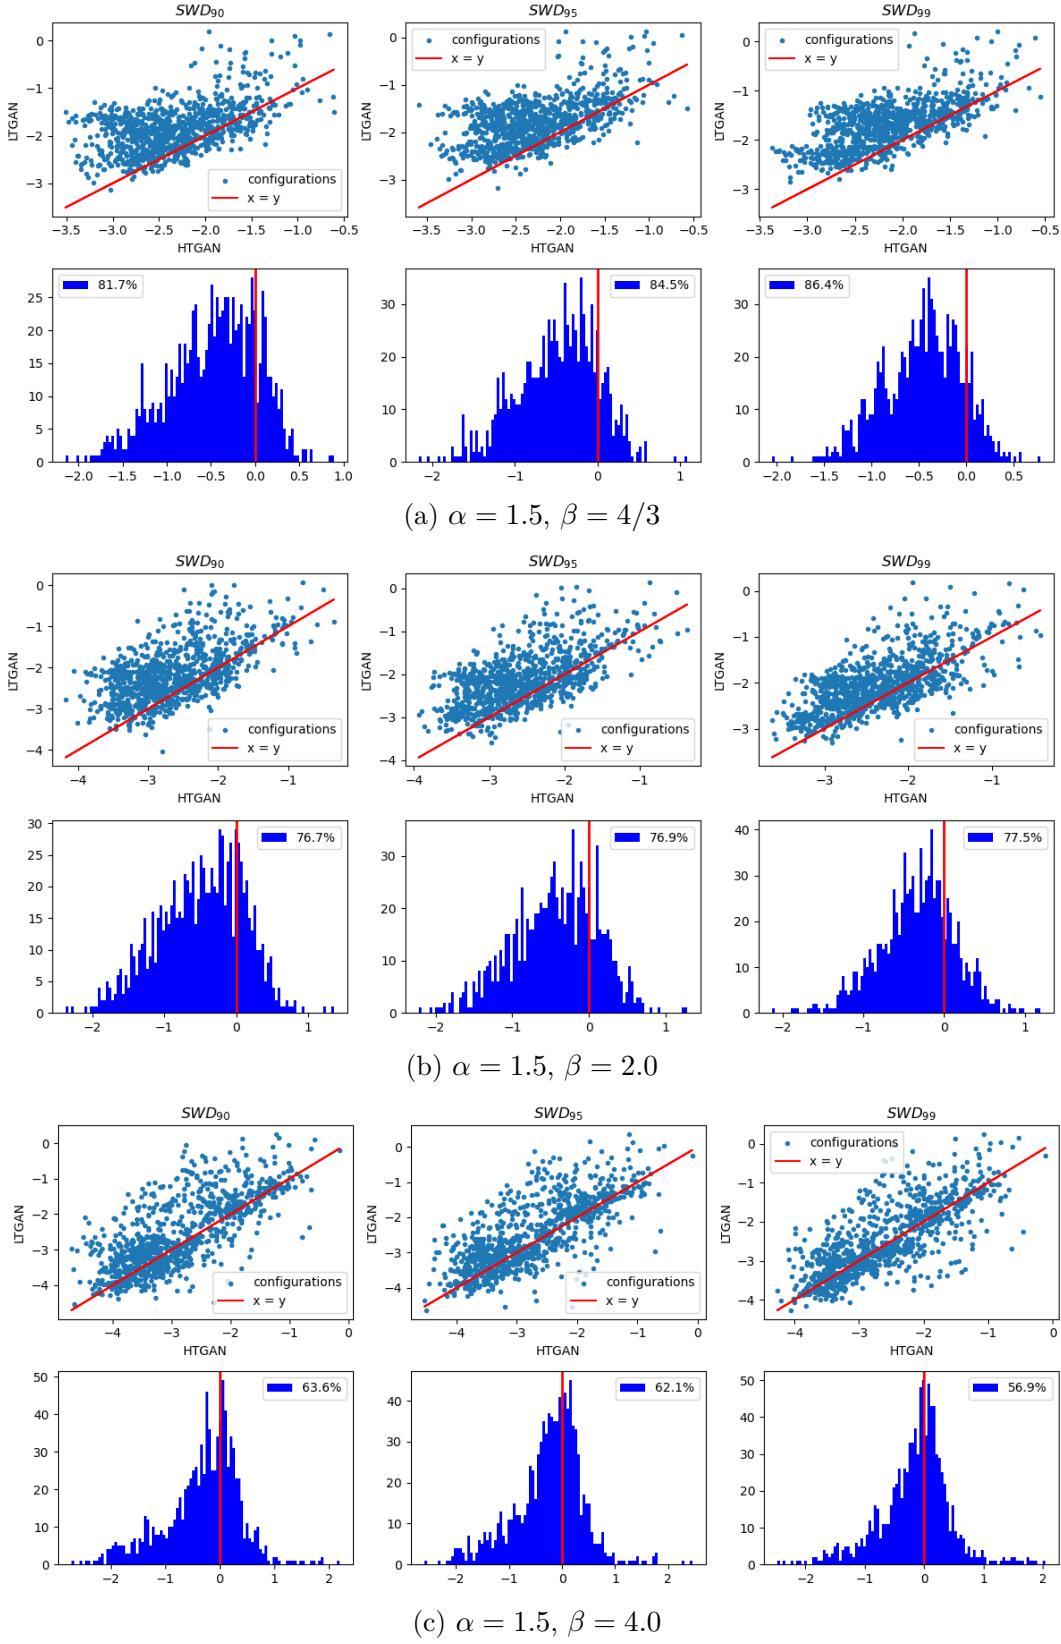


Figure 1.3: Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} = 10,000$ from a Gumbel copula (with dependence parameter β) and Pareto margins (with tail parameter $\alpha = 1.5$). Each group of 2×3 plots corresponds to one experimental setting, i.e. a specification of α and β . Figures are averaged over dimensions $d \in \{2, 5, 10, 20, 50\}$. In each of these three groups: the first row plots the specified metric of HTGAN vs LTGAN in log scale for a given hyperparameter configuration. The second row is a histogram of the difference of the log of the metric for HTGAN vs LHTGAN for each parametrization. The legend corresponds to the proportion of cases where HTGAN performs better.

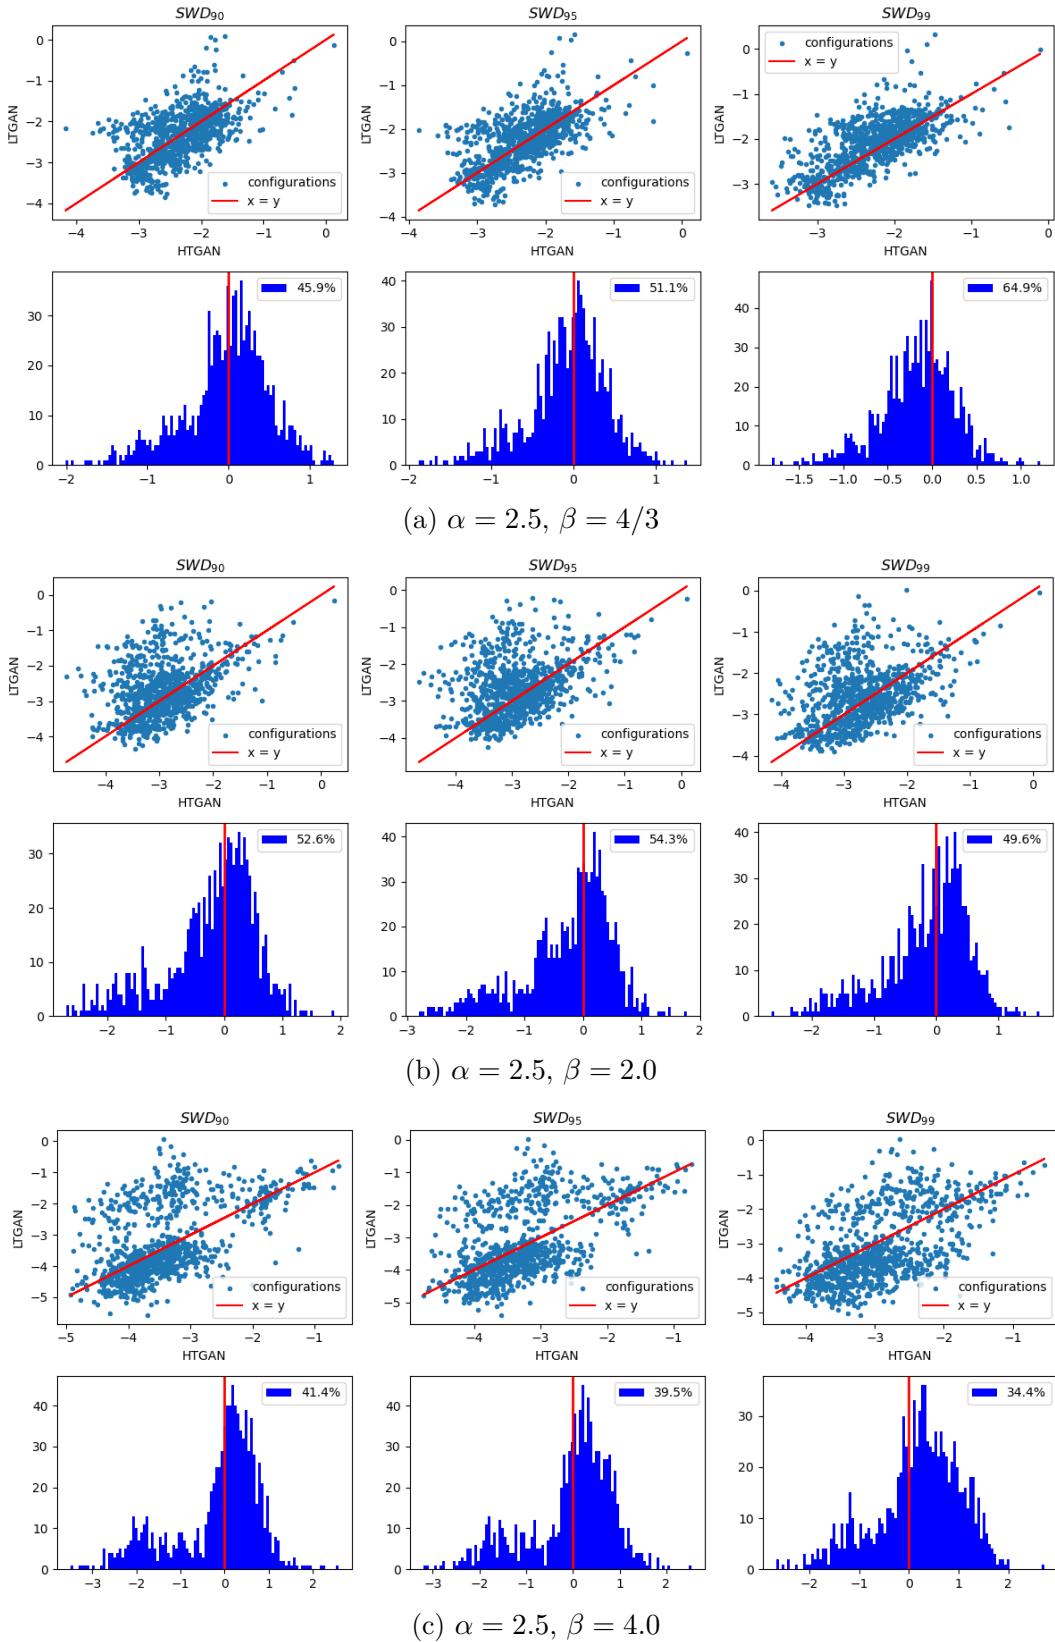


Figure 1.4: Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} = 10,000$ from a Gumbel copula (with dependence parameter β) and Pareto margins (with tail parameter $\alpha = 2.5$), $d \in \{2, 5, 10, 20, 50\}$. See Figure 1.3 for further details.

Table 1.2: Results on simulated data from a Gumbel copula. Statistics (mean, min, max) on three on the top 5% best performing runs of HTGAN (best performing w.r.t. the SWD_90 metric). Statistics are averaged over $\beta \in \{4/3, 2, 4\}$.

α	d	N (latent dimension)			$\hat{\alpha}$			SWD_90		
		mean	min	max	mean	min	max	mean	min	max
1.5	2	75.96	10	195	1.53	1.16	1.72	0.03	0.01	0.06
	5	91.07	17	197	1.61	1.20	1.96	0.04	0.01	0.06
	10	90.69	21	197	1.72	1.30	2.23	0.04	0.01	0.07
	20	91.33	26	182	1.77	1.23	2.64	0.03	0.01	0.07
	50	107.98	22	188	2.05	1.31	4.00	0.02	0.01	0.04
2.0	2	40.71	5	122	2.07	0.70	3.44	0.04	0.01	0.08
	5	55.96	11	164	2.41	1.63	3.68	0.04	0.01	0.12
	10	57.36	10	187	2.83	1.49	4.66	0.04	0.01	0.10
	20	64.60	10	191	3.03	1.06	4.88	0.03	0.01	0.07
	50	69.02	14	187	3.42	1.63	5.27	0.02	0.01	0.04
2.5	2	39.24	2	154	2.82	1.27	4.92	0.04	0.01	0.10
	5	50.20	7	177	3.54	1.33	7.63	0.05	0.01	0.13
	10	57.40	9	174	4.17	1.31	7.40	0.04	0.01	0.09
	20	72.98	12	175	4.32	0.97	8.79	0.03	0.01	0.06
	50	58.02	9	169	4.44	1.07	7.97	0.02	0.01	0.04

1659 upper tail dependence coefficient is

$$\chi_{ij} = 2\{1 - \Phi(\lambda_{ij})\}, \quad (1.20)$$

1660 where Φ is the cumulative distribution function of the standard Gaussian distribution.
1661 Clearly, $\chi_{ij} \downarrow 0$ as $\lambda_{ij} \uparrow \infty$. A convenient spectral construction is via a centered Gaussian
1662 vector $\mathbf{Z} \in \mathbb{R}^d$ with covariance Σ derived from (λ_{ij}) (Brown–Resnick representation): Set
1663 $\mathbf{Y} = \exp \mathbf{Z}$, normalize $\mathbf{S} = \mathbf{Y} / \sum_k Y_k$, draw $E \sim \text{Exp}(1)$, and form a unit-Fréchet
1664 HR vector $\mathbf{X}^{\text{Fr}} = \mathbf{S}/E$. Arbitrary continuous margins are then obtained by monotone
1665 transforms.

1666 **Background: Gaussian copula.** Let $\mathbf{Z} \sim \mathcal{N}_d(0, \Sigma)$ with equicorrelation $\rho \in (-1/(d-1), 1)$, set $U_i = \Phi(Z_i)$ and apply marginal transforms to obtain X from a Gaussian copula.
1667 For $|\rho| < 1$ the Gaussian copula is asymptotically independent ($\chi = 0$), but larger ρ
1668 increases sub-asymptotic dependence.

1669 More details on the specific experimental design for these two models are given in
1670 Subsection 3.I.

1672 **Qualitative summary.** We summarize results in Table 1.3. Across both datasets, the
1673 trends mirror those observed for the Gumbel copula: heavier tails (lower α) consistently
1674 amplify the relative advantage of HTGAN over LTGAN. Therefore, both in asymptotic de-
1675 pendence and in independence settings, HTGAN outperforms LTGAN. Additional figures
1676 are provided in Appendix 3.I.

Table 1.3: Results on simulated data from Hüsler-ReiSS and Gaussian copulas. Percentage (%) of parametrizations for which HTGAN is better than a LTGAN for the six considered metrics. Results are given with precision $\pm 0.1\%$. Values are averaged over data dimensions $d \in \{2, 5, 10, 20, 50\}$.

α	Hüsler-ReiSS			Gaussian		
	1.5	2	2.5	1.5	2	2.5
AKE_90	78.3	81.9	70.8	62.3	57.0	42.9
AKE_95	76.7	78.6	67.4	59.2	56.4	47.2
AKE_99	65.6	73.3	60.2	58.3	53.8	46.2
SWD_90	78.6	63.3	34.4	96.9	96.7	78.3
SWD_95	75.8	68.6	45.6	94.7	93.6	79.4
SWD_99	48.3	44.2	34.2	92.2	90.0	77.8

1677 1.3.3 Experiments on real data

1678 The proposed HTGAN method is tested on the publicly available dataset⁵ consisting of
 1679 daily data for companies of the S&P500. The S&P500 is a stock market index track-
 1680 ing of the 500 most valuable companies in the United States of America. The sig-
 1681 nal of interest is the absolute value of the normalized marginal returns for each stock:
 1682 $r_t = |(X_{t+1} - X_t) / X_t|$, which are bounded below by zero and only have an upper tail. The
 1683 dataset consists of $n_{\text{data}} = 1259$ normalized marginal returns for each ticker. Companies
 1684 are organized using the Global Industry Classification Standard (GICS). This classification
 1685 divides companies into 11 sectors: Energy, Materials, Industrials, Consumer Discretionary,
 1686 Consumer Staples, Health Care, Health Care, Financials, Information Technology, Com-
 1687 munication Services, Utilities and Real Estate. Each sector is divided further into industry
 1688 groups, industry and sub-industry. The returns of the S&P500 present a particular corre-
 1689 lation structure illustrated in Figure 1.5 by a heatmap of the table consisting of estimated
 1690 Kendall's tau for every pair of stock ticker, see the Appendix for more details on Kendall's
 1691 tau coefficient. It appears that returns are more correlated within each sector, especially
 1692 in both the Financials and the Utilities sectors, than between sectors. Therefore, in the
 1693 following, the performance of LTGAN and HTGAN are assessed on the two subsets associ-
 1694 ated with sectors Financials and Utilities, having both strong correlations (see Figure 1.5).
 1695 Note that the number of tickers for the Financials and Utilities sector are not the same:
 1696 $d = 27$ for the Financial sector, and $d = 57$ for the Utilities sectors.

1697 Before moving to the results, let us mention that on the application side of portfolio
 1698 management, it is important to model the dependence of extreme asset returns well, as
 1699 shown in (Mainik et al. 2015), in order to manage risks more accurately. This supports
 1700 the use of HTGAN compared to LTGAN. Their study compares the traditional Markowitz
 1701 approach with the Extreme Risk Index – a method based on multivariate EVT – to optimize
 1702 a portfolio of S&P 500 stocks: their results show the superiority of a model that better
 1703 accounts for dependencies in the tails, compared to the one based only on covariances.

⁵<https://www.kaggle.com/datasets/camnugent/sandp500>

1704 **Sectors and subsectors.** The subsectors for the financial sector are: Asset Management
1705 & Custody Banks, Consumer Finance, Diversified Banks, Financial Exchanges & Data, In-
1706 surance Brokers, Investment Banking & Brokerage, Life & Health Insurance, Multi-Sector
1707 Holdings, Multi-line Insurance, Casualty Insurance, Regional Banks and Processing Ser-
1708 vices. For the Utilities sector, Subsectors are: Electric Utilities, Gas Utilities, Independent
1709 Power Producers & Energy Traders, Multi-Utilities and Water Utilities.

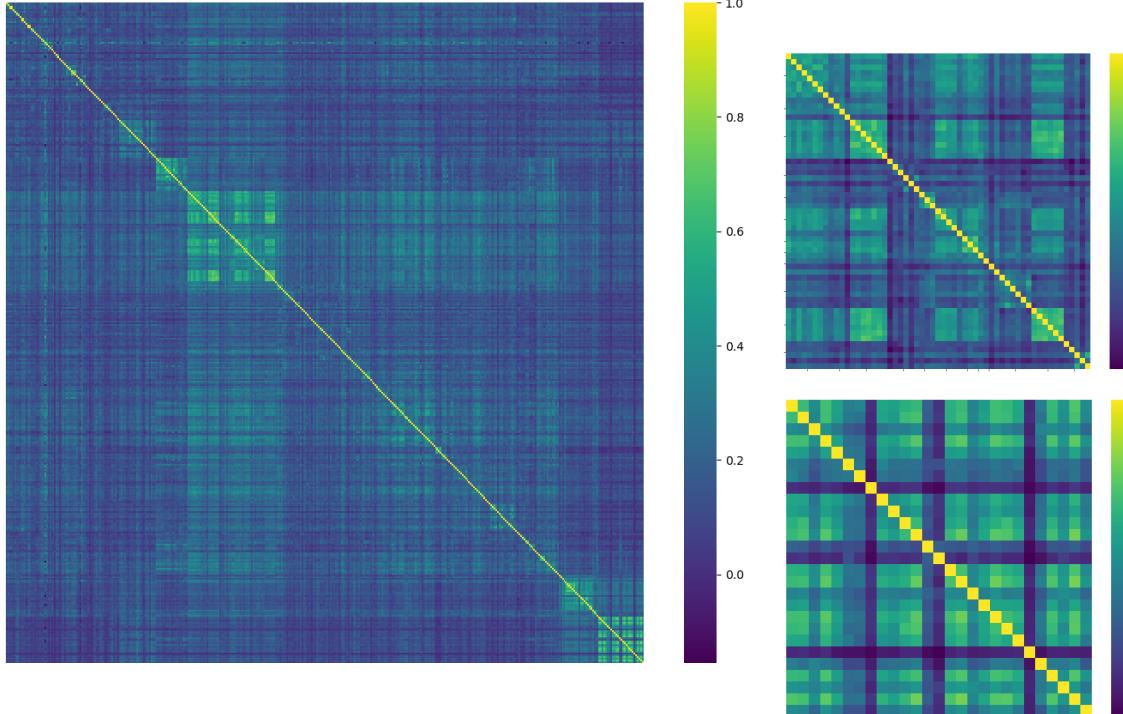


Figure 1.5: Estimated Kendall’s τ for every pair of index considered. Left: every S&P500 ticker. On the top right: Tickers for the Financials sector. On the bottom right: Tickers for the Utilities sector. A comprehensive list of sectors and subsectors is given in Paragraph 1.3.3.

1710 **Data normalization.** In the normalizing step (1.17) of Algorithm 1, two values of $\gamma =$
1711 $1/\alpha$ are investigated: $\alpha = 1.5$ and $\alpha = 2.5$. The cumulative distribution function is
1712 estimated using its empirical counterpart:

$$\widehat{F}_j(x) = \frac{1}{n_{\text{data}} + 1} \sum_{i=1}^{n_{\text{data}}} \mathbb{I} \left\{ X_j^{(i)} \leqslant x \right\}. \quad (1.21)$$

1713 This transformation is not invertible, and therefore it is not possible to proceed to step (1.19)
1714 in the generation step of Algorithm 2. Therefore, we skip this step in the generation pro-
1715 cess. In our evaluation, we solely focus on the good reproduction of dependence in extreme
1716 regions (*i.e.*the difficult task). Our criteria do not assess how well the marginals are fitted.

1717 **Results.** Figure 1.6 presents the results for experiments run on the S&P500 dataset. The
1718 results argue in favor of better performance of the heavy-tailed noise setting when data is

renormalized to Pareto margins with a larger tail index (case $\alpha = 1.5$). This can be seen in the proportion of cases in which HTGAN performs better than LTGAN: 62.5%, 64.3% and 63.7% respectively for SWD_50, SWD_80 and SWD_90 metrics. In the case of a lighter tail, $\alpha = 2.5$, the performance of HTGAN is degraded, with relative performances of 38.7%, 30.4% and 34.5% respectively for SWD_50, SWD_80 and SWD_90 metrics. Therefore, for better accuracy in extreme regions, we argue in favor of a smaller α in the renormalization step. In summary, HTGAN is generally more effective than LTGAN on this dataset. In addition to this conclusion, we provide visual illustrations of the performance of HTGAN on the Financials dataset in Figure 1.11 of the supplementary material.

1.4 Conclusion

We have provided a theoretical framework to analyse how a HTGAN behaves in extreme regions. We have established a bound on the quality of approximation of the stable tail dependence function of the target distribution. This analysis builds upon the deep connections between Extreme Value Theory and the theory of D-norms.

Our numerical experiments support the theoretical results. They show that HTGAN significantly outperforms a standard LTGAN in capturing complex dependencies in the extremes, both on synthetic heavy-tailed datasets with varying dependence structures and on real-world data. We have also demonstrated the key role of the latent dimension in enhancing the model’s ability to reproduce such extreme dependencies.

This work highlights, both theoretically and empirically, how introducing a large latent dimension into generative models can improve the modelling of complex dependence structures in extreme regions. While our application focus here has been on financial data, extending this methodology to other domains – such as image generation – remains an exciting direction for future work.

Some important challenges are left open. In particular, we have not addressed the training instability that may arise when working with heavy-tailed inputs. Furthermore, there is ambiguity in selecting the tail parameter α for the marginal normalization step. We suggest treating α as a hyperparameter to be tuned during training, depending on the application and the data at hand.

Acknowledgements

The authors thank the referees for their constructive remarks which improve the manuscript. They acknowledge the support of the Chair “Stress Test, Risk Management and Financial Steering”, led by the French Ecole Polytechnique and its Foundation and sponsored by BNP Paribas. S. Girard also gratefully acknowledges support from the French National Research Agency under the grant ANR-23-CE40-0009.

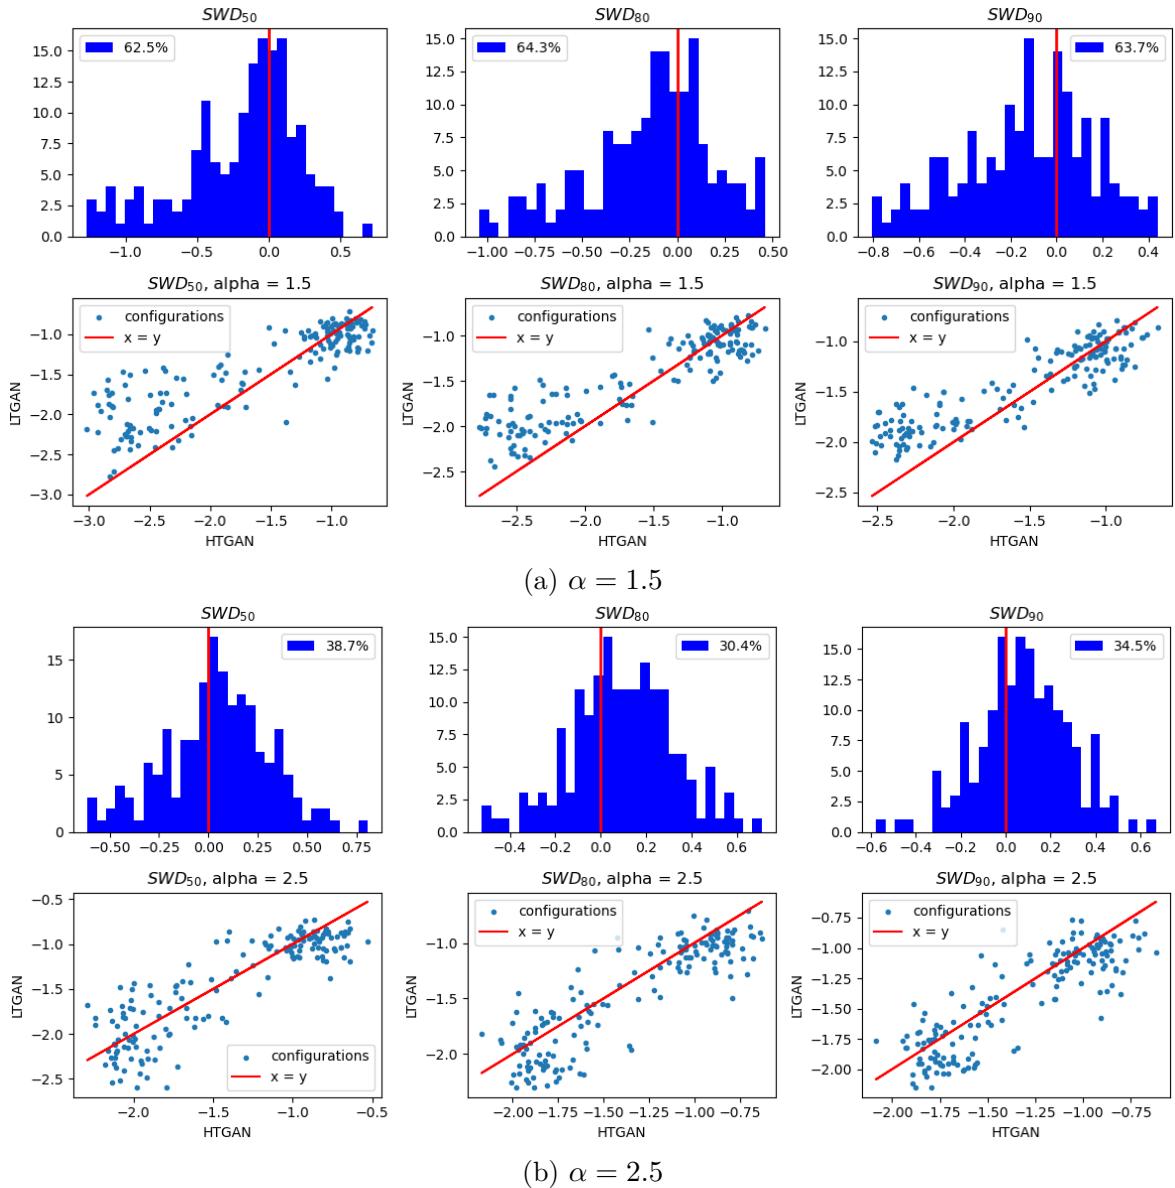


Figure 1.6: Results of the performance of HTGAN vs LTGAN on the Utilities ($d = 27$) and Financials subsectors ($d = 57$) of the S&P500 dataset. Results are aggregated with respect to the dimension d for plot. Each two groups of 2×3 plots corresponds to one experimental setting, *i.e.* $\alpha = 1.5$ (top) and $\alpha = 2.5$ (bottom). In each of these two groups: the first row plots the specified metric of HTGAN vs LTGAN in log scale for a given hyperparametrization configuration. The second row is a histogram of the difference of the log of the metric for HTGAN vs the metric for LTGAN for each parametrization. The legend corresponds to the proportion of cases where HTGAN performs better.

1754 **Appendix**

1755 **3.A Theoretical background**

1756 Section 3.B provides additional background on extreme-value theory to complement Sec-
 1757 tion 1.2.1. Section 3.C is devoted to D-norms, introduced in Section 1.2.2. Section 3.D
 1758 focuses on the theory of quantization. Finally, Section 3.E provides some details on copulas
 1759 and measures of dependence.

1760 **3.B Extreme-value theory**

1761 **General set-up**

1762 Consider a \mathbb{R}^d valued random vector \mathbf{X} with associated cdf $\mathbf{x} \in \mathbb{R}^d \mapsto F(\mathbf{x}) = \mathbb{P}(\mathbf{X} \leq \mathbf{x})$.
 1763 The primary goal of extreme-value theory is to establish the asymptotic distribution of
 1764 well normalized maxima of realizations of \mathbf{X} , see (S. I. Resnick 1987; Beirlant et al. 2004;
 1765 L. d. Haan and Ferreira 2006) for reference textbooks. Let $(\mathbf{X}_i)_{i \in [n]}$ be an i.i.d.sample from
 1766 F and consider $\mathbf{M}_n := \frac{\max_{i \in [n]} \mathbf{X}_i - \mathbf{b}_n}{\mathbf{a}_n}$, where $\mathbf{a}_n > \mathbf{0}$ and \mathbf{b}_n are normalizing sequences in
 1767 \mathbb{R}^d . Clearly, the cdf of \mathbf{M}_n is given by $F_{\mathbf{M}_n}(\mathbf{x}) = F^n(\mathbf{a}_n \mathbf{x} + \mathbf{b}_n)$. This remark gives rise
 1768 to the following definition:

1769 **Definition 1.5** (Max-Domain of Attraction (MDA)), (L. d. Haan and Ferreira 2006, Section
 1770 6.1.2)). A cdf F on \mathbb{R}^d is said to be in the max-domain of attraction of a cdf G , denoted
 1771 by $F \in \text{Dom}(G)$, if there exist sequences $\mathbf{a}_n > 0$ and \mathbf{b}_n , $n \in \mathbb{N}$ such that:

$$\forall \mathbf{x} \in \mathbb{R}^d : F^n(\mathbf{a}_n \mathbf{x} + \mathbf{b}_n) \xrightarrow[n \rightarrow \infty]{} G(\mathbf{x}). \quad (1.22)$$

1772 In other words, the normalized maximum of a sample from F converges in distribution
 1773 to a random vector with cdf G . It can be shown that the limiting distribution function is
 1774 necessarily max-stable, as defined below:

1775 **Definition 1.6** (Max-stable distribution, (L. d. Haan and Ferreira 2006, Section 6.1.2)).
 1776 A cdf G on \mathbb{R}^d is called max-stable if, for any $n \in \mathbb{N}$, there exist sequences $\mathbf{a}_n > 0$ and
 1777 $\mathbf{b}_n \in \mathbb{R}^d$ such that

$$\forall \mathbf{x} \in \mathbb{R}^d : G^n(\mathbf{a}_n \mathbf{x} + \mathbf{b}_n) = G(\mathbf{x}). \quad (1.23)$$

1778 Besides, convergence (1.22) can be extended to a continuous counterpart; one has
 1779 $F \in \text{Dom}(G)$ if and only if there exist two functions $\mathbf{a}_t > 0 \in \mathbb{R}^d$ and $\boldsymbol{\theta} \in \mathbb{R}^d$ such that:

$$\forall \mathbf{x} \in \mathbb{R}^d : F^t(\mathbf{a}_t \mathbf{x} + \boldsymbol{\theta}) \xrightarrow[t \rightarrow \infty]{} G(\mathbf{x}). \quad (1.24)$$

1780 **Margins**

1781 Let us highlight that, if F fulfills (1.22) or (1.24), then its margins F_j satisfy

$$\forall x \in \mathbb{R} : F_j^n(a_{j,n}x + b_{j,n}) \xrightarrow[n \rightarrow \infty]{} G_j(x), \quad (1.25)$$

where G_j denotes the j th margin of G , $j \in [d]$. The *univariate* extreme-value theorem (see for instance (Fisher and Tippett 1928; Gnedenko 1943)) then provides a parametric form for the cdf G_j s which generalizes the parametrizations of Fréchet ($\gamma > 0$), Gumbel ($\gamma = 0$) and reverse-Weibull ($\gamma < 0$):

Theorem 1.4 ((L. d. Haan and Ferreira 2006, Theorem 1.1.3)). *The \mathbb{R} -valued cdf's satisfying (1.22) are of the form $G_\gamma(a \cdot + b)$ with $a > 0$, $b \in \mathbb{R}$ and $\gamma \in \mathbb{R}$, where G_γ is a Generalized Extreme Value Distribution (GEVD) defined as*

$$\forall x \in \mathbb{R}, 1 + \gamma x > 0 : G_\gamma(x) = \exp - (1 + \gamma x)^{-1/\gamma}. \quad (1.26)$$

When $\gamma = 0$, G_0 is interpreted as the pointwise limit in (1.26): $G_0(x) = \exp -e^{-x}$.

Therefore, each margin G_j of G can be written as $G_j = G_{\gamma_j}$ following (1.26) with $\gamma_j \in \mathbb{R}$, $j \in [d]$. Let us note that, for each max-stable cdf G , one can consider G_* , the so-called associated simple max-stable distribution, which is defined as:

$$\forall \mathbf{x} \in \mathbb{R}^d : G_*(\mathbf{x}) = G\left(\frac{\mathbf{x}^\gamma - \mathbf{1}}{\gamma}\right), \quad (1.27)$$

with $\boldsymbol{\gamma} = (\gamma_j)_{j \in [d]} \in \mathbb{R}^d$. The margins of $G_{*,j}$ of G_* are unit Fréchet distributed *i.e.*

$$\forall x > 0, G_{*,j}(x) = \exp -1/x. \quad (1.28)$$

As a consequence, it is possible to characterize a GEVD with two components: On the one hand, the marginal tail indices $\boldsymbol{\gamma}$, which characterize its margins and, on the other hand, its dependence structure encoded in G_* .

Dependence structure

The following result establishes a strong link between the stable tail dependence function (stdf) of G (as introduced in Definition 1.1) and its simple max-stable cdf G_* .

Proposition 1.5 ((L. d. Haan and Ferreira 2006, Section 6.1.5)). *Consider a max-stable cdf G and its associated simple max-stable cdf G_* defined in (1.27). The stdf of G is given by*

$$\forall \mathbf{x} \in (0, \infty)^d : \ell_G(\mathbf{x}) = -\log G_*(1/\mathbf{x}). \quad (1.29)$$

We show in the next paragraph that both ℓ_G and G_* can be interpreted in terms of D-norms, which is a key property for our analysis.

3.C D-norms

Let us first recall classical results from multivariate extreme-value theory through the lens of D-norms. The reference monograph on the matter is (Falk 2019).

1808 **Definition and basic properties**

1809 We start with basic definitions.

1810 **Definition 1.7** (Norm and D-norms). A function $\|\cdot\| : \mathbb{R}^d \rightarrow \mathbb{R}^+$ is a norm on \mathbb{R}^d if it
1811 verifies three conditions: a) *Positive definiteness*, b) *Absolute Homogeneity*, c) *Triangle*
1812 *inequality*.

1813 A subset of norms is the set of D-norms defined by

$$\|\cdot\|_{\boldsymbol{\Gamma}} : \mathbf{x} \in \mathbb{R}^d \mapsto \|\mathbf{x}\|_{\boldsymbol{\Gamma}} = \mathbb{E} \left[\max_{i \in [d]} (|x_i| \boldsymbol{\Gamma}_i) \right] \quad (1.30)$$

1814 for a non-negative random variable $\boldsymbol{\Gamma} \in \mathbb{R}^d$ with unit expectation (Definition 1.2), $\boldsymbol{\Gamma}$ is
1815 also known as the generator of the D-norm.

1816 Let us give a few remarks. First, since D-norms are monotone and radially symmetric,
1817 not all norms are D-norms (see (Falk 2019, Chapter 1)). Second, usual norms (such as
1818 the L_p -norms for $p \in [1, \infty]$) are D-norms. Last, observe that there are infinitely many
1819 generators $\boldsymbol{\Gamma}$ leading to the same D-norm: multiply $\boldsymbol{\Gamma}$ by an independent positive random
1820 variable \mathbf{U} with unit expectation, it readily gives $\|\cdot\|_{\boldsymbol{\Gamma}\mathbf{U}} = \|\cdot\|_{\boldsymbol{\Gamma}}$. Hence, to ease notations,
1821 sometimes we may prefer to write the above norm by $\|\cdot\|_D$ to focus less on the generator
1822 $\boldsymbol{\Gamma}$; in such a situation, we refer to the generator $\boldsymbol{\Gamma}$ associated with the D-norm by writing
1823 $\boldsymbol{\Gamma} \triangleleft \|\cdot\|_D$, which reads “ $\boldsymbol{\Gamma}$ generates $\|\cdot\|_D$ ”. Despite the multiple generators associated with
1824 the same D-norm, the following theorem provides a uniqueness result on the choice of $\boldsymbol{\Gamma}$
1825 when the norm of the random variable $\boldsymbol{\Gamma}$ is constrained.

1826 **Theorem 1.5** (Normed Generators, (Falk 2019, Theorem 1.7.1)). *Let $\|\cdot\|$ be an arbitrary
1827 norm on \mathbb{R}^d . For any D-norm $\|\cdot\|_D$ on \mathbb{R}^d , there exists a generator $\boldsymbol{\Gamma} \triangleleft \|\cdot\|_D$ and a constant c
1828 with the additional property $\mathbb{P}(\|\boldsymbol{\Gamma}\| = c) = 1$. The distribution of the generator is uniquely
1829 defined.*

1830 Let us remark that taking $\|\cdot\| = \|\cdot\|_1$ as the L_1 -norm, the constant c is necessarily d :
1831 indeed, $\|\boldsymbol{\Gamma}\|_1 = c$ implies $c = \mathbb{E}[\|\boldsymbol{\Gamma}\|_1] = \sum_{i \in [d]} \mathbb{E}[\boldsymbol{\Gamma}_i] = d$.

1832 **Relation with stdf**

1833 The next result is fundamental to connect a stdf to a D-norm.

1834 **Theorem 1.6** (Representation of simple max stable distributions). *A cdf G_* on \mathbb{R}^d is
1835 simple max-stable if and only if there exists a D-norm $\|\cdot\|_D$ on \mathbb{R}^d such that:*

$$\forall \mathbf{x} \in (0, \infty)^d : G_*(\mathbf{x}) = \exp - \left\| \frac{\mathbf{1}}{\mathbf{x}} \right\|_D. \quad (1.31)$$

1836 In particular, if G is a max-stable cdf on \mathbb{R}^d , then its stdf is given by

$$\ell_G(\mathbf{x}) = \|\mathbf{x}\|_D. \quad (1.32)$$

1837 The first statement (1.31) is taken from (Falk 2019, Theorem 2.3.3) or (Falk and Fuller
1838 2021, Theorem 4.1), while the second one (1.32) follows from Proposition 1.5.

1839 **3.D Quantization**

1840 **The quantization problem.** The quantization problem is concerned with finding the
1841 best approximation of a random vector $\mathbf{X} \in \mathbb{R}^d$ with a quantized version, *i.e.*a version
1842 that takes a discrete number of values. For a general reference on quantization, see (Graf
1843 and Luschgy 2000). Finding the best approximation of a random vector depends on the
1844 chosen measure. Define the quantization problem (Graf and Luschgy 2000, Section 10.1)
1845 of a probability measure P , for a given norm $\|\cdot\|$ (for example the Euclidean norm), as:

$$e_{N,\infty}(P) := \inf_{f: \mathbb{R}^d \rightarrow \mathbb{R}^d} \{P\text{-esssup } \|f(\mathbf{X}) - f(\mathbf{X})\| : \#f(\mathbb{R}^d) \leq N\} \quad (1.33)$$

$$\begin{aligned} &= \inf_{\boldsymbol{\alpha} \subset \mathbb{R}^d} \sup_{\mathbf{x} \in \text{supp}(P)} \min_{\mathbf{a} \in \boldsymbol{\alpha}} \|\mathbf{x} - \mathbf{a}\|, \\ &\#\boldsymbol{\alpha} \leq N \end{aligned} \quad (1.34)$$

1846 where $f : \mathbb{R}^d \rightarrow \mathbb{R}^d$ is a quantizer. Since $e_{N,\infty}(P)$ depends on the probability measure P
1847 only through its support, sometimes (see Lemma 1.1 below) we will write $e_{N,\infty}(A)$ as a
1848 function of a set $A \subset \mathbb{R}^d$ (that can be taken as the support of P).

1849 Then, if $\text{supp}(P)$ is compact, Jordan measurable with positive volume (Graf and
1850 Luschgy 2000, p.3), the following limit exists:

$$Q_\infty(\text{supp}(P)) = \lim_{N \rightarrow \infty} N^{1/d} e_{N,\infty}(P). \quad (1.35)$$

1851 Note that $Q_\infty(\text{supp}(P))$ does not depend on the law of P but only on its support. There-
1852 fore, we will interchangeably use a probability distribution or a set as input of $Q_\infty(\cdot)$. The
1853 limit $Q_\infty(\text{supp}(P))$ is called the covering coefficient or quantization coefficient of order
1854 ∞ and can be expressed in terms of the covering coefficient $Q_\infty([0, 1]^d)$. Moreover, with
1855 a further constraint, we have the following theorem:

1856 **Theorem 1.7** ((Graf and Luschgy 2000, Theorem 10.7)). *Let $A \subset \mathbb{R}^d$ be a non empty
1857 compact set with $\lambda^d(\partial A) = 0$. Let $Q_\infty([0, 1]^d)$ be the quantization coefficient of $[0, 1]^d$
1858 defined in (1.35). Then $Q_\infty([0, 1]^d) > 0$ and:*

$$\lim_{N \rightarrow \infty} N^{1/d} e_{N,\infty}(A) = Q_\infty([0, 1]^d) \lambda^d(A)^{1/d}, \quad (1.36)$$

1859 where λ^d is the Lebesgue measure in dimension d .

1860 We state the following technical lemma, which will be helpful for the proof of Theo-
1861 rem 1.3:

1862 **Lemma 1.1** ((Graf and Luschgy 2000, Lemma 10.6)).

1863 1. Let $A, B \subset \mathbb{R}^d$ be nonempty compact sets with $A \subset B$. Then,

$$\forall N \in \mathbb{N} : e_{N,\infty}(A) \leq e_{N,\infty}(B). \quad (1.37)$$

1864 2. Let N, m and $\{N_k\}_{k \in [m]}$ be positive integers such that $\sum_{k \in [m]} N_k \leq N$ and consider
1865 m nonempty compact sets $\{A_k\}_{k \in [m]}$. Then,

$$e_{N,\infty}\left(\bigcup_{k \in [m]} A_k\right) \leq \max_{k \in [m]} e_{N_k,\infty}(A_k). \quad (1.38)$$

1866 Here is a particular case of (1.38) when $N \geq 2m$:

$$e_{N,\infty}\left(\bigcup_{k \in [m]} A_k\right) \leq \max_{k \in [m]} e_{\lceil N/m \rceil - 1, \infty}(A_k). \quad (1.39)$$

1867 Proof. The proof of (1.38) can be found in (Graf and Luschgy 2000, Lemma 10.6). Equation
1868 (1.39) is deduced from Equation (1.38) with $N_k = \lceil N/m \rceil - 1$ for all $k \in [m]$. \square

1869 3.E Copulas, measures of dependence

1870 **Copulas.** Let us consider a d -variate random vector \mathbf{X} from a cdf F with continuous
1871 margins F_j . Sklar's Theorem (Sklar 1959) states that there exists a unique function $C : [0, 1]^d \rightarrow [0, 1]$ such that, for all $\mathbf{x} \in \mathbb{R}^d$, $F(x_1, \dots, x_d) = C(F_1(x_1), \dots, F_d(x_d))$.

1873 Introducing for any $j \in [d]$ the uniformly distributed random variable $U_j = F_j(x_j)$,
1874 the copula C is the cdf of the joint random vector (U_1, \dots, U_d) . Copulas are a tool for
1875 studying dependence of probability objects, independently from the margins. See (Nelsen
1876 2006) for a thorough overview. As an example, the Gumbel copula is given by (see (Nelsen
1877 2006, Equation (2.4.2.))):

$$\forall \mathbf{u} \in [0, 1]^d : C_\beta(\mathbf{u}) = \exp - \left(\sum_{i \in [d]} (-\log(u_i))^\beta \right)^{1/\beta}, \quad (1.40)$$

1878 where $\beta \geq 1$.

1879 **Definition 1.8** (Kendall's τ , (Nelsen 2006, Theorem 5.1.3.)). Let C be a bivariate copula.
1880 Its associated Kendall's τ is:

$$\tau = 4\mathbb{E}[C(U_1, U_2)] - 1, \quad (1.41)$$

1881 where $(U_1, U_2) \sim C$.

1882 Kendall's τ is a common measure of dependence for copulas, see (Joe 2014, Section 6.8)
1883 for an estimation procedure. Note that Kendall's τ for the Gumbel copula is given by
1884 $\tau = 1 - 1/\beta$, $\beta \geq 1$.

1885 **Wasserstein distance.** Let μ and ν be two measures on \mathbb{R}^d . Considering a cost function
1886 $c : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}^+$ and $p \geq 1$, the p - Wasserstein distance associated with the cost c

1887 between μ and ν is defined as (Villani 2009, Chap. 6):

$$\mathcal{W}_p(\mu, \nu) = \left(\inf_{\gamma \in \Gamma(\mu, \nu)} \mathbb{E}_{(\mathbf{X}, \mathbf{Y}) \sim \gamma} [c(\mathbf{X}, \mathbf{Y})^p] \right)^{1/p},$$

1888 where $\Gamma(\mu, \nu)$ is the set of *couplings* of μ and ν , *i.e.* the set of distributions on $\mathbb{R}^d \times \mathbb{R}^d$
1889 such that the marginals are respectively μ and ν . The Wasserstein distance suffers from
1890 the fact that there exists no closed formula to compute it, nor very accessible computable
1891 estimates in dimensions higher than one. In order to perform estimation, we therefore rely
1892 on the sliced-Wasserstein distance, which is an efficient proxy, much used for computations
1893 in experimental settings (Nadjahi 2021). It is computed by randomly projecting the two
1894 compared measures on axes defined by directions \mathbf{v} and averaging over the projections.
1895 Specifically, the sliced-Wasserstein distance is defined as:

$$\text{SWD}_p(\mu, \nu) = \mathbb{E}_{\mathbf{v} \sim \mathcal{U}(\mathbb{S}^{d-1})} (\mathcal{W}_p^p(\mathbf{v} \# \mu, \mathbf{v} \# \nu))^{1/p}, \quad (1.42)$$

1896 where \mathbb{S}^{d-1} is the Euclidean unit sphere, $\mathbf{v} \# \mu$ (respectively $\mathbf{v} \# \nu$) stands for the pushforward
1897 of the projection $\mathbf{X} \in \mathbb{R}^d \mapsto \langle \mathbf{X}, \mathbf{v} \rangle \in \mathbb{R}$, where $\mathbf{X} \sim \mu$ (respectively ν).

1898 3.F Proofs of claims

1899 Proof of Proposition 1.3

1900 The following states an interesting corollary of Theorem 1.2, which we are the first to prove
1901 formally as far as we know.

1902 **Corollary 1.1** (of Theorem 1.2). *Let $\mathbf{Z} \in \mathbb{R}^N$ be a positive random vector with continuous
1903 cdf $F_{\mathbf{Z}} \in \text{Dom}(G_{\Gamma_{\mathbf{Z}}})$ where $G_{\Gamma_{\mathbf{Z}}}$ is a simple max-stable distribution *i.e.* $F_{\mathbf{Z}}$ satisfies (1.7)
1904 and (1.8). Consider $\Phi \in \mathbb{H}_1([0, \infty)^N, [0, \infty)^d)$, such that:*

$$\mathbb{E}[\Phi(\Gamma_{\mathbf{Z}})] > 0. \quad (1.43)$$

1905 Let $\mathbf{X} = \Phi(\mathbf{Z})$. Then, $\Phi(\Gamma_{\mathbf{Z}})/\mathbb{E}[\Phi(\Gamma_{\mathbf{Z}})] =: \Gamma_{\mathbf{X}}$ is a valid generator of a D-norm and
1906 $F_{\mathbf{X}} \in \text{Dom}(G_{\Gamma_{\mathbf{X}}})$ where $G_{\Gamma_{\mathbf{X}}}$ is a simple max-stable distribution. Moreover, the normalizing
1907 constants to ensure convergence in (1.24) are $\mathbf{a}_t = t \mathbb{E}[\Phi(\Gamma_{\mathbf{Z}})]$ and $\boldsymbol{\theta} = \mathbf{0}$.

1908 *Proof.* In view of the characterization of Theorem 1.2, consider any $h \in \mathbb{H}_1([0, \infty)^d, [0, \infty))$
1909 and

$$\Psi_h : \begin{cases} [0, \infty)^N & \rightarrow [0, \infty), \\ \mathbf{z} & \mapsto h\left(\frac{\Phi(\mathbf{z})}{\mathbb{E}[\Phi(\Gamma_{\mathbf{Z}})]}\right). \end{cases} \quad (1.44)$$

1910 Clearly, Ψ_h is continuous 1-homogeneous as composition of the continuous 1-homogeneous
1911 functions Φ and h . Since $F_{\mathbf{Z}}$ verifies (1.8), it follows that:

$$t \mathbb{P}\left(h\left(\frac{\Phi(\mathbf{Z})}{\mathbb{E}[\Phi(\Gamma_{\mathbf{Z}})]}\right) > t\right) = t \mathbb{P}(\Psi_h(\mathbf{Z}) > t) \xrightarrow{t \rightarrow \infty} \mathbb{E}[\Psi_h(\Gamma_{\mathbf{Z}})] = \mathbb{E}\left[h\left(\frac{\Phi(\Gamma_{\mathbf{Z}})}{\mathbb{E}[\Phi(\Gamma_{\mathbf{Z}})]}\right)\right]. \quad (1.45)$$

1912 Introducing $\mathbf{Y} = \Phi(\mathbf{Z})/\mathbb{E}[\Phi(\mathbf{\Gamma}_Z)] = \mathbf{X}/\mathbb{E}[\Phi(\mathbf{\Gamma}_Z)]$ and its corresponding distribution
1913 function F_Y , the previous convergence can be rewritten as:

$$\forall h \in \mathbb{H}_1([0, \infty)^d, [0, \infty)) : t \mathbb{P}(h(\mathbf{Y}) > t) \xrightarrow[t \rightarrow \infty]{} \mathbb{E}[h(\mathbf{\Gamma}_X)], \quad (1.46)$$

1914 so that \mathbf{Y} satisfies characterization (1.8) in Theorem 1.2. Moreover, since $\mathbb{E}[\mathbf{\Gamma}_X] = \mathbf{1}$
1915 and the coordinates of $\mathbf{\Gamma}_X$ are non-negative, $\mathbf{\Gamma}_X$ is indeed the generator of a D-norm.
1916 Therefore:

$$\forall \mathbf{x} \in (0, \infty)^d : F_Y^t(t\mathbf{x}) \xrightarrow[t \rightarrow \infty]{} \exp - \left\| \frac{1}{\mathbf{x}} \right\|_{\mathbf{\Gamma}_X}. \quad (1.47)$$

1917 Noting that $F_Y(\mathbf{x}) = F_X(\mathbb{E}[\Phi(\mathbf{\Gamma}_Z)] \mathbf{x})$ finally proves

$$\forall \mathbf{x} \in (0, \infty)^d : F_X^t(t \mathbb{E}[\Phi(\mathbf{\Gamma}_Z)] \mathbf{x}) \xrightarrow[t \rightarrow \infty]{} \exp - \left\| \frac{1}{\mathbf{x}} \right\|_{\mathbf{\Gamma}_X} =: G_{\mathbf{\Gamma}_X}(\mathbf{x}), \quad (1.48)$$

1918 and the result is proved. \square

1919 We now complete the proof of Proposition 1.3. The previous result holds for continuous
1920 1-homogeneous function Φ , which is not exactly the situation with a neural network with
1921 ReLU activation functions because of the bias terms (see Proposition 1.2). However, we
1922 claim that, asymptotically, these bias terms do not play any role on the stdf. To see this,
1923 let us investigate the behavior of a ReLU-neural network when going from a layer to the
1924 next one.

1925 Denoting by σ the ReLU activation function, and supposing that $\varphi_l = \sigma$, $\forall l \in [L]$, we
1926 have

$$\begin{aligned} \mathbf{Z}_l &= \sigma(\mathbf{W}_l \mathbf{Z}_{l-1} + \mathbf{b}_l) \\ &= \sigma(\mathbf{W}_l \mathbf{Z}_{l-1}) + [\sigma(\mathbf{W}_l \mathbf{Z}_{l-1} + \mathbf{b}_l) - \sigma(\mathbf{W}_l \mathbf{Z}_{l-1})]. \end{aligned} \quad (1.49)$$

1927 On the one hand, $\sigma(\mathbf{W}_l \mathbf{Z}_{l-1})$ is a 1-homogeneous function of \mathbf{Z}_{l-1} and is thus a heavy-tailed random variable if $\mathbf{W}_l \neq \mathbf{0}$ and \mathbf{Z}_{l-1} is heavy-tailed itself. On the other hand,
1928 the remaining term $[\dots]$ in (1.49) is a bounded random variable (bounded by $|\mathbf{b}_l|$); the
1929 stdf of the sum of a heavy-tailed random variable and a bounded one is the stdf of the
1930 heavy-tailed random variable. The stdf of (1.49) is therefore the stdf of the transformation
1931 without the bias, and by induction the stdf of the output is the stdf of the input after
1932 iterative applications of $\sigma(\mathbf{W}_l \cdot)$ for $l \in [L]$, which is 1-homogeneous continuous. Let
1933 denote by Φ the output of the L compositions of $\sigma(\mathbf{W}_l \cdot)$: since the matrix weights \mathbf{W}_l
1935 are non-negative, $\Phi \in \mathbb{H}_1([0, \infty)^N, [0, \infty)^d)$. Therefore, Corollary 1.1 applies and so the
1936 stdf associated with the output of the network is $\mathbf{\Gamma}_{G_Z} = \Phi(\mathbf{\Gamma}_Z)/\mathbb{E}\Phi(\mathbf{\Gamma}_Z)$, which has at
1937 most N atoms. Thus, the proof is concluded.

1938 3.G Proof of Proposition 1.4

1939 The next lemma can be interpreted in the following way. If a discrete random vector
1940 generates a D-norm, then it admits a discrete counterpart which can be any other discrete

1941 distribution with modified atoms: it is also a generator of the concerned D-norm.

1942 **Lemma 1.2.** *Let $\Gamma \triangleleft \|\cdot\|_D$ where $\Gamma \in [0, \infty)^d$ is discrete, that is $\Gamma \sim \mu_\Gamma$ where*

$$\mu_\Gamma = \sum_{i \in [N]} p_i \delta_{\gamma^{(i)}}. \quad (1.50)$$

1943 Consider a new discrete probability measure with positive weights $(q_i)_{i \in [N]}$ and a new ran-
1944 dom variable $\tilde{\Gamma} \sim \mu_{\tilde{\Gamma}}$, where

$$\mu_{\tilde{\Gamma}} := \sum_{i \in [N]} q_i \delta_{\tilde{\gamma}^{(i)}}, \quad (1.51)$$

1945 with $\tilde{\gamma}^{(i)} = \frac{p_i}{q_i} \gamma^{(i)}$, $i \in [N]$. Then, $\tilde{\Gamma}$ is a valid generator and $\tilde{\Gamma} \triangleleft \|\cdot\|_D$.

1946 *Proof.* By definition of the D-norm, one has $\forall \mathbf{x} \in \mathbb{R}^N$:

$$\begin{aligned} \|\mathbf{x}\|_D &= \mathbb{E}_\Gamma \left[\max_{j \in [N]} \{|x_j| \Gamma_j\} \right] \\ &= \sum_{i \in [N]} p_i \max_{j \in [N]} \left\{ |x_j| \gamma_j^{(i)} \right\} \\ &= \sum_{i \in [N]} q_i \max_{j \in [N]} \left\{ |x_j| \tilde{\gamma}_j^{(i)} \right\} = \mathbb{E}_{\tilde{\Gamma}} \left[\max_{j \in [N]} \left\{ |x_j| \tilde{\Gamma}_j \right\} \right]. \end{aligned}$$

1947 Since $\mathbb{E}_{\tilde{\Gamma}} [\tilde{\Gamma}] = \mathbb{E}_\Gamma [\Gamma] = \mathbf{1}$ and $\tilde{\Gamma} \in [0, \infty)^d$, $\tilde{\Gamma}$ is a valid generator. \square

1948 In particular, it is possible to choose uniform weights, $q_i = 1/N$, $i \in [N]$. In other
1949 words, if a D-norm admits a generator which admits a finite number of values, it also
1950 admits a generator with the same number of finite values, with uniform probability. This
1951 result is important as it means that the class of generators of D-norms which take a finite
1952 number of values and with uniform weights is as rich as the class of generator which take
1953 finite values with arbitrary weights.

1954 **Proposition 1.6.** *Consider a generator Γ of a D-norm in \mathbb{R}^d . Suppose that it is discrete
1955 uniform with N atoms $(\gamma^{(i)})_{i \in [N]}$:*

$$\mu_\Gamma = \frac{1}{N} \sum_{i \in [N]} \delta_{\gamma^{(i)}}. \quad (1.52)$$

1956 Consider moreover a random vector $\mathbf{Z} \in \mathbb{R}^N$ with i.i.d. unit Fréchet margins: $\mathbf{Z} = (Z_1, \dots, Z_N)$.
1957 Then, there exists a linear mapping $\Phi : \mathbb{R}^N \rightarrow \mathbb{R}^d$ such that:

$$F_{\Phi(\mathbf{Z})} \in \text{Dom}(G_\Gamma). \quad (1.53)$$

1958 *Proof.* Let us first remark that $F_{\mathbf{Z}} \in \text{Dom}(G_{\Gamma_Z})$ with:

$$\mu_{\Gamma_Z} = \frac{1}{N} \sum_{i \in [N]} \delta_{N\mathbf{e}_i},$$

see Example 1.1. Consider $\boldsymbol{\Gamma} \in \mathbb{R}^N$ a valid generator of a D-norm with a discrete uniform distribution (1.52). There exists $\Phi : \mathbb{R}^N \rightarrow \mathbb{R}^d$ a linear function such that $\forall i \in [N] : \Phi(N\mathbf{e}_i) = \gamma^{(i)}$. Let us note that Φ is continuous 1-homogeneous because it is linear. Moreover, remark that $\forall i \in [N] : \gamma^{(i)} \in [0, \infty)^d$ implies $\Phi([0, \infty)^N) \subset \Phi([0, \infty)^d)$ and thus $\Phi(\mathbf{Z}) \in [0, \infty)^d$. Corollary 1.1 ensures that

$$F_{\Phi(\mathbf{Z})} \in \text{Dom}(G_{\Phi(\boldsymbol{\Gamma}_{\mathbf{Z}})/\mathbb{E}(\Phi(\boldsymbol{\Gamma}_{\mathbf{Z}}))}) = \text{Dom}(G_{\boldsymbol{\Gamma}}) \quad (1.54)$$

since, as $\Phi(\mathbf{e}_i) = \gamma^{(i)}$, $\Phi(\boldsymbol{\Gamma}_{\mathbf{Z}})$ is equal in distribution to $\boldsymbol{\Gamma}$ (note that $\mathbb{E}[\Phi(\boldsymbol{\Gamma}_{\mathbf{Z}})] = \mathbf{1}$). The result is thus proved. \square

Remark 1.1. This result ensures that, given any random vector of interest \mathbf{X} for which the associated D-norm generator $\boldsymbol{\Gamma}_{\mathbf{X}}$ has a uniform distribution with N atoms, one can find a linear map Φ transforming a noise with unit Fréchet margins with at least N atoms to a distribution whose D-norm generator is the desired one.

Lemma 1.2 and Proposition 1.6 together prove Proposition 1.4. \square

3.H Proof of Theorem 1.3

The following technical lemma will be useful to prove the result:

Lemma 1.3. $\forall \mathbf{a}, \mathbf{b}, \mathbf{c} \in \mathbb{R}^d$:

$$\left| \max_{i \in [d]} a_i b_i - \max_{i \in [d]} a_i c_i \right| \leq \|\mathbf{a}\|_{\infty} \max_{i \in [d]} |b_i - c_i|. \quad (1.55)$$

Proof. Without loss of generality, suppose that $\max_{i \in [d]} a_i b_i \geq \max_{i \in [d]} a_i c_i$, the argument being symmetric. One has:

$$\begin{aligned} \left| \max_{i \in [d]} a_i b_i - \max_{i \in [d]} a_i c_i \right| &= \max_{i \in [d]} \{a_i c_i + a_i b_i - a_i c_i\} - \max_{i \in [d]} a_i c_i \\ &\leq \max_{i \in [d]} a_i c_i + \max_{i \in [d]} \{a_i b_i - a_i c_i\} - \max_{i \in [d]} a_i c_i \\ &\leq \max_{i \in [d]} |a_i| \max_{i \in [d]} |b_i - c_i| \end{aligned}$$

and the result is proved. \square

We seek to find a upper bound on:

$$\inf_{\theta \in \Theta} \sup_{x \in (0, \infty)^d} |\ell_F(\mathbf{x}) - \ell_{G_{\theta, N}}(\mathbf{x})| / \|\mathbf{x}\|_{\infty}. \quad (1.56)$$

In view of the above, and remembering that the stdf of the output of a neural network is discrete (Proposition 1.3), our goal is to find the best approximation of the D-norm $\|\cdot\|_{\boldsymbol{\Gamma}}$ with a discretized generator with N atoms; we will leverage results from quantization

¹⁹⁸¹ (Section 3.D). Take $\boldsymbol{\Gamma}$ an arbitrary generator of the D-norm related to ℓ_F . One has:

$$\begin{aligned} & \min_{f: \mathbb{R}^d \rightarrow \mathbb{R}^d, \#f(\mathbb{R}^d)=N} \left| \|\mathbf{x}\|_{\boldsymbol{\Gamma}} - \|\mathbf{x}\|_{f(\boldsymbol{\Gamma})} \right| \\ &= \min_{f: \mathbb{R}^d \rightarrow \mathbb{R}^d, \#f(\mathbb{R}^d)=N} \left| \mathbb{E}_{\boldsymbol{\Gamma}} \left[\max_{i \in [d]} (|x_i| \boldsymbol{\Gamma}_i) - \max_{i \in [d]} (|x_i| f(\boldsymbol{\Gamma}_i)) \right] \right|, \end{aligned} \quad (1.57)$$

$$\leq \min_{f: \mathbb{R}^d \rightarrow \mathbb{R}^d, \#f(\mathbb{R}^d)=N} \|\mathbf{x}\|_{\infty} \mathbb{E}_{\boldsymbol{\Gamma}} \left[\max_{i \in [d]} |\boldsymbol{\Gamma}_i - f(\boldsymbol{\Gamma}_i)| \right] \quad (1.58)$$

$$\leq \|\mathbf{x}\|_{\infty} e_{N,\infty}(\boldsymbol{\Gamma}), \quad (1.59)$$

¹⁹⁸² from Lemma 1.3. To bound (1.59), we use a quantizer which has a constant sup-norm,
¹⁹⁸³ which existence is guaranteed by (Theorem 1.5):

$$\|\boldsymbol{\Gamma}\|_{\infty} = c \quad a.s.. \quad (1.60)$$

Observe that the constant c must be related to the stdf using Definition 1.2:

$$\ell_F(1) = \mathbb{E} [\|\boldsymbol{\Gamma}\|_{\infty}] = c.$$

¹⁹⁸⁴ Now, we are in a position to derive a bound on $e_{N,\infty}(\boldsymbol{\Gamma})$.

¹⁹⁸⁵ **Definition 1.9.** For any norm $\|\cdot\|$, let us denote by $\mathcal{B}_{\|\cdot\|}$ the associated unit ball. Besides,
¹⁹⁸⁶ let $\mathcal{B}_{\|\cdot\|}^+ = \mathcal{B}_{\|\cdot\|} \cap [0, \infty)^d$ be the part of the unit ball in the positive orthant.

¹⁹⁸⁷ Owing to (1.60) and Definition 1.2, $\text{supp}(\boldsymbol{\Gamma}) \subset c \cdot [0, 1]^d$. Equation (1.36) shows that
¹⁹⁸⁸ one can get an error bound for (1.59) with rate $N^{-1/d}$. However, it is possible to improve
¹⁹⁸⁹ this result by remarking that $\boldsymbol{\Gamma}$ takes values on the boundary of a hypercube (see (1.60)).
¹⁹⁹⁰ Indeed, now note that:

$$\text{supp}(\boldsymbol{\Gamma}) \subset c \cdot \mathcal{B}_{\|\cdot\|_{\infty}}^+ \subset \bigcup_{i \in [d]} c \cdot \mathcal{B}_{\|\cdot\|_{\infty}}^+(i), \quad (1.61)$$

¹⁹⁹¹ where, for all $i \in [d]$,

$$\mathcal{B}_{\|\cdot\|_{\infty}}^+(i) = \{(x_1, \dots, x_{i-1}, 1, x_{i+1}, \dots, x_d) : (x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_d) \in [0, 1]^{d-1}\}. \quad (1.62)$$

¹⁹⁹² *Remark 1.2.* For all $i \in [d]$, $\mathcal{B}_{\|\cdot\|_{\infty}}^+(i)$ is canonically homeomorphic to $[0, 1]^{d-1}$ in the topo-
¹⁹⁹³ logical space \mathbb{R}^{d-1} endowed with the topology of $\|\cdot\|_{\infty}$.

¹⁹⁹⁴ Using Lemma 1.1 (inequality (1.37) with (1.61) from first to second term and from
¹⁹⁹⁵ second to third and inequality (1.38) from third to fourth), we get

$$e_{N,\infty}(\boldsymbol{\Gamma}) \leq c \cdot e_{N,\infty}(\mathcal{B}_{\|\cdot\|_{\infty}}^+) \leq c \cdot e_{[N/d]-1,\infty}(\mathcal{B}_{\|\cdot\|_{\infty}}^+(1)) = c \cdot e_{[N/d]-1,\infty}([0, 1]^{d-1}). \quad (1.63)$$

¹⁹⁹⁶ Gathering previous arguments, we have proved the following:

¹⁹⁹⁷ **Lemma 1.4.** Consider $\boldsymbol{\Gamma} \in \mathbb{R}^d$ a generator of a D-norm such that $\|\boldsymbol{\Gamma}\|_{\infty} = c$ a.s.. Then,

1998 the following holds true:

$$\forall \mathbf{x} \in \mathbb{R}^d : \min_{f: \mathbb{R}^d \rightarrow \mathbb{R}^d, \#f(\mathbb{R}^d) = N} \left| \|\mathbf{x}\|_{\Gamma} - \|\mathbf{x}\|_{f(\Gamma)} \right| \leq c \cdot \|\mathbf{x}\|_{\infty} e_{\lceil N/d \rceil - 1, \infty} ([0, 1]^{d-1}). \quad (1.64)$$

1999 From (1.36), we get that:

$$e_{\lceil N/d \rceil - 1, \infty} ([0, 1]^{d-1}) \xrightarrow{N \rightarrow \infty} (\lceil N/d \rceil - 1)^{-1/(d-1)} Q_{\infty} ([0, 1]^{d-1}) \lambda([0, 1]^{d-1})^{1/(d-1)} \quad (1.65)$$

$$\xrightarrow{N \rightarrow \infty} d^{1/(d-1)} N^{-1/(d-1)} Q_{\infty} ([0, 1]^{d-1}) \quad (1.66)$$

2000 which proves (1.14), where $C(d) = c \cdot d^{1/(d-1)} Q_{\infty} ([0, 1]^{d-1})$ in (1.15). The proof of 2001 Theorem 1.3 is thus complete. \square

2002 Observe that the quantization upper bound (1.59) is a bit rough since the right-hand 2003 side of (1.58) refers to the L_1 -quantization that we bound using L_{∞} -quantization. One 2004 could use L_1 -quantization estimates but they would depend on the distribution of Γ and we 2005 prefer to provide worst-case estimates. Instead of using the L_{∞} and using L_1 quantization 2006 on (1.58), (Graf and Luschgy 2000, Th. 6.2.) gives that the first order error of the error 2007 is proportional to the $L_{(d-1)/d}$ norm of the density of the generator on the boundaries of 2008 $\mathcal{B}_{\|\cdot\|_{\infty}}^+ (i)$ times $N^{-1/(d-1)}$. In particular, if the generator Γ has a null density with respect 2009 to the Lebesgue measure (e.g. if the distribution is discrete), the multiplicative factor of the 2010 $N^{-1/(d-1)}$ term is null and the convergence of the bound can be quicker than $N^{-1/(d-1)}$.

2011 3.1 Additional experiments

2012 Experiments on Hüsler-ReiSS and Gaussian models

2013 **Experimental setup.** We consider dimensions $d \in \{2, 5, 10, 20, 50\}$. In the Gaussian 2014 copula model, the equicorrelation parameter is fixed at $\rho = 0.5$. In the Hüsler-ReiSS case, 2015 the equi- λ model is adopted with $\lambda = 1.0$ (*i.e.* $\lambda_{ij} = 1$ for $i \neq j$). In both datasets, 2016 the margins are Pareto distributed with tail index $\alpha \in \{1.5, 2.0, 2.5\}$ (implemented by 2017 inverse-cdf transforms after simulating the copula). For each pair (d, α) and dataset type, 2018 we generate 10,000 training and 2,000 validation points. Results are provided in Table 1.3 2019 of the main paper and Figures 1.7, 1.8, 1.9 and 1.10.

2020 3.2 Additional plots for experiments on Real data

2021 We provide additional illustrations for assessing the performance of HTGAN for real data, 2022 on the Financials dataset in Figure 1.11. Real vs generated data are plotted with: (i) 2023 histograms of marginal densities for four tickers, (ii) scatter plots for two tickers and (iii) 2024 histograms of angles of large data points after projection for two tickers. For scatter plots 2025 and angles, two tickers are selected such that they exhibit quite different values of Kendall's 2026 τ , chosen on the relevant heatmap in Figure 1.5. The selected securities are: Ameriprise 2027 Financial (AMP), Franklin Ressources (BEN), BNY Mellon (BK) and BlackRock (BLK). 2028 This contrast in dependence is reflected in the corresponding scatter plots and angular

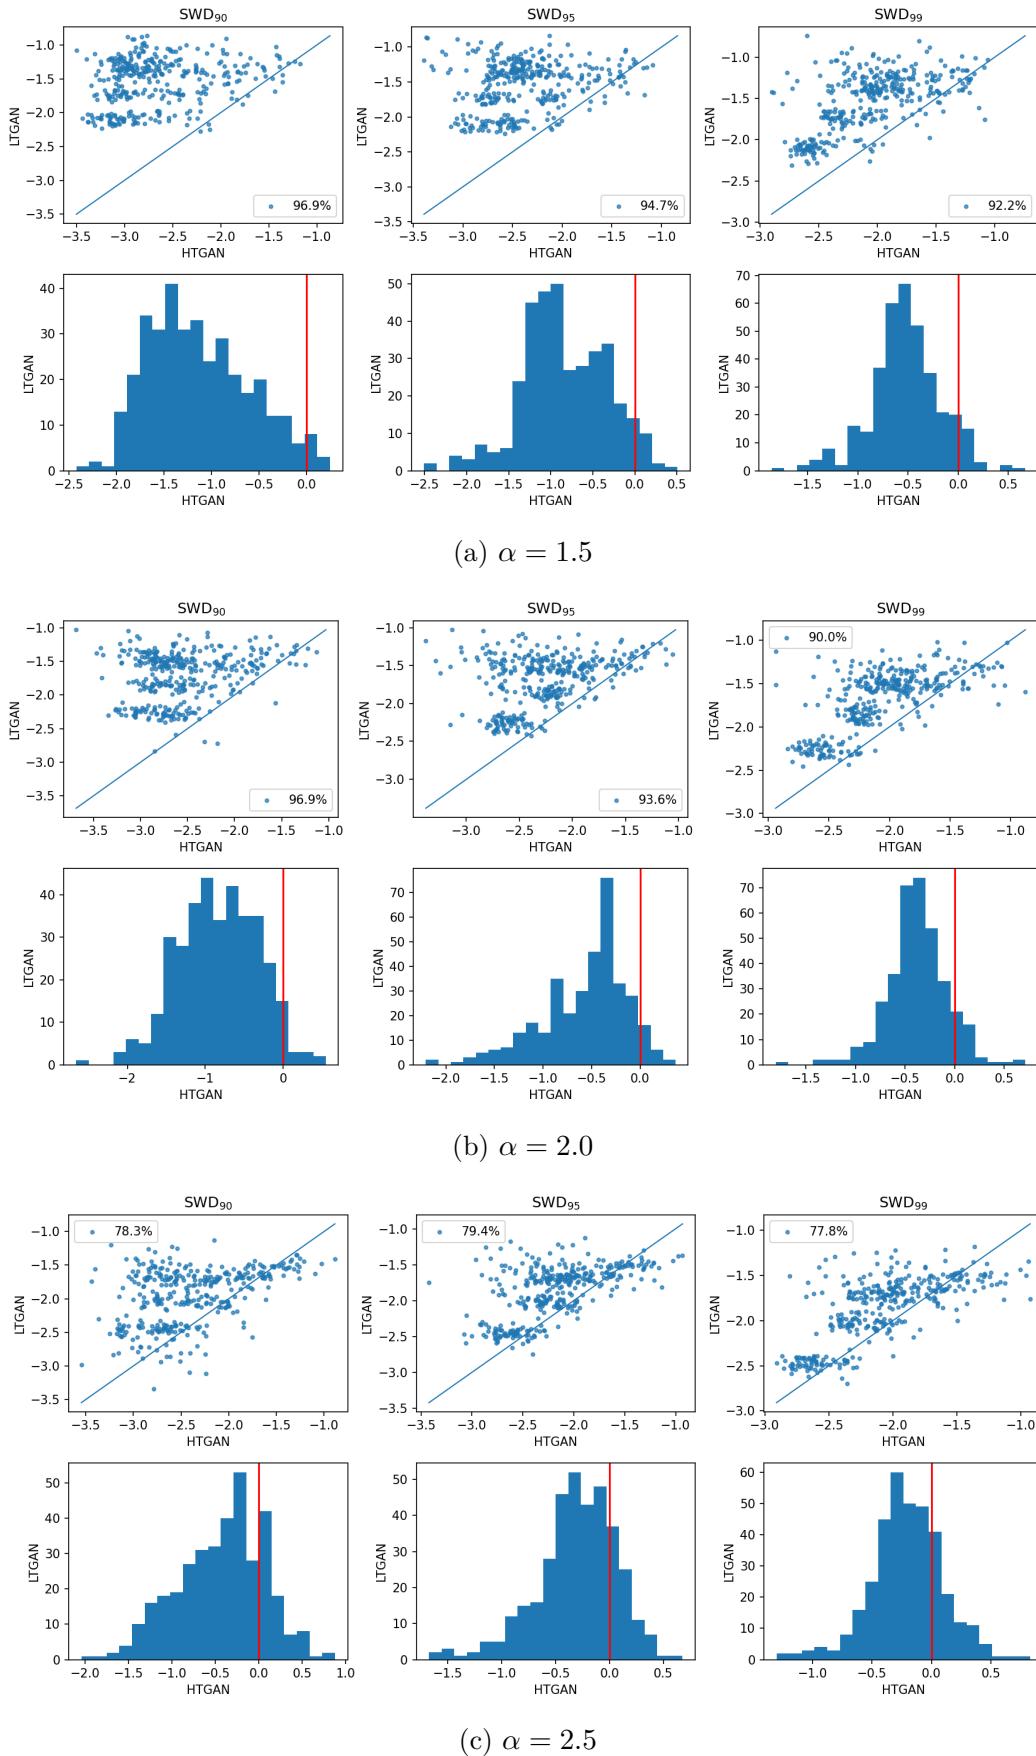


Figure 1.7: Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} = 10,000$ from a Gaussian copula and Pareto margins, $d \in \{2, 5, 10, 20, 50\}$. Results obtained with the SWD metric. See Figure 1.3 for further details.

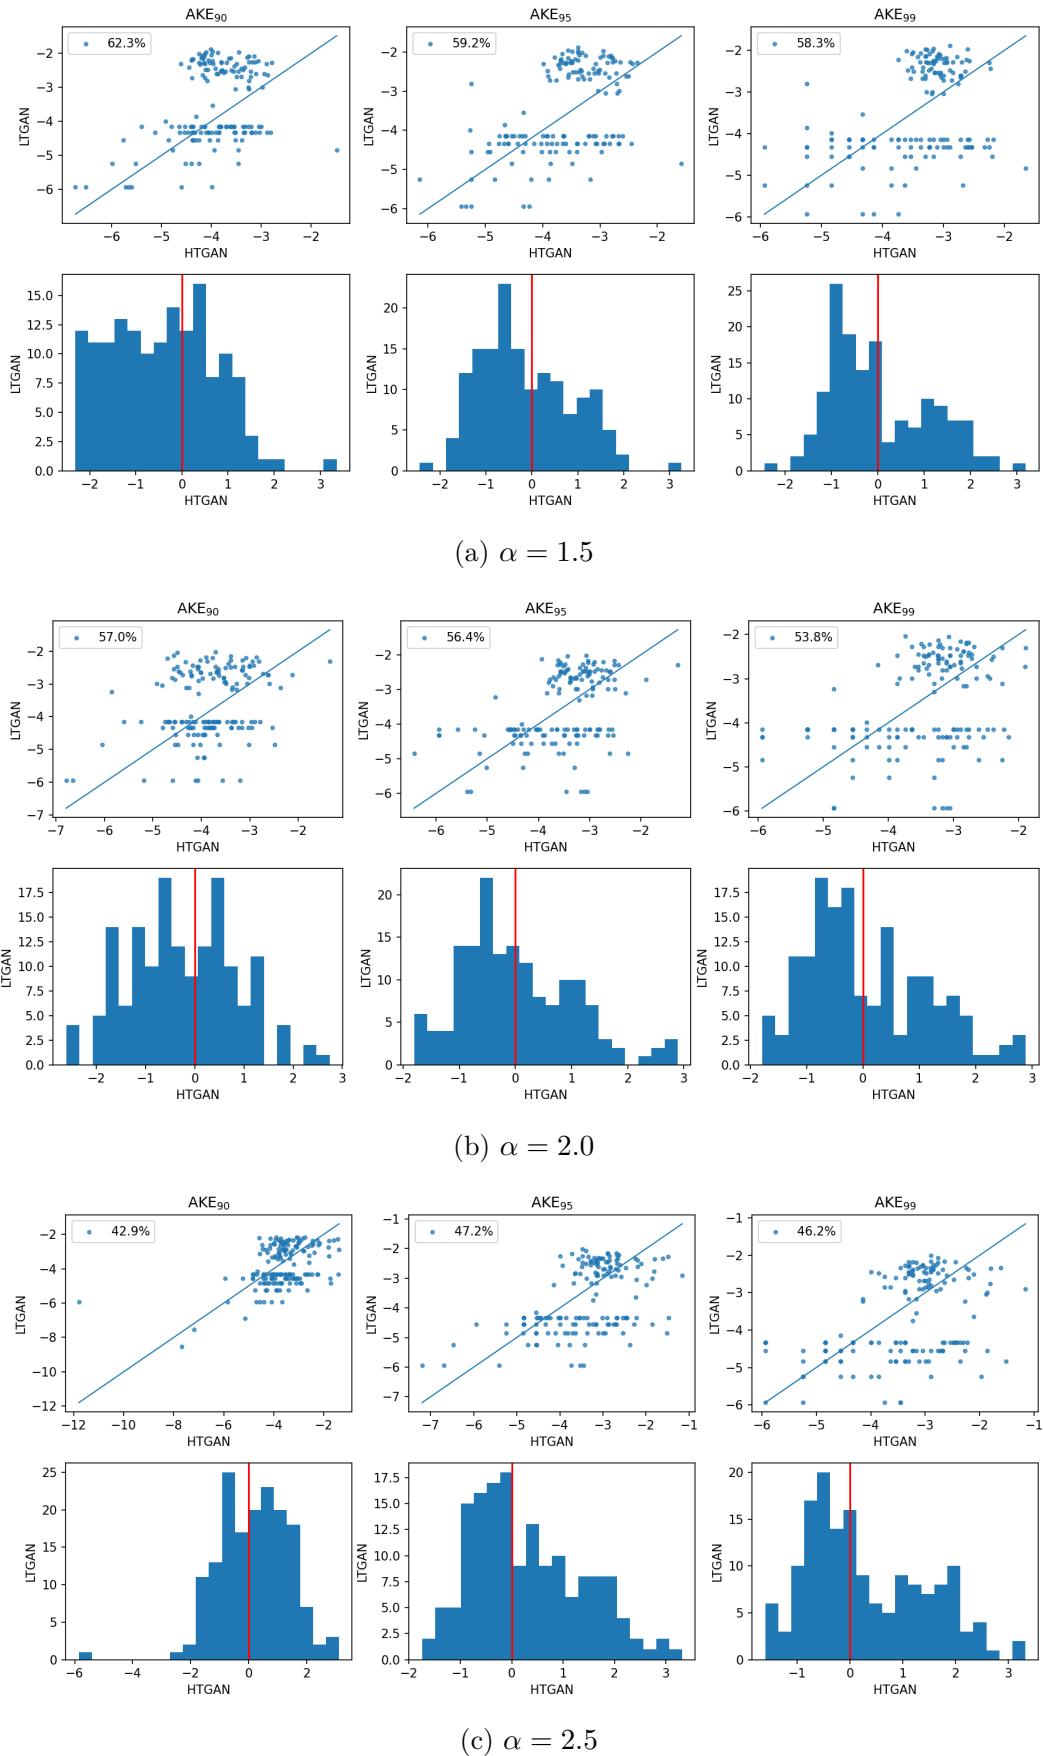


Figure 1.8: Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} = 10,000$ from a Gaussian copula and Pareto margins, $d \in \{2, 5, 10, 20, 50\}$. Results obtained with the AKE metric. See Figure 1.3 for further details.

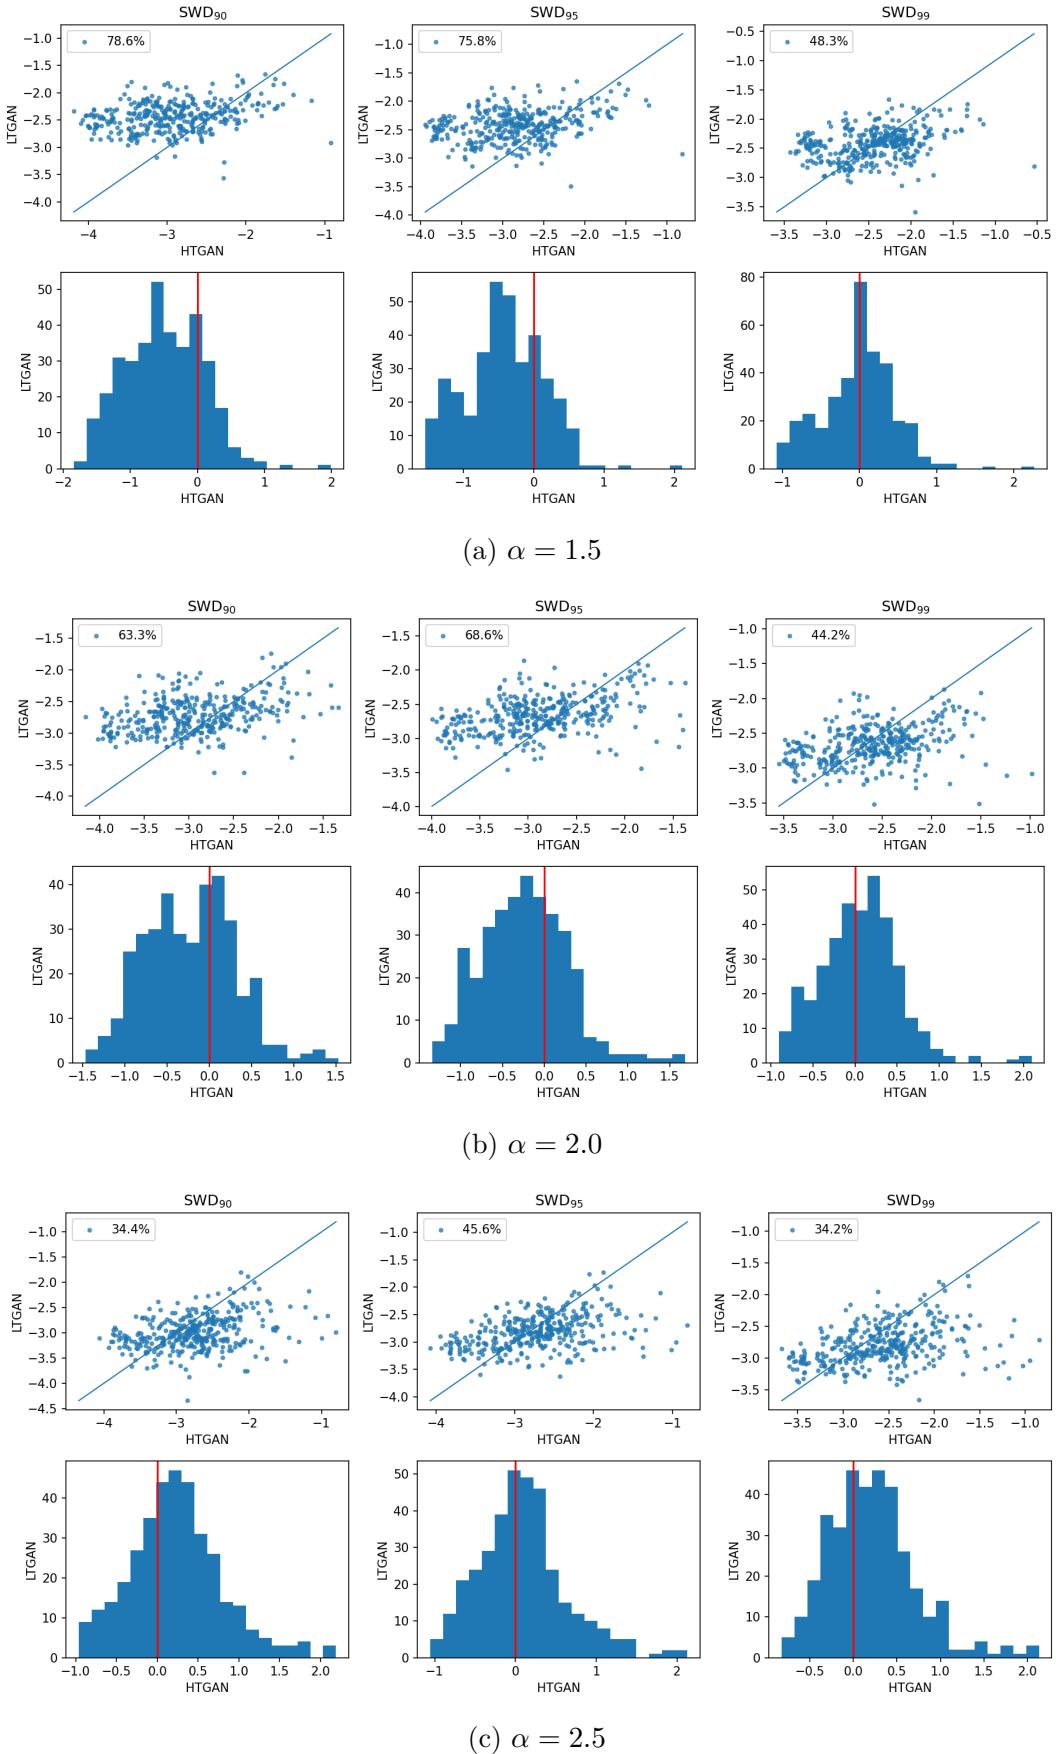


Figure 1.9: Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} = 10,000$ from a Hüsler-ReiSS copula and Pareto margins, $d \in \{2, 5, 10, 20, 50\}$. Results obtained with the SWD metric. See Figure 1.3 for further details.

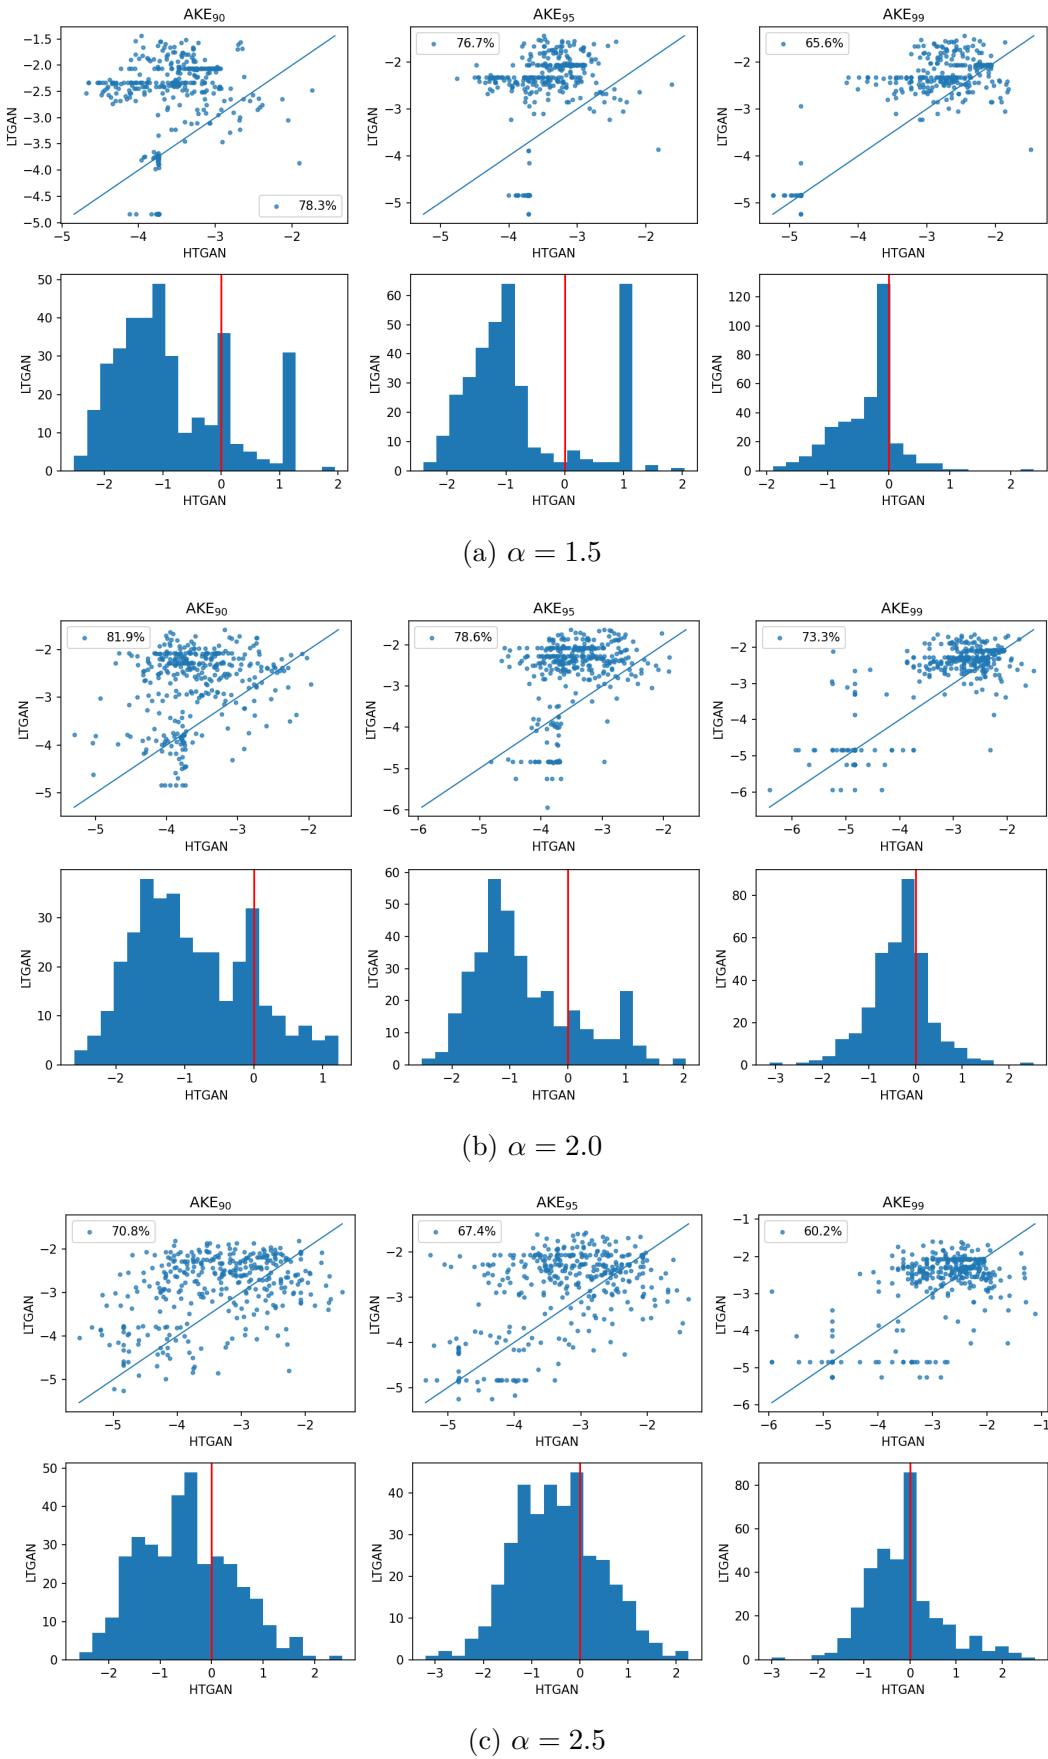


Figure 1.10: Illustration in a d -dimensional setting on simulated data of size $n_{\text{data}} = 10,000$ from a Hüsler-ReiSS copula and Pareto margins, $d \in \{2, 5, 10, 20, 50\}$. Results obtained with the AKE metric. See Figure 1.3 for further details.

2029 distributions. These illustrations highlight HTGAN's capacity to capture various and
2030 complex asymptotic dependence structures - even if not perfectly - despite the challenges
2031 posed by the high-dimensional setting and the limited sample size.

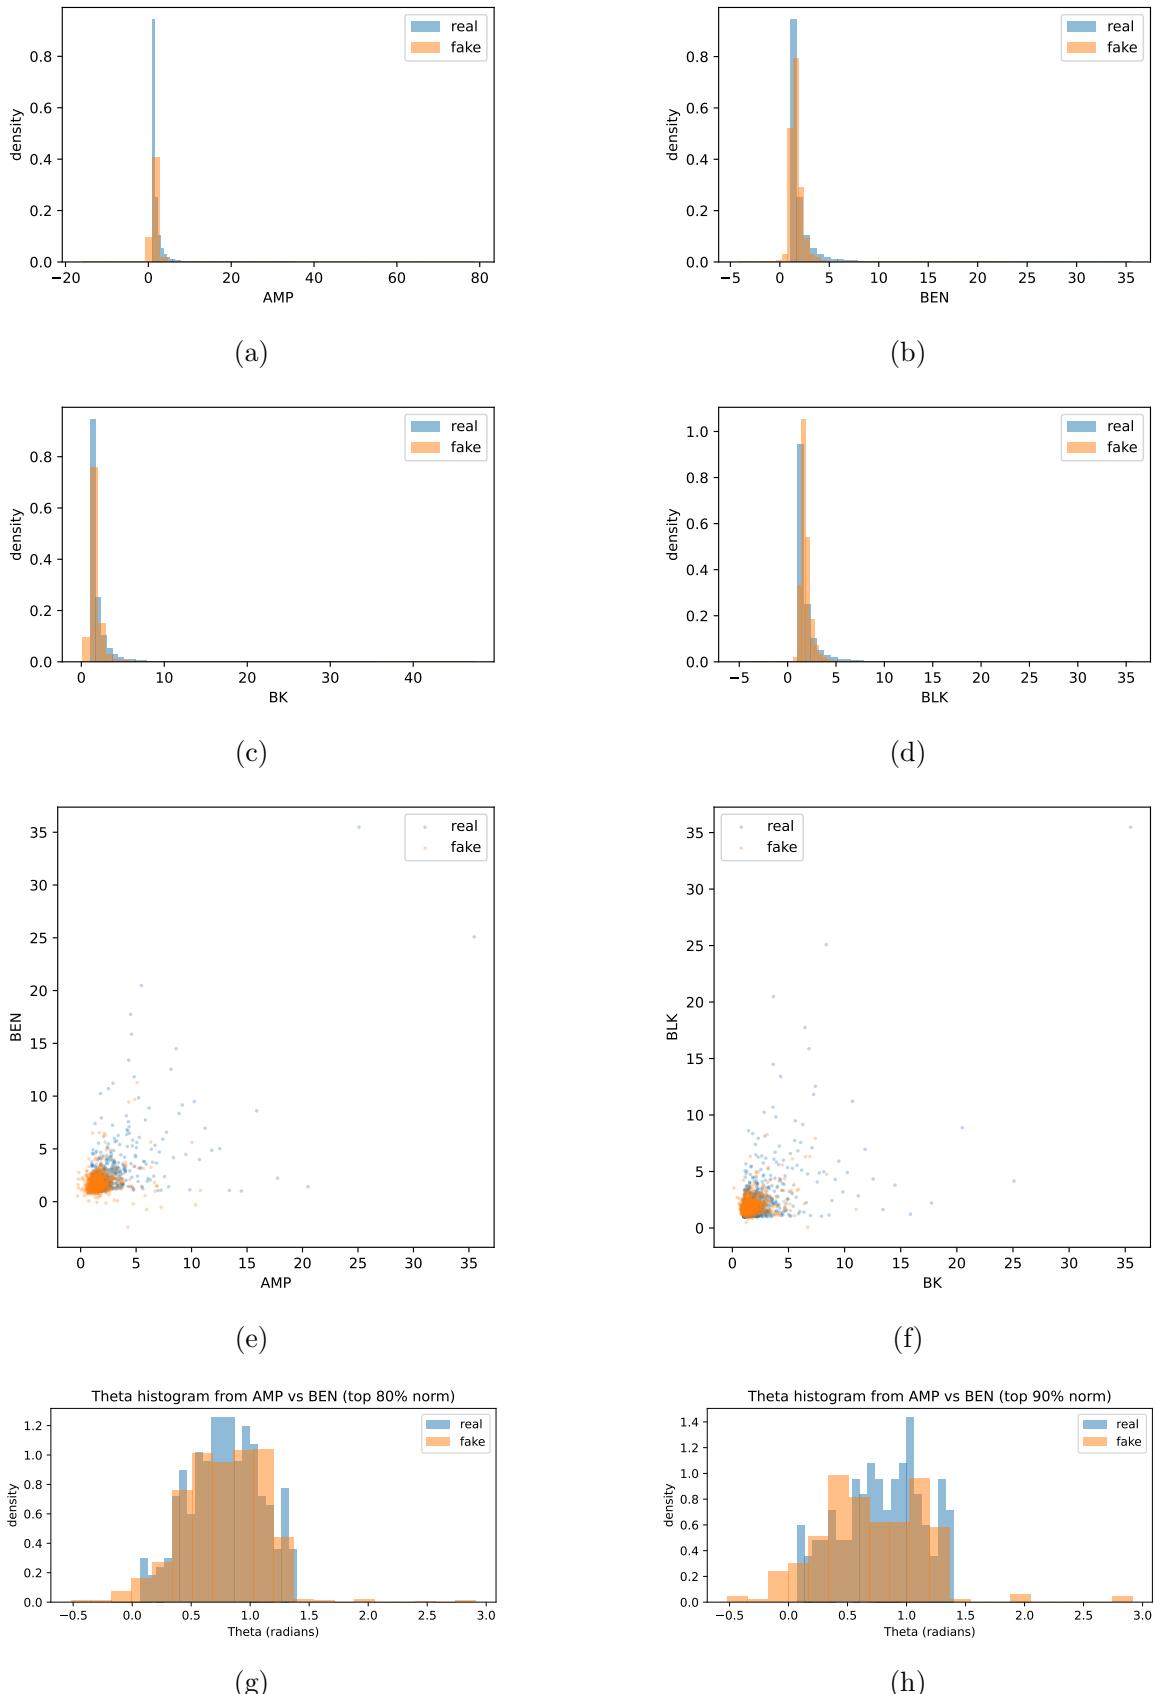


Figure 1.11: (a)-(d): histograms of real vs synthetic marginals across dimensions one to four. (e)-(f): Scatter plot AMF vs BEN and BK vs BLK tickers. (g)-(h): Histogram of θ after selecting largest values, at level 80% and 90%.

2032 **Chapter 2**

2033 **Generative Neural Order
Statistics**

2035 This chapter develops the theory and algorithms for generating order statistics using neural
2036 networks. We analyze the representational capabilities of neural network-based generative
2037 models for sorted data, deriving approximation bounds based on the Sukhatme and Schu-
2038 cany representations.

2039 **2.1 Introduction**

2040 **Sustainable investing.** Sustainable finance refers to the process of taking environmen-
2041 tal, social and governance (ESG) considerations into account when making investment
2042 decisions in the financial sector, leading to more long-term investments in sustainable eco-
2043 nomic activities and projects, as defined by the European Commission ([2024a](#)). Within
2044 sustainable finance, a particular quantitative framework, impact investing, is broadly de-
2045 fined as investments that consider not only financial objectives but also other goals such as
2046 supporting certain social priorities and agendas (Lo and R. Zhang [2021](#)). Impact investing
2047 seeks to optimize both financial objectives and metrics that quantify the compliance level
2048 of assets with sustainable goals. For example, ratings on ESG factors provide information
2049 about the sustainability performance of a company or a financial instrument, by assess-
2050 ing its exposure to sustainability risks and/or its impact on people and the environment
2051 (European Commission [2024b](#)).

2052 **Impact investing, order statistics and induced order statistics.** Recent develop-
2053 ments in Mathematical finance and economics have tackled the challenging task of trans-
2054 lating sustainable investing into a mathematical framework. Among those, the work in
2055 (Lo and R. Zhang [2021](#)) introduces a general framework for quantifying the impact of im-
2056 pact investing. Namely, authors propose to formalize impact investing as the sorting and
2057 selection of an investment universe of N securities based on an *impact factor*, typically the
2058 ESG score, $\mathbf{X} = (X_1, \dots, X_N)^T$ and a *performance metric*, typically the return associated

with an asset, $\Theta = (\Theta_1, \dots, \Theta_N)$. In this framework, \mathbf{X} and Θ are time series evolving in time. It is a common assumption to suppose that both the performance metric and the impact factor are identically distributed throughout time and we will hold this assumption in this work. In this case, the training data consists of the sorted realizations of the time series. Suppose that \mathbf{X} varies with time, denoted \mathbf{X}_t . Under the assumption of identical distributions over time, we drop the time index and treat each observation as a vector. Typically, we are interested in time scales where there are significant changes in the values taken by the impact factor, i.e., periods corresponding to months or years. This excludes high-frequency trading. In this framework, we rank assets by the impact factor and study the associated returns ordered accordingly. This operation fits into the framework of *order statistics* and *induced order statistics*. Formally: suppose that we have a dataset consisting of i.i.d. realizations of the investment universe $\{(X_1, \Theta_1), \dots, (X_N, \Theta_N)\}$, the order statistics of the impact factor \mathbf{X} is defined as a random permutation σ of its elements satisfying:

$$X_{\sigma(1)} \leq \dots \leq X_{\sigma(N)}. \quad (2.1)$$

Recall that throughout the paper, we suppose that the realizations X_1, \dots, X_N are i.i.d.. Note that if the distribution of \mathbf{X} is continuous, such a permutation is unique *almost surely*. Also note that the permutation is random and that each instance depends on the realization X_1, \dots, X_N . We therefore denote, for any $i \in [N]$, $X_{i:N} = X_{\sigma(i)}$, the i^{th} order statistic. Furthermore, note the marginal law of $X_{k:N}$ is different of the marginal law of X_k for each k . The associated *induced order statistics* (IOS) are defined as:

$$\forall i \in [N] : \Theta_{[i:N]} = \Theta_{\sigma(i)}. \quad (2.2)$$

An *impact portfolio strategy*, as defined in (Lo and R. Zhang 2021), is a strategy that considers both the performance metric Θ and the impact factor \mathbf{X} . Typically, an impact driven investor is willing to invest in stocks that have high scores for the impact factor (high values of \mathbf{X}) but also have good values for the performance metric Θ . In this view, it is crucial to perform accurate estimation of the joint distribution of order statistics and the associated induced order statistics counterparts. Modeling order statistics and induced order statistics is an rich, ongoing research field. A thorough treatment of theory of order statistics is given in (H. David and H. Nagaraja 2004). The work (Lo and R. Zhang 2021) proposes a model-based assumption of the impact factor and the performance metric. This approach, combined with a representation theorem of the data, allows for powerful representation theorems for the new portfolio.

Machine Learning and Generative Modeling. Financial datasets are typically scarce and very high dimensional. In a typical setting, \mathbf{X} and Θ are observed over a period of 30 years. If the case where daily data are observable, it amounts to roughly around $30 \text{ days} \times 12 \text{ months} \times 30 \text{ years} \approx 10,000$ data points. Note that it usually makes sense to study data at the granularity of a day for the performance metric (for

example, considering daily stocks returns), but not necessarily for the impact factor. For example, the scale of change for the ESG score is of month or even years. On the other side, the data is of size N , where N corresponds to the size of the investment universe. In the case of the *S&P500* index, the size amounts to roughly 500. Such characteristics make it difficult to build predictive models or perform statistical estimation without a calibration model. To cope with this difficulty, a line of research has focused on augmenting the data available with accurate simulations. Historically, for financial applications, approaches for data augmentation include bootstrap methods, Monte Carlo simulations and parametric models. In portfolio management, Monte Carlo methods are widely used by practitioners (Glasserman 2013). Generally speaking, the simulation approach for data augmentation can be framed as mapping a source of randomness to the data of interest. In our impact investing framework, the problem is defined as finding a generator function \mathcal{G} associated to a source of randomness \mathbf{Z} such that:

$$\mathcal{G}(\mathbf{Z}) \stackrel{d}{=} ((X_{1:N}, \Theta_{[1:N]}), \dots, (X_{N:N}, \Theta_{[N:N]})) \quad (2.3)$$

where $\stackrel{d}{=}$ stands for equality in distribution. The problem of the existence of tractable models achieving exactly (2.3) is an open question. (Lo and R. Zhang 2021, Theorem 1) provides a representation theorem that maps a source of randomness to the data distribution.

Problem statement. It is not known how real (understand finite, approximate) machine learning models are able to correctly approximate the link function Φ . Our analysis is therefore oriented towards finding non-asymptotic guarantees on the quality of approximation of Φ . We build on work from the theory of function approximation with neural network (Cybenko 1989; Hornik et al. 1989; Yarotsky 2017) and properties such as the Markovian property (H. David and H. Nagaraja 2004, Chapter 2, Section 5) of order statistics and induced order statistics. Moreover, the problem of modeling impact portfolios lies in its high dimensional nature. Typically, the number of asset in the universe is large (eg: $N = 500$ for the S&P 500 index and $N = 1,000$ for the Russell 1000 index) and we have only few data at disposition. Typically, 20 to 30 years of data are available. Considering monthly datasets, that yields a dataset with less than 400 data points. It is neither theoretically nor practically feasible to learn a generative model with so few data points. However, in view of impact investing, we are solely interested in modeling assets with best performing performance metrics.

We are focused on modeling best performing assets in regard of the performance metric, in descending order:

$$\mathcal{G}(\mathbf{Z}) \stackrel{d}{=} ((X_{N:N}, \Theta_{[N:N]}), \dots, (X_{N-d+1:N}, \Theta_{[N-K:N]})), \quad (2.4)$$

where typically the dimension of datapoints $d = \lceil \eta N \rceil$, η is a thresholding value, e.g. 5% = 0.05.

2130 3.A Notations

2131 We summarize the main notations used throughout the paper.

- 2132 • **Indices and sets:** For an integer $N \geq 1$, we write $[N] \stackrel{\text{def}}{=} \{1, 2, \dots, N\}$. The symbol
2133 $k \in [N]$ denotes an index.
- 2134 • **Scalars, vectors, matrices:** Lowercase letters (x, y, u, \dots) denote scalars. Bold
2135 ($\mathbf{x}, \mathbf{X}, \mathbf{y}, \mathbf{Y}, \dots$) denote column vectors. The i -th component of \mathbf{x} is x_i .
- 2136 • **Probability:** \mathbb{P} and \mathbb{E} denote probability and expectation. Equality in distribution
2137 is written $X \stackrel{d}{=} Y$. The indicator of an event A is $\mathbb{I}\{A\}$. The cumulative distribution
2138 of X is commonly written F_X , except if it is explicitly stated otherwise.
- 2139 • **Spaces and functions:** \mathbb{R} is the set of real numbers, \mathbb{N} the set of positive integers.
2140 For a cumulative distribution function (cdf) F , F^{-1} denotes the (left-continuous)
2141 quantile function. We use $\ln(\cdot)$ for the natural logarithm.
- 2142 • **Order statistics and induced order statistics (OS/IOS):** For i.i.d. $X_1, \dots, X_N \in$
2143 \mathbb{R} with cdf F_X , the order statistics are $X_{1:N} \leq \dots \leq X_{N:N}$, where $X_{k:N}$ is the k -th
2144 order statistic. For i.i.d. pairs (X_i, Θ_i) , the concomitant of $X_{r:N}$ is $\Theta_{[r:N]}$, i.e., the Θ
2145 attached to the observation whose X -value ranks r th; see Definition 2.1.
- 2146 • **Uniforms and exponentials:** U denotes a standard uniform random variable on
2147 $(0, 1)$. $\mathcal{E}(1)$ denotes the unit exponential law with density $f(z) = e^{-z} \mathbb{I}\{z \geq 0\}$. We
2148 write $Z \sim \mathcal{E}(1)$ and use $F(z) = 1 - e^{-z}$, $F^{-1}(u) = -\ln(1 - u)$.
- 2149 • **Neural networks:** G_{NN} denotes a neural network generator mapping i.i.d. inputs
2150 (e.g., uniforms) to outputs. When a family of networks is used, superscripts indicate
2151 roles, e.g., F_{NN} and F_{NN}^{-1} for approximations of F and F^{-1} .
- 2152 • **Asymptotics:** $a_N \sim b_N$ means $\lim_{N \rightarrow \infty} a_N/b_N = 1$. Big- \mathcal{O} notation is standard.

2153 2.2 A theoretical analysis of generative modeling of 2154 order statistics with Neural Networks

2155 Throughout this section, we let N denote the dimension of the vector to be sorted for
2156 obtaining order statistics. For $k \in [N]$, the k^{th} order statistic is denoted by $X_{k:N}$. We
2157 leverage these representations to analyze the capacity of neural networks to model order
2158 statistics (OS) and induced order statistics (IOS). Formally, we seek a neural network G_{NN}
2159 such that

$$G_{NN}(U_1, \dots, U_N) \stackrel{d}{=} (U_{1:N}, \dots, U_{N:N}), \quad (2.5)$$

2160 where U_1, \dots, U_N are independent standard uniform random variables. To address this,
2161 we consider several possible constructions. We first recall the following result:

2162 *Proposition 2.1* (H. David and H. Nagaraja (2004), eq. (2.3.7), p. 15). Let $X \in \mathbb{R}$ be a
2163 random variable with arbitrary distribution function F_X . Let U be a standard uniform
2164 random variable. Then:

$$(X_{1:N}, \dots, X_{N:N}) \stackrel{d}{=} (F_X^{-1}(U_{1:N}), \dots, F_X^{-1}(U_{N:N})), \quad (2.6)$$

2165 This proposition shows that, up to transformation by the inverse cdf, the joint distri-
2166 bution of order statistics from any i.i.d. sample is determined by the order statistics of the
2167 uniform distribution. Therefore, in this section, we focus solely on approximation results
2168 for the order statistics of the uniform distribution.

2169 3.A Representational results for order statistics

2170 We summarize two key sampling schemes for order statistics, with detailed proofs provided
2171 in Section 2.B.

2172 **Schucany's scheme.** Order statistics exhibit a Markov property. Let U_1, \dots, U_N be
2173 i.i.d. standard uniform random variables. Following (Schucany 1972), the order statistics
2174 of the uniform distribution can be recursively represented as

$$\begin{cases} U_{N:N} = U_1^{1/N}, \\ U_{N-r:N} = U_{N-r+1:N} \cdot U_r^{1/(N-r)} \quad \forall r \in \{1, \dots, N\}. \end{cases} \quad (2.7)$$

2175 **Sukhatme's representation.** Another important representation, due to (Sukhatme
2176 1937), also exhibits a Markov chain-like structure. Let Z_1, \dots, Z_N be i.i.d. standard ex-
2177 ponential random variables, with density

$$f(z) = e^{-z}, \quad z \geq 0. \quad (2.8)$$

2178 Let $F(z) = 1 - e^{-z}$ denote the associated cdf. Then, for all $r = 1, \dots, N$,

$$Z_{r:N} \stackrel{d}{=} \sum_{i=1}^r (Z_{i:N} - Z_{i-1:N}) = \sum_{i=1}^r \frac{Y_i}{N-i+1}, \quad (2.9)$$

2179 where Y_1, \dots, Y_N are i.i.d. standard exponential random variables and $Z_{0:N} = 0$ by con-
2180 vention. Applying Proposition 2.1, and noting that $F^{-1}(u) = -\ln(1-u)$ for $u \in (0, 1)$,
2181 we have

$$U_{r:N} \stackrel{d}{=} F(Z_{r:N}) = 1 - e^{-Z_{r:N}}. \quad (2.10)$$

2182 3.B Generating order statistics: approximation results

2183 To analyze how well neural networks can approximate the mapping from independent
2184 uniforms to order statistics, we establish bounds for the uniform case.

2185 **Learning to generate order statistics based on Schucany's scheme.** We investi-
2186 giate the approximation quality of neural network-based generative schemes for uniform
2187 order statistics, using Schucany's recursive representation. The algorithm iteratively ap-
2188 plies (2.7), with neural networks $\hat{G}_{r:N}^{NN}$ approximating the functions

$$G_{m:N}(u_1, u_2) = u_1 \cdot u_2^{1/m}. \quad (2.11)$$

2189 Thus, (2.7) can be rewritten as

$$U_{N-r:N} = G_{N-r:N}(U_{N-r+1:N}, U_r). \quad (2.12)$$

Algorithm 3: Sampling from order statistics $U_{1:N} \leq \dots \leq U_{N:N}$.

Input : Learned functions $\hat{G}_{r:N}^{NN} : [0, 1] \rightarrow \mathbb{R}$ for $r \in [N]$; N uniform samples

U_1, \dots, U_N

Output: $\hat{U}_{1:N}, \dots, \hat{U}_{N:N}$

2190 1 Set $\hat{U}_{N:N} = U_N^{1/N}$;
2 **for** $r \leftarrow 1$ **to** N **do**
3 $\hat{U}_{N-r:N} \leftarrow \hat{G}_{N-r:N}^{NN}(\hat{U}_{N-r+1:N}, U_r)$;
 // In parallel: $U_{N-r:N} \leftarrow U_{N-r+1:N} U_r^{1/(N-r)}$.
4 **return** $\hat{U}_{r:N}$ and $U_{r:N}$, for $1 \leq r \leq N$.

2191 With this procedure, we obtain the following result:

2192 **Theorem 2.1.** Let $\varepsilon > 0$. There exists a family of ReLU Neural Networks $G_{N-r:N}^{NN}$ for
2193 $1 \leq r \leq N$ such that, using the sampling procedure of Algorithm 3 and setting

$$\Delta V_r = |\hat{U}_{N-r:N} - U_{N-r:N}|, \quad (2.13)$$

2194 we have $\mathbb{E}|\Delta V_r| < \varepsilon$, where the total number of neurons in the family $\{G_{N-r:N}^{NN}\}_{1 \leq r \leq N}$ is

$$\mathcal{C}_{Schucany}(\varepsilon) = \sum_{i=1}^r 3 \left(\left\lceil \left(\frac{2r(N-i)}{(N-r)\varepsilon} \right)^{N-i} \right\rceil + 1 \right) + c_1 \ln \left(\frac{8r(N-i)}{(N-r)\varepsilon} \right) + c_2. \quad (2.14)$$

2195

2196 The proof is given in Section 2.E. This result is pessimistic: the error is asymptotically
2197 polynomial of degree N in ε^{-1} .

2198 **Learning with Sukhatme's representation.** Next, we consider the Sukhatme rep-
2199 resentation for generating order statistics. Let F denote the cdf of the unit exponential
2200 distribution, $F(z) = 1 - e^{-z}$ for $z > 0$, and $F^{-1}(u) = -\ln(1-u)$ for $u \in (0, 1)$. Let F_{NN}
2201 and F_{NN}^{-1} be neural network approximations of F and F^{-1} , respectively. The sampling
2202 procedure is detailed in Algorithm 4:

Algorithm 4: Sampling from order statistics $U_{1:N} \leq \dots \leq U_{N:N}$ based on (2.9)

Input : Learned function \hat{G}^{NN} ; N uniform samples U_1, \dots, U_N

Output: $\hat{Z}_{1:N}, \dots, \hat{Z}_{N:N}$

1 Set $\hat{U}_{0:N} = U_{0:N} = 0$;

2203 2 **for** $r \leftarrow 1$ **to** N **do**

3 $\hat{Z}_{r:N} \leftarrow \hat{Z}_{r-1:N} + \frac{F_{NN}^{-1}(U_r)}{N-r+1};$

 // In parallel: $Z_{r:N} \leftarrow Z_{r-1:N} + \frac{F^{-1}(U_r)}{N-r+1}.$

4 **return** $\hat{U}_{r:N} = F_{NN}(\hat{Z}_{r:N})$ and $U_{r:N} = F(Z_{r:N})$, for $r \in \{1, \dots, N\}$.

2204 **Theorem 2.2.** *With the sampling strategy of Algorithm 4, there exist ReLU Neural Net-*
 2205 *works F_{NN} and F_{NN}^{-1} such that the error in approximating the uniform order statistics*
 2206 *satisfies, in expectation,*

$$\mathbb{E}|U_{r:N} - \hat{U}_{r:N}| \leq \varepsilon \left[1 + 2 \sum_{i=1}^r \frac{1}{N-i+1} \right], \quad (2.15)$$

2207 *with at most*

$$\mathcal{C}_{Sukhatme}(\varepsilon) \leq c\varepsilon^{-2}(2\ln(1/\varepsilon) + 1) + c\varepsilon^{-1}\ln(1/\varepsilon)(\ln(1/\varepsilon) + \ln(\ln(1/\varepsilon)) + 1) \quad (2.16)$$

2208 *neurons involved, for some constant c .*

2209 The proof is provided in Section 2.F. For $r = N$ (the worst case), the error term admits
 2210 the asymptotics, as $N \rightarrow \infty$:

$$\varepsilon \left[1 + 2 \sum_{i=1}^N \frac{1}{N-i+1} \right] \sim \varepsilon \ln(N), \quad (2.17)$$

2211 and the complexity term satisfies the following asymptotic results, as $\varepsilon \rightarrow 0$:

$$\mathcal{C}_{Sukhatme}(\varepsilon) \sim 2c\varepsilon^{-2}\ln(\varepsilon^{-1}). \quad (2.18)$$

2212 This representation is more advantageous than that of Theorem 2.1, as the network com-
 2213 plexity scales asymptotically as a polynomial of order 2 in ε^{-1} times a logarithmic term.

2214 **Discussion.** The two results are not strictly comparable: one yields an error of order
 2215 ε , while the other yields an error of order $\varepsilon \ln(N)$. However, Algorithm 4 presents a key
 2216 advantage: it exploits a structural property linking order statistics to linear combinations
 2217 of i.i.d.exponential variables. As a result, the main complexity in neural network approx-
 2218 imation is concentrated in accurately simulating unit exponential random variables from
 2219 standard uniform inputs. Moreover, it enables computational reuse: we only need to ap-
 2220 proximate the inverse cdf function F^{-1} , where $Z \sim \mathcal{E}(1)$, rather than approximating N
 2221 separate functions as required by Algorithm 3.

2222 2.3 Experiments

2223 In this section, we present a suite of synthetic experiments designed to evaluate the performance
 2224 of generative models in producing samples of order statistics. All results reported
 2225 here are obtained using synthetic data, allowing us to assess model capabilities under
 2226 controlled conditions where the ground truth distribution is known. Our experimental
 2227 framework is specifically tailored to test how well different neural network-based gener-
 2228 ative models—notably various GAN architectures—can approximate the distributional and
 2229 structural properties of order statistics.

2230 We systematically compare several GAN variants, including classic GANs, Wasserstein
 2231 GANs (WGAN), and WGANs with gradient penalty (WGAN-GP). These models are eval-
 2232 uated exclusively on synthetic datasets, which enables precise quantitative measurement of
 2233 model performance. The data used in our experiments are constructed to exhibit the char-
 2234 acteristics of order statistics, as described earlier, providing a rigorous testbed for model
 2235 assessment.

2236 Our experimental design incorporates a range of evaluation metrics that capture both
 2237 standard distributional similarity and the sortedness and dependence properties intrinsic
 2238 to order statistics. This approach allows us to identify model strengths and limitations.

2239 The following subsections detail our experimental protocols, model configurations, and
 2240 results, focusing on synthetic data benchmarks.

2241 3.A Synthetic data

2242 We begin by recalling a setup for a GAN to learn simple nonlinearities across different
 2243 synthetic datasets, and then extend this setup to the case of order statistics. This vanilla
 2244 test case illustrates some design choices for implementation, especially data normalization
 2245 during training and sampling.

2246 We provide evidence that these models can learn simple multivariate distributions,
 2247 commonly used for sanity checks in the literature. We then test whether this architecture
 2248 can learn order statistics. To do so, we quantify performance with respect to two aspects:
 2249 first, the ability of the model to reproduce the marginal distribution of the components;
 2250 second, the ability to reproduce dependence among coordinates. A key structural property
 2251 of order statistics is that they should be ordered. We verify whether the components of the
 2252 model output are sorted by computing a mean 0/1 value, and then proceed to computation
 2253 of a softer metric that quantifies the excess.

2254 GANs, variants and hyperparameters

2255 Generative Adversarial Networks, as introduced by I. J. Goodfellow et al. (2014), are
 2256 trained by playing a minimax game between a generator network \mathcal{G} and a discriminator
 2257 network \mathcal{D} (dependence on parameters is implicit). Namely, both networks optimize the

2258 common objective:

$$\min_{\mathcal{G}} \max_{\mathcal{D}} V(\mathcal{D}, \mathcal{G}) \quad \text{where} \quad V(\mathcal{D}, \mathcal{G}) = \mathbb{E}_{\mathbf{X} \sim p_{\text{data}}} [\log \mathcal{D}(\mathbf{X})] + \mathbb{E}_{\mathbf{Z} \sim p_{\mathbf{Z}}} [\log(1 - \mathcal{D}(\mathcal{G}(\mathbf{Z})))], \quad (2.19)$$

2259 where p_{data} is the underlying distribution of the data and $p_{\mathbf{Z}}$ is the distribution of the
 2260 latent variable, a distribution that is chosen to be easy to sample from, typically Gaussian
 2261 or Uniform. Throughout this work, we use a uniformly distributed latent variable. More
 2262 details on GANs and variants are given in Section 2.G. Regarding architecture, we consider
 2263 both networks to be simple MLP models.

2264 **Hyperparameters.** The hyperparameters to define in an implementation include the
 2265 batch size, the latent dimension, the learning rate, the number of epochs, and of course the
 2266 architectures of both neural networks (generator and discriminator). Moreover, a common
 2267 practice is, at each epoch, to train the generator for `g_step` steps and the discriminator
 2268 for `d_step`. The following choices are common practice: `g_step=1` and `d_steps=5`. We
 2269 provide more details on hyperparameters in Section 2.H.

2270 How accurate is a GAN at approximating order statistics?

2271 In this section, we seek to find the limitations of the representation power of GANs for
 2272 modeling order statistics. For this purpose, we choose a classic MLP architecture for the
 2273 networks, illustrate that it can fit a classical benchmark dataset properly, and then assess
 2274 visually up to which dimension the model can generate proper order statistics. We also
 2275 use this toy-study to discuss the special scaling used.

2276 **Fitting a benchmark dataset.** A simple fitting test can be done on the banana-shaped
 2277 distribution of (Haario et al. 1999), defined as follow: consider $(x_1, x_2) \sim \mathcal{N}(0, I)$ and:

$$\begin{cases} y_1 = x_1 \\ y_2 = x_2 + b(x_1^2 - 1) \end{cases} \quad (2.20)$$

2278 where $b \in \mathbb{R}$ controls the curvature of the distribution. An illustration of the GAN
 2279 performance is given in Figure 2.1.

2280 Fitting naive model on order statistics

2281 We now take the model obtained from the benchmarking on the banana-shaped distribution
 2282 to fit order statistics. We sample *large* order-statistics:

$$(X_{N,N}, \dots, X_{N-d+1:N}). \quad (2.21)$$

2283 **Normalization scheme.** Per-coordinate min–max normalization, while common in GANs,
 2284 applies different affine maps to each coordinate; real data then leave the sorted manifold,

Hyperparameter	Value
Latent dimension (ℓ)	5
Output dimension (d)	4
Batch size	256
Learning rate	$1 \cdot 10^{-4}$
Epochs	5,000
Discriminator steps (n_D)	5
Generator steps (n_G)	1
Normalization scale factor	150
Optimizer	Adam
Activation	LeakyReLU

Table 2.1: Hyperparameters used for training our vanilla model.

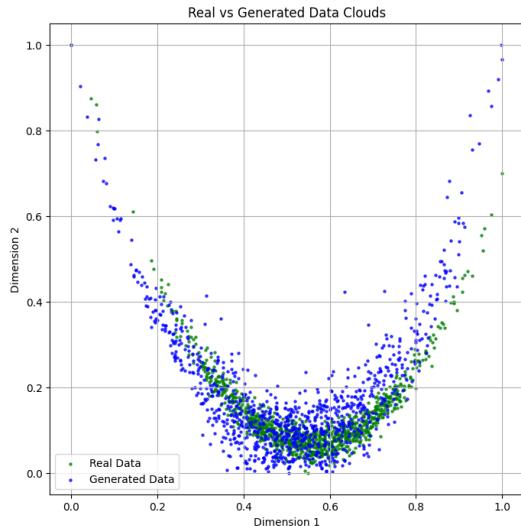


Figure 2.1: Real and generated data after min-max renormalization. n_samples=1,000.

and the generator has no incentive to produce ordered outputs. To respect the combinatorial structure of order statistics during learning, we avoid per-coordinate normalizations (e.g., min–max applied separately to each dimension), which apply different affine maps to each coordinate and can flip pairwise differences within a sample. Instead, we use a single *global* affine transform applied jointly to all coordinates: we subtract the global mean computed over the entire training set and rescale by a positive factor $s > 0$,

$$\tilde{\mathbf{x}} = s(\mathbf{x} - \bar{\mathbf{x}}), \quad \bar{\mathbf{x}} \stackrel{\text{def}}{=} \frac{1}{Nd} \sum_{n=1}^N \sum_{j=1}^d x_j^{(n)}. \quad (2.22)$$

This mean–scale normalization stabilizes optimization *and* exactly preserves the ordering constraint: for any sample \mathbf{x} and indices $i < j$, $\text{sign}(\tilde{x}_i - \tilde{x}_j) = \text{sign}(x_i - x_j)$. The scale s is a hyperparameter chosen for numerical convenience and visualization; adding a constant offset after training can further improve plot readability without affecting sortedness. We apply the same transform to real and generated data for a fair comparison. An illustration of the normalization scheme is provided in Figure 2.2. Empirically, the min–max model yields visually plausible marginals, yet many generated points cross the diagonal, signaling broken sortedness. With mean–scale normalization, generated samples concentrate strictly below the diagonal. Structural metrics corroborate this observation: the proportion of sorted samples is close to 1.0. See Figure 2.2 for a side-by-side comparison.

Ablation study: $d = 2$. This case targets the top two order statistics $(X_{N:N}, X_{N-1:N})$ with $N=500$. An illustration of the generated data cloud is provided in Figure 2.2.

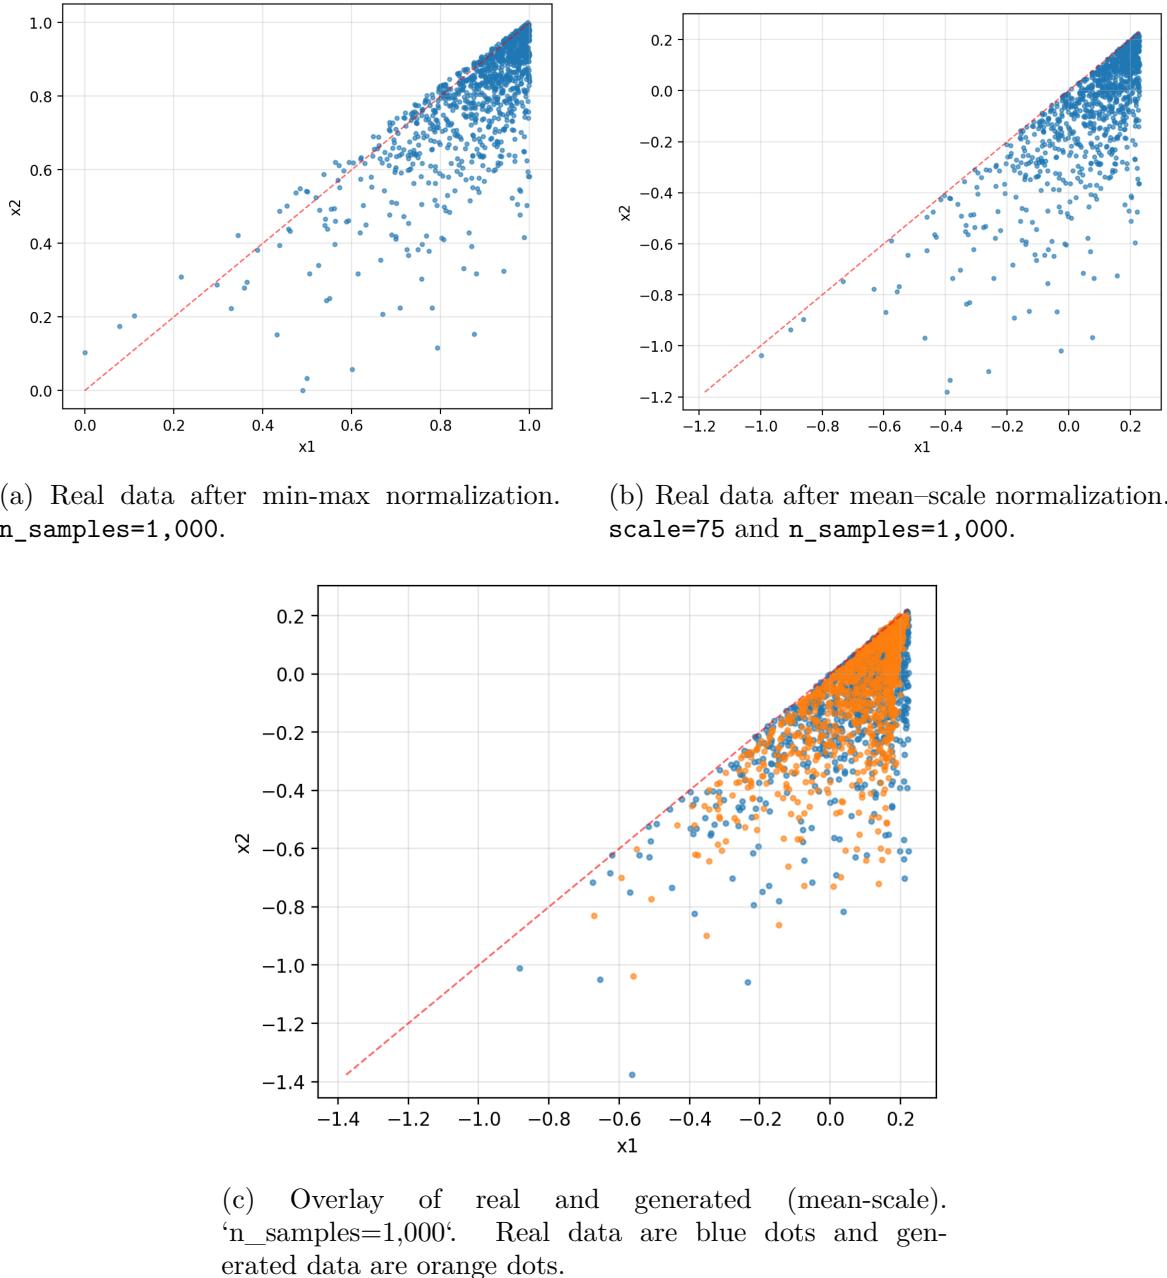


Figure 2.2: Comparison of real and generated 2D order statistics under different normalization schemes, with an additional overlay view. The red dashed lines indicate $x = y$; the mean-scale normalization preserves the sorted structure (no line crossing), while min-max normalization does not.

2303 3.B Benchmarking on synthetic data

2304 In this section, we detail our benchmarking study of GAN architectures to model order
 2305 statistics. We consider three baseline variants of GANs: vanilla GAN (I. J. Goodfellow
 2306 et al. 2014), Wasserstein-GAN (WGAN) (Arjovsky et al. 2017), and WGAN with gradient
 2307 penalty (WGAN-GP) (Wei et al. 2018). We evaluate each of the three models over different
 2308 dimension settings. Once the best model is selected, we enhance it with an auxiliary
 2309 regularization term that proxies the quality of the order-statistics property of the produced
 2310 data points. We then reevaluate the enhanced model against its corresponding base version.

2311 Evaluating the best model out of the three base models

2312 **Scope of the study.** We explore the performance of each model in different settings.
 2313 We proceed to replicate the distribution of (2.21), with $N = 1000$ and various values of k .
 2314 Namely, we first explore dimensions $d = 2, 3, 5, 10, 20$ and 50 for assessing the best model
 2315 out of the three base models. We choose to use a data set of size $n_{\text{data}} = 32,000$, fixed
 2316 across dimensions. For each variant of implementation, we perform $R = 10$ replications
 2317 and compute, for each replicated model, six metrics on generated datasets.

2318 **Metrics.** Among these six metrics, two are based on a Wasserstein criterion: the Wasser-
 2319 stein 1D (abbr. W1D) and Sliced-Wasserstein Distance (SWD). Two are general de-
 2320 pendence metrics: the Spearman’s ρ coefficient difference (Spearman) and the Absolute
 2321 Kendall Error (AKE). Finally, two are specifically designed to assess the order-statistics
 2322 property of the generated samples, the Sortedness and Softsortedness metric. Precise def-
 2323 initions and background on the metrics are detailed in Section 2.I.

2324 **Sweeping with Optuna.** For exploring the best hyperparameters, we use Optuna (Ak-
 2325 iba et al. 2019), a hyperparameter optimization framework. We use the default TPE
 2326 sampler, which models the objective function using a tree-structured Parzen estimator.
 2327 We define a grid of hyperparameters over which we perform optimization. Details on hy-
 2328 perparameters are given in Section 2.H. For each pair of model and dimension, we perform
 2329 200 trials. The target metric for TPE sampling is the W1D criteria. We provide both the
 2330 hyperparameter grid and the hyperparameters selected by the sweep in Section 2.H.

2331 **Replications.** Once the sweep done, we perform $R = 10$ replications for each pair of
 2332 model and dimension. We compute each of the six metrics for each of the replications and
 2333 average over model and dimension. Results are provided in Table 2.2. Best values are in
 2334 bold.

2335 **Results analysis.** The results of Table 2.2 provides us the following insights on the
 2336 performance of models: The WGAN-GP underperforms by a wide margin compared to
 2337 the two other models. On Wasserstein-based metrics, marginal (W1D) and multivariate
 2338 (SWD), the WGAN outperforms the GAN, except in the vanilla 2 dimensional case. This is

expected, as the minimization of the Wasserstein distance is embedded in the model's loss. Looking at classical metrics of dependence, namely Spearman and AKE, the GAN seems to outperform the WGAN, except for $d = 3$ for the Spearman metric and $d = 50$ for the AKE metric. However, looking at the specifically designed Sortedness and Softsortedness metrics, GAN performs better in low dimensional settings $d \in \{2, 3, 5, 10\}$, but WGAN is better for higher dimensional settings $d \in \{20, 50\}$. Furthermore, for higher values of d ($d \in \{20, 50\}$), performances of GAN is especially poor (values of Sortedness of 4.7% and 0% respectively), whereas values for WGAN stay higher (51.9% and 50.7% respectively). Qualitatively, the pairwise projections in Figures 2.3 and 2.4 in Section 2.A exhibit tight alignment between real and generated samples across dimensions. For this reason, we interpret WGAN to be more performant at respecting the order statistics structure in higher dimensions. We henceforth choose WGAN as the best performing model and candidate for enhancement in the latter.

	GAN	WGAN	WGAN OSP		GAN	WGAN	WGAN OSP
dim				dim			
2	0.341	0.328	0.486	2	0.336	0.332	0.475
5	0.22	0.424	0.251	5	0.204	0.396	0.251
10	0.513	0.363	0.488	10	0.473	0.356	0.478
20	1.1	0.55	0.839	20	1.03	0.526	0.815
50	2.33	1.13	1.1	50	2.16	1.1	1.02

(a) W1D			(b) SWD			
	GAN	WGAN	dim	GAN	WGAN	
dim						
2	1.53	1.18	1.17	2	94.8%	88.7%
5	0.315	0.506	0.896	5	89.2%	76.6%
10	0.452	0.989	0.93	10	56.3%	56.2%
20	0.391	0.935	1.33	20	4.0%	60.0%
50	0.713	2.02	1.65	50	0.0%	88.3%

(c) Spearman			(d) Sortedness			
	GAN	WGAN	dim	GAN	WGAN	
dim						
2	1.79	1.84	1.92	2	0.128	0.364
5	2.64	2.71	2.67	5	0.033	0.19
10	3.59	3.56	3.66	10	0.263	0.28
20	4.95	5.05	4.93	20	1.04	0.105
50	7.94	8.08	7.95	50	3.21	0.0807

(e) AKE			(f) Softsortedness			
	GAN	WGAN	dim	GAN	WGAN	
dim						
2	0.128	0.364	0.019	2	0.033	0.19
5	0.033	0.19	0.193	5	0.263	0.000692
10	0.263	0.28	0.0102	10	1.04	0.105
20	3.21	0.0807	0.0832	20	3.21	0.0807

Table 2.2: Summary of metric performances for various models across dimensions.

Order-statistics penalization

In this section, we enforce WGAN to perform better quality samples with an auxiliary regularization term, favoring order statistics samples.

Order statistics penalization (OSP). We define the following penalization term:

$$\mathcal{L}_{\text{OSP}}(G) = \mathbb{E}_{z \sim p(z)} \left[\frac{1}{d-1} \sum_{j=1}^{d-1} (\max \{G(z)_j - G(z)_{j+1} + \varepsilon, 0\})^2 \right] \quad (2.23)$$

where $G(z) \in \mathbb{R}^d$ is the generator output, $\varepsilon > 0$ is a small offset constant, and the sum is over consecutive coordinate pairs. We include this loss in the loss of the generator during training. We name this new version WGAN OSP.

Results. We use the same settings as in the previous section, with the same hyperparameter grid, and we perform a sweep with Optuna to find the best hyperparameters for WGAN OSP. Results for WGAN OSP are given along other models in Table 2.2. We observe that in regard of the Softsortedness metric, WGAN OSP outperforms WGAN in three out of the five dimensional settings. This results also in better performance for the Sortedness metric for four out of the five dimensional settings. Regarding W1D and SWD distances, WGAN outperforms for three dimensions, which is expected, as WGAN is only optimizing for the Wasserstein distance. Regarding the classical dependence metrics, namely Spearman’s ρ and Kendall’s τ , WGAN is consistently better than WGAN OSP, but GAN is the overall best model with respectively four and three dimensions in which it is best performing. Overall, this study shows that adding a Softsortedness regularization term to WGAN improves its performance on metrics related to actually producing order statistics, while still maintaining competitive results across the board.

2.4 Conclusion

This work develops a route to generative modeling of ordered statistics. Starting from classical probability identities, we showed how two representations of order statistics (Schucany’s top-down recursion and Sukhatme’s exponential spacings) can be leveraged to design neural samplers with explicit capacity/error tradeoffs. Concentrating approximation into one scalar inverse-cdf block (with Sukhatme’s representation) yields a parameter growth of order $\mathcal{O}(\varepsilon^{-2} \log(\varepsilon^{-1}))$ and a coordinate-wise expected L_1 error that grows only logarithmically with dimension through a harmonic factor. The theory clarifies what must be learned (and what does not) when the target distribution lives on the ordered manifold.

On the practical side, we argued that respecting the geometry of order statistics during training is as important as the network architecture. A simple global mean-scale normalization preserves pairwise order within each sample and proved crucial to maintain the OS structure under adversarial training. Experiments on synthetic OS benchmarks across dimensions demonstrate that standard GAN/WGAN baselines equipped with this normalization achieve competitive performance on classical metrics while improving sortedness; an additional soft order-violation penalty further tightens structure-aware metrics with small impact on Wasserstein distances.

There are several avenues for future work. First, moving from OS to induced order statistics (IOS) is essential for financial applications in which assets are ranked by an

2391 impact factor and paired with returns. This calls for learning generative IOS models
2392 with the same geometry-preserving principles. Second, a systematic, controlled empirical
2393 study that instantiates architectures mirroring the Schucany and Sukhatme constructions
2394 would enable direct validation of the predicted scalings (e.g., $\varepsilon \log N$ error growth and
2395 $\varepsilon^{-2} \log(\varepsilon^{-1})$ parameter counts). While feasible, such sweeps are tedious and we leave
2396 them to future work. Finally, applying these ideas to real impact-investing datasets will
2397 require robust preprocessing, careful evaluation protocols, and domain-specific stress tests
2398 in which IOS simulators can meaningfully inform portfolio construction.

2399 Taken together, our analysis and experiments show that a small amount of structure—
2400 codified by probabilistic representations and simple normalization—goes a long way for
2401 learning on the ordered manifold. We hope this bridges classical order-statistics theory and
2402 modern deep generative modeling, and serves as a foundation for IOS-aware simulators in
2403 sustainable finance.

2404 Appendix

2405 2.A Additional Figures

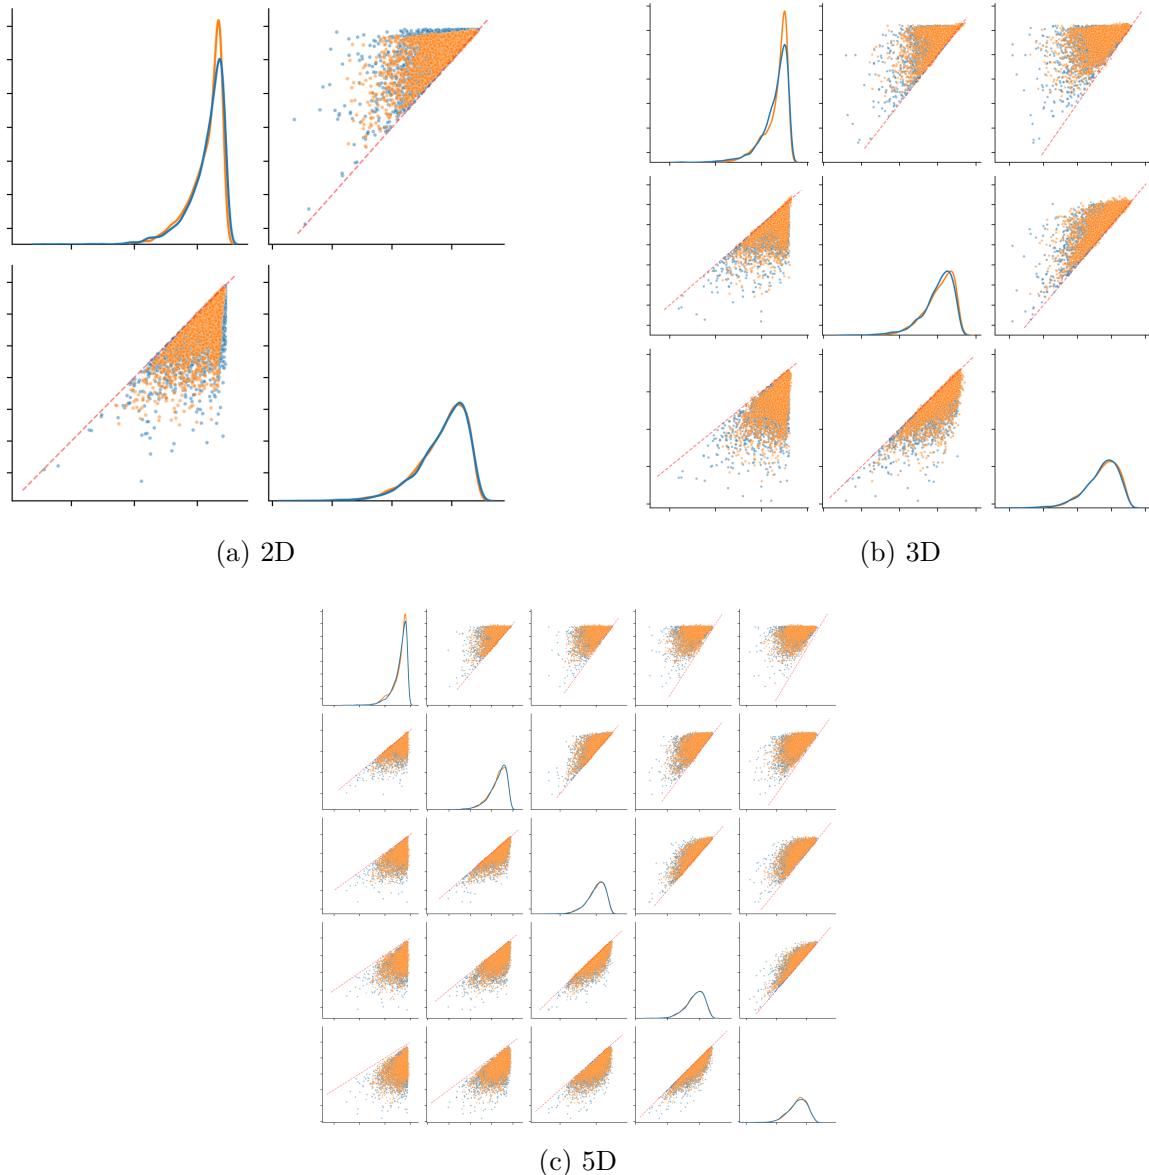


Figure 2.3: Generated vs. real means for order statistics in various dimensions: (a) 2D, (b) 3D, (c) 5D. Every dimension is scattered against one another. Real data points are in blue and generated data points are in orange.

2.B Some theory of concomitants

Representation results for order statistics

Writing order statistics as min and max. Following the results in (Banner and Ghomrasni 2008), order statistics can be written as a combination of min and max. The following results holds:

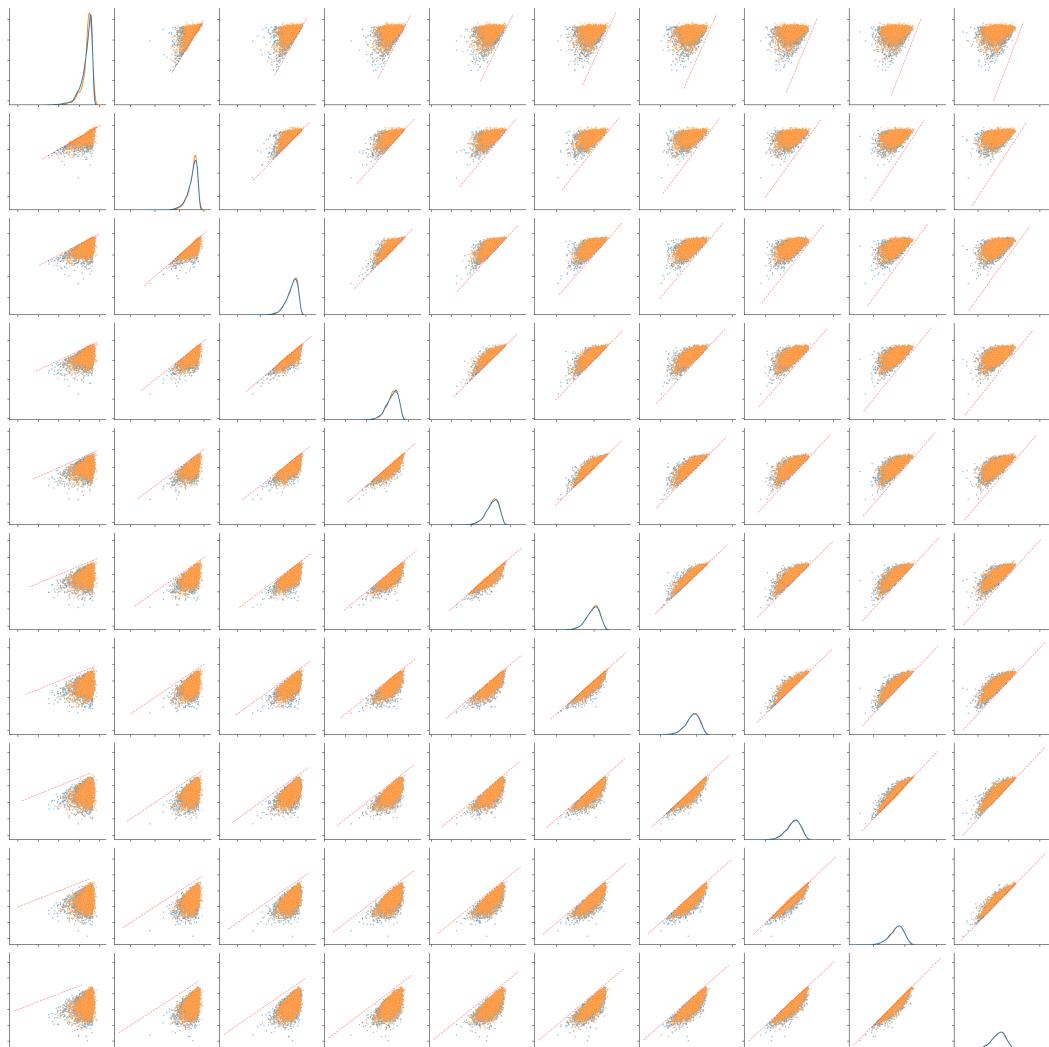


Figure 2.4: Generated vs. real means for order statistics in dimension 10. Every dimension is scattered against one another. Real data points are in blue and generated data points are in orange.

$$X_{k:N} := \max_{1 \leq i_1 < \dots < i_k \leq N} \min \{X_{i_1}, \dots, X_{i_k}\}. \quad (2.24)$$

2411 It is possible to represent min and max with neural network, considering ReLU activation function, that is $\forall x \in \mathbb{R} : \sigma(x) = (x)_+ = \max\{0, x\}$, by writing:

$$\max\{a, b\} = a + (b - a)_+, \quad (2.25)$$

$$\min\{a, b\} = -\max\{-a, -b\}. \quad (2.26)$$

2413 Moreover, it is possible to compute min and max over more than two variables by 2414 computing recursively:

$$\max\{a_1, \dots, a_N\} = \max \{\max\{a_1, \dots, a_{N-1}\}, a_N\}. \quad (2.27)$$

2415 This formula has the benefit, conversely to further representations in this section, of 2416 being fully deterministic. However, this comes at a cost: in practice, it is not reasonably 2417 feasible to perform computation of (2.24), as it involves, for the full sorting of vector 2418 $(X_{1:N}, \dots, X_{N:N})$, 2^N operations.

2419 **Order statistics as a Markov chain (Schucany form).** Order statistics admit a 2420 Markovian structure (H. David and H. Nagaraja 2004, Chapter 1, Section 5). For the 2421 uniform case, the chain is especially simple in the *top-down* direction. Conditionally on 2422 $U_{m+1:N} = x$, the distribution of $U_{m:N}$ on $[0, x]$ is the maximum of m i.i.d. uniforms on 2423 $[0, x]$, hence

$$\forall y \in [0, x] : F_{U_{m:N}|U_{m+1:N}=x}(y) = \left(\frac{y}{x}\right)^m, \quad m = 1, \dots, N-1. \quad (2.28)$$

2424 Equivalently, if $(U_m)_{m=1}^N$ are i.i.d. $\text{Unif}(0, 1)$ independent of $U_{m+1:N}$, then

$$U_{N:N} = U_N^{1/N}, \quad U_{m:N} = U_{m+1:N} U_m^{1/m} \quad (m = N-1, \dots, 1). \quad (2.29)$$

2425 This is precisely Schucany's recursive sampler for uniform order statistics and matches 2426 Equation (2.7) (identify $m = N - r$ and $U_m \equiv U_r$). Defining the bivariate update

$$G_{m:N}(u, w) \stackrel{\text{def}}{=} u w^{1/m}, \quad m \in [N], \quad (2.30)$$

2427 the chain reads $U_{m:N} = G_{m:N}(U_{m+1:N}, U_m)$ with $U_{N:N} = G_{N:N}(1, U_N)$, which is the 2428 formulation used for our neural approximations in Algorithm 3.

2429 **Using standard exponential distributions.** Another method for modeling order statistics stems from Sukhatme's representation (H. David and H. Nagaraja 2004, p.17). Let 2430 $Z_{1:N} \leq \dots \leq Z_{N:N}$ be the order statistics from the standard exponential distribution. 2431 Then the joint distribution of the $Z_{i:N}$ is :

$$N! \exp \left(- \sum_{i=1}^N z_i \right), \quad (2.31)$$

2433 which may be rewritten as (Sukhatme 1937):

$$N! \exp \left(- \sum_{i=1}^N (N-i+1)(z_{i:N} - z_{i-1:N}) \right), \quad (2.32)$$

2434 with convention $z_{0:N} = 0$. Making the transformation

$$y_i = (N-i+1)(z_{i:N} - z_{i+1:N}), \quad (2.33)$$

2435 and noting that the range of each y_i is $(0, \infty)$, we see that the Y_i are statistically
2436 independent variables, each with pdf (2.8). The relation (2.33) allow $Z_{r:N}$ to be expressed
2437 as:

$$Z_{r:N} = \sum_{i=1}^r (Z_{i:N} - Z_{i-1:N}) = \sum_{i=1}^r \frac{Y_i}{N-i+1}. \quad (2.34)$$

2438 Following the probability integral transformation, we can sample from the standard expo-
2439 nential distribution by sampling from an uniform distribution and applying the inverse cdf
2440 $F^{-1} : u \mapsto -\ln(1-u)$:

$$Z_{r:N} = -\ln(1 - U_{r:N}). \quad (2.35)$$

2441 2.C Concomitants of order statistics from a single distribution

2442 Definition

2443 Let $(X_i, Y_i)_{i=1,\dots,n}$ be a random sample from a bivariate distribution with cdf $F(x, y)$. Let
2444 us denote

$$X_{r:n} \quad (2.36)$$

2445 the r -th order statistic (where data have ordered increasingly), for any $r = 1, \dots, n$:

$$X_{1:n} < \dots < X_{r:n} < \dots < X_{n:n}. \quad (2.37)$$

2446 When the marginal cdf of X is continuous, this ordering exists and is almost surely
2447 unique.

Definition 2.1. If $X_{r:n}$ corresponds to the data X_i , its *concomitant* is denoted by

$$Y_{[r:n]}$$

2448 and corresponds to Y_i .

This notion was first introduced of (David, Herbert A 1973) and was called *Induced Order Statistics* (IOS) of (Bhattacharya 1974). See (H. David and H. Nagaraja 2004, Section 6.8) for an introduction.

An important use of concomitants corresponds to select $k < n$ objects Y according to the ranking of their X -values.

Standard properties

When $Y_i = \mathcal{R}(X_i, \varepsilon_i)$ like in a general regression model where X_i and ε_i are independent (see the discussion by H. David and H. Nagaraja (2004, p. 145)), we have

$$Y_{[r:n]} = \mathcal{R}(X_{r:n}, \varepsilon_{[r:n]}). \quad (2.38)$$

When the $(\varepsilon_i)_{1 \leq i \leq n}$ are i.i.d. then the $(\varepsilon_{[r:n]})_{1 \leq r \leq n}$ are i.i.d. too, see (Kim and H. David 1990, Section 4).

We recall the definition of associated random variables.

Definition 2.2 (Kim and H. David (1990, Definition 3.1)). The random variables $X = (X_i)_{1 \leq i \leq n}$ are *associated* if:

$$\text{Cov}(f(X), g(X)) \geq 0 \quad (2.39)$$

for all non-decreasing (or non-increasing) functions f and g for which the covariance exists.

In particular, if X is square integrable, all pairs (X_i, X_j) are non-negatively correlated.

It turns out that the $(Y_{[r:n]})_{1 \leq r \leq n}$ in (2.38) are associated under some mild conditions.

Theorem 2.3 (Kim and H. David (1990), definition 4.1). *Assume that \mathcal{R} in (2.38) is a non-decreasing (or non-increasing) function of both arguments, then $(Y_{[r:n]})_{1 \leq r \leq n}$ are associated.*

Representation theorem under the continuity assumption

Getting the function \mathcal{R} can be achieved through the Rosenblatt transformation (Rosenblatt 1952). (Lo and R. Zhang 2021, Theorem 1) gives an explicit formula for the \mathcal{R} . Their representation theorem formula goes beyond (2.38) by representing all the concomitants using order statistics of uniformly distributed random variables.

Theorem 2.4. *If X and Y have continuous cdf F_X and F_Y , with a copula C having a density, we have*

$$(Y_{[1:n]}, \dots, Y_{[r:n]}, \dots, Y_{[n:n]}) \stackrel{d}{=} (g(U_{1:n}, V_1), \dots, g(U_{r:n}, V_r), \dots, g(U_{n:n}, V_n)) \quad (2.40)$$

2476 where $U_1, \dots, U_n, V_1, \dots, V_n$ are independent random variables uniformly distributed on
2477 $[0, 1]$. The function g is given explicitly by the formula

$$g(u, v) = F_Y^{-1}([\partial_u C(u, \cdot)^{-1}(v)]). \quad (2.41)$$

2478 Note that in the above theorem, C is uniquely defined since the marginal cdf are
2479 continuous.

2480 **Representation theorem without the continuity assumption**

2481 We now state a variant avoiding the continuity assumption of the cdf of X and Y . We
2482 follow the arguments of (Barrera et al. 2025) to well define the quantity under study.

2483 We denote by $\mathcal{B}_{\mathbb{R}}$ the Borel sigma algebra on \mathbb{R} . We fix a conditional distribution
2484 function $\mu : \mathcal{B}_{\mathbb{R}} \times \mathcal{B}_{\mathbb{R}} \rightarrow [0, 1]$ of Y given X (Kallenberg 2002, theorem 5.3, p.84), and we
2485 assume that the function $\mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ defined by $(x, y) \mapsto \mu(x, (-\infty, y])$ is $(\mathcal{B}_{\mathbb{R}} \otimes \mathcal{B}_{\mathbb{R}})/\mathcal{B}_{\mathbb{R}}$
2486 (i.e. Borel)-measurable. With these conventions, we will use implicitly the corresponding
2487 version

$$\mathbb{P}Y \in \cdot | X \stackrel{\text{def}}{=} \mu(X, \cdot). \quad (2.42)$$

2488 The conditional quantile of Y given $X = x$ is defined by

$$q_{Y|X}(x, v) \stackrel{\text{def}}{=} \inf\{y : \mu(x, (-\infty, y]) \geq v\}, \quad v \in [0, 1]. \quad (2.43)$$

2489 **Lemma 2.1.** *The function $(x, v) \in \mathbb{R} \times [0, 1] \mapsto q_{Y|X}(x, v) \in \mathbb{R}$ is $(\mathcal{B}_{\mathbb{R}} \otimes \mathcal{B}_{[0,1]})$ -measurable.*

2490 *Proof.* Given $t \in \mathbb{R}$, since

$$\{q_{Y|X}(x, v) \leq t\} = \{\mu(x, (-\infty, t]) \geq v\}, \quad (2.44)$$

2491 we get the result. □

2492 **Theorem 2.5.** *Denote by F_X and F_Y the respective cdf of X and Y (they can be discontinuous). Then*

$$(Y_{[1:n]}, \dots, Y_{[r:n]}, \dots, Y_{[n:n]}) \quad (2.45)$$

$$\stackrel{d}{=} (q_{Y|X}(F_X^{-1}(U_{1:n}), V_1), \dots, q_{Y|X}(F_X^{-1}(U_{r:n}), V_r), \dots, q_{Y|X}(F_X^{-1}(U_{n:n}), V_n)) \quad (2.46)$$

2494 where $U_1, \dots, U_n, V_1, \dots, V_n$ are independent random variables uniformly distributed on
2495 $[0, 1]$.

2496 *Proof.* The Rosenblatt transformation writes

$$(X_i, Y_i : 1 \leq i \leq n) \stackrel{d}{=} (F_X^{-1}(U_i), q_{Y|X}(X_i, V_i) : 1 \leq i \leq n). \quad (2.47)$$

²⁴⁹⁷ Because the quantile F_X^{-1} is increasing,

$$(X_{r:n}, Y_{[r,n]} : 1 \leq r \leq n) \stackrel{d}{=} (F_X^{-1}(U_{r:n}), q_{Y|X}(X_{r:n}, V_{[r,n]}) : 1 \leq r \leq n) \quad (2.48)$$

$$\stackrel{d}{=} (F_X^{-1}(U_{r:n}), q_{Y|X}(F_X^{-1}(U_{r:n}), V_{[r,n]}) : 1 \leq r \leq n) \quad (2.49)$$

²⁴⁹⁸ Since the V_i 's are i.i.d. we have

$$(V_{[r,n]} : 1 \leq r \leq n) \stackrel{d}{=} (V_r : 1 \leq r \leq n). \quad (2.50)$$

²⁴⁹⁹ In addition, since they are independent from the U_i 's, we get

$$(X_{r:n}, Y_{[r,n]} : 1 \leq r \leq n) \stackrel{d}{=} (F_X^{-1}(U_{r:n}), q_{Y|X}(F_X^{-1}(U_{r:n}), V_r) : 1 \leq r \leq n). \quad (2.51)$$

²⁵⁰⁰

□

²⁵⁰¹ **The case of independent (X_i, Y_i) but with different cdfs $F_i(x, y)$**

²⁵⁰² In that case, the Rosenblatt transformation write

$$(X_i, Y_i : 1 \leq i \leq n) \stackrel{d}{=} (F_{X_i}^{-1}(U_i), q_{Y_i|X_i}(X_i, V_i) : 1 \leq i \leq n). \quad (2.52)$$

²⁵⁰³ **The intermediate case of independent (X_i, Y_i) but with same marginal X_i**

²⁵⁰⁴ In that case, the Rosenblatt transformation writes

$$(X_i, Y_i : 1 \leq i \leq n) \stackrel{d}{=} (F_X^{-1}(U_i), q_{Y_i|X_i}(X_i, V_i) : 1 \leq i \leq n). \quad (2.53)$$

²⁵⁰⁵ **2.D Some elements on Functional approximation with Neural Network**

²⁵⁰⁷ Section 2.B provides different computational methodologies for obtaining order statistics.
²⁵⁰⁸ A natural question is how well a Neural Network-based method can approximate the
²⁵⁰⁹ obtained functions. For quantifying this, we rely on results from the theory of function
²⁵¹⁰ approximation with Neural Networks. (Yarotsky 2017) provides several results on how well
²⁵¹¹ neural networks are able to approximate smooth functions.

²⁵¹² **Lemma 2.2** (Yarotsky (2017), Proposition 3). *Given $M > 0$ and $\varepsilon \in (0, 1)$, there exists
²⁵¹³ a ReLU network η with two input units that implements a function $\tilde{x} : \mathbb{R}^2 \rightarrow \mathbb{R}$ so that:*

- ²⁵¹⁴ • for any input x, y , if $|x| < M$ and $|y| < M$, then $|\tilde{x}(x, y) - xy| \leq \varepsilon$
- ²⁵¹⁵ • if $x = 0$ or $y = 0$ then $\tilde{x}(x, y) = 0$
- ²⁵¹⁶ • the depth and the number of weights and computation units in η is no greater than
²⁵¹⁷ $c_1 \ln(1/\varepsilon) + c_2$ with an absolut constant c_1 and a constant $c_2 = c_2(M)$.

The second proposition gives a result on approximating Sobolev functions with neural network. For a function with sufficient conditions, the Sobolev norm is defined as:

$$\|f\|_{\mathcal{W}^{n,\infty}([0,1]^d)} = \max_{\mathbf{n}:|\mathbf{n}| \leq n} \text{ess}_{\mathbf{x} \in [0,1]^d} |D^{\mathbf{n}} f(\mathbf{x})| \quad (2.54)$$

and its associated unit-norm class is given by:

$$F_{n,d} = \left\{ f \in \mathcal{W}^{n,\infty}([0,1]^d) : \|f\|_{\mathcal{W}^{n,\infty}([0,1]^d)} \leq 1 \right\}. \quad (2.55)$$

The core result of (Yarotsky 2017) is the following:

Theorem 2.6 (Yarotsky (2017), Theorem 1). *For any d, n integers and $\varepsilon \in (0, 1)$, there exists a ReLU network architecture that:*

- *is capable of approximating any function from $F_{d,n}$ with error ε ,*
- *has depth at most $c \ln(1/\varepsilon + 1)$ and at most $c\varepsilon^{-d/n}(\ln(1/\varepsilon) + 1)$ weights and computation units with some constant $c = c(d, n)$.*

2.E Proofs

Proof of Theorem 2.1

Let's investigate the error induced by the sampling procedure of Algorithm 3. Let $r \in [N]$. We are looking to bound the quantity $\Delta V_r = \hat{U}_{N-r:N} - U_{N-r:N}$.

$$\Delta V_0 = 0, \quad (2.56)$$

$$\Delta V_r = \hat{U}_{N-r:N} - U_{N-r:N} = \hat{G}_{N-r:N}^{NN} (\hat{U}_{N-r+1:N}, U_r) - G_{N-r:N} (U_{N-r+1:N}, U_r) \quad (2.57)$$

$$= [\hat{G}_{N-r:N}^{NN} (\hat{U}_{N-r+1:N}, U_r) - G_{N-r:N} (\hat{U}_{N-r+1:N}, U_r)] \quad (2.58)$$

$$+ [G_{N-r:N} (\hat{U}_{N-r+1:N}, U_r) - G_{N-r+1:N} (U_{N-r+1:N}, U_r)]. \quad (2.59)$$

On the one side:

$$G_{N-r:N} (\hat{U}_{N-r+1:N}, U_r) - G_{N-r:N} (U_{N-r+1:N}, U_r) \quad (2.60)$$

$$= (\hat{U}_{N-r+1:N} - U_{N-r+1:N}) U_r^{1/(N-r)} \quad (2.61)$$

$$= \Delta V_{r-1} U_r^{1/(N-r)}. \quad (2.62)$$

On the other:

$$\left| \hat{G}_{N-r:N}^{NN}(\hat{U}_{N-r+1:N}, U_r) - G_{N-r:N}(\hat{U}_{N-r+1:N}, U_r) \right| \quad (2.63)$$

$$\leq \varepsilon_r := \max_{u_1, u_2 \in [0,1]} \left| \hat{G}_{N-r:N}^{NN}(u_1, u_2) - G_{N-r:N}(u_1, u_2) \right|. \quad (2.64)$$

2533 Thus for $1 \leq r \leq N$:

$$|\Delta V_r| \leq |\Delta V_{r-1}| U_r^{1/(N-r)} + \varepsilon_r. \quad (2.65)$$

Lemma 2.3.

$$\text{For } 1 \leq r \leq N : |\Delta V_r| \leq \sum_{i=0}^r \varepsilon_i \prod_{j=i+1}^r U_j^{1/(N-j)}, \quad (2.66)$$

2534 with convention $\varepsilon_0 = 0$.

2535 *Proof.* The proof goes by induction:

2536 • $r = 0$: (2.65) writes:

$$|\Delta V_1| \leq |\Delta V_0| U_1^{1/N} + \varepsilon_1 = \varepsilon_1, \quad (2.67)$$

2537 as $\Delta V_0 = 0$, which is the announced result.

2538 • Suppose that (2.66) holds true for $0 \leq r \leq N-1$. Then:

$$|\Delta V_{r+1}| \leq |\Delta V_r| U_{r+1}^{1/(N-r-1)} + \varepsilon_{r+1} \quad (2.68)$$

$$\leq \left(\sum_{i=0}^r \varepsilon_i \prod_{j=i+1}^r U_j^{1/(N-j)} \right) U_{r+1}^{1/(N-r-1)} + \varepsilon_{r+1} \quad (2.69)$$

$$= \sum_{i=0}^{r+1} \varepsilon_i \prod_{j=i+1}^{r+1} U_j^{1/(N-j)}. \quad (2.70)$$

2539 The result is thus proved. □

2540

2541 A direct computation yields $\mathbb{E}[U^{1/(N-j)}] = \frac{N-j}{N-j+1}$ and thus

$$\mathbb{E} |\Delta V_r| \leq \sum_{i=0}^r \varepsilon_i \prod_{j=i+1}^r \frac{N-j}{N-j+1} \quad (2.71)$$

$$= \sum_{i=0}^r \varepsilon_i \frac{N-r}{N-i}. \quad (2.72)$$

2542 Let's investigate ε_i . We are looking to the best approximation of the function $G_{N-r:N}(u_1, u_2) =$
2543 $u_1 \cdot u_2^{1/(N-r)}$. On the one side, the function

$$\phi_r : u \in [0, 1] \mapsto u^{1/(N-r)} \quad (2.73)$$

2544 is $\zeta_r := 1/(N-r)$ Holder continuous, (the space of Holder continuous functions writes
2545 $C^{0,\zeta_r}([0, 1]))$, with Holder constant 1.

2546 Define the ReLU function as $\sigma^R(x) := \max(0, x)$. Then denote the triangular function
 2547 by:

$$\widehat{\sigma}^R(x) := \sigma^R(x+1) - 2\sigma^R(x) + \sigma^R(x-1) = \begin{cases} 1+x, & \text{if } -1 \leq x \leq 0, \\ 1-x, & \text{if } 0 \leq x \leq 1, \\ 0, & \text{otherwise.} \end{cases}$$

2548 Note that the triangular function can be approximated by a ReLU Neural Network with
 2549 three neurons, see (Michaël Allouche, Girard, et al. 2025). We also have the following
 2550 result:

2551 **Lemma 2.4** (Michaël Allouche, Girard, et al. (2025), Lemma 8). *Let $\widehat{\sigma}^R$ be a triangular
 2552 function, $\tau \in (0, 1)$ and $f \in C^{0,\xi}([0, \tau])$ with $\xi \in (0, 1]$. For all $M \in \mathbb{N} \setminus \{0\}$, let $h = \tau/M$
 2553 and $u_j = jh$ for $j \in \{0, \dots, M\}$. Then,*

$$\sup_{u \in [0, \tau]} \left| f(u) - \sum_{j=0}^M f(u_j) \widehat{\sigma}^R \left(\frac{u - u_j}{h} \right) \right| \leq H_f \left(\frac{\tau}{M} \right)^\xi,$$

2554 where H_f is the Hölder constant associated with f . This construction involves $3(M+1)$
 2555 neurons.

2556 This result thus applies to functions ϕ_r , with $\tau = 1$ and we denote the approximation
 2557 of ϕ_r by $\widehat{\phi}_r$. To obtain the minimal number of neuron sufficient for a given precision, let's
 2558 write the error ε :

$$\varepsilon \leq (1/M)^{\zeta_r} \Rightarrow M \geq 1/\varepsilon^{N-r} \Rightarrow M = \lceil 1/\varepsilon^{N-r} \rceil. \quad (2.74)$$

2559 Thus a neural network with

$$3(M+1) = 3(\lceil 1/\varepsilon^{N-r} \rceil + 1) \quad (2.75)$$

2560 neurons is able to approximate the function ϕ_r with precision ε and a neural network with
 2561 $\lceil (2/\varepsilon)^{N-r} \rceil + 1$ neurons is able to approximate the function ϕ_r with precision $\varepsilon/2$. Now,
 2562 for the multiplicative term, we invoke Lemma 2.2, which writes as follows in this context:

2563 A neural network η_r with $c_1 \ln(8/\varepsilon) + c_2$ weights with an absolute constant c_1 and a
 2564 constant $c_2 = c_2(M)$, with $M = 1$ here, is able to learn $\tilde{x}_r(x, v) = x \cdot v$ with precision
 2565 $\varepsilon/8$.

2566 Now consider

$$\hat{G}_{N-r:N}^{NN}(u_1, u_2) = \tilde{x}_r(u_1, \phi_r(u_2)). \quad (2.76)$$

2567 Then:

$$\left| \hat{G}_{N-r:N}^{NN}(u_1, u_2) - G_{N-r:N}(u_1, u_2) \right| = \left| u_1 \cdot u_2^{1/(N-r)} - \tilde{x}_r(u_1, \phi_r(u_2)) \right| \quad (2.77)$$

$$\leq \left| \tilde{x}_r(u_1, \phi_r(u_2)) - \tilde{x}_r(u_1, u_2^{1/(N-r)}) \right| \quad (2.78)$$

$$+ \left| \tilde{x}_r(u_1, u_2^{1/(N-r)}) - u_1 \cdot u_2^{1/(N-r)} \right|. \quad (2.79)$$

2568 The second term is $\leq \varepsilon/8$ by construction of \tilde{x}_r . Let's look at the first:

$$\left| \tilde{x}_r(u_1, \phi_r(u_2)) - \tilde{x}_r(u_1, u_2^{1/(N-r)}) \right| \leq \left| \tilde{x}_r(u_1, \phi_r(u_2)) - u_1 \cdot \phi_r(u_2) \right| \quad (2.80)$$

$$+ \left| u_1 \phi_r(u_2) - u_1 \cdot u_2^{1/(N-r)} \right| \quad (2.81)$$

$$+ \left| u_1 \cdot u_2^{1/(N-r)} - \tilde{x}_r(u_1, u_2^{1/(N-r)}) \right| \quad (2.82)$$

$$\leq 3\varepsilon/4, \quad (2.83)$$

2569 where (2.80), (2.82) $\leq \varepsilon/8$ by construction of \tilde{x} , (2.81) $\leq \varepsilon/2$ by construction of ϕ_r
2570 and because $0 \leq x \leq 1$. Together with (2.78):

$$\forall u_1, u_2 \in [0, 1] : \left| \hat{G}_{N-r:N}^{NN}(u_1, u_2) - G_{N-r:N}(u_1, u_2) \right| \leq 3\varepsilon/4 + \varepsilon/8 = 7\varepsilon/8 < \varepsilon. \quad (2.84)$$

2571 **Lemma 2.5.** Let $\varepsilon > 0$, $N \in \mathbb{N}^*$ and $r \in [N]$. A feedforward ReLU Neural Network
2572 $\hat{G}_{N-r:N}^{NN}$ with:

$$3(\lceil (2/\varepsilon)^{N-r} \rceil + 1) + c_1 \ln(8/\varepsilon) + c_2 \quad (2.85)$$

2573 neurons is able to approximate $G_{N-r:N}$ with precision ε .

2574 Back to (2.72): choosing for $i \in [r]$, $\varepsilon > 0$, $\varepsilon_i = \frac{(N-r)}{r(N-i)}\varepsilon$, we get $\mathbb{E}|\Delta V_r| \leq \varepsilon$ and thus
2575 we obtain Theorem 2.1 by summing over i .

2. F Proof of Theorem 2.2

2577 Let's investigate the quality of approximation of the scheme in Algorithm 4. For a starter,
2578 let's look into the quality of approximation of the function $F^{-1} : u \mapsto -\ln(1-u)$ with a
2579 neural network. For this sake, consider $\varepsilon \in [0, 1]$ and introduce the function :

$$\tilde{F}_\varepsilon^{-1} : \begin{cases} -\ln(1-u) & \text{if } u \in [0, 1-\varepsilon], \\ F^{-1}(1-\varepsilon) = -\ln(\varepsilon) = \ln(1/\varepsilon) & \text{if } u \in [1-\varepsilon, 1]. \end{cases} \quad (2.86)$$

2580 The function $\varepsilon \cdot \tilde{F}_\varepsilon^{-1}$ is in $F_{1,1}$. Indeed, $\forall u \in [0, 1]$:

$$|\varepsilon \tilde{F}_\varepsilon^{-1}(u)| \leq \varepsilon \tilde{F}_\varepsilon^{-1}(1-\varepsilon) = \varepsilon \ln(1/\varepsilon) \leq \varepsilon \left(\frac{1}{\varepsilon} - 1 \right) = 1 - \varepsilon, \quad (2.87)$$

$$|(\varepsilon \tilde{F}_\varepsilon^{-1})'(u)| = |\varepsilon \tilde{F}'_\varepsilon(u)| I\{u \in [0, 1-\varepsilon]\} = \varepsilon/(1-u) I\{u \in [0, 1-\varepsilon]\} \leq 1. \quad (2.88)$$

Therefore, the assumptions of Theorem 2.6 are met. We apply it to $\varepsilon \tilde{F}_\varepsilon^{-1}$ with precision ε^2 , which writes:

There exists a ReLU network η_{NN} that:

- is capable of approximating $\varepsilon \tilde{F}_\varepsilon^{-1}$ with error ε^2 , i.e.:

$$\max_{u \in [0, 1-\varepsilon]} |\varepsilon \tilde{F}_\varepsilon^{-1}(u) - \eta_{NN}(u)| \leq \varepsilon^2 \quad (2.89)$$

- has depth at most $c(\ln(1/\varepsilon^2 + 1))$ and at most

$$c(\varepsilon^2)^{-1} (\ln(1/\varepsilon^2) + 1) = c\varepsilon^{-2}(2\ln(1/\varepsilon) + 1) \quad (2.90)$$

weights and computation units with some constant $c = c(1, 1)$.

Thus, multiplying the obtained neural network by $1/\varepsilon$, we get that:

$$\max_{u \in [0, 1-\varepsilon]} |\tilde{F}_\varepsilon^{-1}(u) - 1/\varepsilon \cdot \eta_{NN}(u)| \leq \varepsilon \quad (2.91)$$

And thus the following result holds:

There exists a ReLU network architecture η_{NN} (we reuse the notation η_{NN}) such that:

- is capable of approximating $\tilde{F}_\varepsilon^{-1}$ with error ε ,
- has depth at most $c(\ln(1/\varepsilon^2) + 1)$ and at most

$$c(\varepsilon^2)^{-1} (\ln(1/\varepsilon^2) + 1) = c\varepsilon^{-2}(2\ln(1/\varepsilon) + 1) \quad (2.92)$$

weights and computation units with some constant $c = c(1, 1)$.

Note that the Neural Network η_{NN} is defined on $[0, 1 - \varepsilon]$. We artificially set its value on the interval $[1 - \varepsilon, 1]$ to $\eta_{NN}(u) = F^{-1}(1 - \varepsilon) = -\ln(\varepsilon)$. Now consider a standard uniform variable U . We investigate the mean error between $Z = F^{-1}(U)$ (which is $\mathcal{E}(1)$) and its approximation $\hat{Z} = \eta_{NN}(U)$. We have that:

$$\mathbb{E}_U |Z - \hat{Z}| = \int_0^1 |F_Z^{-1}(u) - \eta_{NN}(u)| du \quad (2.93)$$

$$= \int_0^{1-\varepsilon} |F_Z^{-1}(u) - \eta_{NN}(u)| du + \int_{1-\varepsilon}^1 |F_Z^{-1}(u) - \eta_{NN}(u)| du. \quad (2.94)$$

First the LHS of (2.94):

$$\int_0^{1-\varepsilon} |F_Z^{-1}(u) - \eta_{NN}(u)| du \leq \int_0^{1-\varepsilon} \|F_Z^{-1} - \eta_{NN}\|_\infty du \leq \varepsilon, \quad (2.95)$$

2598 by construction of η_{NN} and Theorem 2.6 and second the RHS of (2.94):

$$\int_{1-\varepsilon}^1 |F_Z^{-1}(u) - \eta_{NN}(u)| du = \int_{1-\varepsilon}^1 |- \ln(1-u) + \ln(\varepsilon)| du \quad (2.96)$$

$$= \int_{1-\varepsilon}^1 [\ln(\varepsilon) - \ln(1-u)] du \quad (2.97)$$

$$= \varepsilon \ln(\varepsilon) - \int_{1-\varepsilon}^1 \ln(1-u) du = \varepsilon, \text{ after integration.} \quad (2.98)$$

2599 Together,

$$\mathbb{E}_U |Z - \hat{Z}| \leq 2\varepsilon. \quad (2.99)$$

2600 Therefore,

$$\mathbb{E} |Z_{r:N} - \hat{Z}_{r:N}| = \mathbb{E} \left| \sum_{i=1}^r \frac{Y_i}{N-r+1} - \sum_{i=1}^r \frac{\hat{Y}_i}{N-r+1} \right| \quad (2.100)$$

$$\leq \sum_{i=1}^r \frac{\mathbb{E}|Y_i - \hat{Y}_i|}{N-r+1} \leq 2\varepsilon \sum_{i=1}^r 1/(N-r+1). \quad (2.101)$$

2601 We now consider approximation of the function F . As for the inverse \tilde{F}_ε , we will construct a suitable function to which we can apply Yarotsky's result. For F , the construction 2602 is a bit more subtle: Let $\varepsilon \in [0, 1]$. We consider its restriction to F to $[0, \ln(1/\varepsilon)]$, and 2603 therefore define the function \tilde{F}_ε as follows:

$$\text{for } x \geq 0 : \tilde{F}_\varepsilon(x) = \frac{F(\ln(1/\varepsilon)x)}{\ln(1/\varepsilon)}. \quad (2.102)$$

2605 The function \tilde{F}_ε is in $F_{1,1}$. Indeed $\tilde{F}_\varepsilon(x) \in [0, 1]$ and $\tilde{F}_\varepsilon(x) = e^{-\ln(1/\varepsilon)x} \in [0, 1]$. 2606 Therefore, Theorem 2.6 applies to \tilde{F}_ε . We apply it to F_ε with error $\varepsilon/\ln(1/\varepsilon)$, which 2607 writes out:

2608 There exists a ReLU Neural Network η_{NN} (we reuse the notation η_{NN}) such that:

- 2609 • is capable of approximating \tilde{F}_ε with error $\varepsilon/\ln(1/\varepsilon)$,
- 2610 • has depth at most $c \ln\left(\frac{\ln(1/\varepsilon)}{\varepsilon} + 1\right)$ and at most

$$c(\varepsilon/\ln(\varepsilon))^{-1} \left(\ln\left(\frac{\ln(1/\varepsilon)}{\varepsilon}\right) + 1 \right) = c\varepsilon^{-1} \ln(1/\varepsilon) (\ln(1/\varepsilon) + \ln(\ln(1/\varepsilon)) + 1) \quad (2.103)$$

2611 weights and computational units, for some constant $c = c(1, 1)$.

2612 Finally, doing the converse operation of rescaling on the neural network η_{NN} , that is 2613 considering (we reuse the notation η_{NN}):

$$\eta_{NN}(x) = \ln(1/\varepsilon) \eta_{NN}^{\text{old}} \left(\frac{x}{\ln(1/\varepsilon)} \right) \quad (2.104)$$

2614, we get that, $\forall x > 0$:

$$|F(x) - \ln(1/\varepsilon)\eta_{NN}^{\text{old}}(x/\ln(1/\varepsilon))| = |F(x) - \eta_{NN}(x)| \leq \varepsilon. \quad (2.105)$$

2615 Now, noting $U_{r:N} = F(Z_{r:N})$ and $\hat{U}_{r:N} = \eta_{NN}(\hat{Z}_{r:N})$:

$$|U_{r:N} - \hat{U}_{r:N}| = |F(Z_{r:N}) - \eta_{NN}(\hat{Z}_{r:N})| \leq |\eta_{NN}(\hat{Z}_{r:N}) - F(\hat{Z}_{r:N})| + |F(\hat{Z}_{r:N}) - F(Z_{r:N})| \quad (2.106)$$

$$\leq \varepsilon + |\hat{Z}_{r:N} - Z_{r:N}| \quad (2.107)$$

2616 where the LHS holds in virtue of (2.105) and the RHS because F is a 1-Lipschitz function.

2617 Taking expectations on both sides:

$$\mathbb{E}|U_{r:N} - \hat{U}_{r:N}| \leq \varepsilon + \mathbb{E}|Z_{r:N} - \hat{Z}_{r:N}| \leq \varepsilon + 2\varepsilon \sum_{i=1}^r \frac{1}{N-r+1} \quad (2.108)$$

$$= \varepsilon \left[1 + 2 \sum_{i=1}^r \frac{1}{N-r+1} \right]. \quad (2.109)$$

2618 The total number of neurons involved for the computation of $\hat{U}_{r:N}$ with this precision
2619 is :

$$c\varepsilon^{-2}(2\ln(1/\varepsilon) + 1) + c\varepsilon^{-1}\ln(1/\varepsilon)(\ln(1/\varepsilon) + \ln(\ln(1/\varepsilon)) + 1), \quad (2.110)$$

2620 which achieves the proof of Theorem 2.2.

2621 2.G Additional experimental details

2622 More background on GANs

2623 This appendix summarizes the adversarial training objectives used in our experiments and
2624 highlights practical details for stability.

2625 **GAN objective.** Given a data distribution p_{data} and a latent prior $p_{\mathbf{Z}}$, the vanilla GAN
2626 objective (see Equation (1.12)) defines a two-player minimax game between a discriminator
2627 \mathcal{D} and a generator \mathcal{G} . In practice, the non-saturating generator loss is often used to obtain
2628 stronger gradients early in training:

$$\min_{\mathcal{G}} \mathbb{E}_{\mathbf{Z} \sim p_{\mathbf{Z}}} [-\log \mathcal{D}(\mathcal{G}(\mathbf{Z}))]. \quad (2.111)$$

2629 At equilibrium with idealized capacity, the generator recovers p_{data} and the discriminator
2630 outputs 1/2 (I. J. Goodfellow et al. 2014). Training instabilities and mode collapse are
2631 common; Section Section 2.H records our simple, robust choices.

2632 **Wasserstein GAN (WGAN).** The WGAN replaces the Jensen–Shannon divergence
2633 implicit in the vanilla objective with the Wasserstein-1 distance. Using the Kantorovich–
2634 Rubinstein duality, the critic \mathcal{D} is constrained to be 1-Lipschitz and optimized via

$$\max_{\|\mathcal{D}\|_{\text{Lip}} \leq 1} \mathbb{E}_{\mathbf{X} \sim p_{\text{data}}} [\mathcal{D}(\mathbf{X})] - \mathbb{E}_{\mathbf{Z} \sim p_{\mathbf{Z}}} [\mathcal{D}(\mathcal{G}(\mathbf{Z}))], \quad \min_{\mathcal{G}} -\mathbb{E}_{\mathbf{Z} \sim p_{\mathbf{Z}}} [\mathcal{D}(\mathcal{G}(\mathbf{Z}))]. \quad (2.112)$$

2635 The resulting gradients are more informative when supports are disjoint, improving conver-
2636 gence (Arjovsky et al. 2017). The original implementation enforced the Lipschitz constraint
2637 by weight clipping, which can harm capacity and lead to optimization issues.

2638 **WGAN with Gradient Penalty (WGAN-GP).** To avoid weight clipping, Wei et al.
2639 (2018) propose penalizing the deviation of the gradient norm from 1 along lines interpolat-
2640 ing between real and generated samples. Writing $\tilde{\mathbf{X}} = \mathcal{G}(\mathbf{Z})$, $\hat{\mathbf{X}} = \epsilon \mathbf{X} + (1 - \epsilon) \tilde{\mathbf{X}}$ with
2641 $\epsilon \sim \mathcal{U}[0, 1]$, the critic maximizes

$$\mathbb{E}_{\mathbf{X}} [\mathcal{D}(\mathbf{X})] - \mathbb{E}_{\mathbf{Z}} [\mathcal{D}(\mathcal{G}(\mathbf{Z}))] - \lambda \mathbb{E}_{\hat{\mathbf{X}}} \left(\|\nabla_{\hat{\mathbf{X}}} \mathcal{D}(\hat{\mathbf{X}})\|_2 - 1 \right)^2. \quad (2.113)$$

2642 This penalty enforces the 1-Lipschitz property more effectively and stabilizes training across
2643 architectures and datasets.

2644 2.H Hyperparameters

2645 - Update ratio: several critic steps per generator step (e.g., $n_D=5$, $n_G=1$). - Optimizers:
2646 Adam with small learning rate and $\beta_1 \in [0, 0.5]$ works well for adversarial training. - Reg-
2647 ularization: gradient penalty (WGAN-GP) improves stability; spectral normalization is an
2648 alternative 1-Lipschitz control. - Activations: LeakyReLU for critic; generator activations
2649 as in Section Section 2.H. - Evaluation: rely on distributional metrics (e.g., W1D, SWD)
2650 and structure-aware scores (e.g., sortedness) discussed in Section Section 2.I.

2651 See Appendix Section 2.G for a concise primer on GANs, WGAN, and WGAN-GP,
2652 including objectives, constraints, and training tips.

2653 2.I Metrics

2654 We evaluate generative fidelity and structure with six metrics. Throughout, let real samples
2655 be $\{\mathbf{x}^{(n)}\}_{n=1}^N \subset \mathbb{R}^d$ and generated samples $\{\mathbf{y}^{(n)}\}_{n=1}^N \subset \mathbb{R}^d$. Unless stated otherwise, we
2656 adopt a *descending* order convention when focusing on the largest order statistics.

2657 **Wasserstein 1D (W1D).** For each coordinate $j \in [d]$, we compute the empirical 1-
2658 Wasserstein distance between the marginal distributions by matching order statistics. Let
2659 $x_j^{n:N}$ and $y_j^{n:N}$ denote the n -th order statistics (ascending) of $\{x_j^{(n)}\}_{n=1}^N$ and $\{y_j^{(n)}\}_{n=1}^N$,

2660 respectively. The reported score is the average over coordinates:

$$\widehat{\text{W1D}} = \frac{1}{Nd} \sum_{j=1}^d \sum_{n=1}^N |x_j^{n:N} - y_j^{n:N}|. \quad (2.114)$$

2661 This equals the empirical 1D Wasserstein distance per marginal, averaged across dimensions; it is invariant to permutations within each marginal sample.

2663 **Sliced-Wasserstein Distance (SWD).** We approximate the 1-Wasserstein distance
2664 between the real and generated joint distributions via random 1D projections. Draw unit
2665 vectors $\mathbf{v}_\ell \in \mathbb{S}^{d-1}$, $\ell = 1, \dots, L$, project $u_\ell^{(n)} = \langle \mathbf{x}^{(n)}, \mathbf{v}_\ell \rangle$ and $v_\ell^{(n)} = \langle \mathbf{y}^{(n)}, \mathbf{v}_\ell \rangle$, sort each
2666 set to obtain order statistics $u_\ell^{n:N}$ and $v_\ell^{n:N}$, and average the 1D distances:

$$\widehat{\text{SWD}} = \frac{1}{LN} \sum_{\ell=1}^L \sum_{n=1}^N |u_\ell^{n:N} - v_\ell^{n:N}|. \quad (2.115)$$

2667 As L increases, this Monte Carlo estimator tightens; SWD is sensitive to global geometry
2668 while remaining computationally light.

2669 **Pairwise Spearman difference.** To assess monotone dependence, we compare pairwise
2670 Spearman rank correlations between real and generated data. For each pair (i, j) , let ρ_{ij}^{real}
2671 and ρ_{ij}^{gen} denote Spearman's ρ . The metric is the mean absolute discrepancy over all pairs:

$$\text{Spearman diff} = \frac{2}{d(d-1)} \sum_{1 \leq i < j \leq d} |\rho_{ij}^{\text{real}} - \rho_{ij}^{\text{gen}}|. \quad (2.116)$$

2672 Being rank-based, it is invariant to strictly monotone marginal transforms.

2673 **Sortedness.** Order statistics must be sorted. We report the proportion of generated
2674 samples that satisfy the coordinate-wise descending order $y_1^{(n)} \geq y_2^{(n)} \geq \dots \geq y_d^{(n)}$:

$$\text{Sortedness} = \frac{1}{N} \sum_{n=1}^N \mathbb{I}\{y_k^{(n)} - y_{k+1}^{(n)} \geq 0, \forall k \in [d-1]\}. \quad (2.117)$$

2675 This hard score equals 1 only if all generated vectors respect the ordering constraint.

2676 **Absolute Kendall Error (AKE).** To compare distributions of order statistics, we sort
2677 each sample vector in descending order and compute the mean absolute deviation between
2678 the resulting order-statistic vectors:

$$\text{AKE} = \frac{1}{Nd} \sum_{n=1}^N \sum_{k=1}^d |\text{sort}_{\downarrow}(\mathbf{x}^{(n)})_k - \text{sort}_{\downarrow}(\mathbf{y}^{(n)})_k|. \quad (2.118)$$

2679 This captures alignment of empirical order-statistic marginals; lower is better.

2680 **Softsortedness.** As a differentiable proxy for the sortedness constraint, we average the
2681 positive violations of the descending order:

$$\text{SoftSorted} = \frac{1}{N(d-1)} \sum_{n=1}^N \sum_{k=1}^{d-1} [y_{k+1}^{(n)} - y_k^{(n)}]_+, \quad [t]_+ \stackrel{\text{def}}{=} \max\{0, t\}. \quad (2.119)$$

2682 This penalty is zero if and only if every generated vector is sorted in descending order and
2683 increases linearly with the magnitude of violations.

2684 **Chapter 3**

2685 **Gradient-Free Fine-Tuning of**
2686 **Diffusion Models**

2687 This chapter presents a gradient-free approach to fine-tuning diffusion models for reward-
2688 based generation. Using Fisher’s identity, we derive a score estimator that requires only
2689 forward evaluations of the reward function, avoiding costly backpropagation through the
2690 reward model.

2691 **3.1 Introduction**

2692 Diffusion models are now a standard tool for generative modelling in images, audio, and
2693 related domains (Jascha Sohl-Dickstein et al. 2015; Ho, Jain, et al. 2020; Y. Song, Jascha
2694 Sohl-Dickstein, et al. 2021). They learn a time-dependent score (or denoiser) that inverts
2695 a simple noising process and produces new samples from noise. Following their success, it
2696 has been proposed to fine-tune pretrained diffusion models, taking into account new data
2697 preferences and/or constraints by simply shifting the learnt weights accordingly.

2698 To achieve this, ideas from language-model alignment have recently started to be ap-
2699 plied to diffusion: Either shifting the pre-trained weights using methods based on reward
2700 models (Jiazheng Xu et al. 2023) or reinforcement learning policies (Black et al. 2024;
2701 Fan et al. 2023), or adapting the architecture using lightweight adapters such as LoRA
2702 (Hu et al. 2021). In addition, diffusion-specific techniques like classifier-free guidance (Ho
2703 and Salimans 2022) or improved noise schedulers (Alex Nichol and Dhariwal 2021) provide
2704 complementary control mechanisms.

2705 In this chapter we introduce *iterative tilting*, a gradient-free method for reward-based
2706 fine-tuning that shifts a diffusion model toward a reward-tilted density without differen-
2707 tiating through the reward function. During the preparation of this work, a concurrent
2708 paper was submitted to ICLR (Anonymous 2025). Their method is essentially the same
2709 as ours but formulated using the stochastic interpolants framework and developed in con-
2710 tinuous time. In contrast, our approach is based on a diffusion framework and is derived
2711 in discrete time.

2712 Section 3.2 recalls the theoretical framework of Denoising Diffusion Models (DDM) and
 2713 reviews existing reward fine-tuning formulations, Section 3.3 develops the Iterative Tilting
 2714 method from a theoretical standpoint, and Section 1.3 reports controlled experiments on
 2715 a Gaussian synthetic fine-tuning task, before we conclude on the work of this chapter.

2716 3.2 Preliminaries on Denoising Diffusion Models and 2717 Fine-Tuning

2718 3.3 Diffusion models

2719 Denoising diffusion models (DDMs) (Jascha Sohl-Dickstein et al. 2015; Y. Song and Ermon
 2720 2019; Ho, Jain, et al. 2020) define a generative process for a data distribution p_0 by
 2721 constructing a continuous probability path $(p_t)_{t \in [0,1]}$ of distributions between p_0 and a
 2722 tractable reference $p_1 := N(0, I_d)$.

2723 Forward noising process

2724 The path is defined via the marginals $p_t = \text{Law}(X_t)$, where

$$X_t = \alpha_t X_0 + \sigma_t X_1, \quad X_0 \sim p_0, \quad X_1 \sim N(0, I_d). \quad (3.1)$$

2725 Here X_0 and X_1 are independent, and $(\alpha_t)_{t \in [0,1]}$ and $(\sigma_t)_{t \in [0,1]}$ are deterministic schedules
 2726 such that α_t is non-increasing, σ_t is non-decreasing, with boundary conditions $(\alpha_0, \sigma_0) =$
 2727 $(1, 0)$ and $(\alpha_1, \sigma_1) = (0, 1)$. The path $(p_t)_{t \in [0,1]}$ thus defines an interpolation that gradually
 2728 transforms the clean data distribution p_0 into the Gaussian reference p_1 .

2729 From (3.1), the conditional distribution of X_t given $X_0 = x_0$, denoted $q_{t|0}$, is Gaussian:

$$q_{t|0}(x_t | x_0) = N(x_t; \alpha_t x_0, \sigma_t^2 I_d). \quad (3.2)$$

2730 Regarding the choices of schedules (α_t, σ_t) , two standard choices are the *variance-*
 2731 *preserving (VP)* schedule with $\alpha_t^2 + \sigma_t^2 = 1$ (Ho, Jain, et al. 2020; Dhariwal and Alexander
 2732 Nichol 2021), which ensures $\text{Var}(X_t) = \text{Var}(X_0)$ when X_0 has unit variance, and the *linear*
 2733 (*flow matching*) schedule with $(\alpha_t, \sigma_t) = (1 - t, t)$ (Lipman et al. 2023; Esser et al. 2024),
 2734 corresponding to a straight-line interpolation between data and noise.

2735 Reverse process and DDIM transitions

2736 To generate new samples, DDMs simulate a time-reversed Markov chain. Given a increasing
 2737 sequence $(t_k)_{k=0}^K$ of time steps with $t_K = 1$ and $t_0 = 0$, reverse transitions are
 2738 iteratively applied to map a sample from $p_{t_{k+1}}$ to one from p_{t_k} , progressively denoising
 2739 until convergence to p_0 .

2740 **DDIM framework.** The DDIM framework (J. Song et al. 2021) introduces a general
 2741 family of reverse transitions. It defines a schedule $(\eta_t)_{t \in [0,1]}$ satisfying $\eta_t \leq \sigma_t$ for all

²⁷⁴² $t \in [0, 1]$, along with transition densities given for $s < t$ by

$$p_{s|t}^\eta(x_s | x_t) = \mathbb{E} \left[q_{s|0,1}^\eta(x_s | X_0, X_1) \mid X_t = x_t \right], \quad (3.3)$$

²⁷⁴³ where

$$q_{s|0,1}^\eta(x_s | x_0, x_1) := N\left(x_s; \alpha_s x_0 + \sqrt{\sigma_s^2 - \eta_s^2} x_1, \eta_s^2 I_d\right) \quad (3.4)$$

²⁷⁴⁴ and the random variables (X_0, X_t, X_1) are defined as in (3.1). By construction, this
²⁷⁴⁵ family satisfies the marginalisation property $p_s(x_s) = \int p_{s|t}^\eta(x_s | x_t) p_t(x_t) dx_t$: by the
²⁷⁴⁶ law of total expectation, the right-hand side equals $\mathbb{E}_{(X_0, X_1)}[q_{s|0,1}^\eta(x_s | X_0, X_1)]$, and since
²⁷⁴⁷ $X_1 \sim N(0, I_d)$ is independent of X_0 , sampling from the bridge (3.4) yields $X_s = \alpha_s X_0 +$
²⁷⁴⁸ $\sqrt{\sigma_s^2 - \eta_s^2} X_1 + \eta_s \epsilon$ with $\epsilon \sim N(0, I_d)$; the sum $\sqrt{\sigma_s^2 - \eta_s^2} X_1 + \eta_s \epsilon$ has variance $\sigma_s^2 I_d$, so
²⁷⁴⁹ $X_s \stackrel{d}{=} \alpha_s X_0 + \sigma_s \tilde{X}_1$ with $\tilde{X}_1 \sim N(0, I_d)$, matching the definition of p_s . This ensures that
²⁷⁵⁰ $(p_{t_k|t_{k+1}}^\eta)_{k=0}^{K-1}$ defines a consistent set of reverse transitions.

²⁷⁵¹ **Denoiser.** The transitions (3.3) are intractable and have to be approximated. To this
²⁷⁵² end, it has been proposed to replace X_0 and X_1 by their conditional expectations given
²⁷⁵³ X_t . Let $\hat{x}_0^\theta(\cdot, t)$ and $\hat{x}_1^\theta(\cdot, t)$ denote parametric estimators of

$$\hat{x}_0(x_t, t) := \mathbb{E}[X_0 | X_t = x_t], \quad \hat{x}_1(x_t, t) := \mathbb{E}[X_1 | X_t = x_t]. \quad (3.5)$$

²⁷⁵⁴ From the forward interpolation (3.1), these two quantities are related by $\hat{x}_1(x_t, t) = (x_t -$
²⁷⁵⁵ $\alpha_t \hat{x}_0(x_t, t)) / \sigma_t$, yielding the parametric relationship

$$\hat{x}_1^\theta(x_t, t) = \frac{x_t - \alpha_t \hat{x}_0^\theta(x_t, t)}{\sigma_t}, \quad \hat{x}_0^\theta(x_t, t) = \frac{x_t - \sigma_t \hat{x}_1^\theta(x_t, t)}{\alpha_t}. \quad (3.6)$$

²⁷⁵⁶ The parametric DDIM transition then becomes

$$p_{t_k|t_{k+1}}^{\eta, \theta}(x_{t_k} | x_{t_{k+1}}) := q_{t_k|0,1}^\eta(x_{t_k} | \hat{x}_0^\theta(x_{t_{k+1}}, t_{k+1}), \hat{x}_1^\theta(x_{t_{k+1}}, t_{k+1})). \quad (3.7)$$

²⁷⁵⁷ For $k = 0$, the transition $p_{0|t_1}^{\eta, \theta}(\cdot | x_{t_1})$ is simply defined as the Dirac mass at $\hat{x}_0^\theta(x_{t_1}, t_1)$.

²⁷⁵⁸ **Sampling procedure.** Given timesteps $(t_k)_{k=0}^K$ with $t_K = 1$ and $t_0 = 0$, DDIM sampling
²⁷⁵⁹ is summarised in Algorithm 5. Setting $\eta_t = 0$ yields deterministic sampling; setting
²⁷⁶⁰ $\eta_t = \sigma_t \sqrt{1 - \alpha_t^2 / \alpha_s^2}$ (for the VP schedule) recovers the stochastic DDPM sampler (Ho,
²⁷⁶¹ Jain, et al. 2020).

²⁷⁶² Score function and training

²⁷⁶³ **Score and denoiser relationship.** The *score* of the noised distribution is $\nabla \log p_t(\cdot)$.
²⁷⁶⁴ For the Gaussian forward kernel (3.2), the score admits a closed-form expression in terms of

Algorithm 5: DDIM Sampling

Require: Denoiser \hat{x}_1^θ , timesteps $(t_k)_{k=0}^K$ with $t_K = 1$, $t_0 = 0$, variance schedule $(\eta_{t_k})_{k=0}^{K-1}$

- 1: Sample $x_{t_K} \sim N(0, I_d)$ **for** $k = K - 1, \dots, 1$ **do**
- 2: $\hat{x}_1 \leftarrow \hat{x}_1^\theta(x_{t_{k+1}}, t_{k+1})$ ▷ Predict X_1
- 3: $\hat{x}_0 \leftarrow (x_{t_{k+1}} - \sigma_{t_{k+1}} \hat{x}_1) / \alpha_{t_{k+1}}$ ▷ Estimate clean sample
- 4: $\mu_{t_k} \leftarrow \alpha_{t_k} \hat{x}_0 + \sqrt{\sigma_{t_k}^2 - \eta_{t_k}^2} \hat{x}_1$
- 5: Sample $z \sim N(0, I_d)$
- 6: $x_{t_k} \leftarrow \mu_{t_k} + \eta_{t_k} z$
- 7:
- 8: $x_0 \leftarrow (x_{t_1} - \sigma_{t_1} \hat{x}_1^\theta(x_{t_1}, t_1)) / \alpha_{t_1}$ ▷ Final denoising step
- 9: **return** x_0

2765 the denoiser. Exchanging expectation and gradient under standard regularity assumptions:

$$\nabla_{x_t} \log p_t(x_t) = \mathbb{E} [\nabla_{x_t} \log q_{t|0}(x_t | X_0) \mid X_t = x_t] = -\frac{\hat{x}_1(x_t, t)}{\sigma_t}. \quad (3.8)$$

2766 Thus, the denoiser \hat{x}_1^θ directly provides a score estimator $s_\theta(x_t, t) = -\hat{x}_1^\theta(x_t, t)/\sigma_t$, and
2767 conversely $\hat{x}_1^\theta(x_t, t) = -\sigma_t s_\theta(x_t, t)$.

2768 **Training objective.** The denoiser can be trained by regressing either X_0 or X_1 from
2769 the noised sample $X_t = \alpha_t X_0 + \sigma_t X_1$. The X_0 -prediction loss reads

$$L_{X_0}(\theta) = \int_0^1 \mathbb{E}_{X_0 \sim p_0, X_1 \sim N(0, I_d)} \|\hat{x}_0^\theta(\alpha_t X_0 + \sigma_t X_1, t) - X_0\|^2 dt, \quad (3.9)$$

2770 while the X_1 -prediction loss is

$$L_{X_1}(\theta) = \int_0^1 \mathbb{E}_{X_0 \sim p_0, X_1 \sim N(0, I_d)} \|\hat{x}_1^\theta(\alpha_t X_0 + \sigma_t X_1, t) - X_1\|^2 dt. \quad (3.10)$$

2771 Since $\hat{x}_1^\theta = -\sigma_t s_\theta$ by (3.8), the X_1 -prediction loss is equivalent to denoising score matching
2772 (Hyvärinen 2005; Vincent 2011): minimising (3.10) amounts to learning the score $\nabla \log p_t$.
2773 In practice, the integral is approximated via Monte Carlo: sample $t \sim \text{Unif}[0, 1]$, $x_0 \sim p_0$,
2774 $x_1 \sim N(0, I_d)$, form $x_t = \alpha_t x_0 + \sigma_t x_1$, and regress either x_0 or x_1 .

2775 **Connection to SDEs**

2776 The interpolation (3.1) can be viewed as discretising a continuous-time process. Given a
2777 noise schedule $\beta(t) > 0$, consider the Itô SDE

$$dX_t = -\frac{1}{2}\beta(t)X_t dt + \sqrt{\beta(t)} dW_t, \quad X_0 \sim p_0, \quad (3.11)$$

2778 where $(W_t)_{t \geq 0}$ is a standard Brownian motion in \mathbb{R}^d . For the VP schedule, setting $\alpha_t =$
2779 $\exp(-(1/2) \int_0^t \beta(s) ds)$ and $\sigma_t = \sqrt{1 - \alpha_t^2}$, the marginal Law(X_t) coincides with p_t from

2780 (3.1). The time-reversed process satisfies (Anderson 1982; Y. Song, Jascha Sohl-Dickstein,
2781 et al. 2021)

$$d\bar{X}_t = \left[-\frac{1}{2}\beta(t)\bar{X}_t + \beta(t)\nabla_x \log p_t(\bar{X}_t) \right] dt + \sqrt{\beta(t)} d\bar{W}_t, \quad (3.12)$$

2782 where $\bar{X}_t := X_{1-t}$ and \bar{W}_t is a standard Brownian motion. The deterministic *probability-*
2783 *flow ODE* (Y. Song, Jascha Sohl-Dickstein, et al. 2021) with identical marginals reads

$$dx_t = \left[-\frac{1}{2}\beta(t)x_t + \frac{1}{2}\beta(t)\nabla_x \log p_t(x_t) \right] dt. \quad (3.13)$$

2784 DDIM with $\eta_t = 0$ can be seen as a discretisation of (3.13), while DDPM corresponds to
2785 (3.12).

2786 3.B Fine-tuning formulations

2787 Given a positive reward function r and a pretrained diffusion model sampling from the
2788 base distribution p_0^b , we aim to fine-tune the model to sample from the *tilted distribution*,
2789 defined for $x \in \mathbb{R}^d$:

$$p_0^{\lambda r}(x) = Z_{\lambda r}^{-1} \exp(\lambda r(x)) p_0^b(x), \quad (3.14)$$

2790 where $\lambda > 0$ is the tilt strength and $Z_{\lambda r} = \int \exp(\lambda r(x)) p_0^b(x) dx$ is the normalising con-
2791 stant. This tilted distribution emphasises regions with high reward. Several formulations
2792 have been developed to tackle this goal.

2793 **Direct reward alignment (DRaFT).** DRaFT (Clark et al. 2024) treats the diffu-
2794 sion sampler as a differentiable computation graph. For network parameters θ , the re-
2795 verse sampling process is a deterministic function of θ and the i.i.d. Gaussian noises
2796 $\xi = (Z_K, \dots, Z_1)$, denoted $(\theta, \xi) \mapsto \psi_0(\theta, \xi)$, i.e., $\psi_0(\theta, \xi) = \hat{X}_0$ where \hat{X}_0 is the out-
2797 put of Algorithm 5.

2798 DRaFT maximises the expected reward:

$$J(\theta) = \mathbb{E}_{\xi} [r(\psi_0(\theta, \xi))] \quad (3.15)$$

2799 by backpropagating the reward gradient through the entire chain of reverse steps:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\xi} \left[\left(\frac{\partial \psi_0(\theta, \xi)}{\partial \theta} \right)^{\top} \nabla_x r(\psi_0(\theta, \xi)) \right]. \quad (3.16)$$

2800 Full backpropagation is memory-intensive. DRaFT- K reduces cost by differentiat-
2801 ing only through the final K steps (typically $K \in \{1, 4\}$), detaching earlier gradients.
2802 DRaFT-LV further refines $K = 1$ with a control-variate correction for lower variance
2803 while remaining unbiased.

2804 **Stochastic optimal control and Adjoint Matching.** An alternative approach casts
2805 fine-tuning as stochastic optimal control (SOC) (Uehara et al. 2024; Tang 2024; Domingo-
2806 Enrich, Drozdza, et al. 2024). Rather than backpropagating through sampling (as in

2807 DRaFT), SOC methods add a drift control u_t to shift the generative process toward high-
2808 reward regions. For a controlled SDE $dX_t = [b(X_t, t) + g(t)u_t(X_t)] dt + g(t) dW_t$, where
2809 $g(t)$ denotes the diffusion coefficient (distinct from the noise schedule σ_t in (3.1)), the
2810 objective balances control effort against terminal reward:

$$\min_u \mathbb{E} \left[\frac{1}{2} \int_0^1 \|u_t(X_t)\|^2 dt - \lambda r(X_1) \right]. \quad (3.17)$$

2811 By Girsanov’s theorem, the quadratic control cost equals the KL divergence between con-
2812 trolled and base processes. However, naïvely solving (3.17) yields $p^*(X_0, X_1) \propto p^{\text{base}}(X_0, X_1) \exp(\lambda r(X_1))$
2813 $V(X_0, 0)$), where V is the value function. This does not marginalize to the tilted distri-
2814 bution (3.14) unless the initial noise X_0 and generated sample X_1 are independent.

2815 (Domingo-Enrich, Drozdzal, et al. 2024) show that a *memoryless* diffusion coefficient,
2816 one that ensures $X_0 \perp X_1$, removes the value function bias. Under this schedule, the
2817 optimal control is $u_t^*(x) = g(t)^\top \tilde{a}_t(x)$, where the *lean adjoint* \tilde{a}_t satisfies the backward
2818 ODE

$$\frac{d}{dt} \tilde{a}_t = -[\nabla_x b(x, t)]^\top \tilde{a}_t, \quad \tilde{a}_1 = \lambda \nabla_x r(X_1). \quad (3.18)$$

2819 Adjoint Matching learns a parametric control u_ϕ by minimizing:

$$L_{\text{AM}}(\phi) = \mathbb{E} \left[\int_0^1 \|u_\phi(X_t, t) + g(t)^\top \tilde{a}_t(X_t)\|^2 dt \right]. \quad (3.19)$$

2820 After training, the learned control is absorbed into the drift, yielding a fine-tuned score
2821 or velocity field. Sampling can then use any diffusion coefficient, including deterministic
2822 sampling with $g(t) = 0$.

2823 3.3 Iterative Tilting Method for Fine-Tuning Diffu- 2824 sion Models

2825 The methods reviewed in Section 3.B share a common limitation: they require differen-
2826 tiating through the reward function r . In DRaFT, reward gradients appear explicitly in
2827 (3.16) via backpropagation through the sampling chain. In stochastic optimal control, they
2828 appear in the adjoint terminal condition $a_0 = \lambda \nabla_x r(X_0)$. When r is a large pretrained
2829 model (such as a vision transformer or human preference classifier trained on millions of
2830 examples), computing these gradients becomes prohibitively expensive, both in memory
2831 and compute. Moreover, many reward functions of practical interest are non-differentiable
2832 or defined only through black-box evaluations.

2833 We introduce *iterative tilting*, a method that targets the reward-tilted distribution
2834 (3.14) using only forward evaluations of r , without any gradient computation. The key
2835 idea is to construct a path of intermediate distributions $(p_0^k)_{k=0}^N$ connecting the base p_0^b to
2836 the tilted $p_0^{\lambda r}$, each representing a gentle reweighting of the previous one. At each step,
2837 we reweight samples from the current distribution toward higher rewards, gradually ap-
2838 proaching the target along this path. This sequential approach trades off a single expensive

gradient-based update for multiple cheaper gradient-free updates.

This section develops the method in continuous time, which makes the structure of the approximation explicit and connects naturally to the score-based and control perspectives introduced in Section 3.2.

3.3 Score of the tilted distribution

Suppose we have a data distribution $X_0 \sim p_0^b$ (the base distribution) and wish to sample from the tilted distribution (3.14). We use the forward noising kernel $q_{t|0}$ from (3.2), whose score is available in closed form:

$$\nabla_{x_t} \log q_{t|0}(x_t | x_0) = \frac{\alpha_t x_0 - x_t}{\sigma_t^2}. \quad (3.20)$$

Noising the tilted distribution, the marginal at time t is:

$$p_t^{\lambda r}(x_t) = \int q_{t|0}(x_t | x_0) p_0^{\lambda r}(x_0) dx_0. \quad (3.21)$$

For any prior density p_0° (e.g., base p_0^b or tilted $p_0^{\lambda r}$) and the fixed forward noising kernel $q_{t|0}(x_t | x_0)$, we write $p_{0|t}^\circ(x_0 | x_t)$ for the posterior (inference) density of the clean sample X_0 given its noised version $X_t = x_t$ at time t :

$$p_{0|t}^\circ(x_0 | x_t) \propto p_0^\circ(x_0) q_{t|0}(x_t | x_0). \quad (3.22)$$

Proposition 3.1 (Score of the tilted distribution). Assume that r has at most linear growth, i.e., $\forall x \in \mathbb{R}^d : |r(x)| \leq C(1 + \|x\|)$ for some $C > 0$, and that $q_{t|0}$ is the Gaussian forward kernel (3.2). Then for all $t \in (0, 1]$ and $x_t \in \mathbb{R}^d$,

$$\nabla_{x_t} \log p_t^{\lambda r}(x_t) = \frac{\mathbb{E}_{X_0 \sim p_{0|t}^b(\cdot | x_t)} [\nabla_{x_t} \log q_{t|0}(x_t | X_0) \exp(\lambda r(X_0))]}{\mathbb{E}_{X_0 \sim p_{0|t}^b(\cdot | x_t)} [\exp(\lambda r(X_0))]} \quad (3.23)$$

Proof. Exchanging differentiation and integration in (3.21) (justified by dominated convergence under the stated assumptions) yields

$$\nabla_{x_t} p_t^{\lambda r}(x_t) = \int p_0^{\lambda r}(x_0) \nabla_{x_t} q_{t|0}(x_t | x_0) dx_0. \quad (3.24)$$

Dividing by $p_t^{\lambda r}(x_t)$ and using $\nabla_{x_t} p_t^{\lambda r}(x_t) = p_t^{\lambda r}(x_t) \nabla_{x_t} \log p_t^{\lambda r}(x_t)$, we obtain Fisher's identity:

$$\nabla_{x_t} \log p_t^{\lambda r}(x_t) = \mathbb{E}_{X_0 \sim p_{0|t}^{\lambda r}(\cdot | x_t)} [\nabla_{x_t} \log q_{t|0}(x_t | X_0)]. \quad (3.25)$$

The key observation is that the tilted posterior can be written as a reweighting of the base

2859 posterior:

$$\begin{aligned} p_{0|t}^{\lambda r}(x_0 | x_t) &= \frac{p_0^{\lambda r}(x_0) q_{t|0}(x_t | x_0)}{p_t^{\lambda r}(x_t)} = \frac{Z_{\lambda r}^{-1} \exp(\lambda r(x_0)) p_0^b(x_0) q_{t|0}(x_t | x_0)}{\int p_0^{\lambda r}(x'_0) q_{t|0}(x_t | x'_0) dx'_0} \\ &= \frac{\exp(\lambda r(x_0)) p_{0|t}^b(x_0 | x_t)}{\mathbb{E}_{X'_0 \sim p_{0|t}^b(\cdot | x_t)} [\exp(\lambda r(X'_0))]}, \end{aligned} \quad (3.26)$$

2860 where in the last equality, the marginal $p_t^b(x_t)$ cancels out. The tilted posterior $p_{0|t}^{\lambda r}(\cdot | x_t)$
2861 is thus an exponential tilt of the base posterior $p_{0|t}^b(\cdot | x_t)$. Substituting (3.26) into (3.25)
2862 gives (3.23). \square

2863 **Iterative tilts.** To efficiently approximate the score of the tilted distribution, we decom-
2864 pose the full tilt into a sequence of smaller tilts. Fix $N \in \mathbb{N}^*$ (typically large) and define,
2865 for $k = 0, \dots, N$,

$$p_0^k(x) := Z_k^{-1} \exp\left(\frac{k}{N} \lambda r(x)\right) p_0^b(x), \quad (3.27)$$

2866 where Z_k is the normalising constant. We have $p_0^0 = p_0^b$ and $p_0^N = p_0^{\lambda r}$. The key observation
2867 is that consecutive distributions are related by a small tilt: $p_0^k(x) \propto \exp(\lambda r(x)/N) p_0^{k-1}(x)$.

2868 **Corollary 3.1** (First-order score approximation). *Under the assumptions of Proposi-*
2869 *tion 3.1, for all $x_t \in \mathbb{R}^d$ and $t \in (0, 1]$, define the posterior $\pi_{k-1}(\cdot) := p_{0|t}^{k-1}(\cdot | x_t)$
2870 and the covariance*

$$C_{k-1}(x_t, t) := \text{Cov}_{X_0 \sim \pi_{k-1}} (\nabla_{x_t} \log q_{t|0}(x_t | X_0), r(X_0)). \quad (3.28)$$

2871 Then, for large N ,

$$\nabla_{x_t} \log p_t^k(x_t) = \nabla_{x_t} \log p_t^{k-1}(x_t) + \frac{\lambda}{N} C_{k-1}(x_t, t) + O\left(\frac{1}{N^2}\right). \quad (3.29)$$

2872 *Proof.* Applying Proposition 3.1 with base p_0^{k-1} and tilt strength λ/N :

$$\nabla_{x_t} \log p_t^k(x_t) = \frac{\mathbb{E}_{\pi_{k-1}} [g(X_0) \exp(\frac{\lambda}{N} r(X_0))]}{\mathbb{E}_{\pi_{k-1}} [\exp(\frac{\lambda}{N} r(X_0))]}, \quad (3.30)$$

2873 where $g(X_0) := \nabla_{x_t} \log q_{t|0}(x_t | X_0)$. Taylor-expanding $\exp(\lambda r/N) = 1 + \lambda r/N + O(N^{-2})$
2874 in both numerator and denominator:

$$\begin{aligned} \text{Numerator} &= \mathbb{E}_{\pi_{k-1}} [g] + \frac{\lambda}{N} \mathbb{E}_{\pi_{k-1}} [g r] + O(N^{-2}), \\ \text{Denominator} &= 1 + \frac{\lambda}{N} \mathbb{E}_{\pi_{k-1}} [r] + O(N^{-2}). \end{aligned}$$

2875 Dividing and keeping first-order terms yields $\mathbb{E}_{\pi_{k-1}} [g] + (\lambda/N) C_{k-1}(x_t, t) + O(N^{-2})$. By
2876 Fisher's identity, $\mathbb{E}_{\pi_{k-1}} [g] = \nabla_{x_t} \log p_t^{k-1}(x_t)$. \square

2877 3.B Neural approximation and algorithm

2878 Corollary 3.1 motivates a recursive algorithm for estimating the sequence of score functions
 2879 $\{\nabla_{x_t} \log p_t^k\}_{k=1}^N$ starting from $k = 0$, for which $p_t^0 = p_t^b$. Suppose that at iteration k , we
 2880 have access to an approximation for $\nabla_{x_t} \log p_t^k$, denoted by $s_{\theta^k}^k$. The resulting generative
 2881 models associated to $s_{\theta^k}^k$ is called the teacher at step k . Corollary 3.1 provides a natural
 2882 procedure to obtain an approximation for $\nabla_{x_t} \log p_t^{k+1}$: we train a new network s_{θ}^{k+1} with
 2883 weights θ , so that it matches the target score from (3.29). The updated parameter θ^{k+1} is
 2884 thus obtained as an approximate minimizer of the loss

$$2885 L_{\text{IT}}(\theta; k+1) = \int_0^1 \int_{\mathcal{X}} w(t) \left\| s_{\theta}(x_t, t) - \left(s_{\theta^k}^k(x_t, t) + \frac{\lambda}{N} C_k(x_t, t) \right) \right\|^2 p_t^k(x_t) dx_t dt. \quad (3.31)$$

2885 In practice, we make three simplifications. First, we use a Monte Carlo estimate of the
 2886 loss using a teacher-generated dataset $\{X_0^i\}_{i=1}^N$ at the beginning of each step k . Second,
 2887 for $X_t^i \sim q_{t|0}(\cdot | X_0^i)$, the covariance $C_k(X_t^i, t)$ is approximated by a single-sample estimate
 2888 $g^i = r(X_0^i)(\nabla_{x_t} \log q_{t|0}(X_t^i | X_0^i) - s_{\theta^k}(X_t^i, t))$, where $\nabla_{x_t} \log q_{t|0}$ is computed via (3.20).
 2889 Finally, we multiply by σ_t^2 to stabilize training across noise levels. This leads to the
 2890 following empirical loss for a single sample time t and pair (X_0^i, X_1^i) : setting $\delta = \lambda/N$,

$$2891 L_{\text{prac}}(\theta; k+1) = \sum_{i=1}^N \sigma_t^2 \left\| s_{\theta}^{k+1}(X_t^i, t) - \sigma_t^2 (s_{\theta^k}^k(X_t^i, t) + \delta r(X_0^i) (g^i - s_{\theta^k}^k(X_t^i, t))) \right\|^2, \quad (3.32)$$

2891 Finally, since we expect that $\nabla_{x_t} \log p_t^k$ to be close to $\nabla_{x_t} \log p_t^{k+1}$, the score network
 2892 at step $k+1$, s_{θ}^{k+1} is initialized at $s_{\theta^k}^k$. The iterative tilting procedure is summarised in
 2893 Algorithm 6.

2894 3.4 Experiments

2895 In this section, we validate the iterative tilting method on a two-dimensional Gaussian
 2896 mixture with a reward function. This setup has a key advantage: tilting a Gaussian (or
 2897 mixture of Gaussians) with a quadratic reward (including linear as a special case) yields
 2898 another Gaussian (or mixture), for which we can compute the exact tilted density and
 2899 score analytically.

2900 3.A Gaussian data with quadratic rewards

2901 **One Gaussian.** Suppose the base density is Gaussian:

$$2902 p_0^b(x) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right\} = N(x; \mu, \Sigma). \quad (3.33)$$

2902 and the reward function r is quadratic, i.e.:

$$2903 r(x) = \frac{1}{2} x^T A x + b^T x + c. \quad (3.34)$$

Algorithm 6: Iterative Tilting

Require: Pretrained model s_{θ^0} , number of tilts N , tilt strength λ , samples per tilt S , reverse steps n , batch size B , epochs E , schedule (α_t, σ_t)

- 1: Set step size $\delta \leftarrow \lambda/N$ **for** $k = 1$ to N **do**
- 2: Generate dataset D_k of S samples: $x_0 \leftarrow \text{DDIM}(s_{\theta^k}, n)$
- 3: Initialise $\theta \leftarrow \theta^k$ and freeze teacher s_{θ^k} **for** epoch = 1 to E **do**
- 4: Sample mini-batch $\{X_0^{(i)}\}_{i=1}^B \subset D_k$, $\{t_i\}_{i=1}^B \stackrel{\text{iid}}{\sim} \text{Unif}[0, 1]$, $\{Z_i\}_{i=1}^B \stackrel{\text{iid}}{\sim} N(0, I_d)$ **for**
- 5: *i* = 1 to B **do**
- 6: Noise: $x_{t_i}^{(i)} = \alpha_{t_i} X_0^{(i)} + \sigma_{t_i} Z_i$
- 7: Compute: $g_i = (\alpha_{t_i} X_0^{(i)} - X_{t_i}^{(i)}) / \sigma_{t_i}^2$
- 8: Compute: $s_i^{\text{old}} = s_{\theta^k}(X_{t_i}^{(i)}, t_i)$
- 9: Target: $\text{target}_i = s_i^{\text{old}} + \delta r(X_0^{(i)}) (g_i - s_i^{\text{old}})$
- 10: Update θ by one step of gradient descent on
$$\frac{1}{B} \sum_{i=1}^B \left\| \sigma_{t_i}^2 s_{\theta}(X_{t_i}^{(i)}, t_i) - \sigma_{t_i}^2 \text{target}_i \right\|^2$$
- 11:
- 12: Set $\theta^k \leftarrow \theta$
- 13:
- 14: **return** s_{θ^N}

²⁹⁰³ Then the product $\exp(r(x))p_0^b(x)$ writes:

$$\frac{1}{(2\pi)^{d/2}|\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu) + \left(\frac{1}{2}x^T Ax + b^T x + c \right) \right\}. \quad (3.35)$$

²⁹⁰⁴ Focusing on the term in the exponent:

$$\begin{aligned} & -\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu) + \left(\frac{1}{2}x^T Ax + b^T x + c \right) \\ & = -\frac{1}{2} [x^T (\Sigma^{-1} - A)x - 2(\Sigma^{-1}\mu + b)^T x] + c - \frac{1}{2}\mu^T \Sigma^{-1}\mu. \end{aligned} \quad (3.36)$$

²⁹⁰⁵ With $\Sigma' = (\Sigma^{-1} - A)^{-1}$ and $\mu' = \Sigma'(\Sigma^{-1}\mu + b)$ (we assume $\Sigma^{-1} - A$ is positive definite
²⁹⁰⁶ so that the tilt is normalizable), the exponential can be rewritten as

$$\exp(r(x))p_0^b(x) = \left(\frac{|\Sigma'|}{|\Sigma|} \right)^{1/2} \exp \left\{ \frac{1}{2} [\mu'^T \Sigma'^{-1} \mu' - \mu^T \Sigma^{-1} \mu] + c \right\} N(x; \mu', \Sigma'). \quad (3.37)$$

²⁹⁰⁷ Consequently,

$$\int \exp(r(x))p_0^b(x) dx = \left(\frac{|\Sigma'|}{|\Sigma|} \right)^{1/2} \exp \left\{ \frac{1}{2} [\mu'^T \Sigma'^{-1} \mu' - \mu^T \Sigma^{-1} \mu] + c \right\}. \quad (3.38)$$

²⁹⁰⁸ **Mixture of Gaussians.** Suppose that we now have a mixture of Gaussians with $(\omega_1, \dots, \omega_K) \in$
²⁹⁰⁹ Δ_{K-1} where

$$\Delta_{K-1} = \{(\omega_1, \dots, \omega_K) \in \mathbb{R}_+^K : \sum_{i \in [K]} \omega_i = 1\} \quad (3.39)$$

²⁹¹⁰ is the simplex and $(\mu_1, \Sigma_1), \dots, (\mu_K, \Sigma_K)$ corresponding means and covariances.

$$p_0^b(x) = \sum_{i \in [K]} \omega_i N(x; \mu_i, \Sigma_i). \quad (3.40)$$

²⁹¹¹ Then, we have:

$$\begin{aligned} \exp(r(x))p_0^b(x) &= \sum_{i \in [K]} \omega_i \exp(r(x)) N(x; \mu_i, \Sigma_i) \\ &= \sum_{i \in [K]} \omega'_i N(x; \mu'_i, \Sigma'_i), \end{aligned} \quad (3.41)$$

²⁹¹² where the parameters of the tilted Gaussians are

$$\begin{cases} \Sigma'_i = (\Sigma_i^{-1} - A)^{-1}, \\ \mu'_i = \Sigma'_i(\Sigma_i^{-1}\mu_i + b), \\ \omega'_i = \omega_i \left(\frac{|\Sigma'_i|}{|\Sigma_i|} \right)^{1/2} \exp \left\{ \frac{1}{2} [\mu'^T \Sigma'^{-1} \mu' - \mu^T \Sigma^{-1} \mu] + c \right\}. \end{cases} \quad (3.42)$$

²⁹¹³ We consider the following mixture of two Gaussians in dimension $d = 2$ as base distri-
²⁹¹⁴ bution:

$$\begin{cases} K = 2 \\ \mu_1 = (-2, 0); \mu_2 = (2, 0) \\ \Sigma_1 = \Sigma_2 = \sigma^2 \mathbf{I}; \sigma^2 = 1/2 \\ \omega_1 = \omega_2 = 1/2 \end{cases} \quad (3.43)$$

and we consider the linear reward $r(x) = b^T x$ with $b = (0, 4)$ (corresponding to $A = 0$, $c = 0$ in the general quadratic form). In this case, the resulting distribution is a mixture of two Gaussians with parameters given by Section 3.A:

$$\begin{cases} K' = 2 \\ \mu'_1 = (-2, 2); \mu'_2 = (2, 2) \\ \Sigma'_1 = \Sigma'_2 = \sigma^2 \mathbf{I}; \sigma^2 = 1/2 \\ \omega'_1 = \omega'_2 = 1/2. \end{cases} \quad (3.44)$$

3.B Implementation and results

Experimental setup. We first train a diffusion model on the base distribution using the denoising score matching objective (3.10) with the cosine noise schedule of (Alex Nichol and Dhariwal 2021). The architecture of the base model is a MLP model with depth 5 of width 256. We embed the time t with Fourier embedding (Tancik et al. 2020; Y. Song, Jascha Sohl-Dickstein, et al. 2021) of dimension 128. We train the model on a dataset of 30,000 samples from the base distribution for 200 epochs with early stopping (patience 50). The base model achieves a test negative log-likelihood of 3.278 nats (computed on 5,000 data points) and a mean squared error on the mean of 4.3×10^{-3} .

We then fine-tune the model using the iterative tilting method (Algorithm 6) with $N \in \{20, 50, 100, 200\}$ tilts, $S = 1,000$ samples per tilt, $J = 200$ reverse steps with stochastic DDIM ($\eta = 1.0$), batch size $B = 64$ and $E = 100$ epochs per tilt.

Results. We provide in Figure 3.1 samples obtained with the fine-tuned model at different stages of the tilting process for $N \in \{20, 50, 100, 200\}$.

At each iteration of the tilting procedure, we know the closed form of the score of p_0^k thanks to Section 3.A, and we can thus compute the mean squared error between the learned score and the true score of the tilted distribution. Thus, we monitor the RMSE (Root Mean Squared Error) between the learned score and the true score at each iteration $k = 1, \dots, N$:

$$\text{RMSE}_k = \mathbb{E}_{t \sim \text{Unif}[0,1], X_t \sim p_t^k} [\|s_{\theta^k}(X_t, t) - \nabla_x \log p_t^k(X_t)\|^2]^{1/2}. \quad (3.45)$$

We estimate this RMSE using 5,000 samples from p_t^k obtained with 200 stochastic DDIM reverse steps. We provide in Figure 3.2 the evolution of this RMSE during the tilting procedure for different values of N .

We report running times—total sampling time (s) and total training time (s)—and evaluation metrics (negative log-likelihood and mean squared mean error) in Table 3.1.

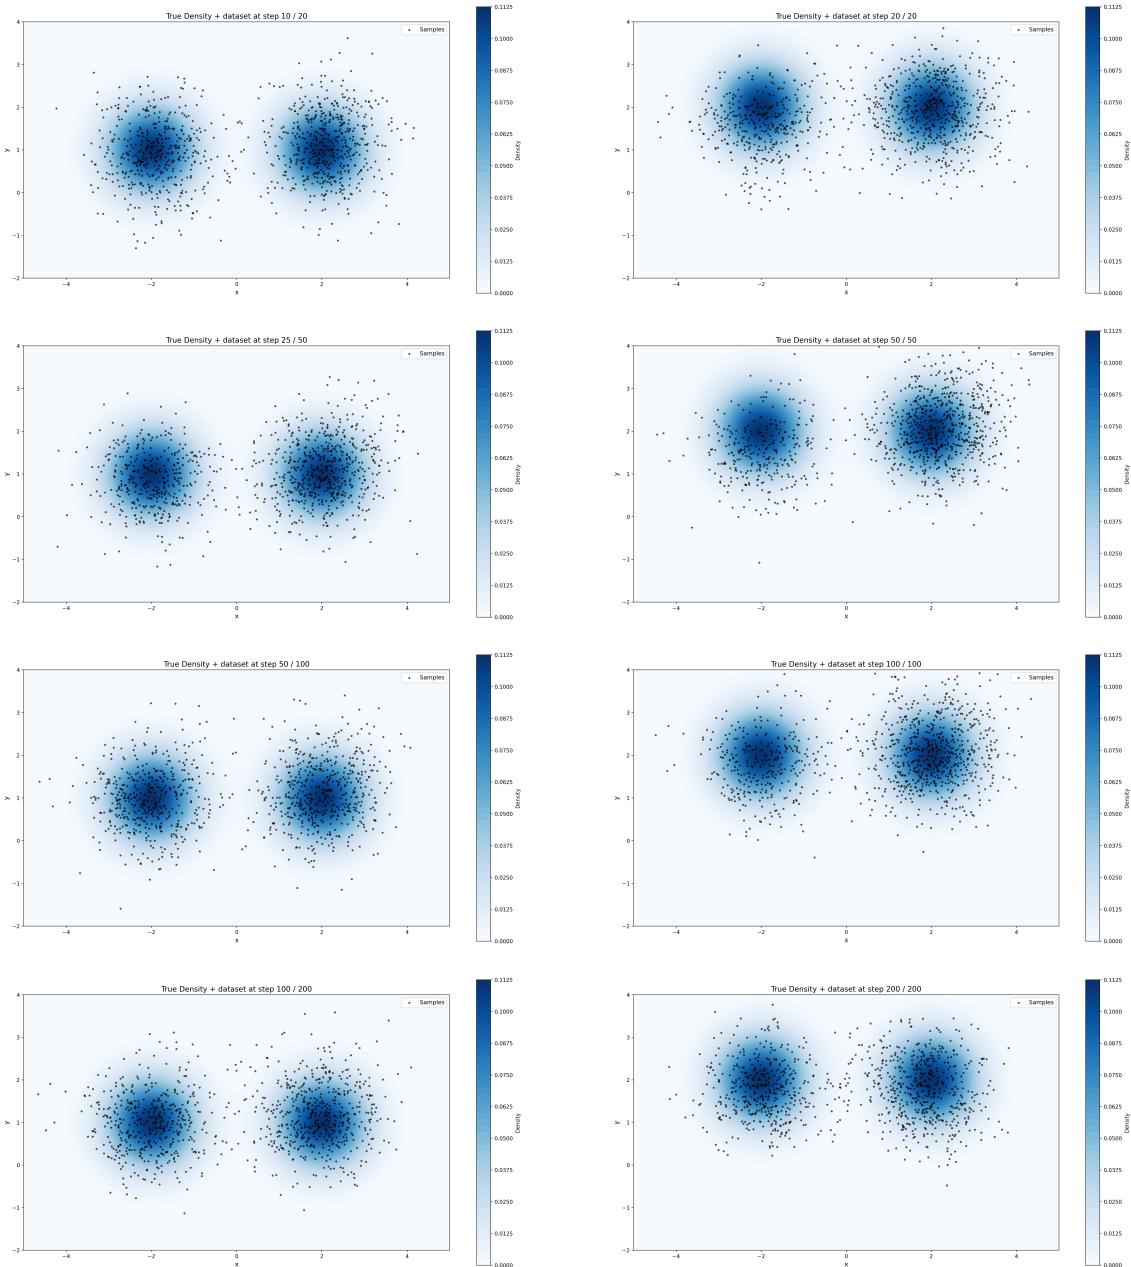


Figure 3.1: Iterative tilting on the 2-D GMM. Each row shows fixed- N settings (top to bottom: $N = 20, 50, 100, 200$) at iterations $N/2$ (left) and N (right). The rightmost plot in each row corresponds to the target distribution.

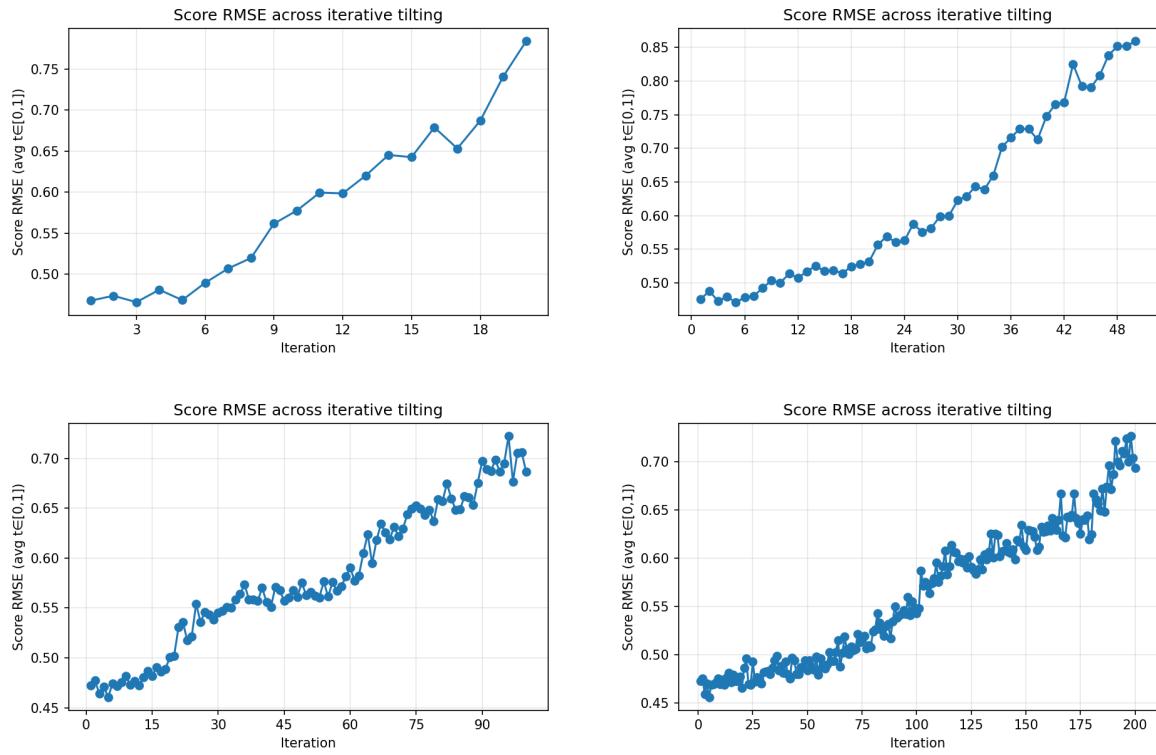


Figure 3.2: Evolution of the score error during iterative tilting for different N . We plot RMSE ($\sqrt{\text{MSE}}$) versus tilt iteration for $N \in \{20, 50, 100, 200\}$.

These controlled experiments validate the mechanics of iterative tilting in a setting where the exact score is observable. They also quantify the computational footprint of increasing the number of tilts, providing practical guidance for further studies on higher-dimensional datasets and richer reward models.

3.5 Conclusion

This chapter introduced iterative tilting, a gradient-free method for fine-tuning diffusion models toward reward-tilted distributions. The core idea is to decompose a single large tilt $\exp(\lambda r)$ into N sequential tilts $\exp(\lambda r/N)$, each producing a tractable score update. At iteration k , the score is approximated by

$$\nabla_x \log p_t^k(x) \approx \nabla_x \log p_t^{k-1}(x) + \frac{\lambda}{N} \text{Cov}_{X_0 \sim p_{0|t}^{k-1}(\cdot|x)} (\nabla_x \log q_{t|0}(x | X_0), r(X_0)),$$

valid for small λ/N via first-order Taylor expansion. Crucially, the method requires only forward evaluations of r , avoiding the cost of backpropagating through sampling chains or solving adjoint equations.

We validated the approach on a two-dimensional Gaussian mixture with linear reward, where the exact tilted distribution is available in closed form. The experiments confirmed that the method successfully recovers the target distribution for $N \in \{20, 50, 100, 200\}$, with score error (RMSE) converging at each tilt and runtime scaling linearly with N .

Number of tilts N	Runtime (s)		Metrics	
	Sampling (s)	Training (s)	Neg. log-lik.	MSE (mean)
20	28.77	34.71	6.48	1.67
50	71.83	85.18	6.94	2.20
100	149.14	182.54	7.23	2.16
200	289.61	355.76	6.34	1.59

Table 3.1: Running times and metrics for iterative tilting on the 2-D GMM. Runtime columns report total sampling and training time in seconds. Metrics report negative log-likelihood and mean squared error on the mean.

2958 Compared to DRaFT (Clark et al. 2024) and stochastic optimal control (Domingo-
 2959 Enrich, Han, et al. 2024), which both require differentiating through the reward, iterative
 2960 tilting trades a single gradient-based update for a sequence of gradient-free updates. This
 2961 makes it well suited to settings where r is a large pretrained model or is only available
 2962 through black-box evaluations.

2963 Several directions remain open: reducing variance in the single-sample covariance ap-
 2964 proximation, establishing principled criteria for selecting N , extending to high-dimensional
 2965 domains with learned reward models, and combining with parameter-efficient methods such
 2966 as LoRA (Hu et al. 2021).

2967 References

- 2968 Akiba, Takuya, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama (July
2969 2019). *Optuna: A Next-generation Hyperparameter Optimization Framework*. Comment:
2970 10 pages, Accepted at KDD 2019 Applied Data Science track. arXiv: [1907.10902](https://arxiv.org/abs/1907.10902)
2971 [cs]. (Visited on 08/03/2025).
- 2972 Allouche, M., S. Girard, and E. Gobet (2022). “EV-GAN: Simulation of extreme events
2973 with ReLU neural networks”. In: *Journal of Machine Learning Research* 23.150, pp. 1–
2974 39. ISSN: 1533-7928. URL: <http://jmlr.org/papers/v23/21-0663.html> (visited on
2975 05/16/2023).
- 2976 — (2024). “Estimation of extreme quantiles from heavy-tailed distributions with neural
2977 networks”. In: *Statistics and Computing* 34, p. 12.
- 2978 Allouche, Michaël, Stéphane Girard, and Emmanuel Gobet (2022). *Robust Estimation for
2979 Extremal Dependence*. arXiv: [2206.10876](https://arxiv.org/abs/2206.10876).
- 2980 — (Feb. 2025). “Learning extreme expected shortfall and conditional tail moments with
2981 neural networks. Application to cryptocurrency data”. In: *Neural Networks* 182, p. 106903.
2982 ISSN: 08936080. DOI: [10.1016/j.neunet.2024.106903](https://doi.org/10.1016/j.neunet.2024.106903). (Visited on 02/21/2025).
- 2983 Allouche, Michaël, Emmanuel Gobet, and Stéphane Girard (2024). “Extreme Event Simu-
2984 lation Using Generative Adversarial Networks”. In: *Handbook of Statistics*. Chapter on
2985 extreme value GAN methods. Elsevier.
- 2986 Anderson, Brian D.O. (1982). “Reverse-time diffusion equation models”. In: *Stochastic
2987 Processes and their Applications* 12.3, pp. 313–326. DOI: [10.1016/0304-4149\(82\)90051-5](https://doi.org/10.1016/0304-4149(82)90051-5).
- 2988 Anonymous (2025). “Tilt Matching for Scalable Sampling and Fine-Tuning”. In: *Submitted
2989 to The Fourteenth International Conference on Learning Representations*. under review.
2990 URL: <https://openreview.net/forum?id=tT7CXL3I9C>.
- 2991 Arjovsky, Martin, Soumith Chintala, and Léon Bottou (2017). “Wasserstein Generative
2992 Adversarial Networks”. In: *Proceedings of the 34th International Conference on Ma-
2993 chine Learning*. Vol. 70. Proceedings of Machine Learning Research. PMLR, pp. 214–
2994 223.
- 2995 Arnold, Barry C., Narayanaswamy Balakrishnan, and Haikady N. Nagaraja (2008). *A First
2996 Course in Order Statistics*. Classics in Applied Mathematics. Philadelphia: SIAM. DOI:
2997 [10.1137/1.9780898719062](https://doi.org/10.1137/1.9780898719062).

- 2999 Asmussen, S. and H. Albrecher (2010). *Ruin probabilities*. second. Advanced Series on
3000 Statistical Science & Applied Probability, 14. World Scientific Publishing Co. Pte. Ltd.,
3001 Hackensack, NJ.
- 3002 Bai, Yuntao et al. (2022). *Training a Helpful and Harmless Assistant with Reinforcement
3003 Learning from Human Feedback*. arXiv: [2204.05862](https://arxiv.org/abs/2204.05862).
- 3004 Banner, Adrian D. and Raouf Ghomrasni (July 2008). “Local times of ranked continuous
3005 semimartingales”. In: *Stochastic Processes and their Applications* 118.7, pp. 1244–1253.
3006 ISSN: 03044149. DOI: [10.1016/j.spa.2007.08.001](https://doi.org/10.1016/j.spa.2007.08.001). (Visited on 12/12/2024).
- 3007 Barrera, D., S. Crépey, E. Gobet, Hoang-Dung Nguyen, and B. Saadeddine (2025). “Statisti-
3008 cal Learning of Value-at-Risk and Expected Shortfall”. In: *Mathematical Finance, to
3009 appear*. arXiv: [2209.06476 \[q-fin, stat\]](https://arxiv.org/abs/2209.06476). URL: <http://arxiv.org/abs/2209.06476>
3010 (visited on 09/20/2024).
- 3011 Barron, Andrew R. (1993). “Universal approximation bounds for superpositions of a sig-
3012 moidal function”. In: *IEEE Transactions on Information Theory* 39.3, pp. 930–945.
3013 DOI: [10.1109/18.256500](https://doi.org/10.1109/18.256500).
- 3014 Beirlant, J., Y. Goegebeur, J. Teugels, and J. Segers (2004). *Statistics of extremes: theory
3015 and applications*. en. Wiley series in probability and statistics. Hoboken, NJ: Wiley.
3016 ISBN: 978-0-471-97647-9.
- 3017 Bellman, Richard (1961). *Adaptive Control Processes: A Guided Tour*. Princeton, NJ:
3018 Princeton University Press.
- 3019 Bevacqua, Emanuele, Carlo De Michele, Colin Manning, Anais Couasnon, Andreia F. S.
3020 Ribeiro, Alexandre M. Ramos, Edoardo Vignotto, Ana Bastos, Suzana Blesic, Fabrizio
3021 Durante, et al. (2021). “Guidelines for studying diverse types of compound weather and
3022 climate events”. In: *Earth's Future* 9.11, e2021EF002340. DOI: [10.1029/2021EF002340](https://doi.org/10.1029/2021EF002340).
- 3023 Bhattacharya, P. K. (1974). “Convergence of Sample Paths of Normalized Sums of Induced
3024 Order Statistics”. In: *The Annals of Statistics* 2.5. DOI: [10.1214/aos/1176342823](https://doi.org/10.1214/aos/1176342823).
- 3025 Black, Kevin, Michael Janner, Yilun Du, Ilya Kostrikov, and Sergey Levine (Jan. 2024).
3026 *Training Diffusion Models with Reinforcement Learning*. DOI: [10.48550/arXiv.2305.13301](https://doi.org/10.48550/arXiv.2305.13301) [cs]. (Visited on 09/18/2025).
- 3027 Blondel, Mathieu, Olivier Teboul, Quentin Berthet, and Josip Djolonga (2020). “Fast Dif-
3028 ferentiable Sorting and Ranking”. In: *Proceedings of the 37th International Conference
3029 on Machine Learning*. PMLR, pp. 950–959.
- 3030 Boulaguiem, Y., J. Zscheischler, E. Vignotto, K. van der Wiel, and S. Engelke (2022).
3031 “Modeling and simulating spatial extremes by combining extreme value theory with
3032 generative adversarial networks”. In: *Environmental Data Science* 1, e5.
- 3033 Boulaguiem, Youssef, Jakob Zscheischler, Edoardo Vignotto, Karin van der Wiel, and
3034 Sebastian Engelke (2022). “Modeling and simulating spatial extremes by combining
3035 extreme value theory with generative adversarial networks”. In: *Environmental Data
3036 Science* 1, e5. DOI: [10.1017/eds.2022.4](https://doi.org/10.1017/eds.2022.4).
- 3037 Bucklew, J. A. (2004). *Introduction to rare event simulation*. Springer Series in Statistics.
3038 Springer-Verlag, New York.

- 3040 Casper, Stephen, Xander Davies, Claudia Shi, Thomas Krendl Gilbert, Jérémie Scheurer,
3041 Javier Rando, Rachel Freedman, Tomasz Korbak, David Lindner, Pedro Freire, et al.
3042 (2023). *Open Problems and Fundamental Limitations of Reinforcement Learning from*
3043 *Human Feedback*. arXiv: [2307.15217](https://arxiv.org/abs/2307.15217).
- 3044 Christiano, Paul F., Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario
3045 Amodei (2017). *Deep Reinforcement Learning from Human Preferences*. NeurIPS 2017.
3046 arXiv: [1706.03741](https://arxiv.org/abs/1706.03741).
- 3047 Clark, Kevin, Paul Vicol, Kevin Swersky, and David J. Fleet (2024). *Directly Fine-Tuning*
3048 *Diffusion Models on Differentiable Rewards*. ICLR 2024. arXiv: [2309.17400](https://arxiv.org/abs/2309.17400).
- 3049 Cont, R., M. Cucuringu, R. Xu, and C. Zhang (2022). “Tail-GAN: Learning to Simulate
3050 Tail Risk Scenarios”. In: *SSRN*.
- 3051 Cont, Rama (2001). “Empirical Properties of Asset Returns: Stylized Facts and Statistical
3052 Issues”. In: *Quantitative Finance* 1.2, pp. 223–236. DOI: [10.1080/713665670](https://doi.org/10.1080/713665670).
- 3053 Cremer, F., B. Sheehan, M. Fortmann, A.N. Kia, M. Mullins, F. Murphy, and S. Materne
3054 (2022). “Cyber risk and cybersecurity: a systematic review of data availability”. In:
3055 *The Geneva Papers on risk and insurance-Issues and practice* 47.3, pp. 698–736.
- 3056 Cuturi, Marco, Olivier Teboul, and Jean-Philippe Vert (2019). “Differentiable Ranking
3057 and Sorting using Optimal Transport”. In: *Advances in Neural Information Processing*
3058 *Systems*. Vol. 32. Curran Associates, Inc.
- 3059 Cybenko, George (1989). “Approximation by superpositions of a sigmoidal function”.
3060 In: *Mathematics of Control, Signals, and Systems* 2.4, pp. 303–314. DOI: [10.1007/BF02551274](https://doi.org/10.1007/BF02551274).
- 3061 David, H.A. and H.N. Nagaraja (2004). *Order statistics*. Wiley series in probability and
3062 statistics. Wiley. ISBN: 978-0-471-65401-8. URL: <https://books.google.fr/books?id=bdhzFXg6xFkC>.
- 3063 David, Herbert A. and Haikady N. Nagaraja (2003). *Order Statistics*. 3rd. Wiley Series in
3064 Probability and Statistics. Hoboken, NJ: Wiley-Interscience. ISBN: 978-0-471-38926-2.
- 3065 David, Herbert A (1973). “Concomitants of order statistics”. In: *Bull. Inst. Internat.*
3066 *Statist.* 45.1, pp. 295–300.
- 3067 Del Moral, P. and J. Garnier (2005). “Genealogical particle analysis of rare events”. In:
3068 *The Annals of Applied Probability* 15.4, pp. 2496–2534.
- 3069 Dhariwal, Prafulla and Alexander Nichol (2021). “Diffusion Models Beat GANs on Image
3070 Synthesis”. In: *Advances in Neural Information Processing Systems*. Ed. by M. Ranzato,
3071 A. Beygelzimer, Y. Dauphin, P.S. Liang, and J. Wortman Vaughan. Vol. 34. Curran
3072 Associates, Inc., pp. 8780–8794.
- 3073 Dinh, Laurent, Jascha Sohl-Dickstein, and Samy Bengio (2017). “Density estimation using
3074 Real-NVP”. In: *International Conference on Learning Representations*.
- 3075 Domingo-Enrich, Carles, Michal Drozdzal, Brian Karrer, and Ricky T. Q. Chen (Jan. 2024).
3076 *Adjoint Matching: Fine-tuning Flow and Diffusion Generative Models with Memoryless*
3077 *Stochastic Optimal Control*. DOI: [10.48550/arXiv.2409.08861](https://doi.org/10.48550/arXiv.2409.08861). arXiv: [2409.08861](https://arxiv.org/abs/2409.08861)
3078 [cs]. (Visited on 09/22/2025).
- 3079
- 3080

- 3081 Domingo-Enrich, Carles, Michal Drozdzal, Brian Karrer, and Ricky T. Q. Chen (2025).
3082 *Adjoint Matching: Fine-tuning Flow and Diffusion Generative Models with Memoryless*
3083 *Stochastic Optimal Control*. arXiv: [2409.08861](https://arxiv.org/abs/2409.08861).
- 3084 Domingo-Enrich, Carles, Jiequn Han, Brandon Amos, Joan Bruna, and Ricky T. Q. Chen
3085 (Oct. 2024). *Stochastic Optimal Control Matching*. DOI: [10.48550/arXiv.2312.02027](https://doi.org/10.48550/arXiv.2312.02027).
3086 arXiv: [2312.02027 \[math\]](https://arxiv.org/abs/2312.02027). (Visited on 09/22/2025).
- 3087 ECB Banking Supervision (2023). “2023 stress test of euro area banks”. In: *Available at*
3088 https://www.bankingsupervision.europa.eu/ecb/pub/pdf/ssm_Report_2023_Stress_Test~96bb5a3af8.en.pdf ECB.
- 3090 Einmahl, John H. J., Laurens de Haan, and Chen Zhou (2016). “Statistics of heteroscedastic
3091 extremes”. In: *Journal of the Royal Statistical Society: Series B* 78.1, pp. 31–51. DOI:
3092 [10.1111/rssb.12099](https://doi.org/10.1111/rssb.12099).
- 3093 Einmahl, John H. J., Andrea Krajina, and Johan Segers (2012). “An M-estimator for tail
3094 dependence in arbitrary dimensions”. In: *The Annals of Statistics* 40.3, pp. 1764–1793.
3095 DOI: [10.1214/12-AOS1023](https://doi.org/10.1214/12-AOS1023).
- 3096 Embrechts, P., C. Klüppelberg, and T. Mikosch (1997). *Modelling Extremal Events for*
3097 *Insurance and Finance*. Springer-Verlag, Berlin.
- 3098 Embrechts, Paul, Claudia Klüppelberg, and Thomas Mikosch (1997). *Modelling Extremal
3099 Events*. Berlin, Heidelberg: Springer. DOI: [10.1007/978-3-642-33483-2](https://doi.org/10.1007/978-3-642-33483-2).
- 3100 Esser, Patrick, Sumith Kulal, Andreas Blattmann, Rahim Entezari, Jonas Müller, Harry
3101 Saini, Yam Levi, Dominik Lorenz, Axel Sauer, Frederic Boesel, Dustin Podell, Tim
3102 Dockhorn, Zion English, Kyle Lacey, Alex Goodwin, Yannik Marek, and Robin Rom-
3103 bach (Mar. 2024). *Scaling Rectified Flow Transformers for High-Resolution Image Syn-
3104 thesis*. DOI: [10.48550/arXiv.2403.03206](https://arxiv.org/abs/2403.03206). arXiv: [2403.03206 \[cs\]](https://arxiv.org/abs/2403.03206). (Visited on
3105 09/22/2025).
- 3106 European Banking Authority (2014). “Guidelines on the revised common procedures and
3107 methodologies for the supervisory review and evaluation process (SREP) and super-
3108 visory stress testing”. In: *Available at* <https://eba.europa.eu/regulation-and-policy/supervisory-review-and-evaluation-srep-and-pillar-2/guidelines-for-common-procedures-and-methodologies-for-the-supervisory-review-and-evaluation-process-srep-and-supervisory-stress-testing>
3112 EBA/GL/2014/13.
- 3113 European Commission (2024a). URL: https://finance.ec.europa.eu/sustainable-finance/tools-and-standards/esg-rating-activities_en.
- 3114 — (2024b). URL: https://finance.ec.europa.eu/sustainable-finance/overview-sustainable-finance_en.
- 3117 Falk, M. (2019). *Multivariate Extreme Value Theory and D-Norms*. en. Springer Series in
3118 Operations Research and Financial Engineering. Cham: Springer International Publishing.
3119 ISBN: 978-3-030-03818-2. DOI: [10.1007/978-3-030-03819-9](https://doi.org/10.1007/978-3-030-03819-9). URL: <http://link.springer.com/10.1007/978-3-030-03819-9> (visited on 05/16/2023).
- 3121 Falk, M. and T. Fuller (2021). “New characterizations of multivariate Max-domain of
3122 attraction and D-Norms”. en. In: *Extremes* 24.4, pp. 849–879. ISSN: 1386-1999, 1572-

- 3123 915X. DOI: [10.1007/s10687-021-00416-4](https://doi.org/10.1007/s10687-021-00416-4). URL: <https://link.springer.com/10.1007/s10687-021-00416-4> (visited on 05/16/2023).
- 3124 Fan, Ying, Olivia Watkins, Yuqing Du, Hao Liu, Moonkyung Ryu, Craig Boutilier, Pieter
3125 Abbeel, Mohammad Ghavamzadeh, Kangwook Lee, and Kimin Lee (2023). *DPOK: Reinforcement Learning for Fine-tuning Text-to-Image Diffusion Models*. NeurIPS 2023.
3126 arXiv: [2305.16381](https://arxiv.org/abs/2305.16381).
- 3127 Feder, R. M., P. Berger, and G. Stein (2020). “Nonlinear 3D cosmic web simulation
3128 with heavy-tailed generative adversarial networks”. In: *Physical Review D*. 102.10,
3129 pp. 103504, 18.
- 3130 Fisher, R. A. and L. H. C. Tippett (1928). “Limiting forms of the frequency distribution
3131 of the largest or smallest member of a sample”. In: *Mathematical Proceedings of the
3132 Cambridge Philosophical Society* 24.2, pp. 180–190. DOI: [10.1017/S0305004100015681](https://doi.org/10.1017/S0305004100015681).
- 3133 Fougères, Anne-Laure and Cécile Mercadier (2013). “Multivariate extremes and max-linear
3134 models”. In: *ESAIM: Probability and Statistics* 17, pp. 109–123.
- 3135 Glasserman, P. (2013). *Monte carlo methods in financial engineering*. Stochastic modelling
3136 and applied probability. Springer New York. ISBN: 978-0-387-21617-1. URL: <https://books.google.fr/books?id=aeAlBQAAQBAJ>.
- 3137 Gnedenko, B. (1943). “Sur La Distribution Limite Du Terme Maximum D’Une Série Aléatoire”. In: *Annals of Mathematics* 44.3, pp. 423–453. ISSN: 0003486X, 19398980. URL:
3138 <http://www.jstor.org/stable/1968974> (visited on 06/25/2024).
- 3139 Gobet, E. and G. Liu (2015). “Rare event simulation using reversible shaking transformations”. In: *SIAM Journal on Scientific Computing* 37.5, A2295–A2316.
- 3140 Goodfellow, I., J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville,
3141 and Y. Bengio (2014). “Generative Adversarial Nets”. In: *Advances in Neural Information
3142 Processing Systems*. Ed. by Z. Ghahramani, M. Welling, C. Cortes, N. Lawrence,
3143 and K.Q. Weinberger. Vol. 27, pp. 2672–2680.
- 3144 Goodfellow, Ian J., Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley,
3145 Sherjil Ozair, Aaron Courville, and Yoshua Bengio (2014). “Generative Adversarial
3146 Nets”. In: *Advances in Neural Information Processing Systems*. Ed. by Z. Ghahramani,
3147 M. Welling, C. Cortes, N. Lawrence, and K.Q. Weinberger. Vol. 27. Curran Associates,
3148 Inc.
- 3149 Graf, S. and H. Luschgy (2000). *Foundations of quantization for probability distributions*.
3150 en. Lecture Notes in Mathematics 1730. Berlin; New York: Springer. ISBN: 978-3-540-
3151 67394-1.
- 3152 Grover, Aditya, Eric Wang, Aaron Zweig, and Stefano Ermon (2019). “Stochastic Optimiza-
3153 tion of Sorting Networks via Continuous Relaxations”. In: *International Conference on
3154 Learning Representations*.
- 3155 Gulrajani, Ishaan, Faruk Ahmed, Martin Arjovsky, Vincent Dumoulin, and Aaron Courville
3156 (2017). “Improved Training of Wasserstein GANs”. In: *Advances in Neural Information
3157 Processing Systems*. Vol. 30. Curran Associates, Inc.
- 3158 Gumbel, Emil J. (1960). “Bivariate exponential distributions”. In: *Journal of the American
3159 Statistical Association* 55.292, pp. 698–707. DOI: [10.1080/01621459.1960.10483368](https://doi.org/10.1080/01621459.1960.10483368).
- 3160

- 3165 Haan, L. de and Ana Ferreira (2006). *Extreme value theory: an introduction*. en. Springer
3166 Series in Operations Research. New York ; London: Springer. ISBN: 978-0-387-23946-0.
- 3167 Haan, Laurens de and Ana Ferreira (2006). *Extreme Value Theory: An Introduction*. Springer
3168 Series in Operations Research. New York: Springer. ISBN: 978-0-387-23946-0.
- 3169 Haario, Heikki, Eero Saksman, and Johanna Tammisen (Sept. 1999). “Adaptive proposal
3170 distribution for random walk Metropolis algorithm”. In: *Computational Statistics* 14.3,
3171 pp. 375–395. ISSN: 0943-4062, 1613-9658. DOI: [10.1007/s001800050022](https://doi.org/10.1007/s001800050022). (Visited on
3172 07/08/2025).
- 3173 Hasan, A., K. Elkhilil, Y. Ng, J.M. Pereira, S. Farsiu, J. Blanchet, and V. Tarokh (2022).
3174 “Modeling extremes with d -max-decreasing neural networks”. In: *Uncertainty in Arti-
3175 ficial Intelligence*. PMLR, pp. 759–768.
- 3176 Hastings, W. K. (1970). “Monte Carlo Sampling Methods Using Markov Chains and Their
3177 Applications”. In: *Biometrika* 57.1, pp. 97–109. DOI: [10.1093/biomet/57.1.97](https://doi.org/10.1093/biomet/57.1.97).
- 3178 Hickling, T. and D. Prangle (2024). *Flexible Tails for Normalizing Flows*. en. arXiv:2406.16971.
3179 URL: <http://arxiv.org/abs/2406.16971> (visited on 06/27/2024).
- 3180 Ho, Jonathan, Ajay Jain, and Pieter Abbeel (2020). *Denoising Diffusion Probabilistic Mod-
3181 els*. arXiv: [2006.11239 \[cs.LG\]](https://arxiv.org/abs/2006.11239).
- 3182 Ho, Jonathan and Tim Salimans (July 2022). *Classifier-Free Diffusion Guidance*. DOI:
3183 [10.48550/arXiv.2207.12598 \[cs\]](https://doi.org/10.48550/arXiv.2207.12598). arXiv: [2207.12598 \[cs\]](https://arxiv.org/abs/2207.12598). (Visited on 09/22/2025).
- 3184 Hofert, M., R. Huser, and A. Prasad (2018). “Hierarchical archimax copulas”. In: *Journal
3185 of Multivariate Analysis* 167, pp. 195–211.
- 3186 Hornik, Kurt, Maxwell Stinchcombe, and Halbert White (1989). “Multilayer feedforward
3187 networks are universal approximators”. In: *Neural Networks* 2.5, pp. 359–366. DOI:
3188 [10.1016/0893-6080\(89\)90020-8](https://doi.org/10.1016/0893-6080(89)90020-8).
- 3189 Hu, Edward J., Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang,
3190 Lu Wang, and Weizhu Chen (Oct. 2021). *LoRA: Low-Rank Adaptation of Large Lan-
3191 guage Models*. DOI: [10.48550/arXiv.2106.09685 \[cs\]](https://doi.org/10.48550/arXiv.2106.09685). (Visited
3192 on 09/22/2025).
- 3193 Hüsler, Jürg and Rolf-Dieter Reiss (1989a). “Maxima of normal random vectors: be-
3194 tween independence and complete dependence”. In: *Statistics & Probability Letters*
3195 7.4, pp. 283–286. DOI: [10.1016/0167-7152\(89\)90106-5](https://doi.org/10.1016/0167-7152(89)90106-5).
- 3196 — (1989b). “Maxima of normal random vectors: between independence and complete de-
3197 pendence”. In: *Statistics & Probability Letters* 7.4, pp. 283–286.
- 3198 Huster, Todd, Jeremy Cohen, Zinan Lin, Kevin Chan, Charles Kamhoua, Nandi Kundu,
3199 Kevin Kwiat, and Noel Leslie (2021). *Pareto GAN: Extending the Representational
3200 Power of GANs to Heavy-Tailed Distributions*. arXiv: [2101.09113](https://arxiv.org/abs/2101.09113).
- 3201 Hyvärinen, Aapo (2005). “Estimation of Non-Normalized Statistical Models by Score Match-
3202 ing”. In: *Journal of Machine Learning Research* 6.24, pp. 695–709.
- 3203 JanSSen, Anja and Phyllis Wan (2020). “K-means clustering of extremes”. In: *Electronic
3204 Journal of Statistics* 14.1, pp. 1211–1233.
- 3205 Joe, H. (2014). *Dependence modeling with copulas*. CRC press.

- 3206 Kallenberg, O. (2002). *Foundations of modern probability*. Probability and its applications.
3207 Springer New York. ISBN: 978-0-387-95313-7. URL: <https://books.google.fr/books?id=L6fhXh130yMC>.
- 3209 Kim, S.H. and H.A. David (Mar. 1990). “On the dependence structure of order statistics
3210 and concomitants of order statistics”. In: *Journal of Statistical Planning and Inference*
3211 24.3, pp. 363–368. ISSN: 0378-3758. DOI: [10.1016/0378-3758\(90\)90055-y](https://doi.org/10.1016/0378-3758(90)90055-y). (Visited
3212 on 07/21/2025).
- 3213 Kingma, D. P. and J. Ba (2014). *Adam: A Method for Stochastic Optimization*. arXiv:1412.6980.
- 3214 Kingma, D. P. and M. Welling (2014). “Auto-Encoding Variational Bayes”. In: *2nd Inter-*
3215 *national Conference on Learning Representation, ICLR*.
- 3216 Kingma, Diederik P. and Max Welling (2014). *Auto-Encoding Variational Bayes*. arXiv:
3217 [1312.6114](https://arxiv.org/abs/1312.6114).
- 3218 Lafon, Marco, Philippe Naveau, and Anthony Leblois (2023). “On the use of variational
3219 autoencoders for extreme values modelling”. In: *arXiv preprint arXiv:2303.11536*.
- 3220 Leshno, Moshe, Vladimir Ya. Lin, Allan Pinkus, and Shimon Schocken (1993). “Multilayer
3221 feedforward networks with a nonpolynomial activation function can approximate any
3222 function”. In: *Neural Networks* 6.6, pp. 861–867. DOI: [10.1016/S0893-6080\(05\)80131-5](https://doi.org/10.1016/S0893-6080(05)80131-5).
- 3224 Liang, F., M. Mahoney, and L. Hodgkinson (2022). “FatTailed Variational Inference with
3225 Anisotropic Tail Adaptive Flows”. en. In: *Proceedings of the 39th International Con-*
3226 *ference on Machine Learning*. PMLR, pp. 13257–13270. URL: <https://proceedings.mlr.press/v162/liang22a.html> (visited on 01/11/2024).
- 3228 Lipman, Yaron, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le (Feb.
3229 2023). *Flow Matching for Generative Modeling*. DOI: [10.48550/arXiv.2210.02747](https://doi.org/10.48550/arXiv.2210.02747).
3230 arXiv: [2210.02747 \[cs\]](https://arxiv.org/abs/2210.02747). (Visited on 09/22/2025).
- 3231 Liu, Tie-Yan (2009). *Learning to Rank for Information Retrieval*. Berlin, Heidelberg: Springer.
3232 DOI: [10.1007/978-3-642-14267-3](https://doi.org/10.1007/978-3-642-14267-3).
- 3233 Lo, Andrew W. and Ruixun Zhang (2021). “Quantifying the Impact of Impact Investing”.
3234 In: *Management Science*. DOI: [10.2139/ssrn.3944367](https://doi.org/10.2139/ssrn.3944367).
- 3235 Maas, A.L., A.Y. Hannun, and A.Y. Ng (2013). “Rectifier Nonlinearities Improve Neural
3236 Network Acoustic Models”. en. In: *Proceedings of the International Conference on Ma-*
3237 *chine Learning (ICML)*.
- 3238 Mainik, Georg, Georgi Mitov, and Ludger Rüschendorf (2015). “Portfolio optimization
3239 for heavy-tailed assets: Extreme Risk Index vs. Markowitz”. In: *Journal of Empirical
3240 Finance* 32, pp. 115–134.
- 3241 McNeil, A.J., R. Frey, and P. Embrechts (2015). *Quantitative risk management*. Princeton
3242 Series in Finance. Princeton University Press, Princeton, NJ, pp. xix+699. ISBN: 978-
3243 0-691-16627-8.
- 3244 McNeil, Alexander J., Rüdiger Frey, and Paul Embrechts (2015). *Quantitative Risk Man-*
3245 *agement: Concepts, Techniques and Tools*. Revised. Princeton Series in Finance. Prince-
3246 ton, NJ: Princeton University Press.

- 3247 Metropolis, Nicholas, Arianna W. Rosenbluth, Marshall N. Rosenbluth, Augusta H. Teller,
3248 and Edward Teller (1953). “Equation of State Calculations by Fast Computing Ma-
3249 chines”. In: *The Journal of Chemical Physics* 21.6, pp. 1087–1092. DOI: [10.1063/1.1699114](https://doi.org/10.1063/1.1699114).
- 3250
- 3251 Murphy, K. P. (2022). *Probabilistic Machine Learning: An introduction*. MIT Press. URL:
3252 probml.ai.
- 3253 Nadjahi, K. (Nov. 2021). “Sliced-Wasserstein distance for large-scale machine learning
3254 : theory, methodology and extensions”. PhD thesis. Institut Polytechnique de Paris.
3255 URL: <https://theses.hal.science/tel-03533097>.
- 3256 Nelsen, R.B. (2006). *An introduction to copulas*. en. second. Springer series in statistics.
3257 New York Berlin Heidelberg: Springer. ISBN: 978-0-387-28659-4.
- 3258 Nichol, Alex and Prafulla Dhariwal (Feb. 2021). *Improved Denoising Diffusion Probabilistic
3259 Models*. DOI: [10.48550/arXiv.2102.09672](https://doi.org/10.48550/arXiv.2102.09672). arXiv: [2102.09672 \[cs\]](https://arxiv.org/abs/2102.09672). (Visited on
3260 10/27/2025).
- 3261 Nikoloulopoulos, Aristidis K., Harry Joe, and Haijun Li (2009). “Extreme value properties
3262 of multivariate t copulas”. In: *Extremes* 12.2, pp. 129–148. DOI: [10.1007/s10687-008-0072-4](https://doi.org/10.1007/s10687-008-0072-4).
- 3263
- 3264 Ouyang, Long, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin,
3265 Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. (2022). *Train-
3266 ing language models to follow instructions with human feedback*. NeurIPS 2022. arXiv:
3267 [2203.02155](https://arxiv.org/abs/2203.02155).
- 3268 Papamakarios, G., E. Nalisnick, D. J. Rezende, S. Mohamed, and B. Lakshminarayanan
3269 (2021). “Normalizing flows for probabilistic modeling and inference”. In: *The Journal
3270 of Machine Learning Research* 22.1, pp. 2617–2680.
- 3271 Prandini, M. and O. J. Watkins (2005). “Probabilistic Aircraft Conflict Detection”. In:
3272 *HYBRIDGE WP3: Reachability analysis for probabilistic hybrid systems*.
- 3273 Resnick, S. I. (1987). *Extreme Values, Regular Variation and Point Processes*. en. Springer
3274 Series in Operations Research and Financial Engineering. New York, NY: Springer.
3275 ISBN: 978-0-387-75952-4. DOI: [10.1007/978-0-387-75953-1](https://doi.org/10.1007/978-0-387-75953-1). URL: <http://link.springer.com/10.1007/978-0-387-75953-1> (visited on 05/16/2023).
- 3276
- 3277 Resnick, Sidney I. (1987). *Extreme Values, Regular Variation and Point Processes*. Springer
3278 Series in Operations Research and Financial Engineering. New York: Springer. DOI:
3279 [10.1007/978-0-387-75953-1](https://doi.org/10.1007/978-0-387-75953-1).
- 3280 Rezende, Danilo Jimenez and Shakir Mohamed (2015). “Variational Inference with Nor-
3281 malizing Flows”. In: *Proceedings of the 32nd International Conference on Machine
3282 Learning*. PMLR, pp. 1530–1538.
- 3283 Robert, Christian P. and George Casella (2004). *Monte Carlo Statistical Methods*. 2nd.
3284 New York: Springer. DOI: [10.1007/978-1-4757-4145-2](https://doi.org/10.1007/978-1-4757-4145-2).
- 3285 Robert, P. (2003). *Stochastic networks and queues*. Vol. 52. Applications of Mathematics.
3286 Springer-Verlag, Berlin.

- 3287 Rosenblatt, Murray (Sept. 1952). “Remarks on a Multivariate Transformation”. In: *The*
3288 *Annals of Mathematical Statistics* 23.3, pp. 470–472. ISSN: 0003-4851. DOI: [10.1214/aoms/1177729394](https://doi.org/10.1214/aoms/1177729394). (Visited on 07/21/2025).
- 3289 Sabourin, Anne, Philippe Naveau, and Anne-Laure Fougères (2013). “Bayesian Dirichlet
3290 mixture model for multivariate extremes: A re-parametrization”. In: *Computational*
3291 *Statistics & Data Analysis* 71, pp. 542–567. DOI: [10.1016/j.csda.2013.04.021](https://doi.org/10.1016/j.csda.2013.04.021).
- 3292 Schucany, William R. (1972). “Order statistics in simulation”. In: *Journal of Statistical*
3293 *Computation and Simulation* 1.3, pp. 281–286. DOI: [10.1080/00949657208810024](https://doi.org/10.1080/00949657208810024).
- 3294 Scott, David W. (2015). *Multivariate Density Estimation: Theory, Practice, and Visualiza-*
3295 *tion*. 2nd. Hoboken, NJ: Wiley. DOI: [10.1002/9781118575574](https://doi.org/10.1002/9781118575574).
- 3296 Silverman, Bernard W. (1986). *Density Estimation for Statistics and Data Analysis*. Lon-
3297 don: Chapman and Hall. DOI: [10.1007/978-1-4899-3324-9](https://doi.org/10.1007/978-1-4899-3324-9).
- 3298 Sklar, M. (1959). “Fonctions de répartition à N dimensions et leurs marges”. fr. In: *Pub-*
3299 *lications de lInstitut Statistique de lUniversité de Paris* 8, pp. 229–231.
- 3300 Sohl-Dickstein, J., E. Weiss, N. Maheswaranathan, and S. Ganguli (2015). “Deep unsuper-
3301 *vised learning using nonequilibrium thermodynamics*”. In: *International conference on*
3302 *machine learning*. PMLR, pp. 2256–2265.
- 3303 Sohl-Dickstein, Jascha, Eric A. Weiss, Niru Maheswaranathan, and Surya Ganguli (Nov.
3304 2015). *Deep Unsupervised Learning using Nonequilibrium Thermodynamics*. DOI: [10.48550/arXiv.1503.03585](https://doi.org/10.48550/arXiv.1503.03585) [cs]. (Visited on 09/22/2025).
- 3305 Song, Jiaming, Chenlin Meng, and Stefano Ermon (2021). “Denoising Diffusion Implicit
3306 *Models*”. In: *International Conference on Learning Representations*.
- 3307 Song, Yang and Stefano Ermon (2019). “Generative Modeling by Estimating Gradients of
3308 the Data Distribution”. In: *Advances in Neural Information Processing Systems*. Ed. by
3309 H. Wallach, H. Larochelle, A. Beygelzimer, F. d’Alché-Buc, E. Fox, and R. Garnett.
3310 Vol. 32. Curran Associates, Inc.
- 3311 Song, Yang, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon,
3312 and Ben Poole (2021). “Score-Based Generative Modeling through Stochastic Differen-
3313 *tial Equations*”. In: *International Conference on Learning Representations*.
- 3314 Sukhatme, Pandurang V. (1937). “Tests of significance for samples of the χ^2 -population
3315 with two degrees of freedom”. In: *Annals of Eugenics* 8.1, pp. 52–56.
- 3316 Tancik, Matthew, Pratul P. Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Ragha-
3317 van, Utkarsh Singhal, Ravi Ramamoorthi, Jonathan T. Barron, and Ren Ng (June
3318 2020). *Fourier Features Let Networks Learn High Frequency Functions in Low Dimen-*
3319 *sional Domains*. DOI: [10.48550/arXiv.2006.10739](https://doi.org/10.48550/arXiv.2006.10739) [cs]. (Visited
3320 on 10/27/2025).
- 3321 Tang, Wenpin (Mar. 2024). *Fine-tuning of diffusion models via stochastic control: entropy*
3322 *regularization and beyond*. DOI: [10.48550/arXiv.2403.06279](https://doi.org/10.48550/arXiv.2403.06279). arXiv: [2403.06279](https://arxiv.org/abs/2403.06279)
3323 [math]. (Visited on 09/22/2025).
- 3324 Uehara, Masatoshi, Yulai Zhao, Kevin Black, Ehsan Hajiramezanali, Gabriele Scalia, Nathaniel
3325 Lee Diamant, Alex M. Tseng, Tommaso Biancalani, and Sergey Levine (Feb. 2024).

- 3328 *Fine-Tuning of Continuous-Time Diffusion Models as Entropy-Regularized Control.*
3329 DOI: [10.48550/arXiv.2402.15194](https://doi.org/10.48550/arXiv.2402.15194). arXiv: [2402.15194 \[cs\]](https://arxiv.org/abs/2402.15194). (Visited on 09/22/2025).
- 3330 Villani, C. (2009). *Optimal transport: old and new*. Vol. 338. Springer.
- 3331 Vincent, Pascal (2011). “A Connection Between Score Matching and Denoising Autoen-
3332 coders”. In: *Neural Computation* 23.7, pp. 1661–1674. DOI: [10.1162/NECO_a_00142](https://doi.org/10.1162/NECO_a_00142).
- 3333 Wei, Xiang, Zixia Liu, Liqiang Wang, and Boqing Gong (2018). “Improving the improved
3334 training of wasserstein GANs”. In: *International conference on learning representa-
3335 tions*. Comment: NIPS camera-ready. URL: <https://openreview.net/forum?id=SJx9GQb0->
- 3336 Wiese, Magnus, Robert Knobloch, Ralf Korn, and Peter Kretschmer (2019). *Deep Hedging:
3337 Learning to Simulate Equity Option Markets*. arXiv: [1911.01700](https://arxiv.org/abs/1911.01700).
- 3338 — (2020). “Quant GANs: Deep generation of financial time series”. In: *Quantitative Fi-
3339 nance* 20.9, pp. 1419–1440.
- 3340 Wu, J., X.-Y. Chen, H. Zhang, L.-D. Xiong, H. Lei, and S.-H. Deng (2019). “Hyperparam-
3341 eter Optimization for Machine Learning Models Based on Bayesian Optimization”. en.
3342 In: 17.1, pp. 26–40.
- 3343 Xu, Jiazheng, Xiao Liu, Yuchen Wu, Yuxuan Tong, Qinkai Li, Ming Ding, Jie Tang, and
3344 Yuxiao Dong (Dec. 2023). *ImageReward: Learning and Evaluating Human Preferences
3345 for Text-to-Image Generation*. DOI: [10.48550/arXiv.2304.05977](https://doi.org/10.48550/arXiv.2304.05977). arXiv: [2304.05977
3346 \[cs\]](https://arxiv.org/abs/2304.05977). (Visited on 04/28/2025).
- 3347 Xu, Jing, Andrew Liu, Andrea Wu, Parisa Saeedi, Jennifer Neville, Mona Diab, and Asli
3348 Celikyilmaz (2024). *ReFL: Reflective Feedback Learning for Language Model Alignment*.
3349 arXiv: [2403.01673](https://arxiv.org/abs/2403.01673).
- 3350 Yarotsky, Dmitry (2017). “Error bounds for approximations with deep ReLU networks”.
3351 In: *Neural Networks* 94, pp. 103–114. DOI: [10.1016/j.neunet.2017.07.002](https://doi.org/10.1016/j.neunet.2017.07.002).
- 3352 Yoon, Sungbin, Juho Park, Jaehyeon Lee, and Gunhee Kim (2023). *Score-based Genera-
3353 tive Modeling through Stochastic Evolution Equations in Hilbert Spaces*. Heavy-tailed
3354 diffusion models using Lévy processes. arXiv: [2305.19241](https://arxiv.org/abs/2305.19241).
- 3355 Zhang, Likun, Andrew Bray, and Joseph Guinness (2024). “Flexible and efficient spatial
3356 extremes emulation via variational autoencoders”. In: *arXiv preprint arXiv:2407.02868*.
- 3357 Zscheischler, J., O. Martius, S. Westra, E. Bevacqua, C. Raymond, R. M. Horton, B. van
3358 den Hurk, A. AghaKouchak, A. Jézéquel, M. D. Mahecha, D. Maraun, A. M. Ramos,
3359 N. N. Ridder, W. Thiery, and E. Vignotto (2020). “A typology of compound weather
3360 and climate events”. In: *Nature Reviews Earth & Environment* 1.7, pp. 333–347.
- 3361 Zscheischler, Jakob, Seth Westra, Bart J. J. M. van den Hurk, Sonia I. Seneviratne, Philip J.
3362 Ward, Andy Pitman, Amir AghaKouchak, David N. Bresch, Michael Leonard, Thomas
3363 Wahl, and Xuebin Zhang (2018). “Future climate risk from compound events”. In:
3364 *Nature Climate Change* 8, pp. 469–477. DOI: [10.1038/s41558-018-0156-3](https://doi.org/10.1038/s41558-018-0156-3).
- 3365

³³⁶⁶ Supplementary Material

³³⁶⁷ .1 Additional Figures

³³⁶⁸ .2 Additional Tables

³³⁶⁹ .3 Research Instruments